

Exploration of Shared SNPs in Thaps

Chr1-qfiltered

June 29, 2017

Rambling exploration of SNP positions shared between two or more of the isolates. Code is included to document it thoroughly, (even if largely uninteresting to anyone else), and I will summarize it as I go.

Contents

1 History	2
2 Preliminaries	2
3 Major Analysis/Performance Parameters.	2
4 Refined SNP Calls	6
4.1 Method	6
4.2 Save them	8
4.3 Examples: Consistent	10
4.4 Examples: Inconsistent	11
4.5 Examples: Homozygous nonref	12
5 Table 1 stats	16
5.1 Table 1 Data	20
6 Shared-SNPs P-Value	20
6.1 SNP Concordance	20
6.2 Notes	23
6.3 P-Value: The Bottom Line	24
7 Sharing	24
7.1 Code	25
7.2 Sanity Checks	26
7.3 Main Analysis	31
8 Trees	35
9 Semi-Automated Tree-Building	45
9.1 Bootstrap	51
10 Notes	54
11 Appendix: Old Trees, etc.	55
11.1 HWE Sharing	55
11.2 Old Tree Stuff	56

1 History

This was added to SVN 1/26/2014; not sure when it was started, but earliest related emails I see are from 1/21/14.

```
r413 | ruzzo | 2014-01-26 08:22:37 -0800 (Sun, 26 Jan 2014) | 2 lines
adding shared-snp analysis.
```

2 Preliminaries

NOTE: Some comments in code and some parts of the text, especially specific numbers and general conclusions, are based on Unqfiltered, Chr1, Medium stringency (i.e., “[[2]]” below) analysis. The broad picture does not appear to change with other choices, but details do, and the text is neither fully parameterized nor fully updated, so proceed with caution.

Load utility R code; do setup:

```
source('.././../R/wlr.R') # load util code; path relative this folder or sibling in scripts/larrys

## Running as: ruzzo @ bicycle.cs.washington.edu; SVN Id, I miss you. $Id: wlr.R 2017-06-26 or later $

setup.my.wd('shared-snps') # set working dir; UPDATE if this file moves, or if COPY/PASTE to new file
setup.my.knitr('f-knitr/')
generic.setup('figs-mine/')
```

3 Major Analysis/Performance Parameters.

Choices here control how this file is processed, what data is analyzed, speed, etc. Set them carefully before running “make.” Major choices are:

1. WHICH SNP TABLES ARE LOADED??? The logical vector `load.tb` selects the desired combination of SNP tables to load, in the order `full.unfiltered`, `chr1.unfiltered`, `full.qfiltered`, `chr1.qfiltered`. E.g., `load.tb=(T, F, T, F)` loads *full* tables for *both* q- and un-qfiltered data. Primary analysis is only performed on one of them, but the others are retained for comparison/debugging.
2. WHICH MAIN ANALYSIS??? If multiple tables are loaded, which is used for the main analysis? Parameter `pri` is a permutation of 1:4, corresponding to `load.tb`; the first loaded table in that order becomes the analysis focus. The default `pri=c(1, 2, 3, 4)` looks at un-q-filtered data in preference to q-filtered, and full tables in preference to Chr1 within each group.
(Choice of data for the “Table 1” coverage summary in section 5 is independent of this; full genome data is preferred over Chr 1 for both q- and unq-filtered reads; change `tset.picker` calls near the end of that section to modify this.)
3. CLEAR CACHE??? `clear.cache=T` forces Knitr cache removal at the start of the run; especially important if the previous parameters have changed since the last run.
4. HOW MANY BOOTSTRAP REPLICATES??? The variable `nboot` is a major performance factor; 1000 reps takes several hours. Set to 5 for debug and quick look; 100 or more for final run.
5. TRUNCATE TABLES TO Chrs ONLY??? I.e., remove mitochondrial-, plastid-, and BD- contigs.

The following code chunk sets the first four parameters based on where it’s run. To prototype/debug on a laptop, faster is better—run on Chr1 with small `nboot`; when run on the linux servers, I typically do full genomes, more replicates. Just override them if these defaults don’t work for you.

```
# for Makefile, params can be command line args, else base on system; see wlr.r for details.
# load.tb order: full.un, chr1.un, full.qfil, chr1.qfil

params <- pick.params(
  mac = list(load.tb=c(F,T,F,F), pri=1:4, clear.cache=F, nboot= 1, trunc.tables=T), # quick on lap
  linux = list(load.tb=c(F,F,F,T), pri=1:4, clear.cache=F, nboot= 5, trunc.tables=T), # quick qfil on server
  linux = list(load.tb=c(T,F,T,F), pri=1:4, clear.cache=T, nboot=101, trunc.tables=T) # full on server
)

# Alternatively, edit/uncomment the following to override the above as needed
#params<-pick.params(default=list(load.tb=c(T,T,T,T),pri=1:4,clear.cache=T,nboot=1000,trunc.tables=T))
print(params)

# $load.tb
# full.unf chr1.unf full.qf chr1.qf
# FALSE TRUE FALSE TRUE
#
# $pri
# [1] 3 4 1 2
#
# $clear.cache
# [1] TRUE
#
# $nboot
# [1] 5
#
# $trunc.tables
# [1] FALSE
```

CLEAR CACHE??!! Some code chunks use the knitr cache, but extent of cache consistency checks unknown. If in doubt, delete “cache/” (knitr’s) directory to force rebuild. T/F set in params above will/won’t force removal (actually, rename):

```
decache(params$clear.cache)

# Rename of 'cache' to 'cache81575' returned TRUE .
```

If still in doubt, also manually remove “00common/mycache/” (mine).

Load the main SNP data file(s) based on the parameters set in section 3.

```
# short names to keep the following chunk compact
tb <- params$load.tb
tset <- list(NULL, NULL, NULL, NULL) # tset = 'table set'

# see wlr.R for load paths
if(tb[1]){tset[[1]] <- load.snp.tables(use.chr1.tables = FALSE, data.name='full.tables.01.26.14')}
if(tb[2]){tset[[2]] <- load.snp.tables(use.chr1.tables = TRUE , data.name='full.tables.01.26.14')}

# Loading ../00common/mycache/snp.tables.chr1.unqfiltered.rda ...Loaded.

if(tb[3]){tset[[3]] <- load.snp.tables(use.chr1.tables = FALSE, data.name='full.tables.02.25.15')}
if(tb[4]){tset[[4]] <- load.snp.tables(use.chr1.tables = TRUE , data.name='full.tables.02.25.15')}

# Loading ../00common/mycache/snp.tables.chr1.qfiltered.rda ...Loaded.
# Bandaidding qfiltered tables...
```

Grrr! I should have excluded non-Chr contigs from full genome runs. Rather than change tons of code below to add mask params, I’m just going to truncate the tables, as follows. (See notes in wlr.r::make.mask for assumptions.)

```
if(params$trunc.tables){
  for(i in 1:4){
    if(!is.null(tset[[i]])){
      first.mito <- match("mitochondria.fasta", tset[[i]][[7]]$Chr)
      if(!is.na(first.mito)){ # will be NA for Chr1 tables
```

```

    for(j in 1:7){
      # hmmm... slow; wonder whether head(tset[[i]][[j]],first.mito-1) is faster;
      # ok, simple tests suggest not: system.time(head(data.frame(1:1e7,1:1e7),5e6))
      tset[[i]][[j]] <- tset[[i]][[j]][1:(first.mito-1),]
    }
  }
}
} else {
  cat('***\n*** DID YOU *REALLY* WANT UNTRUNCATED TABLES???\n***\n')
}

# ***
# *** DID YOU *REALLY* WANT UNTRUNCATED TABLES???
# ***

```

The tersely-named `tset` list is sometimes convenient, but give them more descriptive names, too.

```

snp.tables.full.unfiltered <- tset[[1]]; names(tset)[1] <- 'snp.tables.full.unfiltered'
snp.tables.chr1.unfiltered <- tset[[2]]; names(tset)[2] <- 'snp.tables.chr1.unfiltered'
snp.tables.full.qfiltered <- tset[[3]]; names(tset)[3] <- 'snp.tables.full.qfiltered'
snp.tables.chr1.qfiltered <- tset[[4]]; names(tset)[4] <- 'snp.tables.chr1.qfiltered'

```

Main analysis may just use one of the potentially 4 table sets. Pick it according to the priority specified in section 3, using the shorter name `'snp.tables'` for this default choice.

```
snp.tables <- tset.picker(priority=params$pri, table.set=tset)
```

```

# Sanity check: unlike unfiltered tables, bug in early code gave qfiltered ones different numbers
# of rows per strain, which breaks much code. Verify this is no longer happening.
check.eq.nrows <- function(tables){
  if(!is.null(tables)){
    nrow.snp.tables <- unlist(lapply(tables,nrow))
    print(nrow.snp.tables)
    if(all(nrow.snp.tables == nrow.snp.tables[1])){
      cat('OK, all strains have same number of rows.\n')
    } else {
      cat('***\n*** Warning: Different strains have different numbers of rows! ***\n***\n')
    }
  }
}

dummy<-lapply(tset, check.eq.nrows)

#      1007      1012      1013      1014      1015      3367      1335
# 3042585 3042585 3042585 3042585 3042585 3042585 3042585
# OK, all strains have same number of rows.
#      1007      1012      1013      1014      1015      3367      1335
# 3042585 3042585 3042585 3042585 3042585 3042585 3042585
# OK, all strains have same number of rows.

```

Which tables have we got?:

```

# 'which.snp.tables' return summary of which tables, either as a char string (default), e.g.
# "Chr1-qfiltered", or as vector of 2 strings, e.g. c("full","unfiltered").
cat('This analysis uses: (', paste(unlist(lapply(tset,which.snp.tables)),collapse=', '), ') SNP tables.\n')

# This analysis uses: ( NULL, Chr1-unfiltered, NULL, Chr1-qfiltered ) SNP tables.

cat('Main shared SNP analysis focuses on', which.snp.tables(snp.tables), '\n')

# Main shared SNP analysis focuses on Chr1-qfiltered

```

A \LaTeX hack: I want `which.snp.tables` info in doc title/page headers, but it is unknown until now, so the following writes a command definition `\whichsnptables` into the `.aux` file, which is read during the *next* \LaTeX run, when `\begin{document}` is processed:

```
\makeatletter
\immediate\write\@auxout{\noexpand\gdef\noexpand\whichsnptables{Chr1-qfiltered}}
\makeatother
```

Subsequent analysis was initially all directed at Chr1. In general, I have *not* updated the discussion to reflect genome-wide analysis.

```
if(exists('snp.tables.chr1.qfiltered') && exists('snp.tables.chr1.unqfiltered')){
  # If have both, where is new unequal to old?
  uneq <- snp.tables.chr1.qfiltered[[1]]$Ref[1:chr1.len] != snp.tables.chr1.unqfiltered[[1]]$Ref[1:chr1.len]
  cat('Sum uneq:', sum(uneq, na.rm=T), '\n')
  cat('Sum NA: ', sum(is.na(uneq)), '\n')
  print(which(is.na(uneq)) [1:10])
  seecounts(which(is.na(uneq)) [1:4], snp.tables=snp.tables.qfiltered, debug=F)
}
```

In brief, “snp.tables” will be a list of 7 data frames, one per strain, giving read counts for each nucleotide at each position, SNP calls, etc.:

```
names(snp.tables)

# [1] "1007" "1012" "1013" "1014" "1015" "3367" "1335"

str(snp.tables[[1]])

# 'data.frame': 3042585 obs. of 16 variables:
# $ snp : int 0 0 0 0 0 0 0 0 0 0 ...
# $ Chr : chr "Chr1" "Chr1" "Chr1" "Chr1" ...
# $ Pos : int 1 2 3 4 5 6 7 8 9 10 ...
# $ Ref : chr "T" "C" "C" "A" ...
# $ Cov : num 0 2 3 4 4 4 7 8 9 10 ...
# $ a : num 0 0 0 0 0 0 0 0 0 0 ...
# $ g : num 0 0 0 0 0 0 0 0 0 0 ...
# $ c : num 0 0 0 0 0 0 0 0 0 0 ...
# $ t : num 0 0 0 0 0 0 0 0 0 0 ...
# $ n : num 0 0 0 0 0 0 0 0 0 0 ...
# $ .match: num 0 2 3 4 4 4 7 8 9 10 ...
# $ exon : logi FALSE FALSE FALSE FALSE FALSE FALSE ...
# $ indel : logi FALSE FALSE FALSE FALSE FALSE FALSE ...
# $ chr : Factor w/ 1 level "Chr1": 1 1 1 1 1 1 1 1 1 1 ...
# $ pos : int 1 2 3 4 5 6 7 8 9 10 ...
# $ rawCov: num 1 3 4 5 7 7 10 12 13 15 ...
```

Just for background, also load the desert tables:

```
# from svn+ssh://cegl.ocean.washington.edu/var/svn/7_strains/trunk/code/snpNB/data
#load('.../data/ungit-data/des.rda')
load('.../data/des.rda')
```

What’s the total length of all deserts in each strain? Big deserts (defined as “big.threshold” or longer)?

```
some.desert.stats <- function(big.threshold=0){
  desert.len <- unlist(lapply(des, function(x){sum(unlist(lapply(x, function(y){sum(y[, 'Length'])}))})))
  bigdes.len <- unlist(lapply(des, function(x){sum(unlist(lapply(x, function(y){
    sum(y[y[, 'Length']>=big.threshold, 'Length'])}))})))
  rbind(desert.len, desert.pct=round( desert.len / genome.length.constants()$genome.length.trunc * 100),
        bigdes.len, bigdes.pct=round( bigdes.len / genome.length.constants()$genome.length.trunc * 100))
}
some.desert.stats(big.threshold=50000)
```

	tp1007	tp1012	tp1013	tp1014	tp1015	thapsIT	tp1335
# desert.len	11146526	11332566	5801763	9464213	11251426	6780300	10883723
# desert.pct	36	36	19	30	36	22	35
# bigdes.len	3495805	3936973	55365	3627235	3727061	57119	4046934
# bigdes.pct	11	13	0	12	12	0	13

I.e., looking at all deserts, about 1/3 of L-clade, 1/5 of H-clade are in deserts, whereas, looking at the largest deserts ($> 50k$), only about 12% in L-clade (and none in H-clade). Note that the rough stats above include artifactual “deserts” created by gaps in the reference sequence, large genomic deletions, etc. A more careful analysis of this is found in nc-snp.rnw.

4 Refined SNP Calls

4.1 Method

It is appropriate that SNP calls should be conservative, to avoid many false positives, but, when a position is called a SNP in one isolate, we often see a significant number of reads for the same non-reference nucleotide at that position in other isolates, even if they are not called as SNPs. On the other hand, we sometimes see a position called a SNP in two or more isolates, but with *different* pairs of nucleotides, potentially suggesting technical errors. Analysis in this section attempts to refine the SNP calls by looking for issues such as these by looking at all 7 isolates jointly, at each position called a SNP in any of them.

For a given strain, the following function returns a vector of 0:4 to indicate which nonreference nucleotide has the maximum read count at the corresponding position. The values 1:4 indicate that the max count occurred at A, G, C, T, resp. (Ties are resolved arbitrarily ($a < g < c < t$), which possibly deserves further attention.) The value 0 means all nonreference counts are below threshold, based *either* on absolute count *or* as a fraction of coverage. Default only excludes 0 counts.

```
nref.nuc.new <- function(strain=1, mask=T, thresh.count=0, thresh.rate=0.0){
  # get read count for max nonref nuc
  nref <- apply(snp.tables[[strain]][mask, c('a', 'g', 'c', 't')], 1, max)
  # where does nref count match a (g,c,t, resp) count
  as <- ifelse(nref == snp.tables[[strain]][mask, 'a'], 1, 0)
  gs <- ifelse(nref == snp.tables[[strain]][mask, 'g'], 2, 0)
  cs <- ifelse(nref == snp.tables[[strain]][mask, 'c'], 3, 0)
  ts <- ifelse(nref == snp.tables[[strain]][mask, 't'], 4, 0)
  # most positions will show 3 zeros and one of 1:4, so max identifies max nonref count;
  # ties broken arbitrarily (a<g<c<t)
  merge <- pmax(as, gs, cs, ts)
  # but if max nonref count is zero or below threshold, return 0
  merge[nref == 0 | nref < thresh.count] <- 0
  merge[nref/snp.tables[[strain]][mask, 'Cov'] < thresh.rate] <- 0
  return(merge)
}
```

Get union and intersection of the sets of called SNPs. (“\$snp” is 0/1.) Also, 5-way (L-clade) and 4-way (L-excluding Gyre).

```
# 4-way union/intersection
u4.snps <- snp.tables[[1]]$snp
i4.snps <- snp.tables[[1]]$snp
for(i in c(2,5,7)) {
  u4.snps <- pmax(u4.snps, snp.tables[[i]]$snp)
  i4.snps <- pmin(i4.snps, snp.tables[[i]]$snp)
}
# 5-way: add gyre
u5.snps <- pmax(u4.snps, snp.tables[[4]]$snp)
i5.snps <- pmin(i4.snps, snp.tables[[4]]$snp)
# 7-way
union.snps <- pmax(u5.snps, snp.tables[[3]]$snp, snp.tables[[6]]$snp)
intersect.snps <- pmin(i5.snps, snp.tables[[3]]$snp, snp.tables[[6]]$snp)
nu4snps <- sum(u4.snps)
nu5snps <- sum(u5.snps)
ni4snps <- sum(i4.snps)
ni5snps <- sum(i5.snps)
nusnps <- sum(union.snps)
nisnps <- sum(intersect.snps)
c(n4u=nu4snps, n5u=nu5snps, n7u=nusnps, n4i=ni4snps, n5i=ni5snps, n7i=nisnps)
```

```
#   n4u   n5u   n7u   n4i   n5i   n7i
# 18564 18696 47499 14365 7628 1641
```

There are nusnps=47499 positions called as SNPs in one or more strains (but only nisnps=1641 that are shared among all 7). Note that the 4-way union is only modestly larger (1.2923077 times larger) than the 4-way intersection, emphasizing the inherent similarities among these SNP sets. The corresponding 5-way numbers show that Gyre adds relatively little to the 5-way union vs the 4-way union, whereas it removes a fair bit from the 5-way intersection. However, much of that loss is simply because Gyre has fewer called SNPs: only 8331 vs 14365 in the 4-way intersection, and they are highly concordant:

```
sum(snp.tables[[4]]$snp*i4.snps)/sum(snp.tables[[4]]$snp)

# [1] 0.9156164
```

So, a likely source of the Gyre's difference in called SNPs is technical (lower read coverage, higher read error rate) rather than biological.

Inclusion of the 2 H-clade members, however, causes more dramatic changes in both union and intersection numbers. I examine all these relationships in more detail below, but first I examine what I believe to be a significant source of technical error in these comparisons—erroneous SNP calls, especially false negative calls.

It is appropriate that SNP calls should be conservative, to avoid many false positives, but, if a position is called a SNP in one strain, we often see a significant number of reads for the same non-reference nucleotide at that position in other strains, even if they are not called as SNPs. For my purposes below, these will be considered “shared SNPs,” based on three different levels of permissiveness. Note that, e.g., $\geq 84\%$ of all positions have zero reads for any non-reference nucleotide, and only a small fraction have 2 or more non-reference reads:

```
nonmatch <- rbind(
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 0)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 1)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 2)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 3)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match >= 4)})),
  unlist(lapply(snp.tables,function(x){sum((x$Cov-x$.match)[union.snps==0] >= 4)}))
)/nrow(snp.tables[[1]])*100
rownames(nonmatch) <- c('% ==0', '% ==1', '% ==2', '% ==3', '% >=4', '% >=4, nonSNP')
nonmatch
```

#	1007	1012	1013	1014	1015	3367	1335
# % ==0	97.79158183	97.46508314	95.61527451	97.40214324	97.30475895	96.0503322	96.61320883
# % ==1	1.45497989	1.70969751	2.91324647	2.03304098	1.83551158	2.5290008	2.48311880
# % ==2	0.09754863	0.10606113	0.19447279	0.17521285	0.11904351	0.1896085	0.18402115
# % ==3	0.05166659	0.04680231	0.08019497	0.09935630	0.05324420	0.0796691	0.05416447
# % >=4	0.60422305	0.67235591	1.19681126	0.29024662	0.68744176	1.1513894	0.66548675
# % >=4, nonSNP	0.04486317	0.07924183	0.19759514	0.02635916	0.08581519	0.1951959	0.08811586

Build a table of max non-reference nucleotides at each position in the union.snps set. The three criteria are

- [[1]]: any non-zero count at any coverage is considered significant
- [[2]]: (count ≥ 2 and count/coverage ≥ 0.05) is considered significant
- [[3]]: (count ≥ 4 and count/coverage ≥ 0.10) is considered significant

In all three cases, the nonref nucleotide must also be consistent across all strains passing that threshold; see below.

```
non.refs <- vector('list',4)
for(i in 1:4){
  non.refs[[i]] <- matrix(0, nrow=nusnps, ncol=7)
  colnames(non.refs[[i]]) <- names(snp.tables)
  rownames(non.refs[[i]]) <-
    paste(snp.tables[[1]]$chr[union.snps==1], ': ', snp.tables[[1]]$pos[union.snps==1], sep='')
}
for(j in 1:7){
```

```

non.refs[[1]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=0, thresh.rate=0.00)
non.refs[[2]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=2, thresh.rate=0.05)
non.refs[[3]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=4, thresh.rate=0.10)
}

```

For comparison, I want to look at unfiltered SAMTools SNP calls. In complete opposition to the measures of consistency imposed above, I'm going to simply force this into the “non.refs” structure constructed above by imagining that any position called a SNP in any strain has its max nonref count on “A”, so any given position called a SNP in any strain will automatically be declared “consistent.” This will allow the tree-code, etc. given below to work in a uniform way (even though interpretation of the results is different.) Results will be jammed into a 4th component of the “non.refs” list; i.e., we have a 4th criterion:

- [[4]]: all called SNPs at a given position are considered “consistent.”

As this case was a late addition to the analysis, the commentary throughout this document has not necessarily been updated to reflect that this case is distinct from the first three.

```

for(j in 1:7){
  non.refs[[4]][,j] <- snp.tables[[j]]$snp[union.snps==1]
}

```

```

str(non.refs[[4]])

# num [1:47499, 1:7] 0 0 0 0 0 0 0 0 1 0 ...
# - attr(*, "dimnames")=List of 2
# ..$ : chr [1:47499] "Chr1:333" "Chr1:417" "Chr1:435" "Chr1:438" ...
# ..$ : chr [1:7] "1007" "1012" "1013" "1014" ...

```

“non.refs” indicates, among those positions in the union of all called SNPs having any non-reference read count above the thresholds listed above, the non-ref nucleotide having the highest read count in each strain. If, for a given position, the max of this code is the same as the min (among non-zero values), then every strain having over-threshold nonref reads in that position, in fact has most non-reference reads on the *same* nucleotide. These are defined as the “consistent” SNPs.

```

find.consistent <- function(nr){
  nr.max <- apply(nr,1,max)
  nr.min <- apply(nr,1,function(x){ifelse(max(x)==0,0,min(x[x>0]))})
  return(nr.min == nr.max)
}
consistent <- lapply(non.refs, find.consistent)

```

4.2 Save them

```

# wrap this in a data structure to be cached:
Description <- [2757 chars quoted with '']

filtered.snps <-
  list(Description=Description,

        Data=list(
          based.on.which.snp.tables=which.snp.tables(),
          number.union.snps=nusnps,
          number.intersection.snps=nisnps,
          non.ref.nucleotide=non.refs,
          consistent.snps=consistent),

        Code=list(
          get.snps = function(strain, stringency=2){
            # return nusnps x 1 Bool vector of consistent SNPs @ specified stringency & strain
            return(filtered.snps$Data$consistent.snps[[stringency]] &
              filtered.snps$Data$non.ref.nucleotide[[stringency]][,strain] > 0)
          }
        )

```



```

    }
    ,
    get.snp.locs.char = function(strain, stringency=2){
      # return char vector of locations of consistent SNPs @ specified stringency & strain
      snps <- filtered.snps$Code$get.snps(strain, stringency)
      return(names(snps)[snps])
    }
    ,
    get.snp.locs.df = function(strain, stringency=2){
      # return data frame (Chr/Pos) of locations of consistent SNPs @ specified stringency & strain
      snplist <- strsplit(filtered.snps$Code$get.snp.locs.char(strain, stringency), ':', fixed=TRUE)
      # strsplit returns long list of 2-vectors, 1st=chr, 2nd=char position
      df <- data.frame(Chr=      unlist(lapply(snplist,function(x){return(x[1])})),
                      Pos=as.integer(unlist(lapply(snplist,function(x){return(x[2])}))),
                      stringsAsFactors = FALSE)

      return(df)
    }
  )
)

# dont't clobber existing .rda, but save if absent. (delete to re-save)
rda.filtered <- paste(' ../00common/mycache/filtered.snps', which.snp.tables(), 'rda', sep='.')
if(file.exists(rda.filtered)){
  cat('Pre-existing file', rda.filtered, 'unchanged.\n')
} else {
  cat('Saving', rda.filtered, '...\n')
  save(filtered.snps, file=rda.filtered, compress=TRUE)
  cat('Saved.\n')
}

# Pre-existing file ../00common/mycache/filtered.snps.Chr1-qfiltered.rda unchanged.

```

Knitr seems to be failing to format the long char string above, which says:

```

cat(filtered.snps$Description)

# Contents of this .rda file:
#
# * Description: this text
#
# * Data -- 5 items defining filtered SNPs, at 4 different stringency levels, as defined
#   in shared-snps.rnw:
#
# * based.on.which.snp.tables: {"Chr1","full","trunc"}-{"unfiltered","qfiltered"},
#   depending on which snp tables were used to build this data. ("trunc" = all Chrs.)
#
# * number.union.snps: the total number of SNPs (SAMtools calls) in the union of SNPs
#   across all 7 strains.
#
# * number.intersection.snps: similar, for the 7-way intersection.
#
#   nusnps/nisnps are easily recalculated from the data below, but their inclusion
#   may be convenient, e.g., to quickly see if the .rda represents the full genome
#   (nusnps=488848), or the chr 1 subset (nusnps=47499); (redundant with "based.on...";
#   numbers above are for unfiltered, perhaps slightly different if qfiltered)
#
# * non.ref.nucleotide: 4 arrays, each nusnps x 7, of values 0..4 (0..1 in the 4th
#   array). In the 1st 3 arrays, 0 means the given position in the given strain did
#   not have nonreference read counts above the corresponding filtering threshold,
#   i.e., is NOT a filtered SNP in that strain, whereas 1..4 mean that it did pass
#   threshold, for A,C,G,T resp. In the 4th array, this value is just 1/0,
#   indicating is/is not a called SNP in that strain.
#
# * consistent.snps: 4 Bool vectors of length nusnps flagging positions whose nonref
#   nucs (wrt to the 4 filtering criteria) are deemed *consistent* across
#   all 7 strains. For the 1st 3, this means all nonzero entries of non.ref.nuc
#   are equal, i.e., nonref read counts passing threshold are on the SAME nonref

```

```
# nucleotide in all strains having over-threshold counts. Just for comparison
# and uniformity of data structures, the 4th is all TRUE, i.e., union of SNPs
# across all strains, without any regard for thresholds or consistency.
#
# In short, the filtered SNPs according to our medium filtering criteria are
# strains/positions where consistent.snps[[2]]==TRUE and non.ref.nucleotide[[2]]>0.
#
# Rownames in both non.ref.nucs and consistent define location, e.g. "Chr1:333".
#
# * Code -- simple routines to extract filtered SNPs in (potentially) convenient formats:
#
# * get.snps(strain, stringency=2)
#   returns nusnps x 1 Bool vector of consistent SNPs @ specified stringency in
#   given strain
#
# * get.snp.locs.char(strain, stringency=2)
#   returns n x 1 char vector of locations of consistent SNPs @ specified stringency
#   in given strain, e.g. "Chr1:1234", where n == sum(get.snps(...))
#
# * get.snp.locs.df(strain, stringency=2){
#   As above, but returns data frame (char vector Chr, int vector Pos) with the same info.
```

```
str(consistent[[1]])

# Named logi [1:47499] TRUE FALSE TRUE TRUE TRUE TRUE ...
# - attr(*, "names")= chr [1:47499] "Chr1:333" "Chr1:417" "Chr1:435" "Chr1:438" ...
```

```
consistent.count <- unlist(lapply(consistent, sum)) ; consistent.count

# [1] 44905 47108 47204 47499

inconsistent.count <- consistent.count[4] - consistent.count; inconsistent.count

# [1] 2594 391 295 0

inconsistent.percent <- inconsistent.count/consistent.count[4]*100; inconsistent.percent

# [1] 5.4611676 0.8231752 0.6210657 0.0000000
```

I.e., of the 47499 positions in which a SNP is called, 44905 are consistent by my loose definition, and 47204 are consistent by my tightest definition. The increase in concordance supports the view that the loose definition is too loose. Perhaps misleadingly, these counts include positions that are “consistent SNPs” in only one strain; more below. (*TODO* I suspect, but have not yet systematically checked, that most of the rest are positions with low coverage and/or very low read counts on the mixture of non-reference nucleotides.)

4.3 Examples: Consistent

Here are a few (nonrandomly selected) prototypical consistent SNPs:

```
esnps <- names(consistent[[2]][consistent[[2]])
esnps2 <- as.integer(unlist(lapply(strsplit(esnps[c(7,11:13,92)], ':', fixed=TRUE), function(x){x[2]})))
seecounts(esnps2, snp.tables=snp.tables)
```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	567	T										
# 2				1007	0	0	1	25	0	TRUE	FALSE		
# 3				1012	0	0	14	39	1	TRUE	FALSE		
# 4				1013	0	0	13	87	0	TRUE	FALSE		
# 5				1014	0	1	0	23	0	TRUE	FALSE		
# 6				1015	0	0	8	40	1	TRUE	FALSE		
# 7				3367	0	0	16	38	1	TRUE	FALSE		
# 8				1335	0	0	2	99	0	TRUE	FALSE		
# 9	Chr1	1053	A										
# 10				1007	25	0	0	4	0	TRUE	FALSE		

```

# 11      1012 35   0 0 12   0 TRUE FALSE
# 12      1013  2   1 0 32   0 TRUE FALSE
# 13      1014  5   0 0  5   0 TRUE FALSE
# 14      1015 29   0 0 15   1 TRUE FALSE
# 15      3367  2   0 0  7   0 TRUE FALSE
# 16      1335 56   0 0 39   1 TRUE FALSE
# 17 Chr1 1055   G
# 18      1007  0  23 0  1   0 TRUE FALSE
# 19      1012  0  37 0  6   0 TRUE FALSE
# 20      1013  1  39 0  6   0 TRUE FALSE
# 21      1014  0   6 0  2   1 TRUE FALSE
# 22      1015  0  26 0 14   0 TRUE FALSE
# 23      3367  0  12 0  0   0 TRUE FALSE
# 24      1335  0  54 0 32   1 TRUE FALSE
# 25 Chr1 1176   G
# 26      1007  1  53 0  0   0 FALSE FALSE
# 27      1012  0  54 0  0   0 FALSE FALSE
# 28      1013 19  56 0  0   0 FALSE FALSE
# 29      1014  0  28 0  0   0 FALSE FALSE
# 30      1015  3  85 0  0   0 FALSE FALSE
# 31      3367  9   2 0  0   1 FALSE FALSE
# 32      1335  0 156 0  0   0 FALSE FALSE
# 33 Chr1 8685   G
# 34      1007  6  15 0  0   0 TRUE FALSE
# 35      1012 10  23 0  0   0 TRUE FALSE
# 36      1013 18  21 0  0   1 TRUE FALSE
# 37      1014  4   8 0  0   0 TRUE FALSE
# 38      1015 10  24 0  0   1 TRUE FALSE
# 39      3367  0   4 0  0   0 TRUE FALSE
# 40      1335  5  32 0  0   0 TRUE FALSE

```

4.4 Examples: Inconsistent

Here is a brief look at some *in*-consistent positions. E.g., Chr1:2013 shows nontrivial counts on 3 alleles in Wales, as do 2319, 3286, 5002, 5433, whereas 7878 shows a different alternate allele in Italy than in Wales.

```

unc <- names(consistent[[2]][!consistent[[2]]])
unc2 <- as.integer(unlist(lapply(strsplit(unc[1:10], ':', fixed=TRUE), function(x){x[2]})))
seecounts(unc2, snp.tables=snp.tables)

```

```

#      chr   pos Ref Strain  A   G   C   T SNP  exon indel nrf rat
# 1  Chr1  2013   T
# 2      1007  4   0 0 15   0 TRUE FALSE
# 3      1012  6   0 0 21   0 TRUE FALSE
# 4      1013  7  10 0  6   1 TRUE FALSE
# 5      1014  1   0 0  6   0 TRUE FALSE
# 6      1015 13   0 0 13   1 TRUE FALSE
# 7      3367  7   0 0 25   0 TRUE FALSE
# 8      1335 16   0 0 42   1 TRUE FALSE
# 9  Chr1  2319   C
# 10     1007  0  28 10  0   1 TRUE FALSE
# 11     1012  0  43 17  0   1 TRUE FALSE
# 12     1013 13  15  9  0   1 TRUE FALSE
# 13     1014  0  18  6  0   1 TRUE FALSE
# 14     1015  0  53 20  0   1 TRUE FALSE
# 15     3367  4   0 24  0   0 TRUE FALSE
# 16     1335  0 118 28  0   1 TRUE FALSE
# 17 Chr1  3286   T
# 18     1007  4   0 1 10   0 TRUE FALSE
# 19     1012  7   0 3 32   0 TRUE FALSE
# 20     1013 34   0 30  1   1 TRUE FALSE
# 21     1014  4   0 4 10   0 TRUE FALSE
# 22     1015 11   0 6 31   0 TRUE FALSE
# 23     3367  5   0 29  0   0 TRUE FALSE
# 24     1335 14   0 3 55   0 TRUE FALSE

```

```

# 25 Chr1 5002 T
# 26 1007 0 14 0 7 0 TRUE FALSE
# 27 1012 0 20 0 19 1 TRUE FALSE
# 28 1013 18 10 0 22 0 TRUE FALSE
# 29 1014 0 5 0 2 0 TRUE FALSE
# 30 1015 0 18 0 12 1 TRUE FALSE
# 31 3367 0 0 0 31 0 TRUE FALSE
# 32 1335 0 46 0 44 0 TRUE FALSE
# 33 Chr1 5178 G
# 34 1007 0 20 0 0 0 TRUE FALSE
# 35 1012 0 32 0 0 0 TRUE FALSE
# 36 1013 47 9 0 0 1 TRUE FALSE
# 37 1014 0 13 0 0 0 TRUE FALSE
# 38 1015 0 30 0 0 0 TRUE FALSE
# 39 3367 32 19 0 0 1 TRUE FALSE
# 40 1335 0 38 0 2 0 TRUE FALSE
# 41 Chr1 5433 G
# 42 1007 0 40 0 3 0 TRUE FALSE
# 43 1012 0 53 0 5 0 TRUE FALSE
# 44 1013 16 29 0 7 1 TRUE FALSE
# 45 1014 9 8 0 0 1 TRUE FALSE
# 46 1015 6 53 0 2 0 TRUE FALSE
# 47 3367 8 37 0 0 0 TRUE FALSE
# 48 1335 6 72 0 2 0 TRUE FALSE
# 49 Chr1 7858 C
# 50 1007 0 0 42 0 0 TRUE FALSE
# 51 1012 0 0 35 0 0 TRUE FALSE
# 52 1013 0 0 81 8 0 TRUE FALSE
# 53 1014 0 0 12 0 0 TRUE FALSE
# 54 1015 0 0 71 0 0 TRUE FALSE
# 55 3367 20 0 2 0 1 TRUE FALSE
# 56 1335 0 0 83 0 0 TRUE FALSE
# 57 Chr1 8974 C
# 58 1007 0 1 5 0 0 TRUE FALSE
# 59 1012 0 2 13 0 0 TRUE FALSE
# 60 1013 9 15 2 0 1 TRUE FALSE
# 61 1014 0 1 1 0 0 TRUE FALSE
# 62 1015 0 1 9 0 0 TRUE FALSE
# 63 3367 2 0 1 0 0 TRUE FALSE
# 64 1335 0 11 30 0 0 TRUE FALSE
# 65 Chr1 10099 T
# 66 1007 16 0 0 24 0 TRUE FALSE
# 67 1012 45 0 0 26 0 TRUE FALSE
# 68 1013 0 2 6 55 0 TRUE FALSE
# 69 1014 32 0 0 11 0 TRUE FALSE
# 70 1015 38 0 0 37 0 TRUE FALSE
# 71 3367 0 1 0 7 0 TRUE FALSE
# 72 1335 52 0 0 61 1 TRUE FALSE
# 73 Chr1 15154 A
# 74 1007 13 0 0 0 0 FALSE FALSE
# 75 1012 37 0 0 1 0 FALSE FALSE
# 76 1013 2 0 35 7 1 FALSE FALSE
# 77 1014 10 0 0 0 0 FALSE FALSE
# 78 1015 24 0 0 0 0 FALSE FALSE
# 79 3367 3 0 0 12 1 FALSE FALSE
# 80 1335 47 0 0 3 0 FALSE FALSE

```

4.5 Examples: Homozygous nonref

And at some *homozygous nonreference* positions (defined to be those with nonref fraction > 0.75):

```

hnr <- lapply(snp.tables, function(x){x$.match/x$Cov < 0.25}) # find them
hnr <- lapply(hnr, function(x){ifelse(is.na(x), FALSE, x)}) # remove NA
unlist(lapply(hnr, sum)) # count per strain

# 1007 1012 1013 1014 1015 3367 1335

```

```
# 316 247 12082 938 236 13863 167
```

Hmm, in L-clade, excluding the ref isolate (1335) this tracks time-in culture to some degree; Maybe many of these are in hemizygous regions. Next two chunks lifted from nc-snps to get tables for hemi-deletion.

```
cnv.chrononly <- load.cnv.tables('.././../data/cnv.txt', chrs.only=TRUE)

str(cnv.chrononly)

# 'data.frame': 1956 obs. of 11 variables:
# $ strain : Factor w/ 7 levels "IT","tp1007",...: 3 3 3 3 3 3 3 3 3 ...
# $ chr : Factor w/ 65 levels "BD1_7","BD10_65",...: 38 38 38 38 38 38 38 38 38 ...
# $ start : int 10601 112001 215001 358901 536501 554801 673401 781801 806901 853201 ...
# $ end : int 13500 116500 221100 370300 538600 559300 685000 787400 811100 855600 ...
# $ length : int 2900 4500 6100 11400 2100 4500 11600 5600 4200 2400 ...
# $ filtered : logi FALSE FALSE FALSE TRUE FALSE FALSE ...
# $ type : Factor w/ 1 level "CNVnator": 1 1 1 1 1 1 1 1 1 1 ...
# $ cov_ratio: num 0.63738 1.54893 1.65381 0.00204 0.68486 ...
# $ dup_frac : num 0.41188 0.00908 0.01178 0.97997 0.0211 ...
# $ iStart : num 10601 112001 215001 358901 536501 ...
# $ iEnd : num 13500 116500 221100 370300 538600 ...

cnv.chrononly[c(1:4,nrow(cnv.chrononly)+c(-1,0)),] ## first/last few rows

# strain chr start end length filtered type cov_ratio dup_frac iStart iEnd
# 1 tp1012 Chr1 10601 13500 2900 FALSE CNVnator 0.63738000 0.41187900 10601 13500
# 2 tp1012 Chr1 112001 116500 4500 FALSE CNVnator 1.54893000 0.00907677 112001 116500
# 3 tp1012 Chr1 215001 221100 6100 FALSE CNVnator 1.65381000 0.01178470 215001 221100
# 4 tp1012 Chr1 358901 370300 11400 TRUE CNVnator 0.00204431 0.97997300 358901 370300
# 1955 tp1335 Chr24 259901 278000 18100 FALSE CNVnator 1.41458000 0.38091100 31264334 31282433
# 1956 tp1335 Chr24 286901 289800 2900 FALSE CNVnator 1.74941000 0.74228100 31291334 31294233
```

```
get.cnv.dels <- function(cov.thresh.lo = 0.0,
                        cov.thresh.hi = 0.8,
                        cnv,
                        snp.tables = NULL,
                        DEBUG = FALSE)
{
  # build list of 7 Bool vectors of genome length, with i-th == T iff
  # * i-th pos is 'NA' in genome seq (if snp.tables are provided), or
  # * in CNVnator call for coverage in half-open [cov.thresh.lo, hi), and
  # * not marked 'filtered' by CNVnator
  cnv.deletions <- vector(mode='list',7) # make list of bool vectors
  if(is.null(snp.tables)){
    # if no tables, assume full
    t.len <- genome.length.constants()$genome.length.trunc
  } else {
    t.len <- nrow(snp.tables[[1]])
  }
  for(st in 1:7){
    if(is.null(snp.tables)){
      cnv.deletions[[st]] <- logical(t.len) # all F
    } else {
      cnv.deletions[[st]] <- is.na(snp.tables[[st]]$Pos[1:t.len]) # NA positions in genome
    }
  }
  strain.names <- c(paste('tp10',c('07',12:15),sep=''),'IT','tp1335')
  names(cnv.deletions) <- strain.names
  for(i in 1:nrow(cnv)){
    if(!cnv$filtered[i] &&
        cnv$cov_ratio[i] >= cov.thresh.lo &&
        cnv$cov_ratio[i] < cov.thresh.hi)
    {
      if(DEBUG){
        print(cnv[i,])
        print(as.character(cnv$strain[i]))
      }
    }
  }
}
```

```

    }
    # following ASSUMES no CNVnator call crosses a chromosome bdry, & that
    # t.len ends at chr end (typically chr1 or chr24)
    if(cnv$iEnd[i] <= t.len){
      cnv.deletions[[as.character(cnv$strain[i])]][cnv$iStart[i]:cnv$iEnd[i]] <- TRUE
    }
  }
}
return(cnv.deletions)
}

# sanity check:
cnv.dels.38 <- get.cnv.dels(0.3, 0.8, cnv.chrononly, snp.tables = NULL)
unlist(lapply(cnv.dels.38, sum)) # does it match low.length.38 in tic ?

# tp1007 tp1012 tp1013 tp1014 tp1015 IT tp1335
# 1672500 1781500 1383600 1313700 988400 320900 1453000

# 1672500 1781500 1399400 1313700 988400 336500 1453000 <== low.length.38 from tic (circa page 8)
# 1672500 1781500 1399400 1313700 988400 336500 1453000 <== low.length.38 from tic (pg9, 6/28/17)
rm(cnv.dels.38)

```

Slight discrepancy in H-clade that I should hunt down, but basically OK. (hmm; maybe untrunc tbls.)

```

# the ones we want for the current analysis:
hemi.masks <- get.cnv.dels(0.3, 0.8, cnv.chrononly, snp.tables=snp.tables)

rbind(
  homnr      = unlist(lapply(hnr, sum)),
  hemi       = unlist(lapply(hemi.masks, sum)),
  homnr.unhemi = unlist(lapply(list(1,2,3,4,5,6,7), function(i){sum(hnr[[i]] & !hemi.masks[[i]]))}))
)

#
#      1007  1012  1013  1014  1015  3367  1335
# homnr      316   247 12082   938   236 13863   167
# hemi      11761 16653 24337 11603 19824 49254 15367
# homnr.unhemi 316   246 11890   938   219 13474   167

# based on the thought that hnr in 1335 may reflect errors in the ref seq,
# are they shared with others?
unlist(lapply(hnr, function(x){sum(x & hnr[[7]])})) # hnr shared with 1335

# 1007 1012 1013 1014 1015 3367 1335
#   43   55   66   30   56   68  167

# answer: around 300 in each strain, of 558 in NY, genomewide,
# so that seems like a plausibly important factor.

hnr.lclade <- hnr[[1]] | hnr[[2]] | hnr[[4]] | hnr[[5]] | hnr[[7]] # union over L-clade
sum(hnr.lclade) # count all in L-clade

# [1] 1496

sum(hnr[[3]] | hnr[[6]]) # present in H-clade

# [1] 18363

sum(hnr[[3]] & hnr[[6]]) # shared in H-clade

# [1] 7582

# look at a few in L-clade
w.hnr.l <- which(hnr.lclade)
seecounts(w.hnr.l[1:10], snp.tables=snp.tables)

```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	1559	A										
# 2				1007	7	0	0	24	0	TRUE	FALSE		
# 3				1012	11	0	0	37	0	TRUE	FALSE		
# 4				1013	9	0	0	5	0	TRUE	FALSE		
# 5				1014	4	0	0	16	0	TRUE	FALSE		
# 6				1015	47	0	0	35	0	TRUE	FALSE		
# 7				3367	0	0	0	0	0	TRUE	FALSE		
# 8				1335	60	0	0	50	0	TRUE	FALSE		
# 9	Chr1	1575	G										
# 10				1007	24	7	0	0	0	TRUE	FALSE		
# 11				1012	42	13	0	0	0	TRUE	FALSE		
# 12				1013	17	16	0	0	0	TRUE	FALSE		
# 13				1014	15	4	0	0	0	TRUE	FALSE		
# 14				1015	43	31	0	0	1	TRUE	FALSE		
# 15				3367	0	2	0	0	0	TRUE	FALSE		
# 16				1335	34	74	0	0	0	TRUE	FALSE		
# 17	Chr1	1893	C										
# 18				1007	0	0	14	32	0	TRUE	FALSE		
# 19				1012	0	0	38	52	0	TRUE	FALSE		
# 20				1013	0	0	95	14	0	TRUE	FALSE		
# 21				1014	0	0	5	31	0	TRUE	FALSE		
# 22				1015	0	0	47	44	0	TRUE	FALSE		
# 23				3367	0	0	29	0	0	TRUE	FALSE		
# 24				1335	0	0	68	85	0	TRUE	FALSE		
# 25	Chr1	2223	A										
# 26				1007	25	13	0	0	0	TRUE	FALSE		
# 27				1012	13	12	1	0	0	TRUE	FALSE		
# 28				1013	5	24	0	0	0	TRUE	FALSE		
# 29				1014	0	4	0	0	0	TRUE	FALSE		
# 30				1015	19	22	0	0	1	TRUE	FALSE		
# 31				3367	15	3	0	0	0	TRUE	FALSE		
# 32				1335	33	22	0	0	0	TRUE	FALSE		
# 33	Chr1	2319	C										
# 34				1007	0	28	10	0	1	TRUE	FALSE		
# 35				1012	0	43	17	0	1	TRUE	FALSE		
# 36				1013	13	15	9	0	1	TRUE	FALSE		
# 37				1014	0	18	6	0	1	TRUE	FALSE		
# 38				1015	0	53	20	0	1	TRUE	FALSE		
# 39				3367	4	0	24	0	0	TRUE	FALSE		
# 40				1335	0	118	28	0	1	TRUE	FALSE		
# 41	Chr1	2502	A										
# 42				1007	14	2	0	0	0	FALSE	FALSE		
# 43				1012	17	6	0	0	0	FALSE	FALSE		
# 44				1013	6	13	0	0	0	FALSE	FALSE		
# 45				1014	1	6	0	0	0	FALSE	FALSE		
# 46				1015	20	7	0	0	0	FALSE	FALSE		
# 47				3367	3	3	0	0	0	FALSE	FALSE		
# 48				1335	29	17	0	0	0	FALSE	FALSE		
# 49	Chr1	2573	C										
# 50				1007	0	0	11	28	1	TRUE	FALSE		
# 51				1012	0	0	30	50	1	TRUE	FALSE		
# 52				1013	0	0	231	12	0	TRUE	FALSE		
# 53				1014	0	0	4	18	1	TRUE	FALSE		
# 54				1015	0	0	50	38	1	TRUE	FALSE		
# 55				3367	0	0	71	0	0	TRUE	FALSE		
# 56				1335	0	0	62	75	1	TRUE	FALSE		
# 57	Chr1	3938	G										
# 58				1007	12	20	0	0	0	TRUE	FALSE		
# 59				1012	9	22	0	0	0	TRUE	FALSE		
# 60				1013	35	19	0	0	0	TRUE	FALSE		
# 61				1014	8	2	0	0	0	TRUE	FALSE		
# 62				1015	25	53	0	0	0	TRUE	FALSE		
# 63				3367	14	13	0	0	0	TRUE	FALSE		
# 64				1335	59	42	0	0	0	TRUE	FALSE		
# 65	Chr1	4876	G										
# 66				1007	0	1	0	0	0	FALSE	FALSE		

```
# 67      1012  1  4  0  0  0 FALSE FALSE
# 68      1013  0  0  0  0  0 FALSE FALSE
# 69      1014  1  0  0  0  0 FALSE FALSE
# 70      1015  0  3  0  0  0 FALSE FALSE
# 71      3367  4  4  0  0  0 FALSE FALSE
# 72      1335  2  2  0  0  0 FALSE FALSE
# 73 Chr1 4938  T
# 74      1007  0  43  0  23  1 FALSE FALSE
# 75      1012  0  63  0  48  1 FALSE FALSE
# 76      1013  0  83  0  2  0 FALSE FALSE
# 77      1014  0  27  0  4  1 FALSE FALSE
# 78      1015  0  75  0  47  1 FALSE FALSE
# 79      3367  0  19  0  12  1 FALSE FALSE
# 80      1335  0  57  0  59  1 FALSE FALSE

# one of those is a little weird:
xx<-snp.tables[[1]][149457,]
for (i in 2:7){xx <- rbind(xx, snp.tables[[i]][149457,])}
row.names(xx)<-names(snp.tables)
# My guess is that Chr/Pos/Ref are left as NA if coverage is zero.
xx

#      snp  Chr    Pos  Ref  Cov  a  g  c  t  n .match  exon  indel  chr    pos  rawCov
# 1007  0 <NA>    NA <NA>    0  0  0  0  0  0  0 FALSE FALSE <NA>    NA    0
# 1012  0 <NA>    NA <NA>    0  0  0  0  0  0  0 FALSE FALSE <NA>    NA    0
# 1013  0 <NA>    NA <NA>    0  0  0  0  0  0  0 FALSE FALSE <NA>    NA    0
# 1014  0 Chr1 149457  G    0  0  0  0  0  0  0 FALSE FALSE Chr1 149457    1
# 1015  0 <NA>    NA <NA>    0  0  0  0  0  0  0 FALSE FALSE <NA>    NA    0
# 3367  0 <NA>    NA <NA>    0  0  0  0  0  0  0 FALSE FALSE <NA>    NA    0
# 1335  0 Chr1 149457  G    0  0  0  0  0  0  0 FALSE FALSE Chr1 149457    1
```

5 Table 1 stats

Here is a brief summary of per-strain SNP counts, pairwise overlaps, and other conveniently available stats, such as those shown in Table 1 of the paper.

```
snp.counts <- matrix(NA,7,4)
snp.pctofny <- matrix(NA,7,4)
snp.pctofself <- matrix(NA,7,4)
snp.inter <- matrix(NA,7,7)
snp.union <- matrix(NA,7,7)
rownames(snp.counts) <- names(snp.tables)
rownames(snp.pctofny) <- names(snp.tables)
rownames(snp.pctofself) <- names(snp.tables)
rownames(snp.inter) <- names(snp.tables)
colnames(snp.inter) <- names(snp.tables)
rownames(snp.union) <- names(snp.tables)
colnames(snp.union) <- names(snp.tables)
for(stringency in 1:4){
  cat('\nStringency', stringency, 'ifelse(stringency==4, '(i.e. raw SAMTools SNP calls)', ''),
      '\n-----\n')
  for(i in 1:7){
    f.snps.i <- filtered.snps$Code$get.snps(i, stringency)
    snp.counts[i,stringency] <- sum(f.snps.i)
    for(j in i:7){
      f.snps.j <- filtered.snps$Code$get.snps(j, stringency)
      snp.inter[i,j] <- sum(f.snps.i & f.snps.j)
      snp.union[i,j] <- sum(f.snps.i | f.snps.j)
    }
  }
  snp.pctofny[,stringency] <- snp.inter[,7]/snp.counts[7,stringency]
  snp.pctofself[,stringency] <- snp.inter[,7]/snp.counts[,stringency]
  cat('Union Counts:\n'); print(snp.union)
  cat('Intersect Counts:\n'); print(snp.inter)
```



```

cat('Intersect as percent of union:\n'); print(snp.inter/snp.union*100,digits=3)
}

#
# Stringency 1 :
# -----
# Union Counts:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 18262 18723 36431 18614 18906 35671 18816
# 1012   NA 18475 36501 18769 19016 35729 18925
# 1013   NA   NA 29970 35615 36655 39685 36480
# 1014   NA   NA   NA 15827 18929 34748 18774
# 1015   NA   NA   NA   NA 18651 35878 19063
# 3367   NA   NA   NA   NA   NA 28699 35711
# 1335   NA   NA   NA   NA   NA   NA 18403
# Intersect Counts:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 18262 18014 11801 15475 18007 11290 17849
# 1012   NA 18475 11944 15533 18110 11445 17953
# 1013   NA   NA 29970 10182 11966 18984 11893
# 1014   NA   NA   NA 15827 15549  9778 15456
# 1015   NA   NA   NA   NA 18651 11472 17991
# 3367   NA   NA   NA   NA   NA 28699 11391
# 1335   NA   NA   NA   NA   NA   NA 18403
# Intersect as percent of union:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 100  96.2  32.4  83.1  95.2  31.7  94.9
# 1012   NA 100.0  32.7  82.8  95.2  32.0  94.9
# 1013   NA   NA 100.0  28.6  32.6  47.8  32.6
# 1014   NA   NA   NA 100.0  82.1  28.1  82.3
# 1015   NA   NA   NA   NA 100.0  32.0  94.4
# 3367   NA   NA   NA   NA   NA 100.0  31.9
# 1335   NA   NA   NA   NA   NA   NA 100.0
#
# Stringency 2 :
# -----
# Union Counts:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 17996 18521 37541 18222 18729 36692 18479
# 1012   NA 18326 37654 18474 18794 36815 18579
# 1013   NA   NA 30826 35649 37844 41411 37498
# 1014   NA   NA   NA 12861 18694 34625 18147
# 1015   NA   NA   NA   NA 18563 37002 18768
# 3367   NA   NA   NA   NA   NA 29507 36612
# 1335   NA   NA   NA   NA   NA   NA 17867
# Intersect Counts:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 17996 17801 11281 12635 17830 10811 17384
# 1012   NA 18326 11498 12713 18095 11018 17614
# 1013   NA   NA 30826  8038 11545 18922 11195
# 1014   NA   NA   NA 12861 12730  7743 12581
# 1015   NA   NA   NA   NA 18563 11068 17662
# 3367   NA   NA   NA   NA   NA 29507 10762
# 1335   NA   NA   NA   NA   NA   NA 17867
# Intersect as percent of union:
#      1007  1012  1013  1014  1015  3367  1335
# 1007 100  96.1  30.0  69.3  95.2  29.5  94.1
# 1012   NA 100.0  30.5  68.8  96.3  29.9  94.8
# 1013   NA   NA 100.0  22.5  30.5  45.7  29.9
# 1014   NA   NA   NA 100.0  68.1  22.4  69.3
# 1015   NA   NA   NA   NA 100.0  29.9  94.1
# 3367   NA   NA   NA   NA   NA 100.0  29.4
# 1335   NA   NA   NA   NA   NA   NA 100.0
#
# Stringency 3 :
# -----
# Union Counts:
#      1007  1012  1013  1014  1015  3367  1335

```

```

# 1007 16801 18054 36539 17040 18269 35673 17872
# 1012    NA 17738 36954 17864 18437 36089 18190
# 1013    NA    NA 30064 33057 37184 40928 36622
# 1014    NA    NA    NA 7895 18141 31952 17244
# 1015    NA    NA    NA    NA 18035 36335 18388
# 3367    NA    NA    NA    NA    NA 28785 35724
# 1335    NA    NA    NA    NA    NA    NA 17020
# Intersect Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 16801 16485 10326 7656 16567 9913 15949
# 1012    NA 17738 10848 7769 17336 10434 16568
# 1013    NA    NA 30064 4902 10915 17921 10462
# 1014    NA    NA    NA 7895 7789 4728 7671
# 1015    NA    NA    NA    NA 18035 10485 16667
# 3367    NA    NA    NA    NA    NA 28785 10081
# 1335    NA    NA    NA    NA    NA    NA 17020
# Intersect as percent of union:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 100 91.3 28.3 44.9 90.7 27.8 89.2
# 1012    NA 100.0 29.4 43.5 94.0 28.9 91.1
# 1013    NA    NA 100.0 14.8 29.4 43.8 28.6
# 1014    NA    NA    NA 100.0 42.9 14.8 44.5
# 1015    NA    NA    NA    NA 100.0 28.9 90.6
# 3367    NA    NA    NA    NA    NA 100.0 28.2
# 1335    NA    NA    NA    NA    NA    NA 100.0
#
# Stringency 4 (i.e. raw SAMTools SNP calls) :
# -----
# Union Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 16530 17707 35005 16864 17989 34289 17382
# 1012    NA 17019 35294 17276 18074 34563 17577
# 1013    NA    NA 25412 30445 35599 39448 34479
# 1014    NA    NA    NA 8331 17634 29704 16078
# 1015    NA    NA    NA    NA 17397 34876 17881
# 3367    NA    NA    NA    NA    NA 24613 33699
# 1335    NA    NA    NA    NA    NA    NA 15582
# Intersect Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 16530 15842 6937 7997 15938 6854 14730
# 1012    NA 17019 7137 8074 16342 7069 15024
# 1013    NA    NA 25412 3298 7210 10577 6515
# 1014    NA    NA    NA 8331 8094 3240 7835
# 1015    NA    NA    NA    NA 17397 7134 15098
# 3367    NA    NA    NA    NA    NA 24613 6496
# 1335    NA    NA    NA    NA    NA    NA 15582
# Intersect as percent of union:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 100 89.5 19.8 47.4 88.6 20.0 84.7
# 1012    NA 100.0 20.2 46.7 90.4 20.5 85.5
# 1013    NA    NA 100.0 10.8 20.3 26.8 18.9
# 1014    NA    NA    NA 100.0 45.9 10.9 48.7
# 1015    NA    NA    NA    NA 100.0 20.5 84.4
# 3367    NA    NA    NA    NA    NA 100.0 19.3
# 1335    NA    NA    NA    NA    NA    NA 100.0

vs.stringency <- cbind(snp.counts, matrix(NA,7,1), round(snp.counts[,1:3]/snp.counts[,4]*100,1))
colnames(vs.stringency) <- c('[[1]]', '[[2]]', '[[3]]', '[[4]]', '----', '[[1]]%', '[[2]]%', '[[3]]%')

# SNPs vs filtering stringency (raw counts and as % of [[4]]). Medium filter
# adds 10-20% in most cases. Big exception is Gyre, where low coverage,
# high err rate and SAMTools conservatism seemed to seriously undercall:
print(vs.stringency)

#      [[1]] [[2]] [[3]] [[4]] ---- [[1]]% [[2]]% [[3]]%
# 1007 18262 17996 16801 16530    NA 110.5 108.9 101.6
# 1012 18475 18326 17738 17019    NA 108.6 107.7 104.2
# 1013 29970 30826 30064 25412    NA 117.9 121.3 118.3

```

```
# 1014 15827 12861 7895 8331 NA 190.0 154.4 94.8
# 1015 18651 18563 18035 17397 NA 107.2 106.7 103.7
# 3367 28699 29507 28785 24613 NA 116.6 119.9 117.0
# 1335 18403 17867 17020 15582 NA 118.1 114.7 109.2

# Intersect NY as % of self (vs stringency):
print(snp.pctofself*100, digits=3)

#      [,1] [,2] [,3] [,4]
# 1007 97.7 96.6 94.9 89.1
# 1012 97.2 96.1 93.4 88.3
# 1013 39.7 36.3 34.8 25.6
# 1014 97.7 97.8 97.2 94.0
# 1015 96.5 95.1 92.4 86.8
# 3367 39.7 36.5 35.0 26.4
# 1335 100.0 100.0 100.0 100.0

# Intersect NY as % of NY (vs stringency):
print(snp.pctofny*100, digits=3)

#      [,1] [,2] [,3] [,4]
# 1007 97.0 97.3 93.7 94.5
# 1012 97.6 98.6 97.3 96.4
# 1013 64.6 62.7 61.5 41.8
# 1014 84.0 70.4 45.1 50.3
# 1015 97.8 98.9 97.9 96.9
# 3367 61.9 60.2 59.2 41.7
# 1335 100.0 100.0 100.0 100.0
```

Quick look at coverage. Are there any NA?:

```
nacount <- NULL
for(i in 1:4){
  if(!is.null(tset[[i]])){
    nacount <- rbind(nacount,
                     unlist(lapply(tset[[i]], function(x){sum(is.na(x$Cov))}))
    rownames(nacount)[nrow(nacount)] <- names(tset)[i]
  }
}
nacount

#      1007 1012 1013 1014 1015 3367 1335
# snp.tables.chr1.unfiltered 0 0 0 0 0 0 0
# snp.tables.chr1.qfiltered 0 0 0 0 0 0 0
```

Seemingly no. What's average in unq- vs q-filtered:

```
snp.tables.unqfil <- tset.picker(c(1,2), table.set = tset)
snp.tables.qfil <- tset.picker(c(3,4), table.set = tset)
cov.unqfil <- unlist(lapply(snp.tables.unqfil, function(x){mean(x$Cov)}))
cov.qfil <- unlist(lapply(snp.tables.qfil, function(x){mean(x$Cov, na.rm=T)}))
cov.both <- rbind(cov.unqfil, cov.qfil, cov.qfil/cov.unqfil)
i <- 1
if(!is.null(snp.tables.unqfil)){
  rownames(cov.both)[i] <- which.snp.tables(snp.tables.unqfil)
  i <- i+1
}
if(!is.null(snp.tables.qfil)){
  rownames(cov.both)[i] <- which.snp.tables(snp.tables.qfil)
  i <- i+1
}
if(i==3){
  rownames(cov.both)[i] <- 'Ratio'
}
cat('Mean Coverage:\n'); cov.both
```

```
# Mean Coverage:
#           1007      1012      1013      1014      1015      3367      1335
# Chr1-unfiltered 36.2816326 68.2005811 66.6908911 31.2663216 59.4704151 62.3834535 103.9124774
# Chr1-qfiltered  27.5849516 49.2557296 43.2293645 12.4319866 46.9874722 43.4403699  78.7789843
# Ratio           0.7603007  0.7222186  0.6482049  0.3976159  0.7900983  0.6963444  0.7581282
```

5.1 Table 1 Data

Throw together the conveniently-available Table 1 data, in Table 1 row order:

```
# if coverage unavailable, build NA vector
if(!is.null(cov.unqfil)){cov.unqfilv <- cov.unqfil} else {cov.unqfilv <- rep(NA,times=7)}
if(!is.null(cov.qfil )){cov.qfilv  <- cov.qfil } else {cov.qfilv  <- rep(NA,times=7)}
tldata.df <- data.frame(
  id      = st.locs(1:7, id=T, loc=F, date=F),
  loc     = st.locs(1:7, id=F, loc=T, date=F),
  date    = st.locs(1:7, id=F, loc=F, date=T),
  cov.unq = cov.unqfilv,
  cov.q    = cov.qfilv,
  SNPs.4   = snp.counts[,4],
  SNPs.2   = snp.counts[,2],
  olap.ny.4 = snp.pctofny[,4]*100,
  olap.ny.2 = snp.pctofny[,2]*100
)
tlrow.order <- c(7,1,2,5,3,6,4)
print(tldata.df[tlrow.order,],digits=3)

#           id      loc date cov.unq cov.q SNPs.4 SNPs.2 olap.ny.4 olap.ny.2
# 1335 CCMP1335    New York 1958  103.9  78.8  15582  17867  100.0  100.0
# 1007 CCMP1007    Virginia 1964   36.3  27.6  16530  17996   94.5   97.3
# 1012 CCMP1012 W. Australia 1965   68.2  49.3  17019  18326   96.4   98.6
# 1015 CCMP1015    Puget Sound 1985   59.5  47.0  17397  18563   96.9   98.9
# 1013 CCMP1013      Wales 1973   66.7  43.2  25412  30826   41.8   62.7
# 3367 CCMP3367      Italy 2007   62.4  43.4  24613  29507   41.7   60.2
# 1014 CCMP1014 N. Pacific Gyre 1971   31.3  12.4   8331  12861   50.3   70.4
```

6 Shared-SNPs P-Value

Text of the main paper quotes a “p-value” for the observed degree of SNP sharing in L-clade (and/or L-clade excluding Gyre) under a null model that these isolates were sampled from a population globally in Hardy-Weinberg equilibrium. Details of this analysis are as follows.

6.1 SNP Concordance

Arbitrarily pick one isolate, say, A , as the “template”. Arbitrarily pick a heterozygous (aka “SNP”) position in A . Let p_1 , and $q_1 = 1 - p_1$ be the frequencies in the overall population of the two nucleotides observed at that position in A . (Positions having 3 or 4 nucleotide variants segregating in the population are assumed to be negligibly rare.) Under the HWE null model, a second isolate B will also be heterozygous at the same position with probability $2p_1q_1 \leq 1/2$. Similarly, this position will be heterozygous in a third isolate C with the same probability, *independently*, and so on for isolates D and E . Overall, (assuming HWE) the probability that a heterozygous position in A is simultaneously heterozygous in the other 4 isolates is at most $1/2^4 = 1/16$. Continuing, suppose we pick a second heterozygous position in A , on a different chromosome with allele frequencies $p_2, q_2 = 1 - p_2$, say. Again assuming HWE, this position will be a SNP in all of B, C, D and E with probability $(2p_2q_2)^4 \leq 1/16$, and this is independent of the first position, since segregation on different chromosomes is unlinked. Repeat this at 24 heterozygous positions in A , one per chromosome. Then, the number of five-way concordant positions observed should be dominated by the number observed when sampling from a binomial distribution with parameters $n = 24$ and $p = 1/16$, i.e., we expect at most $1/16 = 6.25\%$ of positions to agree, or at most $24/16 = 1.5$ five-way concordant positions in total. In sharp contrast,

choosing CCMP 1014 (North Pacific Gyre) as the template, we see many more five-way concordant positions than predicted under these assumptions:

```
gyre.count <- sum(snp.tables[[4]]$snp)
# 'unfil.' => unfiltered for consistency; see below.
unfil.fiveway.count <- sum(snp.tables[[4]]$snp * i4.snps)
unfil.fiveway.percent <- unfil.fiveway.count / gyre.count * 100
unfil.p.value <- pbinom(floor(unfil.fiveway.count/gyre.count*24)-1, 24, 1/16, lower.tail = FALSE)
consistency.comparison <-
  data.frame(
    fiveway.count = unfil.fiveway.count,
    fiveway.percent = unfil.fiveway.percent,
    p.value = unfil.p.value
  )
consistency.comparison

#   fiveway.count fiveway.percent      p.value
# 1           7628          91.56164 8.700771e-23
```

Namely, 8331 positions are called as SNPs in CCMP1014, of which 7628 or 91.5616373% are also called as SNPs in *all four* other L-clade isolates. 91.5616373% of 24 is 21.9747929, and the probability of seeing 21 or more “Heads” in 24 flips of a biased coin with $P(\text{Heads}) \leq 1/16$, i.e., our p-value under the HWE null hypothesis, is at most: 8.700771×10^{-23} based on this simple binomial model. This is obviously strong evidence against the null hypothesis.

This analysis is potentially overly-simplistic in four respects, addressed below.

1. “ $2pq \leq 1/2$ ” is conservative. Neutral theory predicts that most variant nucleotides are rare in the population, so $2pq \ll 1/2$ is to be expected. This should make our quoted p-value very conservative.
2. Effect of Erroneous SNP calls. We base our analysis on *predicted* (by SAMTOOLS) heterozygous positions, not absolute-truth, which may affect our conclusions. However,
 - False negatives in *A* are irrelevant, since we never examine those positions. (This is the motivation for using CCMP1014 as the template; it has the lowest predicted SNP rate, likely due to a high false negative rate in that sequencing run. As noted elsewhere, it had the lowest coverage and lowest sequence quality of the 7 isolates, both of which impare SNP calling.)
 - False negatives in *BCDE* make such positions appear *non*-concordant. For our purpose, this makes our statistic more conservative since it can only deflate a statistic that we argue is nevertheless unexpectedly large.
 - False positive calls in *A* are conservatively treated, as well: barring simultaneous false-positive calls in all of *BCDE*, such a position will appear non-concordant, again deflating the statistic. The *false* positive rates in *B, C, D* and *E* are unknown, but cannot exceed SAMTOOLS *total* positive rate, which is below 1% in all 7 isolates, suggesting a simultaneous *BCDE* false positive rate $< 10^{-8}$, which will have a negligible effect.
 - A potentially more serious issue is a true positive in *A* aligned to false positives in *BCD* and/or *E*. (I.e., a position that is polymorphic in the population and heterozygous in *A*, under the HWE null model is likely to be homozygous for one of the two alleles in one or more of *BCDE*; false positive SNP calls in all of those isolates would make the site appear concordant, i.e., provide evidence against the null model.) However, (a) my impression is that SAMTOOLS is more prone to false negative calls than to false positive calls (see Section 4), and (b) we would need a high rate of false positives to turn a truly heterozygous but non-concordant *A* call into a false “concordant” call—I’d expect at most half (especially given point 1 above) of *BCDE* to be heterozygous, but all would need to be falsely declared heterozygous. Such a high false positive rate on *BCDE* seems unlikely (see previous bullet), and would likely be counterbalanced by a similarly increased rate of false positives on *A*, which, as noted, tend to deflate our statistic (previous bullet again).
 - Systematic errors. If there were, say, a sequence-context-dependent bias in the DNA sequencing, mapping and/or SNP-calling that tended to suggest (or hide) a SNP at some position, we’re going to systematically over- (or under-) estimate concordant SNPs across isolates. The discordance of called SNPs between the

L- and H-clades and within the H-clade suggests that this is not a major problem, but it is worth noting as a possibility.

- Discordant nucleotides at “concordant” SNP positions. A “shared” SNP at a given position might be, say, G/C in one isolate vs T/C in another, reflecting an unexpected tri-allelic position in the population or a technical sequencing error. It is inappropriate to count such a “shared” SNP position as evidence against the null hypothesis, since it isn’t clear that it is truly shared. Instead, I will identify such inconsistent positions, based on the “stringency [[2]]” criteria established above, and treat each as non-concordant. I.e., a position will be considered to be a “5-way concordant SNP” if and only if it was called as a SNP by SAMTOOLS (independently) in all 5 L-clade isolates, *and* shows the same dominant non-reference nucleotide in all 5, according to criteria [[2]] above. As it turns out, this correction has a very minor effect on the resulting p-value:

```
# 'unfil.' => Ignoring "consistency"; 'fil.' => Filtering for "consistency":
fil.fiveway.count <- sum((snp.tables[[4]]$snp * i4.snps)[union.snps == 1] & consistent[[2]])
fil.fiveway.percent <- fil.fiveway.count / gyre.count * 100
fil.p.value <- pbinom(floor(fil.fiveway.count/gyre.count*24)-1, 24, 1/16, lower.tail = FALSE)
# append new stats to previous table for easy comparison
consistency.comparison <-
  rbind(consistency.comparison,
        data.frame(
          fiveway.count = fil.fiveway.count,
          fiveway.percent = fil.fiveway.percent,
          p.value = fil.p.value
        )
  )
rownames(consistency.comparison) <- c('unfiltered', 'consistency.filtered')
consistency.comparison

#           fiveway.count fiveway.percent      p.value
# unfiltered           7628          91.56164 8.700771e-23
# consistency.filtered    7537          90.46933 8.700771e-23
```

In particular, it removes 1.1% of five-way consistent positions (only 91 of 7628 positions), and still shows a highly significant p-value.

- $P(E[X]) \neq E[P(X)]$. I’m expressing this poorly, but finding the p-value based on the *expected* number of concordant positions is somewhat non-standard. A more typical set-up would use the *actual* value of some statistic, then calculate the probability of observing a value that extreme (or more extreme) under the null model. The fundamental problem is that we have thousands of SNPs, but I don’t see an easy way to use more than 24 of them at a time, because potential genetic linkage seemingly destroys statistical independence, which is key to most simple analyses. A somewhat more formal, but still non-standard, approach is the following. Suppose we randomly sample one SNP per chromosome and count the number X of them that are 5-way concordant. What I outlined above calculated the p-value based on $E[X]$, the expected value of X , i.e., $P(E[X])$. Alternatively, we can calculate $E[P(X)]$, the expected p-value. (They are not the same.) In effect, this averages the p-values that would be seen over many different randomly-sampled sets of 24 SNPs. This is not difficult to calculate. First, the probability that we would observe $0 \leq i \leq 24$ concordant positions in a sample of 24, given that 90.47% of positions are concordant follows this binomial distribution:

```
x.equals.i.distribution <- dbinom(0:24, 24, fil.fiveway.percent/100)
print(x.equals.i.distribution, digits=3)

# [1] 3.15e-25 7.19e-23 7.85e-21 5.46e-19 2.72e-17 1.03e-15 3.11e-14 7.58e-13 1.53e-11 2.58e-10
# [11] 3.68e-09 4.44e-08 4.57e-07 4.00e-06 2.98e-05 1.89e-04 1.01e-03 4.50e-03 1.66e-02 4.98e-02
# [21] 1.18e-01 2.14e-01 2.77e-01 2.28e-01 9.04e-02
```

Second, the p-value corresponding to $0 \leq i \leq 24$ observed concordant positions also follows a different binomial distribution:

```
p.val.of.x.equals.i <- c(1, pbinom(0:23, 24, 1/16, lower.tail = F))
print(p.val.of.x.equals.i, digits=3)

# [1] 1.00e+00 7.88e-01 4.48e-01 1.87e-01 5.95e-02 1.49e-02 3.01e-03 4.99e-04 6.90e-05 8.02e-06
# [11] 7.89e-07 6.60e-08 4.72e-09 2.87e-10 1.49e-11 6.59e-13 2.46e-14 7.66e-16 1.98e-17 4.14e-19
# [21] 6.88e-21 8.70e-23 7.88e-25 4.56e-27 1.26e-29
```

Finally, the expected (or “average”) p-value is just the weighted average of the latter values, weighted by the former:

```
e.of.p.of.x <- sum(x.equals.i.distribution * p.val.of.x.equals.i)
e.of.p.of.x

# [1] 1.33456e-14
```

This is still highly significant, but weaker than the $P(E[X])$ analysis, basically because $X < E[X]$ has a fair probability of occurring, and the corresponding p-value $P(X)$ rises rapidly as X declines.

Another way to look at the numbers:

```
pvdof <- data.frame(x.density=x.equals.i.distribution,
                    x.cdf=cumsum(x.equals.i.distribution),
                    pval.of.x=p.val.of.x.equals.i)
print(pvdof, digits=4)

#   x.density    x.cdf pval.of.x
# 1 3.155e-25 3.155e-25 1.000e+00
# 2 7.187e-23 7.219e-23 7.875e-01
# 3 7.846e-21 7.918e-21 4.476e-01
# 4 5.461e-19 5.541e-19 1.869e-01
# 5 2.722e-17 2.777e-17 5.950e-02
# 6 1.033e-15 1.061e-15 1.490e-02
# 7 3.106e-14 3.213e-14 3.010e-03
# 8 7.583e-13 7.904e-13 4.994e-04
# 9 1.530e-11 1.609e-11 6.899e-05
# 10 2.581e-10 2.742e-10 8.015e-06
# 11 3.675e-09 3.949e-09 7.887e-07
# 12 4.440e-08 4.835e-08 6.603e-08
# 13 4.566e-07 5.049e-07 4.716e-09
# 14 4.001e-06 4.506e-06 2.875e-10
# 15 2.984e-05 3.434e-05 1.493e-11
# 16 1.888e-04 2.232e-04 6.590e-13
# 17 1.008e-03 1.231e-03 2.456e-14
# 18 4.504e-03 5.735e-03 7.662e-16
# 19 1.663e-02 2.236e-02 1.977e-17
# 20 4.984e-02 7.220e-02 4.143e-19
# 21 1.183e-01 1.905e-01 6.877e-21
# 22 2.139e-01 4.043e-01 8.701e-23
# 23 2.768e-01 6.811e-01 7.884e-25
# 24 2.285e-01 9.096e-01 4.556e-27
# 25 9.037e-02 1.000e+00 1.262e-29
```

E.g., row 9 in that table says that the concordance rate (90%) is so high that a sample of 24 SNPs will almost always have 9 or more five-way concordant positions (probability of fewer is only 1.609e-11), while under the null model, seeing 9 or more is very unlikely (probability at most 6.899e-05). ***AM I OFF-BY-ONE INTERPRETING ROW 9 HERE??***

6.2 Notes

In earlier drafts, an analog of the above analysis was based on the concordance of *refined* SNPs. This now seems to me to be questionable, since the “refined” SNP calling makes SNPs called across L-clade non-independent. OTOH,

the above analysis seems valid: SAMTOOLS was run on each isolate independently, and likewise “criterion [[2]]” is evaluated independently in each strain, and is being used here solely to remove SNP predictions, not to add them. “Systematic errors” as outlined above remain a potential problem, but again discordance with/within H-clade suggests that this is of limited concern.

For completeness, I did a similar analysis including a sample of H-clade comparisons: Gyre vs Italy, NY vs Italy, NY vs Italy+Wales, and of Italy vs Wales. As expected, none of these show a statistically significant p-value, although the $\approx 40\%$ concordance in the 2-way comparisons, while $< 1/2$ as predicted, is a bit higher than I expected based on “neutral theory implies many rare variants.” (I did not bother to include “criterion[[2]] filtering” in these calculations.)

```
# 'gi.twoway' => gyre vs italy 2-way concordance;
# 'ni.twoway' => new york vs italy 2-way concordance;
# not bothering with criterion[[2]] filtering
gi.twoway.count <- sum(snp.tables[[4]]$snp * snp.tables[[6]]$snp)
gi.twoway.percent <- gi.twoway.count / gyre.count * 100
gi.p.value <- pbinom(floor(gi.twoway.count/gyre.count*24)-1, 24, 1/2, lower.tail = FALSE)
ny.count <- sum(snp.tables[[7]]$snp)
ni.twoway.count <- sum(snp.tables[[7]]$snp * snp.tables[[6]]$snp)
ni.twoway.percent <- ni.twoway.count / ny.count * 100
ni.p.value <- pbinom(floor(ni.twoway.count/ny.count*24)-1, 24, 1/2, lower.tail = FALSE)
niw.threeway.count <- sum(snp.tables[[7]]$snp * snp.tables[[6]]$snp * snp.tables[[3]]$snp)
niw.threeway.percent <- niw.threeway.count / ny.count * 100
niw.p.value <- pbinom(floor(niw.threeway.count/ny.count*24)-1, 24, 1/4, lower.tail = FALSE)
it.count <- sum(snp.tables[[6]]$snp)
iw.twoway.count <- sum(snp.tables[[6]]$snp * snp.tables[[3]]$snp)
iw.twoway.percent <- iw.twoway.count / it.count * 100
iw.p.value <- pbinom(floor(iw.twoway.count/it.count*24)-1, 24, 1/2, lower.tail = FALSE)
consistency.comparison <-
  rbind(consistency.comparison,
    data.frame(
      fiveway.count = c(gi.twoway.count, ni.twoway.count, niw.threeway.count, iw.twoway.count),
      fiveway.percent = c(gi.twoway.percent, ni.twoway.percent, niw.threeway.percent, iw.twoway.percent),
      p.value = c(gi.p.value, ni.p.value, niw.p.value, iw.p.value)
    )
  )
colnames(consistency.comparison)[1:2] <- c('552232way.count', '552232way.percent') # old col names misleading
rownames(consistency.comparison)[3:6] <- c('gyre.vs.italy', 'new.york.vs.italy', # new rows
      'ny.vs.it.plus.wales', 'it.vs.wales')

consistency.comparison

#           552232way.count 552232way.percent      p.value
# unfiltered              7628          91.56164 8.700771e-23
# consistency.filtered    7537          90.46933 8.700771e-23
# gyre.vs.italy           3240          38.89089 9.242052e-01
# new.york.vs.italy       6496          41.68913 8.462719e-01
# ny.vs.it.plus.wales     3804          24.41278 7.533516e-01
# it.vs.wales             10577          42.97323 8.462719e-01
```

6.3 P-Value: The Bottom Line

So, what to say in the body of the paper? $E[P(X)]$ is highly significant, and conservative, but complex to explain. $P(E[X])$ is simpler to explain, but may be criticized as misleading if we aren’t very careful in that explanation. I’m slightly leaning towards the last option, but want to sleep on it and draft the key sentence or two before settling.

7 Sharing

The following analysis looks at the sharing patterns among the consistent SNPs. I assume that shared SNPs reflect shared ancestry, and that SNPs accumulate slowly over time. Then, in outline, the story is consistent with what we have seen in other analyses—there seem to be 3 groups: 1013 (Wales) in one, 3367 (Italy) in another, and the other 5 in a third, with some hints as to the order of divergence. A caveat is that in a sexual population, non-shared SNPs do not immediately imply non-shared ancestry; they may merely reflect Hardy-Weinberg capturing a homozygous state

in one isolate vs the other. (Or read errors, etc.) Thus, if we are right that the H-isolates retain sex, then the large number of “private” SNPs in H may be at least partially due to HWE.

Analysis is broken into cases based on how many strains share a particular SNP.

7.1 Code

To categorize SNPs by sharing patterns, first convert the 7-way consistent sharing pattern into a 7-bit binary number, and tabulate based on that:

```
# convert (n x 7) 0-1 matrix to n vector of 0-127
tobin <- function(x){
  bin <- integer(nrow(x)) # initialized to 0
  for(i in 1:7){
    bin <- bin*2 + as.integer(x[,i]>0)
  }
  return(bin)
}

# get full set of patterns
snp.pattern.all <- lapply(non.refs,tobin)
# prune to just the consistent ones
snp.pattern <- snp.pattern.all
for(i in 1:3){
  snp.pattern[[i]][!consistent[[i]]] <- NA
}

# analogous to built-in ``table'' but simpler. Count entries in an integer
# vector sharing values in a (smallish) range. Result is a 2-column matrix with
# the shared values in col 1 and count of occurrences of that value in col 2.
# Out-of-range values cause subscript error.
mytable <- function(vec, therange=range(vec,na.rm=T)){
  counts <- matrix(0,nrow=therange[2]-therange[1]+1,ncol=2,dimnames=list(NULL,c('val','count')))
  counts[1:nrow(counts),1] <- therange[1]:therange[2]
  for(i in 1:length(vec)){
    if(!is.na(vec[i])){
      counts[vec[i]-therange[1]+1,2] <- counts[vec[i]-therange[1]+1,2] + 1
    }
  }
  return(counts)
}

pattern.counts <- lapply(snp.pattern, function(x){mytable(x,c(0,127))})
```

To display the results, build a data frame whose i -th row, $0 \leq i \leq 127$ shows one of the 128 possible sharing patterns, with counts of the numbers of consistent, shared SNPs with that pattern according to criteria c1-c3.

```
tobitvec <- function(x){
  bitvec <- integer(7)
  for(i in 1:7){
    bitvec[i] <- x %% 2
    x <- x %% 2
  }
  return(bitvec)
}

flg <- function(x){
  return(ifelse(x==1, 'X', ''))
}

pat.summary <- function(listOfTbls){
  mydf <- data.frame(pat=0:127,sharedBy=NA,
    tp1007='',tp1012='',tp1013='',tp1014='',tp1015='',tp3367='',tp1335='',
    count1=NA,count2=NA,count3=NA,count4=NA,stringsAsFactors=F)

  for(i in 1:128){
```

```

bvec <- tobitvec(i-1)
mydf[i, 'sharedBy'] = sum(bvec)
mydf[i, 'tp1007'] = flg(bvec[1])
mydf[i, 'tp1012'] = flg(bvec[2])
mydf[i, 'tp1013'] = flg(bvec[3])
mydf[i, 'tp1014'] = flg(bvec[4])
mydf[i, 'tp1015'] = flg(bvec[5])
mydf[i, 'tp3367'] = flg(bvec[6])
mydf[i, 'tp1335'] = flg(bvec[7])
}

for(i in 1:length(listOfTbls)){
  tbl <- listOfTbls[[i]]
  if(!is.null(tbl)){
    mydf[,9+i] <- tbl[,2] ## count1/2/3/4 are columns 10/11/12/13 in mydf
    #for(j in 1:length(tbl)){
    #  k <- as.integer(rownames(tbl)[j]);
    #  mydf[k+1,9+i] <- tbl[j] ## count1/2/3 are columns 10/11/12
    #}
  }
}

mydf$pat <- as.octmode(mydf$pat) # display bit pattern in octal
return(mydf)
}

pat.summaries <- pat.summary(pattern.counts)

```

7.2 Sanity Checks

Some sanity checking: table sums equal to number of consistent positions?

```

all(consistent.count == apply(pat.summaries[,10:13],2,sum))

# [1] TRUE

```

More sanity checking: visually inspect a pattern with small counts, specifically pattern 12, i.e., consistent SNPs shared by only strains 1014 and 1015 (2nd and 3 rows from bottom, binary code $12 = 2^3 + 2^2$). There are only 10 such positions on Chr1. Chr1 2524239 has pattern 12 under criteria c1 and c2 but not c3; Chr1 1088766 has in c2 only. Both look good. Neither position is a *called* SNP except in 1015. However, all but 1 nonreference read agree with the called SNP (the exception being one read in Wales). Both 1014 and 1015 have at least 2 non-reference reads, comprising at least 5% of coverage, and in both strains, those reads are on the same non-reference base, satisfying criterion c2. The other strains have higher coverage and/or lower non-reference counts, so they do not satisfy c2. Position 2524239 also satisfies c1, but not c3, since 2 reads out of 35 is below the 10% threshold. (It is pattern 4 under c3, i.e., a SNP private to 1015.) Position 1088766 is also pattern 4 under c3 (2 reads out of 56 in 1335 is below both thresholds), and it is not consistent under c1, since the single A read in 1013 is discordant with the other non-reference reads.

```

unlist(lapply(snp.pattern,function(x){sum(x==12,na.rm=T)}))

# [1] 4 1 6 12

sp1 <- snp.pattern[[1]]==12
sp2 <- snp.pattern[[2]]==12
sp3 <- snp.pattern[[3]]==12
sp4 <- snp.pattern[[4]]==12
c(sum(sp1,na.rm=T), sum(sp2,na.rm=T), sum(sp3,na.rm=T), sum(sp4,na.rm=T))

# [1] 4 1 6 12

r1 <- rownames(non.refs[[1]])[which(sp1)]
r2 <- rownames(non.refs[[2]])[which(sp2)]
r3 <- rownames(non.refs[[3]])[which(sp3)]
r4 <- rownames(non.refs[[4]])[which(sp4)]

r2

```

```
# [1] "Chr1:1799155"

c1 <- as.integer(unlist(lapply(strsplit(r1[1:min(20,length(r1))],':',fixed=TRUE),function(x){x[2]})))
c2 <- as.integer(unlist(lapply(strsplit(r2[1:min(20,length(r2))],':',fixed=TRUE),function(x){x[2]})))
c3 <- as.integer(unlist(lapply(strsplit(r3[1:min(20,length(r3))],':',fixed=TRUE),function(x){x[2]})))
c4 <- as.integer(unlist(lapply(strsplit(r4[1:min(20,length(r4))],':',fixed=TRUE),function(x){x[2]})))

c1

# [1] 614335 914018 1317406 2388286

c2

# [1] 1799155

c3

# [1] 371484 518347 1210354 2209068 2264683 2898352

c4

# [1] 518347 691730 767408 1049906 1390437 2072951 2254059 2254789 2264683 2823796 2898352
# [12] 2998868

seecounts(c2,snp.tables=snp.tables)

# chr pos Ref Strain A G C T SNP exon indel nrf rat
# 1 Chr1 1799155 C
# 2 1007 0 0 10 1 0 TRUE FALSE
# 3 1012 0 0 16 1 0 TRUE FALSE
# 4 1013 0 0 10 0 0 TRUE FALSE
# 5 1014 0 0 8 2 0 TRUE FALSE
# 6 1015 0 0 12 3 1 TRUE FALSE
# 7 3367 1 0 1 1 1 TRUE FALSE
# 8 1335 0 0 7 1 0 TRUE FALSE
```

Position 1088766, however, is a good example of the situation that motivated this analysis—one strain has a G/C SNP and 5 of the other 6 strains have nonreference reads consistent with that SNP. Although, excluding 1015, the nonreference read counts are not high enough to justify a SNP call in any strain considered in isolation, the fact that they *consistently* agree with the 1015 SNP suggests that they are real. One alternative hypothesis is that there is some sequence-dependent bias at this locus that favors misreading a G as a C. On the other hand, one could equally well posit a shared SNP, and a locus-dependant bias that *supresses* C reads, explaining the unbalanced readout that we observe. However, it is hard to reconcile either view with the significant strain-specific patterns that we see in the shared SNPs (as seen below). I think a more likely explanation is that (a) there are some number of relatively rare SNPs present in each of the sampled populations, (b) some of these SNPs happened to be present in one or two cells of the roughly 5-10 cells that we believe constituted the founding population of the culture grown for sequencing, and (c) stochastic effects during culture growth and during sequencing may have further perturbed the apparent frequency of each variant, but the bottom line is that the above-threshold presence of consistent non-reference reads is evidence for shared SNPs at the population level (and the proportions of such reads represent estimates of the population-level frequencies of the variants, albeit a noisy estimate at any specific position).

An aside: I was curious to see whether there is any consistent pattern to positions that are called consistent SNPs in all but Italy, so I repeated the above, basically. My summary is that coverage in Italy tends to be below average in these positions, but otherwise they don't stand out. For the record:

```
abit <- snp.pattern[[2]]==125
abit[is.na(abit)]<-F
sum(abit)

# [1] 1352

rabit <- rownames(non.refs[[2]])[which(abit)]
rabits <- rabbit[1:20]
cabit <- as.integer(unlist(lapply(strsplit(rabits,':',fixed=TRUE),function(x){x[2]})))
cabit
```

```
# [1] 1244 1575 6485 7181 7220 7661 8144 8208 8518 8552 8567 8670 8685 14361 15254
# [16] 15280 16103 25546 30784 33852
```

```
seecounts(cabit,snp.tables=snp.tables)
```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	1244	G										
# 2				1007	2	25	0	0	0	TRUE	FALSE		
# 3				1012	3	32	0	0	0	TRUE	FALSE		
# 4				1013	10	24	0	0	1	TRUE	FALSE		
# 5				1014	3	17	0	0	0	TRUE	FALSE		
# 6				1015	15	43	0	0	1	TRUE	FALSE		
# 7				3367	0	1	0	0	0	TRUE	FALSE		
# 8				1335	82	65	0	0	1	TRUE	FALSE		
# 9	Chr1	1575	G										
# 10				1007	24	7	0	0	0	TRUE	FALSE		
# 11				1012	42	13	0	0	0	TRUE	FALSE		
# 12				1013	17	16	0	0	0	TRUE	FALSE		
# 13				1014	15	4	0	0	0	TRUE	FALSE		
# 14				1015	43	31	0	0	1	TRUE	FALSE		
# 15				3367	0	2	0	0	0	TRUE	FALSE		
# 16				1335	34	74	0	0	0	TRUE	FALSE		
# 17	Chr1	6485	G										
# 18				1007	24	19	0	0	0	TRUE	FALSE		
# 19				1012	29	29	0	0	0	TRUE	FALSE		
# 20				1013	49	33	0	0	0	TRUE	FALSE		
# 21				1014	6	5	0	0	0	TRUE	FALSE		
# 22				1015	31	32	0	0	1	TRUE	FALSE		
# 23				3367	0	37	0	0	0	TRUE	FALSE		
# 24				1335	62	52	0	0	0	TRUE	FALSE		
# 25	Chr1	7181	G										
# 26				1007	0	30	29	0	0	TRUE	FALSE		
# 27				1012	0	52	34	0	0	TRUE	FALSE		
# 28				1013	0	19	72	0	0	TRUE	FALSE		
# 29				1014	0	13	7	0	0	TRUE	FALSE		
# 30				1015	0	40	33	0	1	TRUE	FALSE		
# 31				3367	0	29	0	0	0	TRUE	FALSE		
# 32				1335	0	78	73	0	0	TRUE	FALSE		
# 33	Chr1	7220	C										
# 34				1007	16	0	19	6	0	TRUE	FALSE		
# 35				1012	38	0	22	11	0	TRUE	FALSE		
# 36				1013	82	1	30	9	0	TRUE	FALSE		
# 37				1014	12	0	6	2	0	TRUE	FALSE		
# 38				1015	55	0	22	5	1	TRUE	FALSE		
# 39				3367	0	0	8	0	0	TRUE	FALSE		
# 40				1335	55	0	32	20	0	TRUE	FALSE		
# 41	Chr1	7661	T										
# 42				1007	0	0	9	9	0	TRUE	FALSE		
# 43				1012	0	0	5	19	0	TRUE	FALSE		
# 44				1013	0	0	24	14	1	TRUE	FALSE		
# 45				1014	0	0	6	3	0	TRUE	FALSE		
# 46				1015	0	0	5	34	0	TRUE	FALSE		
# 47				3367	0	0	0	4	0	TRUE	FALSE		
# 48				1335	0	0	4	24	0	TRUE	FALSE		
# 49	Chr1	8144	G										
# 50				1007	8	9	0	0	0	TRUE	FALSE		
# 51				1012	12	10	0	0	1	TRUE	FALSE		
# 52				1013	38	29	0	0	0	TRUE	FALSE		
# 53				1014	5	4	0	0	0	TRUE	FALSE		
# 54				1015	15	16	0	0	0	TRUE	FALSE		
# 55				3367	0	0	0	0	0	TRUE	FALSE		
# 56				1335	12	15	0	0	1	TRUE	FALSE		
# 57	Chr1	8208	G										
# 58				1007	0	6	0	7	1	TRUE	FALSE		
# 59				1012	0	19	0	11	0	TRUE	FALSE		
# 60				1013	0	1	0	48	0	TRUE	FALSE		
# 61				1014	0	5	0	3	0	TRUE	FALSE		
# 62				1015	0	19	0	11	1	TRUE	FALSE		
# 63				3367	0	1	0	0	0	TRUE	FALSE		
# 64				1335	0	28	0	16	1	TRUE	FALSE		
# 65	Chr1	8518	T										
# 66				1007	0	0	20	15	1	FALSE	FALSE		
# 67				1012	0	0	40	20	1	FALSE	FALSE		
# 68				1013	0	0	45	56	1	FALSE	FALSE		
# 69				1014	0	0	10	16	0	FALSE	FALSE		
# 70				1015	0	0	36	13	1	FALSE	FALSE		
# 71				3367	0	0	0	2	0	FALSE	FALSE		
# 72				1335	0	0	113	53	1	FALSE	FALSE		
# 73	Chr1	8552	G										
# 74				1007	3	9	0	0	0	TRUE	FALSE		

# 75				1012	20	21	0	0	0	TRUE	FALSE
# 76				1013	28	16	0	0	1	TRUE	FALSE
# 77				1014	6	2	0	0	0	TRUE	FALSE
# 78				1015	14	13	0	0	0	TRUE	FALSE
# 79				3367	0	12	0	0	0	TRUE	FALSE
# 80				1335	24	47	0	0	0	TRUE	FALSE
# 81	Chr1	8567	A								
# 82				1007	14	18	0	0	1	TRUE	FALSE
# 83				1012	26	30	0	0	1	TRUE	FALSE
# 84				1013	50	66	0	0	1	TRUE	FALSE
# 85				1014	1	3	0	0	0	TRUE	FALSE
# 86				1015	12	31	0	0	1	TRUE	FALSE
# 87				3367	22	0	0	0	0	TRUE	FALSE
# 88				1335	51	40	0	0	1	TRUE	FALSE
# 89	Chr1	8670	A								
# 90				1007	7	0	0	5	0	TRUE	FALSE
# 91				1012	16	0	0	10	0	TRUE	FALSE
# 92				1013	16	0	0	11	0	TRUE	FALSE
# 93				1014	2	0	0	4	0	TRUE	FALSE
# 94				1015	14	0	0	10	1	TRUE	FALSE
# 95				3367	5	0	0	0	0	TRUE	FALSE
# 96				1335	7	0	0	6	0	TRUE	FALSE
# 97	Chr1	8685	G								
# 98				1007	6	15	0	0	0	TRUE	FALSE
# 99				1012	10	23	0	0	0	TRUE	FALSE
# 100				1013	18	21	0	0	1	TRUE	FALSE
# 101				1014	4	8	0	0	0	TRUE	FALSE
# 102				1015	10	24	0	0	1	TRUE	FALSE
# 103				3367	0	4	0	0	0	TRUE	FALSE
# 104				1335	5	32	0	0	0	TRUE	FALSE
# 105	Chr1	14361	A								
# 106				1007	20	7	0	0	0	FALSE	FALSE
# 107				1012	35	5	0	0	0	FALSE	FALSE
# 108				1013	1	11	0	0	1	FALSE	FALSE
# 109				1014	6	2	0	0	0	FALSE	FALSE
# 110				1015	35	7	0	0	0	FALSE	FALSE
# 111				3367	2	1	0	0	0	FALSE	FALSE
# 112				1335	50	8	0	0	0	FALSE	FALSE
# 113	Chr1	15254	T								
# 114				1007	11	0	0	16	1	FALSE	FALSE
# 115				1012	26	0	0	38	1	FALSE	FALSE
# 116				1013	37	0	0	48	1	FALSE	FALSE
# 117				1014	3	0	0	8	1	FALSE	FALSE
# 118				1015	18	0	0	32	1	FALSE	FALSE
# 119				3367	0	0	0	73	0	FALSE	FALSE
# 120				1335	13	0	0	32	1	FALSE	FALSE
# 121	Chr1	15280	T								
# 122				1007	0	13	0	20	1	FALSE	FALSE
# 123				1012	0	27	0	28	1	FALSE	FALSE
# 124				1013	0	5	0	64	0	FALSE	FALSE
# 125				1014	0	2	0	8	0	FALSE	FALSE
# 126				1015	0	19	0	29	1	FALSE	FALSE
# 127				3367	0	0	0	42	0	FALSE	FALSE
# 128				1335	0	21	0	70	1	FALSE	FALSE
# 129	Chr1	16103	A								
# 130				1007	10	0	11	0	1	FALSE	FALSE
# 131				1012	44	0	19	0	1	FALSE	FALSE
# 132				1013	21	0	13	0	1	FALSE	FALSE
# 133				1014	14	0	2	0	0	FALSE	FALSE
# 134				1015	29	0	10	0	1	FALSE	FALSE
# 135				3367	33	0	0	0	0	FALSE	FALSE
# 136				1335	47	0	11	0	0	FALSE	FALSE
# 137	Chr1	25546	A								
# 138				1007	23	0	0	14	1	FALSE	FALSE
# 139				1012	46	0	0	19	1	FALSE	FALSE
# 140				1013	6	0	0	42	1	FALSE	FALSE
# 141				1014	7	0	0	15	1	FALSE	FALSE
# 142				1015	52	0	0	17	1	FALSE	FALSE
# 143				3367	60	0	0	0	0	FALSE	FALSE
# 144				1335	67	0	0	5	0	FALSE	FALSE
# 145	Chr1	30784	C								
# 146				1007	16	0	13	0	1	TRUE	FALSE
# 147				1012	33	0	32	0	1	TRUE	FALSE
# 148				1013	19	0	33	0	1	TRUE	FALSE
# 149				1014	4	0	11	0	1	TRUE	FALSE
# 150				1015	39	0	29	0	1	TRUE	FALSE
# 151				3367	0	0	55	0	0	TRUE	FALSE
# 152				1335	46	0	50	0	1	TRUE	FALSE
# 153	Chr1	33852	C								
# 154				1007	0	24	25	0	1	FALSE	FALSE

```
# 155      1012  0 18  26  0  1 FALSE FALSE
# 156      1013  0 28  33  0  1 FALSE FALSE
# 157      1014  0  9   4  0  1 FALSE FALSE
# 158      1015  0 19  28  0  1 FALSE FALSE
# 159      3367  0  0  26  0  0 FALSE FALSE
# 160      1335  0 30  53  0  1 FALSE FALSE
```

More sanity: there are 83 sites on Chr1 shared by zero strains in the tightest condition. (I.e., SAMTOOLS called it a SNP, but the read counts/proportions fall below our 3rd threshold). Are they due to low coverage? Seemingly yes:

```
zp3 <- snp.pattern[[3]] == 0
zr3 <- rownames(non.refs[[3]])[which(zp3)]
zc3 <- as.integer(unlist(lapply(strsplit(zr3[1:min(100,length(zr3))],':',fixed=TRUE),function(x){x[2]})))
zc3
```

```
# [1] 16115 16615 19117 25748 43500 55857 56591 65787 66879 68328 80862 81001 90622
# [14] 90721 91284 110754 116443 116453 120183 126702 127986 129056 147698 153874 159756 160912
# [27] 161271 170686 180314 181477 182139 196862 196864 199166 206132 206143 221888 234931 242276
# [40] 242914 244505 268954 274655 282391 282511 283646 289363 311952 312625 314132 326217 371008
# [53] 376784 387078 387091 389263 395153 406158 410771 431788 438958 438976 443898 447253 448223
# [66] 452774 488812 495476 498133 501830 501975 504462 506422 515441 515595 530113 530114 532320
# [79] 534149 541667 543095 575081 585297 586276 612732 622585 651159 652889 655373 655380 657704
# [92] 657955 658216 685697 687653 692115 692139 700484 700845 701061
```

```
seecounts(zc3[1:5], snp.tables=snp.tables)
```

```
#      chr   pos Ref Strain  A G  C  T SNP  exon indel nrf rat
# 1  Chr1 16115   T
# 2      1007  0 0  0  5  0 FALSE FALSE
# 3      1012  0 0  0  9  0 FALSE FALSE
# 4      1013  0 0  0  6  0 FALSE FALSE
# 5      1014  0 0  0  3  0 FALSE FALSE
# 6      1015  0 0  0 10  0 FALSE FALSE
# 7      3367  0 0  3  3  1 FALSE FALSE
# 8      1335  0 0  0  6  0 FALSE FALSE
# 9  Chr1 16615   C
# 10     1007  0 0 39  0  0 FALSE FALSE
# 11     1012  0 0 54  0  0 FALSE FALSE
# 12     1013  0 0  4  2  1 FALSE FALSE
# 13     1014  0 0 19  0  0 FALSE FALSE
# 14     1015  0 0 46  0  0 FALSE FALSE
# 15     3367  0 0 13  0  0 FALSE FALSE
# 16     1335  0 0 40  0  0 FALSE FALSE
# 17 Chr1 19117   A
# 18     1007 16 0  0  0  0 TRUE FALSE
# 19     1012 21 0  0  0  0 TRUE FALSE
# 20     1013  1 0  0  1  0 TRUE FALSE
# 21     1014  6 0  0  0  0 TRUE FALSE
# 22     1015 21 0  0  0  0 TRUE FALSE
# 23     3367  0 0  0  1  1 TRUE FALSE
# 24     1335 24 0  0  0  0 TRUE FALSE
# 25 Chr1 25748   C
# 26     1007  0 0 17  0  0 FALSE FALSE
# 27     1012  0 0 36  0  0 FALSE FALSE
# 28     1013  3 0  7  0  1 FALSE FALSE
# 29     1014  1 0  4  0  0 FALSE FALSE
# 30     1015  0 0 32  0  0 FALSE FALSE
# 31     3367  0 0  1  0  0 FALSE FALSE
# 32     1335  1 0 34  0  0 FALSE FALSE
# 33 Chr1 43500   A
# 34     1007 10 0  0  3  1 FALSE FALSE
# 35     1012 10 0  0  3  1 FALSE FALSE
# 36     1013 10 0  1  1  0 FALSE FALSE
# 37     1014  5 0  0  0  0 FALSE FALSE
# 38     1015 11 0  0  2  0 FALSE FALSE
# 39     3367  6 0  0  3  0 FALSE FALSE
# 40     1335 13 0  0  1  0 FALSE FALSE
```

7.3 Main Analysis

Turning to the main analysis, there is a large increase in the number of consistent positions between the loose and medium stringency levels; medium and tight are similar in most respects. The likely interpretation is that the loose criterion is including many “SNPs” induced by read errors, and that either of the tighter criteria are successfully filtering them out. In the interest of simplicity, the narrative below will focus on the shared SNPs at the medium stringency level (the “count2” column in the data frame), although the numbers for all three (sometimes all 4) are displayed. Also note that the prose and some comments in the code were based on the Chr1 analysis, and so may occasionally be off-target for the whole-genome data.

```
# Show a subset of pat.summaries, optionally with totals of count_i in last row, and optionally
# aggregating low-count rows as ``Other``
#
#   sharedBy=c(2,4) selects SNPs shared by 2 or 4 strains,
#   subset=as.octmode('35') select those with sharing pattern a subset (optionally proper) of this
#   split=as.octmode('14') additionally restricts to patterns straddling split/subset minus split
#   c2.thresh=42 suppresses printout of rows with count2 < 42
#   restrict.to=c(0,42,127) restrict to these 3 rows
showgroup <- function(p.summ=pat.summaries, sharedBy=0:7, subset=127, split=NULL, proper.subset=F,
                      total=T, c2.thresh=0, fourteenth=F, restrict.to=NULL){
  # pick just those bit patterns that are subsets of 'subset'
  pick <- bitwAnd(0:127,bitwNot(subset))==0
  if(proper.subset){
    pick[subset+1] <- F
  }
  if(!is.null(split)){ # AND that straddle left/right subtrees
    cosplit <- bitwAnd(subset,bitwNot(split))
    pick <- pick & bitwAnd(0:127,split)!=0 & bitwAnd(0:127,cosplit)!=0
  }
  # and have desired shareBy counts
  pick <- pick & (p.summ$sharedBy %in% sharedBy)
  # and are among the set of interest
  if(!is.null(restrict.to)){
    pick <- pick & (0:127 %in% restrict.to)
  }
  # find rows with low counts
  pick.low <- pick & (p.summ$count2 < c2.thresh)
  # now show them
  show <- p.summ[pick & ! pick.low,]
  # rename columns just to narrow the printouts
  colnames(show) <- c('Pat','ShrBy','1007', '1012', '1013', '1014', '1015', '3367', '1335',
                      'count1', 'count2', 'count3','count4')
  show[,1] <- format(show[,1]) # convert octal col to char so can override in last row(2)
  nlow <- sum(pick.low)
  if(nlow > 0){
    n <- nrow(show)+1
    lows <- apply(p.summ[pick.low,10:13],2,sum)
    show[n,10:13] <- lows
    show[n,1:9] <- ''
    row.names(show)[n] <- 'Other'
    if(fourteenth){
      # do this: add 14th col just to hold this comment:
      show <- cbind(show, ' ', stringsAsFactors=F)
      show[n,14] <- paste('(', nlow, 'rows w/ c2 < ', c2.thresh, ')')
    } else {
      ## or this (looks a bit funky, but fits across page without line-wrap):
      show[n,1:8] <- c('(', nlow, 'rows', 'w/', 'c2', '<', c2.thresh, ')')
    }
  }
  if(total){
    n <- nrow(show)+1
    tots <- apply(show[,10:13],2,sum)
    show[n,10:13] <- tots
    show[n,1:9] <- ''
    row.names(show)[n] <- 'Total'
    if(ncol(show)==14){show[n,14]<-''}
  }
}
```

```
}
  return(show)
}
```

First, are there any SNPs that are not “consistent SNPs?” Yes, a few in c3. As noted above, they seem to be mainly low-coverage positions.

```
showgroup(pat.summaries,0,total=F) # chr1 totals: 0 0 83

#   Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 1    0     0                                9    51    468     0
```

Next, look at completely shared SNPs, those found in all 7 strains.

```
showgroup(pat.summaries,7,total=F) # Chr1 count1 = 8593, count2 = 7054, count3 = 4790 c4=1641

#   Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 128 177    7    X    X    X    X    X    X    X    8132    6449    3873    1641
```

I.e., of the 47108 consistent positions, 6449 or 13.7% are shared by all 7 strains.

Next look at singletons, aka private SNPs—SNPs that are called in one strain and no other strain has a significant number of non-ref reads at that position. Presumably these are variants that arose in a given population after it separated from the others.

```
showgroup(pat.summaries,1) # chr1 totals: 9669 18865 19670 23574

#   Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 2    001    1                                X    35    42    62    135
# 3    002    1                                X    7818  8912  9095  10949
# 5    004    1                                X    176    218    264    385
# 9    010    1                                X    42    62    53    113
# 17   020    1                                X    8529  9752  9961  11677
# 33   040    1                                X    53    62    103    174
# 65   100    1    X                                25    33    38    141
# Total                                16678  19081  19576  23574
```

The import of shared/private SNPs changes between sexual and asexual populations. Presumably asexuals slowly gain and rarely lose private SNPs; shared ones predate separation of the lineages. In sexual lineages, however, SNPs may be rather freely “gained” or “lost,” merely by recombination (converting between homo- and heterozygous in the sample we sequenced). Thus, the low private counts for the 5 L-isolates compared to the large count of het positions overall suggest that (a) they are asexual, and (b) none of them has been isolated from the others for very long (if at all). Conversely, the high counts for Italy and Wales suggest that (a) if asexual, they have been separated from each other and from the rest for a long time, but (b) if sexual, there is little surprise: we have $\approx 160\text{K}$ SNPs shared between the two (90K just in those two (below), plus 70K shared by all 7), and $\approx 90\text{K}$ additional positions that are het in one but not the other. These are close to, but not exactly equal to, the 1:2:1 ratios we would naively expect from two samples of a single HWE population. The most parsimonious explanation seems to be that the H-clade is sexual, but perhaps some het positions private to each population separates them.

Aside: counts of “consistent” SNPs minus these singletons yeilds count of shared SNPs:

```
singlets <- apply(pat.summaries[pat.summaries$sharedBy==1,10:13],2,sum)
rbind(consistent=consistent.count,singlets=singlets,shared=consistent.count-singlets)

#           count1 count2 count3 count4
# consistent 44905 47108 47204 47499
# singlets   16678 19081 19576 23574
# shared     28227 28027 27628 23925
```

The slightly higher count of shared positions in the medium case further supports this choice for subsequent analysis.

Next look at consistent SNPs shared between just a pair of isolates.


```
showgroup(pat.summaries,2) # chr 1 counts: 7641 9549 9472 6924
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 4	003	2						X	X	87	15	37	31
# 6	005	2					X		X	17	23	40	52
# 7	006	2					X	X		65	9	20	41
# 10	011	2				X			X	3	5	4	14
# 11	012	2				X		X		41	4	2	5
# 13	014	2				X	X			4	1	6	12
# 18	021	2			X				X	102	7	24	9
# 19	022	2			X			X		8822	9349	8911	6177
# 21	024	2			X		X			62	15	21	46
# 25	030	2			X	X				65	5	1	9
# 34	041	2		X					X	2	4	23	36
# 35	042	2		X				X		47	2	20	27
# 37	044	2		X			X			6	17	80	155
# 41	050	2		X		X				5	1	3	7
# 49	060	2		X	X					59	8	24	28
# 66	101	2	X						X	3	5	7	20
# 67	102	2	X					X		34	4	11	25
# 69	104	2	X				X			2	10	40	107
# 73	110	2	X			X				2	1	3	7
# 81	120	2	X		X					49	5	7	31
# 97	140	2	X	X						14	20	25	85
# Total										9491	9510	9309	6924

I.e., of the 9510 paired SNPs, 9349 or 98.3% are found between Italy and Wales, with comparatively few shared between any other pairs (only).

SNPs shared among exactly 3 isolates are relatively rare. (The 5 trios containing both Italy and Wales predominate in the loose set, probably because they share many pairs that become triples with the addition of a few read errors.)

```
showgroup(pat.summaries,3) # chr 1 counts: 1438 294 671 1034
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 8	007	3					X	X	X	4	3	19	10
# 12	013	3				X		X	X	3	2	2	2
# 14	015	3				X	X		X	6	4	4	7
# 15	016	3				X	X	X		1	0	1	2
# 20	023	3			X			X	X	89	23	35	17
# 22	025	3			X		X		X	8	4	23	21
# 23	026	3			X		X	X		84	31	37	32
# 26	031	3			X	X			X	3	1	0	0
# 27	032	3			X	X		X		66	9	6	5
# 29	034	3			X	X	X			2	5	1	1
# 36	043	3		X				X	X	8	10	14	6
# 38	045	3		X			X		X	12	44	184	131
# 39	046	3		X			X	X		4	12	41	55
# 42	051	3		X		X			X	0	2	4	4
# 43	052	3		X		X		X		1	0	0	1
# 45	054	3		X		X	X			1	7	12	18
# 50	061	3		X	X				X	1	8	11	12
# 51	062	3		X	X			X		86	21	29	36
# 53	064	3		X	X		X			2	17	52	60
# 57	070	3		X	X	X				3	0	0	2
# 68	103	3	X					X	X	2	3	6	8
# 70	105	3	X				X		X	9	14	37	63
# 71	106	3	X				X	X		4	10	9	27
# 74	111	3	X			X			X	0	1	1	1
# 75	112	3	X			X		X		0	1	0	0
# 77	114	3	X			X	X			2	4	4	8
# 82	121	3	X		X				X	2	0	2	4
# 83	122	3	X		X			X		45	6	12	26
# 85	124	3	X		X		X			2	7	21	35
# 89	130	3	X		X	X				1	1	1	1
# 98	141	3	X	X					X	5	9	20	40
# 99	142	3	X	X				X		3	3	9	15

# 101	144	3	X	X		X		18	74	159	355
# 105	150	3	X	X		X		0	1	0	6
# 113	160	3	X	X	X			6	3	7	23
# Total								483	340	763	1034

Four-way sharing is more common, but dominated by the coastal (i.e., non-Gyre) L-clade isolates. This is likely a reflection of the strong 5-way sharing among the L-clade, from which the Gyre commonly drops out due to the lower coverage/higher error rate in that sequencing run.

showgroup(pat.summaries,4) # chr 1 counts: 564 1346 2552 3479

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 16	017	4				X	X	X	X	1	2	1	2
# 24	027	4			X		X	X	X	15	16	37	24
# 28	033	4			X	X		X	X	4	2	4	6
# 30	035	4			X	X	X		X	2	4	0	0
# 31	036	4			X	X	X	X		5	0	3	3
# 40	047	4		X			X	X	X	9	26	68	60
# 44	053	4		X		X		X	X	0	1	1	2
# 46	055	4		X		X	X		X	8	15	24	36
# 47	056	4		X		X	X	X		2	2	2	5
# 52	063	4		X	X			X	X	9	12	34	21
# 54	065	4		X	X		X		X	8	21	68	48
# 55	066	4		X	X		X	X		15	43	102	76
# 58	071	4		X	X	X			X	0	2	0	0
# 59	072	4		X	X	X		X		4	2	1	2
# 61	074	4		X	X	X	X			1	2	1	7
# 72	107	4	X				X	X	X	6	10	6	16
# 76	113	4	X			X		X	X	0	0	0	1
# 78	115	4	X			X	X		X	1	4	6	9
# 79	116	4	X			X	X	X		1	0	1	5
# 84	123	4	X		X			X	X	5	9	13	8
# 86	125	4	X		X		X		X	3	4	14	16
# 87	126	4	X		X		X	X		10	17	14	43
# 90	131	4	X		X	X			X	0	0	0	2
# 91	132	4	X		X	X		X		1	1	1	1
# 93	134	4	X		X	X	X			6	3	2	3
# 100	143	4	X	X				X	X	1	3	4	20
# 102	145	4	X	X			X		X	598	1356	2429	2585
# 103	146	4	X	X			X	X		9	34	69	140
# 106	151	4	X	X		X			X	2	2	4	14
# 107	152	4	X	X		X		X		1	3	1	4
# 109	154	4	X	X		X	X			24	45	34	103
# 114	161	4	X	X	X				X	3	8	10	18
# 115	162	4	X	X	X			X		8	11	20	33
# 117	164	4	X	X	X		X			19	51	71	163
# 121	170	4	X	X	X	X				0	1	1	3
# Total										781	1712	3046	3479

Five-way sharing is much more common, and is strongly dominated by the 5 L-clade isolates.

showgroup(pat.summaries,5) # chr 1 counts: 3969 5047 4624 6125

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 32	037	5			X	X	X	X	X	12	11	8	5
# 48	057	5		X		X	X	X	X	4	9	8	17
# 56	067	5		X	X		X	X	X	48	109	257	104
# 60	073	5		X	X	X		X	X	3	3	3	3
# 62	075	5		X	X	X	X		X	8	7	10	11
# 63	076	5		X	X	X	X	X		9	10	12	7
# 80	117	5	X			X	X	X	X	1	1	3	7
# 88	127	5	X		X		X	X	X	13	27	49	47
# 92	133	5	X		X	X		X	X	2	3	0	0
# 94	135	5	X		X	X	X		X	4	2	0	7
# 95	136	5	X		X	X	X	X		5	3	0	5
# 104	147	5	X	X			X	X	X	205	421	740	1160

# 108	153	5	X	X		X		X	X	4	0	0	7
# 110	155	5	X	X		X	X		X	4136	3560	2135	3228
# 111	156	5	X	X		X	X	X		11	7	9	43
# 116	163	5	X	X	X			X	X	15	14	21	33
# 118	165	5	X	X	X		X		X	318	591	957	1140
# 119	166	5	X	X	X		X	X		46	154	220	254
# 122	171	5	X	X	X	X			X	4	4	2	7
# 123	172	5	X	X	X	X		X		3	6	3	5
# 125	174	5	X	X	X	X	X			5	14	17	35
# Total										4856	4956	4454	6125

Six-way sharing is also common, with the sets *excluding* Gyre, Italy, or Wales having the most mutually-shared SNPs.

```
showgroup(pat.summaries,6) # chr 1 counts: 4166 4741 5312 4722
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 64	077	6		X	X	X	X	X	X	43	46	62	26
# 96	137	6	X		X	X	X	X	X	8	6	8	14
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 120	167	6	X	X	X		X	X	X	1305	2445	4098	1852
# 124	173	6	X	X	X	X		X	X	15	5	5	3
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# 127	176	6	X	X	X	X	X	X		57	79	43	68
# Total										4475	5009	5715	4722

8 Trees

So, overall, the picture looks like a long shared history (6449 7-way shared positions), followed by a split of the 5 L-isolates from the 2 H-isolates, then a long shared history in the 5 (3560 quintuples), in parallel with a long shared history in H- (9349 pairs), then separate histories in Italy and Wales (>8912 “private” SNPs in each, although again if they are sexual, many of these just reflect HWE), and very limited differentiation among the 5 L-isolates.

Branch lengths of course depend on filtering criteria used (and, of course, full vs Chr1 differ by about a factor of 10), but the tree *topology* appears to be fairly stable. Various versions are drawn below, exactly to explore how robust this story is. I think we should go with “medium stringency” SNP filtering (based on un-qfiltered reads).

NOTE: Much of this analysis make less sense for q-filtered read data, since (a) the point of the SNP filtering was to try to correct for noise in the raw reads, which may (or may not; haven’t looked closely, yet) be largely fixed by qfiltering (e.g., “loose” or no SNP filtering may be more appropriate, post-q-filtering, esp. if we had re-run SAMTools to call SNPs based on the q-filtered reads), and (b) tree topology *does* appear to change, in that Gyre’s coverage has been so sharply reduced by qfiltering that it clearly stands aside from the others (and that’s confirmed by bootstrap), but this also seems to be clearly a technical rather than a biological artifact. SO, code below will run on q-filtered data, but *is not tuned to it*. Likewise, most comments in the prose below were made to describe the un-q-filtered data, and *are misleading and in some cases flatly wrong* for qfiltered data, but it doesn’t seem worthwhile to bother with a rewrite...

Trees are coded in newick format, which doesn’t seem to tolerate line-breaks; print with line-wrap:.

```
# wrap a long char string across multiple lines in printout
cat.hardwrap <- function(str,width=80){
  while(nchar(str)>width){
    cat(substr(str,1,width),'\n')
    str <- substr(str,width+1,nchar(str))
  }
  cat(str,'\n')
}
```

Trees are built as follows. Code for drawing, especially, is specific to the topology of the medium tree, and placement of some of the figure elements have been hand-optimized for this case; drawings for the other variants will not be as pretty.

```

# set up for tree figs

# the newick parser in ape seems to be confused by commas and parens in
# tip names, and blanks are not allowed, so replace by *, <, >, _, resp.
newick.name <- function(name){
  name <- gsub(' ', '_', name, fixed=TRUE)
  name <- gsub(',', '*', name, fixed=TRUE)
  name <- gsub('(', '<', name, fixed=TRUE)
  name <- gsub(')', '>', name, fixed=TRUE)
  return(name)
}

# undo above changes
newick.name.undo <- function(name){
  #name <- gsub('_', ' ', name, fixed=TRUE) # unnecessary; ape plot routine handles this one
  name <- gsub('*', ',', name, fixed=TRUE)
  name <- gsub('<', '(', name, fixed=TRUE)
  name <- gsub('>', ')', name, fixed=TRUE)
  return(name)
}

# make a newick string from tree; see it below
# 'pre' is prefixed to ccmpid; 'nb' optionally included;
# 'alt' can be used instead of pre/ccmp/nb/where for less formal labeling
# 'newstyle'=T => new node label: [nb_]where[(pre-less-id)]
# 'newstyle'=F => old node label: [nb_][pre id]where
newickize <- function(tree,pre='CCMP',nb=TRUE,alt=F,newstyle=TRUE){
  if(is.null(tree$where)){
    # not a leaf; paste together newick from subtrees
    sub1 <- newickize(tree$sub1,pre=pre,nb=nb,alt=alt,newstyle=newstyle)
    sub2 <- newickize(tree$sub2,pre=pre,nb=nb,alt=alt,newstyle=newstyle)
    new <- paste('(', sub1, ',', sub2, ')', sep='')
    if(!is.null(tree$length)){
      # internal node, add length
      return(paste(new, ':', tree$length, sep=''))
    } else {
      # top level; escape blanks and add trailing ';'
      return(paste(gsub(' ', '_', new), ';', sep=''))
    }
  } else {
    # a leaf; build label and branch length
    if(alt){
      # label is just alt; if alt omitted, default to where
      new <- newick.name(ifelse(is.null(tree$alt), tree$where, tree$alt))
    } else {
      if(newstyle){
        # new node label = [nb_]where[(pre-less-id)]
        new <- ifelse(nb && !is.null(tree$nb), paste(tree$nb, '_', sep=''), '')
        new <- newick.name(paste(new, tree$where, sep=''))
        new <- ifelse(is.null(tree$id), new, paste(new, '_', tree$id, ')', sep=''))
        new <- newick.name(new)
      } else {
        # old style node label = [nb_][pre id]where
        new <- ifelse(nb && !is.null(tree$nb), paste(tree$nb, '_', sep=''), '')
        new <- ifelse(is.null(tree$id), new, paste(new, pre, tree$id, '_', sep=''))
        new <- newick.name(paste(new, tree$where, sep=''))
      }
    }
    #add length to either
    new <- paste(new, ':', tree$length, sep='')
  }
  return(new)
}

# Make a tree as nested lists, **based on the chr1, count2 topology**, but using any of the counts.
# Internal nodes have subtrees sub1/2 and length
# Root has sub1/2, but no length
# Leaves have where, length, optionally, id, alt, nb. (Omit id for 'outgroup'. Use 'alt' for less formal
# labeling in cartoon version; it defaults to 'where'. Use 'nb' to add abcde annotations for legend.)
# The single parameter v is any of the 4 count vectors contained in pat.summaries (most conveniently
# indexed in octal). E.g., make.tree(pat.summaries['count2']) reproduces the count2 tree.
# (This was previously built by hand-pasting the edge lengths; tree.by.hand is retained in appendix
# for comparison, & its counts are in comments below).
#
make.tree <- function(v){
  pat.count <- function(pat, pat.counts=v){return(pat.counts[1+strtoi(pat,8)])}
  thetree <-
    list(
      sub1 = list(
        sub1 = list(
          sub1 = list(id=3367, length=pat.count('002'), where='Venice, Italy', alt='Venice'), #8813

```

```

sub2 = list(id=1013, length=pat.count('020'), where='Wales, UK'), #9652
length=pat.count('022')), #9365
sub2 = list(
  sub1 = list(
    sub1 = list(
      sub1 = list(id=1007, length=pat.count('100'), nb='e', where='Virginia, USA'), #30
      sub2 = list(id=1012, length=pat.count('040'), nb='d', where='Perth, W. Australia', alt='Perth'), #61
      length=pat.count('140')), #19
    sub2 = list(
      sub1 = list(id=1015, length=pat.count('004'), nb='c', where='Washington, USA', alt='Puget Sound'), #207
      sub2 = list(id=1335, length=pat.count('001'), nb='b', where='New York, USA', alt='NY'), #41
      length=pat.count('005')), #18
      length=pat.count('145')), #1005
      sub2 = list(id=1014, length=pat.count('010'), nb='a', where='N. Pacific Gyre'), #61
      length=pat.count('155')), #3912
      length=pat.count('177')), #7054
    sub2 = list(length=0, where='outgroup')
  )
)
return(thetree)
}

```

Code to plot a tree given newick description. Again, code is somewhat general, but has some specializations tied to the medium-stringency, full-genome, un-qfiltered data.

```

# run following 2 lines after an R upgrade
# update.packages()
# install.packages("ape")
library(ape)
show.tree <- function(newick.str=newick.medium,
  col.edge = 'darkblue', lwd.edge = 2,
  col.elabel = 'darkblue', cex.elabel = 0.8, font.elabel = 3,
  col.arrow = 'red', lwd.arrow = 1.5, cex.arrow = 0.9, font.arrow = 4,
  col.clade = 'black', lwd.clade = 1, cex.clade = 1.0, font.clade = 3,
  col.legend = 'beige', cex.legend = 0.8,
  col.tip = 'darkblue', font.tip = 4,
  plusx = FALSE, pltdebug = FALSE, total.snps = consistent.count[2]){

####
#
# ADJUST NEWICK & GET LENGTHS, COORDINATES
#
newick.str.noout <- sub('outgroup', '_', newick.str) # Hide outgroup ('_' prints as blank)
the.tree <- read.tree(text=newick.str.noout)

## nasty hack: ape's newick parser seems to be confused by commas, () in tip labels, so
## newickize replaced them by '*<>'; before plotting, I want to convert them back, and hope
## this doesn't break anything else... And if a revised version of ape changes the internal
## representation of a tree, this may need to be redone.
the.tree$tip.label <- newick.name.undo(the.tree$tip.label)

# extract branch lengths as char string of comma-separated numbers via pattern matching hack:
# lengths always preceded by colon
lengths.ch <- strsplit(paste(newick.str, ','), '[^0-9][^:]*:')[[1]]

# then convert to ints, dropping empty string at front
lengths.int <- scan(what=integer(), quiet=T, sep=',', text=lengths.ch[-1])

# then to data frame with named rows; a..g are terminal branches; others are internal.
# a..e match legend in plot; f/g = wales/italy. lengths appear in postfix order of
# newick tree, and ape draws the 1st of them at the bottom of the plot.
lmed <- data.frame(lengths=lengths.int,
  row.names=c('g', 'f', 'fg', 'e', 'd', 'de', 'c', 'b', 'bc', 'bcde', 'a', 'abcde', 'all', 'out'))

# extract counts needed for legend:
#leg.counts <- c( 61, 41, 207, 61, 30, 1005, 18, 19) #by hand, medium chr1
leg.counts <- lmed[c('a', 'b', 'c', 'd', 'e', 'bcde', 'bc', 'de'), 1]
discord <- total.snps - sum(lmed$lengths)

#tree.labels <- list( ## x,y,text; coords are all picked by eye
# 3000, 3.62, paste(lmed['all', 1], 'shared by 7', sep='\n'), # 7054
# 8900, 5.75, paste(lmed['abcde', 1], 'by 5', sep='\n'), # 3912
# 12000, 1.50, paste(lmed['fg', 1], 'shared by 2', sep='\n'), # 9365
# 21000, 2.00, paste(lmed['f', 1], 'only\nin Wales'), # 9652
# 21000, 1.00, paste(lmed['g', 1], 'only\nin Italy'), # 8813
# 11500, 4.50, '*')
# automating x-placement, below; retain above for comparison...
tip <- integer(7) # x coords of tree tips
tip[1] <- sum(lmed[c('all', 'fg', 'g'), 1])
tip[2] <- sum(lmed[c('all', 'fg', 'f'), 1])

```

```

tip[3] <-sum(lmed[c('all','abcde','bcde','de','e'),1])
tip[4] <-sum(lmed[c('all','abcde','bcde','de','d'),1])
tip[5] <-sum(lmed[c('all','abcde','bcde','bc','c'),1])
tip[6] <-sum(lmed[c('all','abcde','bcde','bc','b'),1])
tip[7] <-sum(lmed[c('all','abcde','a'),1])

inode <- integer(5) # x coords of (some) internal nodes
inode[1] <- 0 # root
inode[2] <- lmed['all',1] # lca of all
inode[3] <- sum(lmed[c('all','fg'),1]) # lca H-clade
inode[4] <- sum(lmed[c('all','abcde'),1]) # lca L-clade
inode[5] <- sum(lmed[c('all','abcde','bcde'),1]) # lca L-clade, nonGyre
tree.labels <- list( ## x,y,text; y coords partially picked by eye
  sum(inode[c(1,2)])/2, 3.62, paste(lmed['all',1], 'shared by 7', sep='\n'), # 7054
  sum(inode[c(2,4)])/2, 5.75, paste(lmed['abcde',1], 'by 5', sep='\n'), # 3912
  sum(inode[c(2,3)])/2, 1.50, paste(lmed['fg',1], 'shared by 2', sep='\n'), # 9365
  (inode[3]+tip[2])/2, 2.00, paste(lmed['f',1], 'only\nin 1013'), # 9652
  (inode[3]+tip[1])/2, 1.00, paste(lmed['g',1], 'only\nin 3367'), # 8813
  sum(inode[c(4,5)])/2, 4.35, '* ')

tree.labels <- list( ## x,y,text; y coords partially picked by eye
  sum(inode[c(1,2)])/2, 3.62, paste(lmed['all',1], 'in 7', sep='\n'), # 7054
  sum(inode[c(2,4)])/2, 5.75, paste(lmed['abcde',1], 'in 5', sep='\n'), # 3912
  sum(inode[c(2,3)])/2, 1.50, paste(lmed['fg',1], 'in 2', sep='\n'), # 9365
  (inode[3]+tip[2])/2, 2.00, paste(lmed['f',1], 'only\nin 1013'), # 9652
  (inode[3]+tip[1])/2, 1.00, paste(lmed['g',1], 'only\nin 3367'), # 8813
  sum(inode[c(4,5)])/2, 4.35, '* ')

####
#
# BOGUS PLOT
#
# a messy bit: need string widths to set xlim; but strwidth needs x-scale so must plot first.
# M plot completely invisible, overlay 2nd plot via par(new=F...) .
#
# PROVISIONALLY set x.lim here at about 30% wider than tree; fine tune it for the real plot
# based on strwidth(tip labels) below.
#
provisional.tree.x.lim <- 1.3 * max(tip) # <= PROVISIONAL plot width
plot(0,0, type='n', bty='n', xaxt='n', yaxt='n', xlab='', ylab='', xlim=c(0,provisional.tree.x.lim), ylim=c(0,7))

tiplabel.x <- integer(7)
for(i in 1:7){
  # see warning above about internals of the tree; labels have '_', printed as ' '.
  tiplabel.x[i] <- tip[i]+strwidth(gsub('_', ' ', the.tree$tip.label[i], fixed=T), font=font.tip)
}

# visually show tip coords & max x to debug placement issues
plt.debug <- function(tree.x.lim, tip, tiplabel.x, spx=NULL, spy=NULL){
  if(plt.debug){ # F to hide/T to show debug
    cat('Tip labels:', paste(the.tree$tip.label, sep='', collapse='/'), '\n')
    axis(2) # useful only for placing labels
    for(i in 1:7){
      points(c(tip[i], tiplabel.x[i]), c(i,i)) # debug: do I have right tip coordinates?
    }
    lines(rep(tree.x.lim, 2), c(0, 7)) # where is right edge?
    if(!is.null(spx)){
      points(spx, spy) # show spline control points, for tweaking
    }
  }
}

plt.debug(provisional.tree.x.lim, tip, tiplabel.x)

label.end.H <- max(tiplabel.x[1:2])
label.end.L <- max(tiplabel.x[3:7])
clade.dx <- strwidth('x') # space between clade marker line and its label
xdel <- 3*clade.dx # space between labeled clade tips and marker line

tree.x.lim <- 1.03*(max(tiplabel.x)+xdel) # <= FINAL plot width
if(plt.debug){cat('Plot width hacking:', provisional.tree.x.lim, tree.x.lim, tree.x.lim/1.03/max(tip), clade.dx)}

par(new=T) # I.e., NOT starting a new plot

####
#
# REAL PLOT
#

```

```

plot(the.tree,
     x.lim = tree.x.lim,
     y.lim = c(0,7),
     font=font.tip, label.offset=100,           # bold-italic, nudged slightly right
     tip.color=col.tip, edge.color=col.edge,
     edge.width=lwd.edge,
     edge.lty=c(1,1,1,1, 1 ,1,1,1,1,1,1,1,0)   # 5th is bottleneck edge; 14th is outgroup
)
lines(00+c(0,0),c(3.5,6),col='white',lwd=6)     # Hide vertical line to outgroup
axis(1, pos=0.25, at=seq(0,25,by=5)*10^round(log10(max(tip)/25)))

if(pltdebug){text(tip[1]+100, 1.0, 'Venice, Italy (3367)', adj=0, font=font.tip)}

####
#
# BOTTLENECK ANNOTATION
#
# spline/ellipse control points (spy/y) & tweaks thereto (dx/y)
dx <- 0.01 * tree.x.lim
dy <- .04
spx <- c(7400, 7400, 9900, 10500) # by eye, chr1, for comparison
spx <- c(inode[2]+dx,inode[2]+dx,inode[4]-3*dx,inode[4]-dx)
spy <- c( 3.8,  3.9,  5.6-dy,  5.6-dy)

plt.debug(tree.x.lim, tip, tiplabel.x, spx, spy)

if(T){
  #ellipse version, defined by rect thru 2 middle pts of spx/y
  spf<-function(x){
    ifelse(x <= spx[2], spy[1],
           ifelse(x >= spx[3], spy[4],
                  spy[2]+(spy[3]-spy[2])*sqrt(pmax(0,1-((x-spx[3])/(spx[3]-spx[2]))^2))))
  }
} else {
  # spline version
  spf <- splinefun(spx,spy,method='hyman')
}
serx <- seq(spx[1],spx[length(spx)],length.out=50)
sery <- spf(serx)
tailx <- spx[1]
taily <- spy[1]
headx <- spx[4]
heady <- spy[4]
arrows(headx,heady,headx+tree.x.lim*1e-3,heady, length=.1,col=col.arrow,lwd=lwd.arrow)
lines(rev(serx), rev(sery), lty=c(5,1),col=col.arrow, lwd=lwd.arrow)
bottle.txt <- "inbreeding\nLoH / LoS"
if(T){
  text((headx+tailx)/2+(headx-tailx)*(-.01), (heady+taily)/2+(heady-taily)*(-.10),
       bottle.txt, srt=66, font=font.arrow, cex=cex.arrow, col=col.arrow)
} else {
  # experiment at wrapping text along curved path; not too pretty, but retain for now, maybe revisit
  bottlec <- strsplit(bottle,split=NULL)[[1]]
  for(i in 1:length(bottlec)){
    text(xser[i],yser[i],bottlec[i], srt=65, font=4, cex=.7, col=col.arrow)
  }
}

####
#
# CLADE ANNOTATION
#
clade.L.x <- label.end.L + xdel
clade.H.x <- label.end.H + xdel
dy <- .33
lines(rep(clade.L.x,2),c(3-dy,7+dy),lwd=lwd.clade,col=col.clade)
lines(rep(clade.H.x,2),c(1-dy,2+dy),lwd=lwd.clade,col=col.clade)
text(clade.L.x+clade.dx,5.0,'L-clade',srt=90,font=font.clade,cex=cex.clade,col=col.clade)
text(clade.H.x+clade.dx,1.5,'H-clade',srt=90,font=font.clade,cex=cex.clade,col=col.clade)

####
#
# LEGEND
#
# parameter plusx controls whether we try to annotate b/c (+) and d/e (x) sharing in tree; I think
# it looks cluttered, rather than adding clarity, so I vote no, but code is here, in case. "Logic,"
# if any, for my symbol choice is that + overlaid on x looks like the * at the next level; this
# analogy is more visible if we use pch 3/4/8 rather than Courier or Helvetica chars, but probably
# should use same in both tree & legend, which will take a modicum of additional work.
legend.text <- c('a: only in 1014 ',
                'b: only in 1335 ',

```

```

        'c: only in 1015 ',
        'd: only in 1012 ',
        'e: only in 1007 ',
        '*: shared by bcde',
        paste(ifelse(plusx,'+:',' '), 'shared by b/c '),
        paste(ifelse(plusx,'x:',' '), 'shared by d/e ')
    )

    legend.text <- c('a: only in 1014 ',
                    'b: only in 1335 ',
                    'c: only in 1015 ',
                    'd: only in 1012 ',
                    'e: only in 1007 ',
                    '*: in bcde ',
                    paste(ifelse(plusx,'+:',' '), 'in bc '),
                    paste(ifelse(plusx,'x:',' '), 'in de '),
                    'Discordant SNPs ')
  )
  legend.text <- paste(legend.text, format(c(leg.counts, discord), width=4), sep=' - ')
  legend.text <- paste(legend.text, ' ') # add a little more right margin in box
  opar <- par(family="mono", cex=cex.legend)
  legend("topright", legend=legend.text, cex=cex.legend, inset=c(0.05, 0), bg=col.legbox, box.col=col.legbox)
  par(opar)
  if(plusx){
    points(tree.labels[[16]], tree.labels[[17]]+.14, pch=8, col=col.elabel)
    points(tree.labels[[16]]+200, tree.labels[[17]]+1, pch=3, col=col.elabel)
    points(tree.labels[[16]]+200, tree.labels[[17]]-1, pch=4, col=col.elabel)
  }

  ####
  #
  # EDGE LENGTHS
  #
  for(i in seq(1, length(tree.labels)-ifelse(plusx, 5, 2), by=3)){
    if(F){ # T for \n in edge labels; F to remove (except "by 5")
      text(tree.labels[[i]], tree.labels[[i+1]], tree.labels[[i+2]])
    } else {
      # points(tree.labels[[i]], tree.labels[[i+1]], pch=3, col='green') # for debugging
      text(tree.labels[[i]], tree.labels[[i+1]], sub('\n([`z])', ' \\1', tree.labels[[i+2]]),
           pos=3, offset=.4, font=font.elabel, col=col.elabel, cex=cex.elabel)
    }
  }
}

caption <- function(stringency, which.tables=which.snp.tables(string.val=F)){
  caption.where <- '(UNKNOWN genome subset).'
  if(which.tables[1]=='Chr1') {caption.where <- 'on Chr1.'}
  if(which.tables[1]=='full') {caption.where <- 'genome-wide.'}
  if(which.tables[1]=='trunc'){caption.where <- 'all Chrs.'}
  cap.stringency <- c(
    'loose SNP filters.',
    'medium SNP filters.',
    'strict SNP filters.',
    'unfiltered SNPs.')
  cap <- paste('Tree based on', which.tables[2], 'reads and', cap.stringency[stringency],
              'Lengths\\' are numbers of shared/private SNPs', caption.where)
  return(cap)
}

```

Trees based on all four SNP filtering criteria are shown below. Their topologies are exactly the same, although the branch lengths are different. In all four, the length of the branch labeled “*” is probably inflated by lower coverage and higher error rate in 1014, which may mask further legitimate sharing between it and the other L-isolates. The branch lengths among the other 4 are too short for their topology to be convincing without a more rigorous analysis (e.g., a bootstrap test), but detail there is irrelevant to the story.

My sense is that the “medium” version is the best for the paper, made here and shown in Fig 1. In theory, this should look exactly like Fig 3, but something is apparently different between Knitr and direct-to-pdf. (Increasing fig.width in Knitr’s chunk headers from 8 (as in the pdf call below) to 9 helps somewhat, but probably still best to make the paper fig directly rather than via Knitr.)

```

###
#
# MAKE PDF FOR PAPER
#
if(which.snp.tables() == 'trunc-unfiltered'){
  paperfig.path <- paste('figs-mine/Fig3-paperfig-medium-tree-', which.snp.tables(), '.pdf', sep='')
}

```

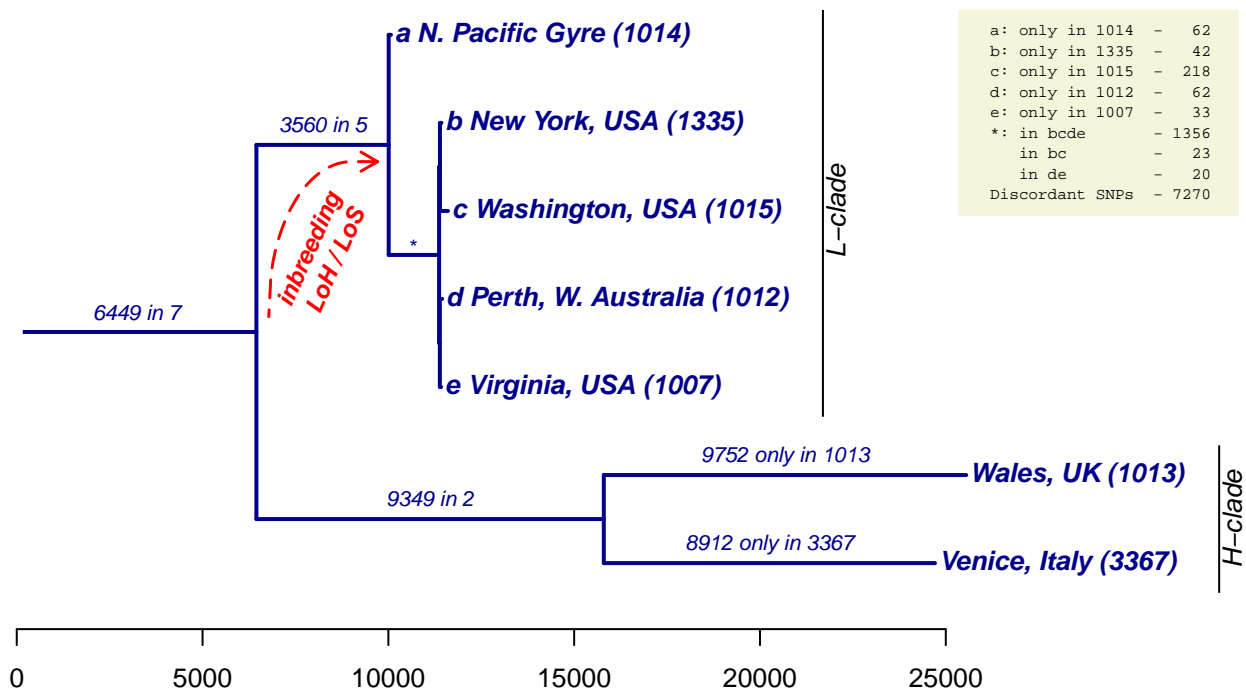



Figure 1: Proposed fig. for paper: Tree based on qfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1.

```

} else {
  paperfig.path <- paste('figs-mine/paperfig-medium-tree-', which.snp.tables(), '.pdf', sep='')
}
pdf(paperfig.path, width=8,height=5,onfile=TRUE,family='Helvetica',fonts='Courier',pointsize=10)
newick.medium <- newickize(make.tree(pat.summaries[, 'count2']))
show.tree(newick.medium, total.snps=consistent.count[2], pltdebug=F)
dev.off()

# pdf
# 2

```

```

# fig.paths for knitr chunks below; .h for "hand-made" trees; plain for automatic chr1/full versions
myfigpath <- paste(getwd(), '/figs-knitr/newick-', which.snp.tables(), '-', sep='')
myfigpath.h <- paste(getwd(), '/figs-knitr/newick-', sep='')

```

Figure 2, i.e., criteria [[1]]:

```

newick.loose <- newickize(make.tree(pat.summaries[, 'count1']))
show.tree(newick.loose, total.snps=consistent.count[1])

```

Figure 3, i.e. [[2]]:




Figure 2: Tree based on qfiltered reads and loose SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1.

```
# newick.medium <- newickize(tree.by.hand)
# simple.newick.medium <- newickize(tree.by.hand, alt=TRUE)
newick.medium <- newickize(make.tree(pat.summaries[, 'count2']))
simple.newick.medium <- newickize(make.tree(pat.summaries[, 'count2']), alt=TRUE)
show.tree(newick.medium, total.snps=consistent.count[2])
```

Figure 4, i.e. [[3]]:

```
newick.strict <- newickize(make.tree(pat.summaries[, 'count3']))
show.tree(newick.strict, total.snps=consistent.count[3])
```

Figure 5, i.e. [[4]]:

```
newick.unfiltered <- newickize(make.tree(pat.summaries[, 'count4']))
show.tree(newick.unfiltered, total.snps=consistent.count[4])
```

Some other versions of the trees are included in the appendix.

Counts for all tree edges in the medium tree:

```
#pat.summaries[c(128,110,102,6,97,19,9,2,5,33,65,17,3),]
tree.edges <- c(128,110,102,6,97,19,9,2,5,33,65,17,3)-1
non.edges <- setdiff(0:127, tree.edges)
sg.edges <- showgroup(restrict.to=tree.edges) ; sg.edges
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 2	001	1							X	35	42	62	135
# 3	002	1						X		7818	8912	9095	10949
# 5	004	1					X			176	218	264	385
# 6	005	2					X		X	17	23	40	52
# 9	010	1				X				42	62	53	113
# 17	020	1			X					8529	9752	9961	11677
# 19	022	2			X			X		8822	9349	8911	6177
# 33	040	1		X						53	62	103	174
# 65	100	1	X							25	33	38	141
# 97	140	2	X	X						14	20	25	85
# 102	145	4	X	X			X		X	598	1356	2429	2585
# 110	155	5	X	X		X	X		X	4136	3560	2135	3228
# 128	177	7	X	X	X	X	X	X	X	8132	6449	3873	1641
# Total										38397	39838	36989	37342




Figure 3: Tree based on qfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1.

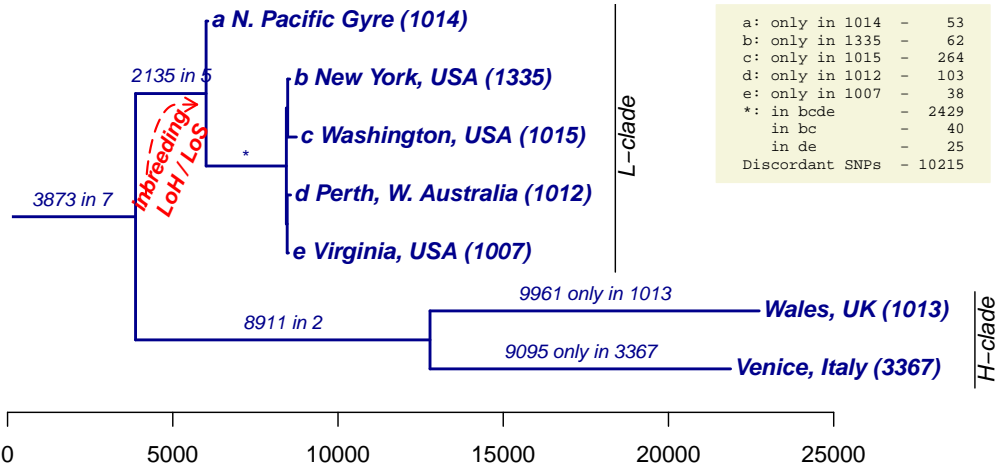


Figure 4: Tree based on qfiltered reads and strict SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1.




Figure 5: Tree based on qfiltered reads and unfiltered SNPs. “Lengths” are numbers of shared/private SNPs on Chr1.

Counts for the top 10 discordant patterns, i.e., SNPs whose sharing pattern does not match any of the bifurcations in the tree:

```
tenth <- sort(showgroup(restrict.to=non.edges)[- (length(non.edges)+1), 'count2'], decreasing=T)[10]
sg.non.edges <- showgroup(restrict.to=non.edges, c2.thresh = tenth) ; sg.non.edges
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 1	000	0								9	51	468	0
# 56	067	5		X	X		X	X	X	48	109	257	104
# 101	144	3	X	X			X	X	X	18	74	159	355
# 104	147	5	X	X			X	X	X	205	421	740	1160
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 117	164	4	X	X	X		X			19	51	71	163
# 118	165	5	X	X	X		X		X	318	591	957	1140
# 119	166	5	X	X	X		X	X		46	154	220	254
# 120	167	6	X	X	X		X	X	X	1305	2445	4098	1852
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# 127	176	6	X	X	X	X	X	X		57	79	43	68
# Other	(104 rows	w/	c2	<	51)			1436	867	1703	2302
# Total										6508	7270	10215	10157

And percent of discordant SNPs:

```
nsge <- nrow(sg.edges)
discordv <- consistent.count - sg.edges[nsge, c('count1', 'count2', 'count3', 'count4')] ; discordv
```

#	count1	count2	count3	count4
# Total	6508	7270	10215	10157

```
discordv.pct <- round(discordv/consistent.count*100,1) ; discordv.pct
```

#	count1	count2	count3	count4
# Total	14.5	15.4	21.6	21.4

In short, the sharing pattern observed at 7270 or 15.4% of the 47108 medium-stringency consistent SNPs positions observed across all 7 isolates are discordant with the medium tree. (The strict tree has slightly more.)

A majority of the discordant SNPs fall into one of three patterns: 6-way sharing excluding Gyre (likely a technical artifact since the low coverage in Gyre reduces our power to detect SNPs there), or 6-way sharing excluding one of

the two H-isolates (likely a reflection of sexuality in the H-clade—SNP positions in a population in Hardy-Weinberg equilibrium are fairly likely to be homozygous for the reference allele in a given individual).

```
third.biggest <- sort(showgroup(pat.summaries,6)[-8,'count2'],decreasing=T)[3]
big.three <- showgroup(pat.summaries,6,c2.thresh = third.biggest); big.three

#      Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 112 157    6    X    X    X    X    X    X    1338 1076    665 1343
# 120 167    6    X    X    X    X    X    X    1305 2445 4098 1852
# 126 175    6    X    X    X    X    X    X    1709 1352  834 1416
# Other (    4 rows w/  c2  < 1076 )      123    136    118    111
# Total                                4475  5009  5715  4722

big.three.frac <- sum(big.three[1:3,'count2'])/discordv$count2; big.three.frac

# [1] 0.6702889
```

I.e., 67% of discordant SNPs fall into one of these three categories.

Out of curiosity: what is the ratio of full genome to Chr 1 branch lengths. Except for the shortest few, generally $\approx 10x$, as expected given the length of Chr 1:

```
# (vectors derived by editing Newick strings, and in that order)
print(
  c(Italy=86155, Wales=95697, IW=89598, Virg=330, Aust=632, VA=1296,
    Puget=2113, NY=658, PNY=480, four=10059, Gyre=568, five=39517, all=69526) /
  c(Italy=8813, Wales=9652, IW=9365, Virg=30, Aust=61, VA=19,
    Puget=207, NY=41, PNY=18, four=1005, Gyre=61, five=3912, all= 7054),
  digits=3)

# Italy Wales IW Virg Aust VA Puget NY PNY four Gyre five all
# 9.78 9.91 9.57 11.00 10.36 68.21 10.21 16.05 26.67 10.01 9.31 10.10 9.86

round(genome.length.constants()$genome.length.trunc / genome.length.constants()$chr1.length, digits=4)

# [1] 10.2879
```

9 Semi-Automated Tree-Building

Slightly formalizing the process above: Look for the bifurcation of the 7 strains that maximizes the number of shared SNPs *within* each side of the partition while minimizing the number and fraction of SNPs that are shared by subsets that include at least one strain on each side of the partition. The 2/5 split is the winner, with 6418 SNPs in conflict with that partition (16% of the 39842 SNPs not shared by all 7; Chr1 data). The runner-up places the Gyre in a group by itself (7079 = 18% in conflict).

```
treepart <- function(p.summ=pat.summaries, root=127, verbose=T, stringency='count2'){
  root.shared <- p.summ[root+1,stringency]
  df<-NULL
  for(i in 1:floor(root/2)){
    if(bitwAnd(i,root)==i && i < root-i){
      l1 <- showgroup(p.summ,subset=i,split=NULL,proper.subset=F,total=T)
      l <- l1[nrow(l1),stringency]
      r1 <- showgroup(p.summ,subset=root-i,split=NULL,proper.subset=F,total=T)
      r <- r1[nrow(r1),stringency]
      c1 <- showgroup(p.summ,subset=root,split=i,proper.subset=T,total=T)
      c <- c1[nrow(c1),stringency]
      df <- rbind(df, data.frame(pat=i, left=l, right=r, both=l+r, cross=c, all=l+r+c, ratio=c/(l+r+c),
                                best=' ',stringsAsFactors=F))
    }
  }
  df$pat<-as.octmode(df$pat)
  maxl <- which.max(df$left)
  maxr <- which.max(df$right)
  maxb <- which.max(df$both)
  minc <- which.min(df$cross)
  minr <- which.min(df$ratio)
```

```

df$best[c(maxl,maxr,maxb,minc,minr)] <- '<'
df$best[maxl] <- paste(df$best[maxl], 'L') # max Left
df$best[maxr] <- paste(df$best[maxr], 'R') # max Right
df$best[maxb] <- paste(df$best[maxb], 'B') # max Both (L+R)
df$best[minc] <- paste(df$best[minc], 'C') # min Cross
df$best[minr] <- paste(df$best[minr], 'O') # min ratio (Cross/(Left+Right+Cross))
if(verbose){
  same <- all(maxl==c(maxr,maxb,minc,minr))
  cat('root:', format(as.octmode(root),width=3),
      '; shared:', root.shared,
      '. max l', format(as.octmode(df$pat[maxl]),width=3),
      ', max r', format(as.octmode(df$pat[maxr]),width=3),
      ', max both', format(as.octmode(df$pat[maxb]),width=3),
      ', min cross', format(as.octmode(df$pat[minc]),width=3),
      ', min ratio', format(as.octmode(df$pat[minr]),width=3),
      '. \nAll the same?:', same,
      '\n')
  cat('\n')
}
return(df)
}

```

treepart()

```

# root: 177 ; shared: 6449 . max l 077 , max r 010 , max both 010 , min cross 010 , min ratio 010 .
# All the same?: FALSE
#   pat  left right  both cross  all      ratio      best
# 1   01    93 29241 29334 11376 40710 0.2794399
# 2   02   8963 17601 26564 14146 40710 0.3474822
# 3   03   9020 10496 19516 21194 40710 0.5206092
# 4   04    269 28545 28814 11896 40710 0.2922132
# 5   05    334 28340 28674 12036 40710 0.2956522
# 6   06   9190 10106 19296 21414 40710 0.5260133
# 7   07   9273 10006 19279 21431 40710 0.5264309
# 8   10    113 34247 34360  6350 40710 0.1559813 < R B C O
# 9   11    160 28961 29121 11589 40710 0.2846721
# 10  12   9029 12483 21512 19198 40710 0.4715795
# 11  13   9093 10343 19436 21274 40710 0.5225743
# 12  14    332 28414 28746 11964 40710 0.2938836
# 13  15    406 28242 28648 12062 40710 0.2962908
# 14  16   9257 10017 19274 21436 40710 0.5265537
# 15  17   9353  9934 19287 21423 40710 0.5262343
# 16  20   9803 16282 26085 14625 40710 0.3592483
# 17  21   9852  9610 19462 21248 40710 0.5219356
# 18  22 28064  5697 33761  6949 40710 0.1706952
# 19  23 28151   607 28758 11952 40710 0.2935888
# 20  24 10036  9264 19300 21410 40710 0.5259150
# 21  25 10112  9160 19272 21438 40710 0.5266028
# 22  26 28337   301 28638 12072 40710 0.2965365
# 23  27 28470   231 28701 12009 40710 0.2949889
# 24  30   9870 11459 21329 19381 40710 0.4760747
# 25  31   9925  9471 19396 21314 40710 0.5235569
# 26  32 28144  1982 30126 10584 40710 0.2599853
# 27  33 28241   485 28726 11984 40710 0.2943748
# 28  34 10109  9178 19287 21423 40710 0.5262343
# 29  35 10199  9087 19286 21424 40710 0.5262589
# 30  36 28423   226 28649 12061 40710 0.2962663
# 31  37 28587   166 28753 11957 40710 0.2937116
# 32  40    113 28782 28895 11815 40710 0.2902235
# 33  41    159 28529 28688 12022 40710 0.2953083
# 34  42  9027 10293 19320 21390 40710 0.5254237
# 35  43  9098 10173 19271 21439 40710 0.5266274
# 36  44    348 28314 28662 12048 40710 0.2959469
# 37  45    461 28196 28657 12053 40710 0.2960698
# 38  46  9283  9971 19254 21456 40710 0.5270450
# 39  47  9450  9910 19360 21350 40710 0.5244412

```

```
# 40 50 176 28634 28810 11900 40710 0.2923115
# 41 51 229 28429 28658 12052 40710 0.2960452
# 42 52 9094 10190 19284 21426 40710 0.5263080
# 43 53 9175 10091 19266 21444 40710 0.5267502
# 44 54 419 28216 28635 12075 40710 0.2966102
# 45 55 558 28112 28670 12040 40710 0.2957504
# 46 56 9360 9895 19255 21455 40710 0.5270204
# 47 57 9567 9841 19408 21302 40710 0.5232621
# 48 60 9873 9454 19327 21383 40710 0.5252518
# 49 61 9934 9320 19254 21456 40710 0.5270450
# 50 62 28157 478 28635 12075 40710 0.2966102
# 51 63 28278 380 28658 12052 40710 0.2960452
# 52 64 10140 9141 19281 21429 40710 0.5263817
# 53 65 10293 9068 19361 21349 40710 0.5244166
# 54 66 28519 200 28719 11991 40710 0.2945468
# 55 67 28886 147 29033 11677 40710 0.2868337
# 56 70 9941 9362 19303 21407 40710 0.5258413
# 57 71 10012 9247 19259 21451 40710 0.5269221
# 58 72 28240 396 28636 12074 40710 0.2965856
# 59 73 28379 312 28691 12019 40710 0.2952346
# 60 74 10223 9065 19288 21422 40710 0.5262098
# 61 75 10416 9000 19416 21294 40710 0.5230656
# 62 76 28629 131 28760 11950 40710 0.2935397
# 63 77 29112 84 29196 11514 40710 0.2828298 < L
```

Comparing the 5/2 split to the second-place NPG/rest split (below), the former has fewer pattern instances in conflict with the split (6418 vs 7079), as well as somewhat more random distribution of the conflicting patterns (92 vs 62 rows), whereas the 1/6 split has the majority of its conflicts (3912 of 7079, or 55%) concentrated in one pattern—the 5 NE strains. Collectively, these seem to favor the 5/2 split as the correct “history.”

```
showgroup(pat.summaries,split=strtoi('022'), subset=127, proper.subset=T, c2.thresh=100)
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 56	067	5		X	X		X	X	X	48	109	257	104
# 104	147	5	X	X			X	X	X	205	421	740	1160
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 118	165	5	X	X	X		X		X	318	591	957	1140
# 119	166	5	X	X	X		X	X		46	154	220	254
# 120	167	6	X	X	X		X	X	X	1305	2445	4098	1852
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# Other	(85 rows	w/	c2	<	100)			1415	801	1317	1735
# Total										6384	6949	9088	9004

```
showgroup(pat.summaries,split=strtoi('010'), subset=127, proper.subset=T, c2.thresh=100)
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 110	155	5	X	X		X	X		X	4136	3560	2135	3228
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# Other	(59 rows	w/	c2	<	100)			470	362	335	590
# Total										7653	6350	3969	6577

Below is the full summary of shared SNPs that do *not* directly correspond to tree splits, e.g. deep coalescence, independent coincident mutations, false positives/false negatives in the shared SNP calls, loss of SNPs in hemizygous regions, etc. (Additionally, SAMTools’ SNP calls exclude positions it judges to be homozygous, and I think it operates without regard to the reference sequence, so homozygous nonreference positions, while rare except in IT/Wales, often are not called SNPs by SAMTools, but are relevant for this analysis. Provided the position is called a SNP in some other isolate, the consistency filtering we’ve done above should recover it, but this is still worth keeping in mind when examining the data.)

First, here are SNPs that “coalesce” on the branch from the LCA of bcde, i.e., shared among some nonempty, proper subset of bcde other than bc or de. There are 8 such patterns: any of the 4 choose 3 trios plus any of the 4 pairs having exactly one of bc.

```
sg4 <- showgroup(pat.summaries, subset=strtoi('0145'), split=5, proper.subset = F)
sg4
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 34	041	2		X					X	2	4	23	36
# 37	044	2		X			X			6	17	80	155
# 38	045	3		X			X		X	12	44	184	131
# 66	101	2	X						X	3	5	7	20
# 69	104	2	X				X			2	10	40	107
# 70	105	3	X				X		X	9	14	37	63
# 98	141	3	X	X					X	5	9	20	40
# 101	144	3	X	X			X			18	74	159	355
# 102	145	4	X	X			X		X	598	1356	2429	2585
# Total										655	1533	2979	3492

```
sg4n <- nrow(sg4)
sg4pct <- round(sg4$count2[sg4n-1]/sg4$count2[sg4n]*100,1)
sg4pct
```

```
# [1] 88.5
```

So, of the 1533 SNPs found only in bcde, 88.5% have a sharing pattern consistent with the given tree structure.

Similarly, we analyze patterns relative to the root of the L-clade (14 patterns—any nonempty proper subset of bcde together with a):

```
sg5 <- showgroup(pat.summaries, subset=strtoi('0155'), split=8, proper.subset = F)
sg5
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 10	011	2				X			X	3	5	4	14
# 13	014	2				X	X			4	1	6	12
# 14	015	3				X	X		X	6	4	4	7
# 41	050	2		X		X				5	1	3	7
# 42	051	3		X		X			X	0	2	4	4
# 45	054	3		X		X	X			1	7	12	18
# 46	055	4		X		X	X		X	8	15	24	36
# 73	110	2	X			X				2	1	3	7
# 74	111	3	X			X			X	0	1	1	1
# 77	114	3	X			X	X			2	4	4	8
# 78	115	4	X			X	X		X	1	4	6	9
# 105	150	3	X	X		X				0	1	0	6
# 106	151	4	X	X		X			X	2	2	4	14
# 109	154	4	X	X		X	X			24	45	34	103
# 110	155	5	X	X		X	X		X	4136	3560	2135	3228
# Total										4194	3653	2244	3474

```
sg5n <- nrow(sg5)
sg5pct <- round(sg5$count2[sg5n-1]/sg5$count2[sg5n]*100,1)
```

I.e., of the 3653 SNPs found only in abcde, 97.5% have a sharing pattern consistent with the given tree structure.

Finally, how many SNPs have patterns inconsistent with the 5-2 split, i.e., include at least one strain on each side of the 5-2 split, but not shared by all 7?

```
sg7 <- showgroup(pat.summaries, subset=127, split=strtoi('022'), proper.subset=F)
sg7
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 4	003	2						X	X	87	15	37	31
# 7	006	2					X	X		65	9	20	41
# 8	007	3					X	X	X	4	3	19	10
# 11	012	2				X		X		41	4	2	5
# 12	013	3				X		X	X	3	2	2	2
# 15	016	3				X	X	X		1	0	1	2
# 16	017	4				X	X	X	X	1	2	1	2
# 18	021	2			X				X	102	7	24	9

# 20	023	3			X			X	X	89	23	35	17
# 21	024	2			X		X			62	15	21	46
# 22	025	3			X		X		X	8	4	23	21
# 23	026	3			X		X	X		84	31	37	32
# 24	027	4			X		X	X	X	15	16	37	24
# 25	030	2			X	X				65	5	1	9
# 26	031	3			X	X			X	3	1	0	0
# 27	032	3			X	X		X		66	9	6	5
# 28	033	4			X	X		X	X	4	2	4	6
# 29	034	3			X	X	X			2	5	1	1
# 30	035	4			X	X	X		X	2	4	0	0
# 31	036	4			X	X	X	X		5	0	3	3
# 32	037	5			X	X	X	X	X	12	11	8	5
# 35	042	2		X					X	47	2	20	27
# 36	043	3		X					X	8	10	14	6
# 39	046	3		X			X	X		4	12	41	55
# 40	047	4		X			X	X	X	9	26	68	60
# 43	052	3		X		X		X		1	0	0	1
# 44	053	4		X		X		X	X	0	1	1	2
# 47	056	4		X		X	X	X		2	2	2	5
# 48	057	5		X		X	X	X	X	4	9	8	17
# 49	060	2		X	X					59	8	24	28
# 50	061	3		X	X				X	1	8	11	12
# 51	062	3		X	X			X		86	21	29	36
# 52	063	4		X	X			X	X	9	12	34	21
# 53	064	3		X	X		X			2	17	52	60
# 54	065	4		X	X		X		X	8	21	68	48
# 55	066	4		X	X		X	X		15	43	102	76
# 56	067	5		X	X		X	X	X	48	109	257	104
# 57	070	3		X	X	X				3	0	0	2
# 58	071	4		X	X	X			X	0	2	0	0
# 59	072	4		X	X	X		X		4	2	1	2
# 60	073	5		X	X	X		X	X	3	3	3	3
# 61	074	4		X	X	X	X			1	2	1	7
# 62	075	5		X	X	X	X		X	8	7	10	11
# 63	076	5		X	X	X	X	X		9	10	12	7
# 64	077	6		X	X	X	X	X	X	43	46	62	26
# 67	102	2	X						X	34	4	11	25
# 68	103	3	X						X	2	3	6	8
# 71	106	3	X				X	X		4	10	9	27
# 72	107	4	X				X	X	X	6	10	6	16
# 75	112	3	X			X		X		0	1	0	0
# 76	113	4	X			X		X	X	0	0	0	1
# 79	116	4	X			X	X	X		1	0	1	5
# 80	117	5	X			X	X	X	X	1	1	3	7
# 81	120	2	X		X					49	5	7	31
# 82	121	3	X		X				X	2	0	2	4
# 83	122	3	X		X			X		45	6	12	26
# 84	123	4	X		X			X	X	5	9	13	8
# 85	124	3	X		X		X			2	7	21	35
# 86	125	4	X		X		X		X	3	4	14	16
# 87	126	4	X		X		X	X		10	17	14	43
# 88	127	5	X		X		X	X	X	13	27	49	47
# 89	130	3	X		X	X				1	1	1	1
# 90	131	4	X		X	X			X	0	0	0	2
# 91	132	4	X		X	X		X		1	1	1	1
# 92	133	5	X		X	X		X	X	2	3	0	0
# 93	134	4	X		X	X	X			6	3	2	3
# 94	135	5	X		X	X	X		X	4	2	0	7
# 95	136	5	X		X	X	X	X		5	3	0	5
# 96	137	6	X		X	X	X	X	X	8	6	8	14
# 99	142	3	X	X				X		3	3	9	15
# 100	143	4	X	X				X	X	1	3	4	20
# 103	146	4	X	X			X	X		9	34	69	140
# 104	147	5	X	X			X	X	X	205	421	740	1160
# 107	152	4	X	X		X		X		1	3	1	4
# 108	153	5	X	X		X		X	X	4	0	0	7

# 111	156	5	X	X		X	X	X		11	7	9	43
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 113	160	3	X	X	X					6	3	7	23
# 114	161	4	X	X	X				X	3	8	10	18
# 115	162	4	X	X	X			X		8	11	20	33
# 116	163	5	X	X	X			X	X	15	14	21	33
# 117	164	4	X	X	X		X			19	51	71	163
# 118	165	5	X	X	X		X		X	318	591	957	1140
# 119	166	5	X	X	X		X	X		46	154	220	254
# 120	167	6	X	X	X		X	X	X	1305	2445	4098	1852
# 121	170	4	X	X	X	X				0	1	1	3
# 122	171	5	X	X	X	X			X	4	4	2	7
# 123	172	5	X	X	X	X		X		3	6	3	5
# 124	173	6	X	X	X	X		X	X	15	5	5	3
# 125	174	5	X	X	X	X	X			5	14	17	35
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# 127	176	6	X	X	X	X	X	X		57	79	43	68
# 128	177	7	X	X	X	X	X	X	X	8132	6449	3873	1641
# Total										14516	13398	12961	10645

```
sg7n <- nrow(sg7)
sg7pct <- round(sg7$count2[sg7n-1]/sg7$count2[sg7n]*100,1)
sg7pct

# [1] 48.1
```

A more compact version of that table, showing only the larger counts:

```
thresh <- signif(.02 * sg7$count2[sg7n],1)
thresh

# [1] 300

showgroup(pat.summaries, subset=127, split=strtoi('022'), proper.subset=F, c2.thresh = thresh)
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 104	147	5	X	X			X	X	X	205	421	740	1160
# 112	157	6	X	X		X	X	X	X	1338	1076	665	1343
# 118	165	5	X	X	X		X		X	318	591	957	1140
# 120	167	6	X	X	X		X	X	X	1305	2445	4098	1852
# 126	175	6	X	X	X	X	X		X	1709	1352	834	1416
# 128	177	7	X	X	X	X	X	X	X	8132	6449	3873	1641
# Other	(87 rows	w/	c2	<	300)			1509	1064	1794	2093
# Total										14516	13398	12961	10645

So, of the 13398 SNPs found both in the L- and H-clades, 48.1% have a sharing pattern consistent with the given tree structure, i.e., are found in all 7 isolates. Among the others, three patterns dominate—(i) the 6-way pattern excluding the Gyre is the largest, plausibly explained by 7-way sharing from which the Gyre drops out due to low coverage/high error rate, (ii) the 6-way excluding Italy, and (iii) ditto for Wales. Origin of the later two cases is unclear, but may partly reflect Hardy-Weinberg—some positions that are *population-level* SNPs in those isolates will be homozygous-reference in the CCMP founder cell for IT or Wales. If I take the 7-way shared SNP count (69526) as a surrogate approximating the number of population-level SNPs in either IT or Wales that are shared with the L-clade, then I might expect, based on HWE, roughly half that number to be lost (become homozygous) in IT, and a similar number in Wales. However, the observed counts of these positions are lower by $\approx 20K$ than I might have guessed from HWE, perhaps suggesting that IT and Wales are distinct populations, each with a pool of many thousand private polymorphisms.

In aggregate:

```
untreelike <-
  sg7$count2[sg7n]-sg7$count2[sg7n-1] +
  sg5$count2[sg5n]-sg5$count2[sg5n-1] +
  sg4$count2[sg4n]-sg4$count2[sg4n-1]
untreelike
```

```
# [1] 7219

consistent.count[2]

# [1] 47108

unpct <- round(untreelike/consistent.count[2]*100,1)
unpct

# [1] 15.3
```

I.e., 7219 or 15.3% of the 47108 consistent SNPs identified (by criterion 2) across all 7 isolates are discordant with the assumed tree.

Overall, based on this data, I take the following to be obvious: (a) separation of the the H-isolates from the L-isolates (and from each other??), and (b) near-identity of the L-isolates. Due to the small counts, the exact topology among the L-isolates (esp. bcde) is uncertain, but *any* topology there is consistent with the asexual/clonal/global-expansion hypothesis, so there is little point in examining this subtree more carefully. Again, we believe the (apparent) slight separation of the Gyre from the other L-isolates is largely driven by technical artifacts (lower coverage/higher error rates) in the sequencing rather than by biological effects. However, the discord between Gyre SNPs and others is the major substantive ambiguity in the offered tree. Nevertheless, in the next section we show by a bootstrap analysis that the offered placement of Gyre with respect to the other 4 L-isolates is strongly supported by the data.

9.1 Bootstrap

How robust is the inferred tree? Italy/Wales seem clearly related to each other but separate from the other 5. Likewise, the 4 coastal L-isolates seem to be closely related, with little data to separate them (and perhaps little sense in trying). So, the key question here is whether the top level bifurcation is 2/5 or NPG/6. Here, we do a simple bootstrap test (on c2 numbers only) to see whether the 2/5 split is consistently the most parsimonious.

```
n2 <- sum(pattern.counts[[2]][,2]); n2

# [1] 47108
```

Conceptually, we sample, with replacement, $n_2=47108$ SNP positions from among the 47108 positions declared consistent SNPs according to criterion c2, and recalculate the statistics examined above to see whether the 2/5 split again minimizes conflicting sharing patterns. This resampling/calculation is repeated `nboot` times (set near front of file). Since all that matters is the sharing pattern, this procedure is expedited by actually sampling 47108 independent integers in the range 0:127 with probabilities proportional to the pattern counts given in column 2 of `pattern.counts[[2]]`. The sample is then tabulated in a 128 row table analogous to `pattern.summaries`, for analysis by `showgroups/treepart`, as above.

```
boot.sample <- sample(0:127,n2,replace=T,prob=pattern.counts[[2]][,2])
str(boot.sample)

# int [1:47108] 127 18 125 18 18 16 127 109 2 101 ...

boot.count <- mytable(boot.sample,c(0,127))
boot.count[c(1:4,125:128),] # show a few rows

#      val count
# [1,]    0    51
# [2,]    1    46
# [3,]    2  8744
# [4,]    3    14
# [5,]  124    15
# [6,]  125 1333
# [7,]  126    68
# [8,]  127 6555

boot.counts <- list(NULL,boot.count,NULL) # dummy list with just c2 summaries
cor(pattern.counts[[2]][,2],boot.counts[[2]][,2]) # just curious - how correlated are they?
```

```
# [1] 0.9998476

boot.summaries <- pat.summary(boot.counts)
showgroup(boot.summaries, c2.thresh=400) #show a few rows
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 3	002	1						X		NA	8744	NA	NA
# 17	020	1			X					NA	9838	NA	NA
# 19	022	2			X			X		NA	9397	NA	NA
# 102	145	4	X	X			X		X	NA	1403	NA	NA
# 110	155	5	X	X		X	X		X	NA	3638	NA	NA
# 112	157	6	X	X		X	X	X	X	NA	1103	NA	NA
# 118	165	5	X	X	X		X		X	NA	594	NA	NA
# 120	167	6	X	X	X		X	X	X	NA	2259	NA	NA
# 126	175	6	X	X	X	X	X		X	NA	1333	NA	NA
# 128	177	7	X	X	X	X	X	X	X	NA	6555	NA	NA
# Other	(118 rows	w/	c2	<	400)			NA	2244	NA	NA
# Total										NA	47108	NA	NA

Tree partition analysis (and how to pluck out only the best rows based on 3 smallest cross counts and “best” criteria):

```
tp <- treepart(boot.summaries, root=127) ; tp
```

```
# root: 177 ; shared: 6555 . max l 077 , max r 010 , max both 010 , min cross 010 , min ratio 010 .
# All the same?: FALSE
# pat left right both cross all ratio best
# 1 01 97 29198 29295 11309 40604 0.2785194
# 2 02 8795 17813 26608 13996 40604 0.3446951
# 3 03 8855 10594 19449 21155 40604 0.5210078
# 4 04 289 28500 28789 11815 40604 0.2909812
# 5 05 350 28293 28643 11961 40604 0.2945769
# 6 06 9038 10198 19236 21368 40604 0.5262536
# 7 07 9116 10083 19199 21405 40604 0.5271648
# 8 10 104 34045 34149 6455 40604 0.1589745 < R B C O
# 9 11 154 28933 29087 11517 40604 0.2836420
# 10 12 8851 12630 21481 19123 40604 0.4709635
# 11 13 8916 10441 19357 21247 40604 0.5232736
# 12 14 345 28380 28725 11879 40604 0.2925574
# 13 15 413 28208 28621 11983 40604 0.2951187
# 14 16 9097 10115 19212 21392 40604 0.5268446
# 15 17 9185 10020 19205 21399 40604 0.5270170
# 16 20 9889 16279 26168 14436 40604 0.3555315
# 17 21 9941 9461 19402 21202 40604 0.5221653
# 18 22 28030 5851 33881 6723 40604 0.1655748
# 19 23 28116 625 28741 11863 40604 0.2921633
# 20 24 10142 9106 19248 21356 40604 0.5259580
# 21 25 10217 8985 19202 21402 40604 0.5270909
# 22 26 28316 311 28627 11977 40604 0.2949709
# 23 27 28437 222 28659 11945 40604 0.2941828
# 24 30 9946 11346 21292 19312 40604 0.4756182
# 25 31 10002 9328 19330 21274 40604 0.5239385
# 26 32 28100 2053 30153 10451 40604 0.2573884
# 27 33 28192 506 28698 11906 40604 0.2932223
# 28 34 10206 9030 19236 21368 40604 0.5262536
# 29 35 10291 8926 19217 21387 40604 0.5267215
# 30 36 28393 240 28633 11971 40604 0.2948232
# 31 37 28541 167 28708 11896 40604 0.2929761
# 32 40 112 28742 28854 11750 40604 0.2893804
# 33 41 162 28504 28666 11938 40604 0.2940104
# 34 42 8859 10392 19251 21353 40604 0.5258841
# 35 43 8935 10275 19210 21394 40604 0.5268939
# 36 44 365 28272 28637 11967 40604 0.2947247
# 37 45 477 28153 28630 11974 40604 0.2948971
# 38 46 9131 10060 19191 21413 40604 0.5273618
# 39 47 9300 9990 19290 21314 40604 0.5249237
# 40 50 166 28600 28766 11838 40604 0.2915476
# 41 51 224 28412 28636 11968 40604 0.2947493
# 42 52 8916 10293 19209 21395 40604 0.5269185
# 43 53 9001 10198 19199 21405 40604 0.5271648
# 44 54 429 28183 28612 11992 40604 0.2953404
# 45 55 570 28079 28649 11955 40604 0.2944291
# 46 56 9200 9991 19191 21413 40604 0.5273618
# 47 57 9415 9930 19345 21259 40604 0.5235691
# 48 60 9961 9298 19259 21345 40604 0.5256871
```

```
# 49 61 10023 9167 19190 21414 40604 0.5273865
# 50 62 28122 492 28614 11990 40604 0.2952911
# 51 63 28242 397 28639 11965 40604 0.2946754
# 52 64 10244 8978 19222 21382 40604 0.5265984
# 53 65 10392 8895 19287 21317 40604 0.5249975
# 54 66 28500 206 28706 11898 40604 0.2930253
# 55 67 28857 142 28999 11605 40604 0.2858093
# 56 70 10019 9213 19232 21372 40604 0.5263521
# 57 71 10090 9102 19192 21412 40604 0.5273372
# 58 72 28197 414 28611 11993 40604 0.2953650
# 59 73 28333 335 28668 11936 40604 0.2939612
# 60 74 10319 8912 19231 21373 40604 0.5263767
# 61 75 10508 8839 19347 21257 40604 0.5235199
# 62 76 28605 144 28749 11855 40604 0.2919663
# 63 77 29098 89 29187 11417 40604 0.2811792 < L
```

```
otp <- order(tp[, 'cross'])[1:3] # 3 smallest 'cross' counts
btp <- which(tp[, 'best'] != '') # 'best' by Left/Right/Both/Cross/ratio
toptp <- unique(c(otp, btp, 18, 8)) # above, plus 5/2, 6/1 splits
print(tp[toptp,]) # show the winners
```

```
# pat left right both cross all ratio best
# 8 10 104 34045 34149 6455 40604 0.1589745 < R B C O
# 18 22 28030 5851 33881 6723 40604 0.1655748
# 26 32 28100 2053 30153 10451 40604 0.2573884
# 63 77 29098 89 29187 11417 40604 0.2811792 < L
```

Now repeat the above nboot times, and summarize results:

```
nboot <- params$nboot # default from params set in section 2
nboot <- ((nboot+2) %/% 4) * 4 + 1 # summary is cleaner if n mod 4 == 1, so int median/quartiles
cat('***\n*** Doing', nboot, 'bootstrap replicates.\n***\n')

# ***
# *** Doing 5 bootstrap replicates.
# ***

bcor <- numeric(nboot)
b52cross <- integer(nboot)
b61cross <- integer(nboot)
brev <- logical(nboot)
for(i in 1:nboot){
  boot.sample <- sample(0:127, n2, replace=T, prob=pattern.counts[[2]][,2])
  boot.count <- mytable(boot.sample, c(0, 127))
  boot.counts <- list(NULL, boot.count, NULL) # dummy list with just c2 summaries
  boot.summaries <- pat.summary(boot.counts)
  tp <- treepart(boot.summaries, root=127, verbose=F)
  bcor[i] <- cor(pattern.counts[[2]][,2], boot.counts[[2]][,2]) # just curious - how correlated are they?
  b52cross[i] <- tp[18, 'cross']
  b61cross[i] <- tp[ 8, 'cross']
  brev[i] <- (b52cross[i] > b61cross[i])
  if(brev[i]){
    # show the unexpected ones; probably breaks w/ cache
    otp <- order(tp[, 'cross'])[1:3]
    btp <- which(tp[, 'best'] != '')
    toptp <- unique(c(otp, btp, 18, 8))
    print(tp[toptp,])
  }
}

# pat left right both cross all ratio best
# 8 10 116 34263 34379 6341 40720 0.1557220 < R B C O
# 18 22 28026 5683 33709 7011 40720 0.1721758
# 26 32 28116 2000 30116 10604 40720 0.2604126
# 63 77 29060 79 29139 11581 40720 0.2844057 < L
# pat left right both cross all ratio best
# 8 10 94 34228 34322 6341 40663 0.1559403 < R B C O
# 18 22 28053 5631 33684 6979 40663 0.1716302
```

```
# 26 32 28124 1969 30093 10570 40663 0.2599415
# 63 77 29123 80 29203 11460 40663 0.2818287 < L
# pat left right both cross all ratio best
# 8 10 112 34274 34386 6239 40625 0.1535754 < R B C O
# 18 22 28009 5760 33769 6856 40625 0.1687631
# 26 32 28098 2058 30156 10469 40625 0.2576985
# 63 77 29074 94 29168 11457 40625 0.2820185 < L
# pat left right both cross all ratio best
# 8 10 103 34244 34347 6282 40629 0.1546186 < R B C O
# 18 22 28010 5767 33777 6852 40629 0.1686480
# 26 32 28090 1980 30070 10559 40629 0.2598883
# 63 77 29080 82 29162 11467 40629 0.2822368 < L
# pat left right both cross all ratio best
# 8 10 110 34422 34532 6212 40744 0.1524642 < R B C O
# 18 22 28225 5667 33892 6852 40744 0.1681720
# 26 32 28299 1973 30272 10472 40744 0.2570194
# 63 77 29283 88 29371 11373 40744 0.2791331 < L

# summarize:
corsummary <- t(as.matrix(c(summary(bcor),sd=sd(bcor))))
row.names(corsummary) <- 'bcor'
bdelta <- b61cross-b52cross
brevp <- 100*brev # make it percent reversed instead of logical
thesummary <- rbind(summary(b52cross), summary(b61cross), summary(c(bdelta)), summary(brevp))
row.names(thesummary) <- c('b52cross', 'b61cross', 'b61-b52', '% rev')
thesummary <- cbind(thesummary, sd=c(sd(b52cross),sd(b61cross),sd(bdelta),sd(brevp)))
```

SUMMARY: In 5 bootstrap replicates, we saw 5 samples with the 6/1 split having fewer conflicts than the 5/2 split, and the minimum separation between them was ≈ -18 sigma, hence highly statistically significant.

```
# 'opt' hacking is trying to force knitr to show more digits of bcor in summary, as Rstudio does, but
# it still fails... Bottom line is the correlation seems to be > .999 in all samples, rounds to 1.0,
# as seen in this run of 1001 samples cut/paste from Rstudio:
#      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.      sd
# bcor   " 0.9998" " 0.9999" " 0.9999" " 0.9999" " 1" " 1" " 0.00003462"
# > max(bcor)
# [1] 0.9999915
o.opts <- options(digits=7,width=127)
format(rbind(corsummary,thesummary),scientific=F,digits=4,drop0trailing=T)

#      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.      sd
# bcor   " 0.9998723" " 0.9999124" " 0.9999239" " 0.9999219" " 0.9999498" " 0.999951" " 0.0000323
# b52cross "6852"      "6852"      "6856"      "6910"      "6979"      "7011"      " 78.4314988
# b61cross "6212"      "6239"      "6282"      "6283"      "6341"      "6341"      " 58.5363135
# b61-b52  "-670"      "-640"      "-638"      "-627"      "-617"      "-570"      " 37.0405184
# % rev    " 100"      " 100"      " 100"      " 100"      " 100"      " 100"      " 0"

options(o.opts)
```

Based on this, it is reasonable to claim that we are confident that the tree topology is as shown in the earlier figures, with the exception of the exact order of the splits with the 4 NE coastal isolates.

10 Notes

This section is a random brain dump of limitations of the current analysis, ideas for improvements, etc. In the main, these may not be worth doing, unless we see significant holes or get pushed by reviewers, etc, but I wanted to catalog before we forget them.

Noise: Various sources of “noise” in the data:

1. Read errors, low read depth — perhaps fixed by medium/strict thresholding
2. Deep coalescence

3. Skew because 1335 is the reference. (Julie notes we could partially fix this by remapping based on discovered SNPs, tho that wouldn't fix gross misassembly in 1335, e.g. collapsed or misordered tandem duplicates, or segments missing in 1335 that are present in one or more other strains, etc.; much harder to fix those, let's just hope they are rare...)
4. Varying error rates and sequencing depth among the 7. E.g., plausibly the 1000 SNPs shared by 4 but not by Gyre are a result of lower read depth (we missed a SNP that is actually present) and/or higher error rates (causing a position to appear inconsistent in gyre) in the gyre data. I can't think of a way to correct for this effect. It might be possible, perhaps by simulation, to estimate the size of the effect and see whether it could explain ≈ 1000 SNPs.
5. Varying numbers of founder cells in the sequencing cultures. (Again, I made some attempts at modeling this, but nothing very satisfactory yet.)
6. Tri-allelic positions where stochastic fluctuation in sequence sampling promotes the rare allele to prominence. (Julie replies: "isn't this the same as more than one founder cell? If they are diploid there should only ever be two alleles, unless there were random and very rare, thus unlikely, trisomy events?" I agree, but it is a concrete example of an effect of multiple founders that might be important. Not sure this is the most important such effect...)
7. Gaps/indels - alignments are likely to be of lower quality in the vicinity of an indel, so, maybe lower coverage/more SNPs. We ignored them. Does this add any systematic bias? e.g. if one strain had more indels than another, would this confound other analyses? unclear. Julie suggested a paper titled "Barking up the wrong tree-length: yada yada yada gap penalties"; maybe relevant?

Other Items/Potential To Dos:

1. any spacial structure to various sub-classes?
2. after top level split, should I reanalyze halves of partition in isolation? said another way, I think the tree-building is sensible, but not sure it's optimal.
3. if we believe no sex, then I think gain of SNP should be more common than loss of SNP, since the later can only happen by (a) mutation reverting to reference, (b) second mutation matching nonreference, (c) homologous repair (look for blocks of LOH), or (d) false negative e.g. from low read depth. Does tree-building appropriately weight the gain vs loss cases? (Does it even care?)
4. should we weight coding and/or nonsynonymous SNPs more heavily? Julie says "you do not want to weight the coding or nonsynonymous/coding SNPs because for time you want the more clock-like neutral mutations." I.e., I got this backwards. Maybe should redo tree based on noncoding SNPs only.
5. We could also do an actual parsimony analysis based on 2-state model (homozygous-ref vs not), but I'm not quite sure how to handle this in a mixed sex/nosex case.
6. Might be interesting to look at sharing just within (shared?) deserts. Given tree model above and expectation that bottleneck followed split of H- from L-clades, I would expect little or no sharing of L-clade desert SNPs with H-clade; sharing between It/Wales might suggest "desert" is actually a region under strong purifying selection (e.g. a gene); sharing/non-sharing within L-clade deserts might suggest more about evo history of the 5.

11 Appendix: Old Trees, etc.

Tangents, old stuff of historical interest at best, etc..

11.1 HWE Sharing

Tangent: As a function of nonref allele freq, assuming HWE, what is probability that nonref allele will be seen in k strains, $0 \leq k \leq 4$ (Fig 6).




Figure 6: Sharing Probability

```
myfigpath.h <- paste(getwd(), '/figs-knitr/', sep='')

```

```
p <- (0:20)/20
q <- 1-p
r <- 2*p*q+p^2
plot(p, 1*q^0*r^4, type='b', pch='4', ylab="share prob")
points(p, 4*q^2*r^3, type='b', pch='3')
points(p, 6*q^4*r^2, type='b', pch='2')
points(p, 4*q^6*r^1, type='b', pch='1')
points(p, 1*q^8*r^0, type='b', pch='0')

```

11.2 Old Tree Stuff

All based on un-q-filtered reads.

The first pass at the tree analysis was the Chr1 tree, *loose criteria* (c1); it is rendered via <http://iubio.bio.indiana.edu/treeapp/treeprint-form.html> as Fig 7, and in newick format is:

```
newick.chr1.loose <- '(((tp3367_Italy:4551,tp1013_Wales:4954):5920,(((tp1007_Virginia:10,tp1012_Australia:29):9,(
cat.hardwrap(newick.chr1.loose)

# (((tp3367_Italy:4551,tp1013_Wales:4954):5920,(((tp1007_Virginia:10,tp1012_Austra
# lia:29):9,(tp1015_Puget_Sound:90,tp1335_NY:13):11):320,tp1014_Gyre:22):3484):859
# 3,outgroup:0);

```

Chr 1 tree based on *medium criteria* (c2) has exactly the same topology is, although the branch lengths are different. As noted earlier, the length of the branch labeled “*” is probably inflated by lower coverage and higher error rate in 1014, which may mask further legitimate sharing between it and the other L-isloates. The branch lengths among the other 4 are too short for its topology to be convincing without a more rigorous analysis (e.g., a bootstrap test).




Figure 7: Inferred Tree, based on Chr1, un-q-filtered reads, loose criteria. (Note: to visually resolve the edges among the 5, their lengths were scaled by 5x – 10x in this figure, but not in the newick description shown in the text.)

Chr1 tree, medium criteria, in newick format:

```
newick.chr1.med <- '(((tp3367_Italy:8813,tp1013_Wales:9652):9365,((e_tp1007_Virginia:30,d_tp1012_Australia:61):19,(c_tp1015_Puget_Sound:207,b_tp1335_NY:41):18):1005,a_tp1014_Gyre:61):3912):7054,outgroup:0);
cat.hardwrap(newick.chr1.med)

# (((tp3367_Italy:8813,tp1013_Wales:9652):9365,((e_tp1007_Virginia:30,d_tp1012_Australia:61):19,(c_tp1015_Puget_Sound:207,b_tp1335_NY:41):18):1005,a_tp1014_Gyre:61):3912):7054,outgroup:0);
```

NOTE: In early code, tree was not being recalculated; it was defined by constants in the following code chunk, hand-copied from the analysis above.

```
# tree parameters as nested lists
# Internal nodes have subtrees sub1/2 and length
# Root has sub1/2, but no length
# Leaves have where, length, optionally, id, alt, nb. (Omit id for 'outgroup'. Use 'alt' for less formal labeling in cartoon version; it defaults to 'where'. Use 'nb' to add abode annotations for legend.)
# This hand-made version is now subsumed by make.tree; retained for comparison
tree.by.hand <-
list(
  sub1 = list(
    sub1 = list(
      sub1 = list(id=3367, length=8813, where='Venice, Italy', alt='Venice'),
      sub2 = list(id=1013, length=9652, where='Wales, UK'),
      length=9365),
    sub2 = list(
      sub1 = list(
        sub1 = list(
          sub1 = list(id=1007, length=30, nb='e', where='Virginia, USA'),
          sub2 = list(id=1012, length=61, nb='d', where='Perth, W. Australia', alt='Perth'),
          length=19),
        sub2 = list(
          sub1 = list(id=1015, length=207, nb='c', where='Washington, USA', alt='Puget Sound'),
          sub2 = list(id=1335, length=41, nb='b', where='New York, USA', alt='NY'),
          length=18),
        length=1005),
      sub2 = list(id=1014, length=61, nb='a', where='N. Pacific Gyre'),
      length=3912),
      length=7054),
    sub2 = list(length=0, where='outgroup')
  )
)

# historical, format example, and debug help:
oldwick.medium <- '(((CCMP3367_Italy:8813,CCMP1013_Wales:9652):9365,((e_CCMP1007_Virginia:30,d_CCMP1012_Australia:61):19,(c_CCMP1015_Puget_Sound:207,b_CCMP1335_NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
# with simpler labeling for cartoon
simple.oldwick.medium <- '(((Italy:8813,Wales:9652):9365,((Virginia:30,Australia:61):19,(Puget:207,NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
cat.hardwrap(oldwick.medium)

# (((CCMP3367_Italy:8813,CCMP1013_Wales:9652):9365,((e_CCMP1007_Virginia:30,d_CCMP1012_Australia:61):19,(c_CCMP1015_Puget_Sound:207,b_CCMP1335_NY:41):18):1005,a_CCMP1014_NPG:61):3912):7054,outgroup:0);
cat.hardwrap(simple.oldwick.medium)

# (((Italy:8813,Wales:9652):9365,((Virginia:30,Australia:61):19,(Puget:207,NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
```

Two other versions of the tree, for possible use in figs in the main paper.

Figure 8: [** as of 10/4/2015, this fig and next have stray stars on virginia, wales labels; probably due to hacking with commas in newick; not worth fixing unless we resurrect these trees for some purpose, but if so, see use of newick.name.undo in show.tree as probable fix. **]

```
tree.scale <- ifelse(which.snp.tables(string.val=F)[1]=='Chr1', 1, 10)
tree.x.lim <- 3e4 * tree.scale
the.simple.tree <- read.tree(text=simple.newick.medium)
plot(the.simple.tree, x.lim = tree.x.lim)
axis(1)
```

Figure 9:

```
plot(the.simple.tree, x.lim = tree.x.lim)
axis(1, (0:4)*7000*tree.scale, (0:4)*7000*tree.scale)
```




Figure 8: Tree based on qfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1. (no edge labels, nolegend)




Figure 9: Tree based on qfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs on Chr1. (no edge labels, no legend, short scale bar)

At some much earlier point, Tony ran the whole-genome version of the then-current code above, and manually entered tree branch lengths/legend for the resulting tree, shown in Fig 10. Code above can now automatically generate such a tree, but retain the following for comparison. The basic story seems clear—same topology and branch lengths scaled by about 10x, which is completely reasonable given that Chr1 is about 10% of the genome. Note that this tree is not being recalculated; it is defined by constants in the following code chunk.

```
fullgenome.newick.medium <- '(((3367_Italy:86155,1013_Wales:95697):89598,((e_1007_VA:330,d_1012_Australia:632):1296,(c_1015_WA:2
cat.hardwrap(fullgenome.newick.medium)

# (((3367_Italy:86155,1013_Wales:95697):89598,((e_1007_VA:330,d_1012_Australia:63
# 2):1296,(c_1015_WA:2113,b_1335_NY:658):480):10059,a_1014_NPG:568):39517):69526,o
# utgroup:0);

legend.text <- c('a: only in 1014 ',
                 'b: only in 1335 ',
                 'c: only in 1015 ',
                 'd: only in 1012 ',
                 'e: only in 1007 ',
                 '*: shared by bcde',
                 '  shared by b/c ',
                 '  shared by d/e ')

fullgenome.tree.x.lim <- 300000
fullgenome.counts <- c( 568, 658, 2113, 632, 330, 10059, 480, 1296 )
fullgenome.legend.text <- paste(legend.text, format(fullgenome.counts,width=5), sep=' - ')
fullgenome.tree.labels <- list( ## x,y,text
  41000,3.63,'69526\nshared by 7',
  90000,5.75,'39517\nby 5 (*)',
  115000,1.5, '89598\nshared by 2',
  210000,2.0, '95697 only\nin Wales',
  210000,1.0, '86155 only\nin Italy',
  113500,4.6, '*' )
```

Figure 10:

```
library(ape)
the.fullgenome.tree <- read.tree(text=fullgenome.newick.medium)
plot(the.fullgenome.tree, x.lim = fullgenome.tree.x.lim)
axis(1) # ; axis(2) useful only for placing labels
opar <- par(family='mono',cex=.8)
legend('topright', legend=fullgenome.legend.text)
par(opar)
for(i in seq(1,length(fullgenome.tree.labels)-2,by=3)){
  text(fullgenome.tree.labels[[i]], fullgenome.tree.labels[[i+1]], fullgenome.tree.labels[[i+2]])
}
```




Figure 10: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs genome-wide. (By-hand legacy version)

