

Exploration of Shared SNPs in Thaps

trunc-unfiltered

July 26, 2017

Rambling exploration of SNP positions shared between two or more of the isolates. Code is included to document it thoroughly, (even if largely uninteresting to anyone else), and I will summarize it as I go.

Contents

1 History	2
2 Preliminaries	2
3 Major Analysis/Performance Parameters.	2
4 Refined SNP Calls	6
4.1 Method	6
4.2 Save them	8
4.3 Examples: Consistent	10
4.4 Examples: Inconsistent	11
4.5 Examples: Homozygous nonref	13
5 Table 1 stats	16
5.1 Table 1 Data	20
6 Shared-SNPs P-Value	20
6.1 SNP Concordance	20
6.2 Notes	24
6.3 P-Value: The Bottom Line	24
7 Sharing	24
7.1 Code	25
7.2 Sanity Checks	26
7.3 Main Analysis	33
8 Trees	38
9 Semi-Automated Tree-Building	49
9.1 Bootstrap	56
10 Notes	59
11 Appendix: Old Trees, etc.	60
11.1 HWE Sharing	60
11.2 Old Tree Stuff	61

1 History

This was added to SVN 1/26/2014; not sure when it was started, but earliest related emails I see are from 1/21/14.

```
r413 | ruzzo | 2014-01-26 08:22:37 -0800 (Sun, 26 Jan 2014) | 2 lines
adding shared-snp analysis.
```

2 Preliminaries

NOTE: Some comments in code and some parts of the text, especially specific numbers and general conclusions, are based on Unqfiltered, Chr1, Medium stringency (i.e., “[[2]]” below) analysis. The broad picture does not appear to change with other choices, but details do, and the text is neither fully parameterized nor fully updated, so proceed with caution.

Load utility R code; do setup:

```
source('.././../R/wlr.R') # load util code; path relative this folder or sibling in scripts/larrys

## Running as: ruzzo @ bicycle.cs.washington.edu; SVN Id, I miss you. $Id: wlr.R 2017-07-21 or later $

setup.my.wd('shared-snps') # set working dir; UPDATE if this file moves, or if COPY/PASTE to new file
setup.my.knitr('figs-knitr/')
generic.setup('figs-mine/')
```

3 Major Analysis/Performance Parameters.

Choices here control how this file is processed, what data is analyzed, speed, etc. Set them carefully before running “make.” Major choices are:

1. WHICH SNP TABLES ARE LOADED??? The logical vector `load.tb` selects the desired combination of SNP tables to load, in the order `full.unfiltered`, `chr1.unfiltered`, `full.qfiltered`, `chr1.qfiltered`. E.g., `load.tb=(T, F, T, F)` loads *full* tables for *both* q- and un-qfiltered data. Primary analysis is only performed on one of them, but the others are retained for comparison/debugging.
2. WHICH MAIN ANALYSIS??? If multiple tables are loaded, which is used for the main analysis? Parameter `pri` is a permutation of 1:4, corresponding to `load.tb`; the first loaded table in that order becomes the analysis focus. The default `pri=c(1, 2, 3, 4)` looks at un-q-filtered data in preference to q-filtered, and full tables in preference to Chr1 within each group.
(Choice of data for the “Table 1” coverage summary in section 5 is independent of this; full genome data is preferred over Chr 1 for both q- and unq-filtered reads; change `tset.picker` calls near the end of that section to modify this.)
3. CLEAR CACHE??? `clear.cache=T` forces Knitr cache removal at the start of the run; especially important if the previous parameters have changed since the last run.
4. HOW MANY BOOTSTRAP REPLICATES??? The variable `nboot` is a major performance factor; 1000 reps takes several hours. Set to 5 for debug and quick look; 100 or more for final run.
5. TRUNCATE TABLES TO Chrs ONLY??? I.e., remove mitochondrial-, plastid-, and BD- contigs.

The following code chunk sets the first four parameters based on where it’s run. To prototype/debug on a laptop, faster is better—run on Chr1 with small `nboot`; when run on the linux servers, I typically do full genomes, more replicates. Just override them if these defaults don’t work for you.

```

# for Makefile, params can be command line args, else base on system; see wlr.r for details.
# load.tb order: full.un, chr1.un, full.qfil, chr1.qfil

params <- pick.params(
  mac = list(load.tb=c(F,T,F,F), pri=1:4, clear.cache=F, nboot= 1, trunc.tables=T), # quick on lap
  linux = list(load.tb=c(F,F,F,T), pri=1:4, clear.cache=F, nboot= 5, trunc.tables=T), # quick qfil on server
  linux = list(load.tb=c(T,F,T,F), pri=1:4, clear.cache=T, nboot=101, trunc.tables=T) # full on server
)

# Alternatively, edit/uncomment the following to override the above as needed
#params<-pick.params(default=list(load.tb=c(T,T,T,T),pri=1:4,clear.cache=T,nboot=1000,trunc.tables=T))
print(params)

# $load.tb
# full.unf chr1.unf full.qf chr1.qf
# TRUE FALSE TRUE FALSE
#
# $pri
# [1] 1 2 3 4
#
# $clear.cache
# [1] TRUE
#
# $nboot
# [1] 101
#
# $trunc.tables
# [1] TRUE

```

CLEAR CACHE??!! Some code chunks use the knitr cache, but extent of cache consistency checks unknown. If in doubt, delete “cache/” (knitr’s) directory to force rebuild. T/F set in params above will/won’t force removal (actually, rename):

```

decache(params$clear.cache)

# Rename of 'cache' to 'cache86293' returned TRUE .

```

If still in doubt, also manually remove “00common/mycache/” (mine).

Load the main SNP data file(s) based on the parameters set in section 3.

```

# short names to keep the following chunk compact
tb <- params$load.tb
tset <- list(NULL, NULL, NULL, NULL) # tset = 'table set'

# see wlr.R for load paths
if(tb[1]){tset[[1]] <- load.snp.tables(use.chr1.tables = FALSE, data.name='full.tables.01.26.14')}

# Loading full tables from ../../../../data/ungit-data/full.tables.01.26.14.rda ...Loaded.
# ../00common/mycache/snp.tables.chr1.unqfiltered.rda saved.

if(tb[2]){tset[[2]] <- load.snp.tables(use.chr1.tables = TRUE, data.name='full.tables.01.26.14')}
if(tb[3]){tset[[3]] <- load.snp.tables(use.chr1.tables = FALSE, data.name='full.tables.02.25.15')}

# Loading full tables from ../../../../data/ungit-data/full.tables.02.25.15.rda ...Loaded.
# ../00common/mycache/snp.tables.chr1.qfiltered.rda saved.
# Bandaiding qfiltered tables...

if(tb[4]){tset[[4]] <- load.snp.tables(use.chr1.tables = TRUE, data.name='full.tables.02.25.15')}

```

Grrr! I should have excluded non-Chr contigs from full genome runs. Rather than change tons of code below to add mask params, I’m just going to truncate the tables, as follows. (See notes in wlr.r::make.mask for assumptions.)

```

if(params$trunc.tables){
  for(i in 1:4){
    if(!is.null(tset[[i]])){
      first.mito <- match("mitochondria.fasta", tset[[i]][[7]]$Chr)
      if(!is.na(first.mito)){ # will be NA for Chr1 tables
        for(j in 1:7){
          # hmmm... slow; wonder whether head(tset[[i]][[j]],first.mito-1) is faster;
          # ok, simple tests suggest not: system.time(head(data.frame(1:1e7,1:1e7),5e6))
          tset[[i]][[j]] <- tset[[i]][[j]][1:(first.mito-1),]
        }
      }
    }
  }
} else {
  cat('***\n*** DID YOU *REALLY* WANT UNTRUNCATED TABLES???\n***\n')
}

```

The tersely-named `tset` list is sometimes convenient, but give them more descriptive names, too.

```

snp.tables.full.unfiltered <- tset[[1]]; names(tset)[1] <- 'snp.tables.full.unfiltered'
snp.tables.chr1.unfiltered <- tset[[2]]; names(tset)[2] <- 'snp.tables.chr1.unfiltered'
snp.tables.full.qfiltered <- tset[[3]]; names(tset)[3] <- 'snp.tables.full.qfiltered'
snp.tables.chr1.qfiltered <- tset[[4]]; names(tset)[4] <- 'snp.tables.chr1.qfiltered'

```

Main analysis may just use one of the potentially 4 table sets. Pick it according to the priority specified in section 3, using the shorter name 'snp.tables' for this default choice.

```

snp.tables <- tset.picker(priority=params$pri, table.set=tset)

```

```

# Sanity check: unlike unfiltered tables, bug in early code gave qfiltered ones different numbers
# of rows per strain, which breaks much code. Verify this is no longer happening.
check.eq.nrows <- function(tables){
  if(!is.null(tables)){
    nrow.snp.tables <- unlist(lapply(tables,nrow))
    print(nrow.snp.tables)
    if(all(nrow.snp.tables == nrow.snp.tables[1])){
      cat('OK, all strains have same number of rows.\n')
    } else {
      cat('***\n*** Warning: Different strains have different numbers of rows! ***\n***\n')
    }
  }
}

dummy<-lapply(tset, check.eq.nrows)

#      1007      1012      1013      1014      1015      3367      1335
# 31301782 31301782 31301782 31301782 31301782 31301782 31301782
# OK, all strains have same number of rows.
#      1007      1012      1013      1014      1015      3367      1335
# 31301782 31301782 31301782 31301782 31301782 31301782 31301782
# OK, all strains have same number of rows.

```

Which tables have we got?:

```

# 'which.snp.tables' return summary of which tables, either as a char string (default), e.g.
# "Chr1-qfiltered", or as vector of 2 strings, e.g. c("full","unfiltered").
cat('This analysis uses: (', paste(unlist(lapply(tset,which.snp.tables)),collapse=', '), ') SNP tables.\n')

# This analysis uses: ( trunc-unfiltered, NULL, trunc-qfiltered, NULL ) SNP tables.

cat('Main shared SNP analysis focuses on', which.snp.tables(snp.tables), '\n')

# Main shared SNP analysis focuses on trunc-unfiltered

```

A \LaTeX hack: I want which.snp.tables info in doc title/page headers, but it is unknown until now, so the following writes a command definition `\whichsnp tables` into the .aux file, which is read during the *next* \LaTeX run, when `\begin{document}` is processed:

```
\makeatletter
\immediate\write\@auxout{\noexpand\gdef\noexpand\whichsnp tables{trunc-unfiltered}}
\makeatother
```

Subsequent analysis was initially all directed at Chr1. In general, I have *not* updated the discussion to reflect genome-wide analysis.

```
if(exists('snp.tables.chr1.qfiltered') && exists('snp.tables.chr1.unqfiltered')){
  # If have both, where is new unequal to old?
  uneq <- snp.tables.chr1.qfiltered[[1]]$Ref[1:chr1.len] != snp.tables.chr1.unqfiltered[[1]]$Ref[1:chr1.len]
  cat('Sum uneq:', sum(uneq, na.rm=T), '\n')
  cat('Sum NA: ', sum(is.na(uneq)), '\n')
  print(which(is.na(uneq)) [1:10])
  seecounts(which(is.na(uneq)) [1:4], snp.tables=snp.tables.qfiltered, debug=F)
}
```

In brief, “snp.tables” will be a list of 7 data frames, one per strain, giving read counts for each nucleotide at each position, SNP calls, etc.:

```
names(snp.tables)

# [1] "1007" "1012" "1013" "1014" "1015" "3367" "1335"

str(snp.tables[[1]])

# 'data.frame': 31301782 obs. of 15 variables:
# $ chr : Factor w/ 66 levels "BD10_65","BD11_74",...: 39 39 39 39 39 39 39 39 39 39 ...
# $ pos : int 1 2 3 4 5 6 7 8 9 10 ...
# $ snp : int 0 0 0 0 0 0 0 0 0 0 ...
# $ Chr : chr "Chr1" "Chr1" "Chr1" "Chr1" ...
# $ Pos : int 1 2 3 4 5 6 7 8 9 10 ...
# $ Ref : chr "T" "C" "C" "A" ...
# $ Cov : num 1 3 4 5 7 7 10 12 13 15 ...
# $ a : num 0 0 1 0 0 0 0 0 1 0 ...
# $ g : num 0 0 0 0 0 0 0 0 0 0 ...
# $ c : num 0 0 0 0 0 0 0 0 0 0 ...
# $ t : num 0 0 0 0 0 0 0 0 0 0 ...
# $ n : num 0 0 0 0 0 0 0 0 0 0 ...
# $ .match: num 1 3 3 5 7 7 10 12 12 15 ...
# $ exon : logi FALSE FALSE FALSE FALSE FALSE FALSE ...
# $ indel : logi FALSE FALSE FALSE FALSE FALSE FALSE ...
```

Just for background, also load the desert tables:

```
# from svn+ssh://cegl.ocean.washington.edu/var/svn/7_strains/trunk/code/snpNB/data
#load('../.../data/ungit-data/des.rda')
load('../.../data/des.rda')
```

What’s the total length of all deserts in each strain? Big deserts (defined as “big.threshold” or longer)?

```
some.desert.stats <- function(big.threshold=0){
  desert.len <- unlist(lapply(des, function(x){sum(unlist(lapply(x, function(y){sum(y[, 'Length'])}))})))
  bigdes.len <- unlist(lapply(des, function(x){sum(unlist(lapply(x, function(y){
    sum(y[, 'Length']>=big.threshold, 'Length'])}))})))
  rbind(desert.len, desert.pct=round( desert.len / genome.length.constants()$genome.length.trunc * 100),
        bigdes.len, bigdes.pct=round( bigdes.len / genome.length.constants()$genome.length.trunc * 100))
}
some.desert.stats(big.threshold=50000)

#           tp1007  tp1012  tp1013  tp1014  tp1015 thapsIT  tp1335
# desert.len 11146526 11332566 5801763 9464213 11251426 6780300 10883723
```

# desert.pct	36	36	19	30	36	22	35
# bigdes.len	3495805	3936973	55365	3627235	3727061	57119	4046934
# bigdes.pct	11	13	0	12	12	0	13

I.e., looking at all deserts, about 1/3 of L-clade, 1/5 of H-clade are in deserts, whereas, looking at the largest deserts ($> 50k$), only about 12% in L-clade (and none in H-clade). Note that the rough stats above include artifactual “deserts” created by gaps in the reference sequence, large genomic deletions, etc. A more careful analysis of this is found in nc-snps.rnw.

4 Refined SNP Calls

4.1 Method

It is appropriate that SNP calls should be conservative, to avoid many false positives, but, when a position is called a SNP in one isolate, we often see a significant number of reads for the same non-reference nucleotide at that position in other isolates, even if they are not called as SNPs. On the other hand, we sometimes see a position called a SNP in two or more isolates, but with *different* pairs of nucleotides, potentially suggesting technical errors. Analysis in this section attempts to refine the SNP calls by looking for issues such as these by looking at all 7 isolates jointly, at each position called a SNP in any of them.

For a given strain, the following function returns a vector of 0:4 to indicate which nonreference nucleotide has the maximum read count at the corresponding position. The values 1:4 indicate that the max count occurred at A, G, C, T, resp. (Ties are resolved arbitrarily ($a < g < c < t$), which possibly deserves further attention.) The value 0 means all nonreference counts are below threshold, based *either* on absolute count *or* as a fraction of coverage. Default only excludes 0 counts.

```
nref.nuc.new <- function(strain=1, mask=T, thresh.count=0, thresh.rate=0.0){
  # get read count for max nonref nuc
  nref <- apply(snp.tables[[strain]][mask, c('a', 'g', 'c', 't')], 1, max)
  # where does nref count match a (g,c,t, resp) count
  as <- ifelse(nref == snp.tables[[strain]][mask, 'a'], 1, 0)
  gs <- ifelse(nref == snp.tables[[strain]][mask, 'g'], 2, 0)
  cs <- ifelse(nref == snp.tables[[strain]][mask, 'c'], 3, 0)
  ts <- ifelse(nref == snp.tables[[strain]][mask, 't'], 4, 0)
  # most positions will show 3 zeros and one of 1:4, so max identifies max nonref count;
  # ties broken arbitrarily (a<g<c<t)
  merge <- pmax(as, gs, cs, ts)
  # but if max nonref count is zero or below threshold, return 0
  merge[nref == 0 | nref < thresh.count] <- 0
  merge[nref/snp.tables[[strain]][mask, 'Cov'] < thresh.rate] <- 0
  return(merge)
}
```

Get union and intersection of the sets of called SNPs. (“\$snp” is 0/1.) Also, 5-way (L-clade) and 4-way (L-excluding Gyre).

```
# 4-way union/intersection
u4.snp <- snp.tables[[1]]$snp
i4.snp <- snp.tables[[1]]$snp
for(i in c(2,5,7)) {
  u4.snp <- pmax(u4.snp, snp.tables[[i]]$snp)
  i4.snp <- pmin(i4.snp, snp.tables[[i]]$snp)
}
# 5-way: add gyre
u5.snp <- pmax(u4.snp, snp.tables[[4]]$snp)
i5.snp <- pmin(i4.snp, snp.tables[[4]]$snp)
# 7-way
union.snp <- pmax(u5.snp, snp.tables[[3]]$snp, snp.tables[[6]]$snp)
intersect.snp <- pmin(i5.snp, snp.tables[[3]]$snp, snp.tables[[6]]$snp)
nu4.snp <- sum(u4.snp)
nu5.snp <- sum(u5.snp)
ni4.snp <- sum(i4.snp)
```

```

ni5snps <- sum(i5.snps)
nusnps <- sum(union.snps)
nisnps <- sum(intersect.snps)
c(n4u=nusnps, n5u=nusnps, n7u=nusnps, n4i=ni4snps, n5i=ni5snps, n7i=nisnps)

#      n4u      n5u      n7u      n4i      n5i      n7i
# 196296 197799 474613 128683  70687  15186

```

There are `nusnps=474613` positions called as SNPs in one or more strains (but only `nisnps=15186` that are shared among all 7). Note that the 4-way union is only modestly larger (1.5254229 times larger) than the 4-way intersection, emphasizing the inherent similarities among these SNP sets. The corresponding 5-way numbers show that Gyre adds relatively little to the 5-way union vs the 4-way union, whereas it removes a fair bit from the 5-way intersection. However, much of that loss is simply because Gyre has fewer called SNPs: only 89184 vs 128683 in the 4-way intersection, and they are highly concordant:

```

sum(snp.tables[[4]]$snp*i4.snps)/sum(snp.tables[[4]]$snp)

# [1] 0.7925973

```

So, a likely source of the Gyre's difference in called SNPs is technical (lower read coverage, higher read error rate) rather than biological.

Inclusion of the 2 H-clade members, however, causes more dramatic changes in both union and intersection numbers. I examine all these relationships in more detail below, but first I examine what I believe to be a significant source of technical error in these comparisons—erroneous SNP calls, especially false negative calls.

It is appropriate that SNP calls should be conservative, to avoid many false positives, but, if a position is called a SNP in one strain, we often see a significant number of reads for the same non-reference nucleotide at that position in other strains, even if they are not called as SNPs. For my purposes below, these will be considered “shared SNPs,” based on three different levels of permissiveness. Note that, e.g., $\geq 84\%$ of all positions have zero reads for any non-reference nucleotide, and only a small fraction have 2 or more non-reference reads:

```

nonmatch <- rbind(
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 0)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 1)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 2)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match == 3)})),
  unlist(lapply(snp.tables,function(x){sum(x$Cov-x$.match >= 4)})),
  unlist(lapply(snp.tables,function(x){sum((x$Cov-x$.match)[union.snps==0] >= 4)}))
)/nrow(snp.tables[[1]])*100
rownames(nonmatch) <- c('% ==0', '% ==1', '% ==2', '% ==3', '% >=4', '% >=4, nonSNP')
nonmatch

```

#	1007	1012	1013	1014	1015	3367	1335
# % ==0	92.4326481	88.6711338	84.5965383	86.25661312	90.4442629	86.3165937	84.9608722
# % ==1	6.3661967	9.4731028	12.3462396	11.94901300	7.9893279	10.9701135	11.8031619
# % ==2	0.4436073	0.9107788	1.4480773	1.14379111	0.6581766	1.1767158	1.8092900
# % ==3	0.0908830	0.1568633	0.2616369	0.19303054	0.1231208	0.2238563	0.4539230
# % >=4	0.6666649	0.7881213	1.3475079	0.45755222	0.7851119	1.3127208	0.9727529
# % >=4, nonSNP	0.1041826	0.1892065	0.3693783	0.09624053	0.1760091	0.3790263	0.3876297

Build a table of max non-reference nucleotides at each position in the `union.snps` set. The three criteria are

- `[[1]]`: any non-zero count at any coverage is considered significant
- `[[2]]`: (count ≥ 2 and count/coverage ≥ 0.05) is considered significant
- `[[3]]`: (count ≥ 4 and count/coverage ≥ 0.10) is considered significant

In all three cases, the nonref nucleotide must also be consistent across all strains passing that threshold; see below.

```

non.refs <- vector('list',4)
for(i in 1:4){
  non.refs[[i]] <- matrix(0, nrow=nusnps, ncol=7)
}

```

```

colnames(non.refs[[i]]) <- names(snp.tables)
rownames(non.refs[[i]]) <-
  paste(snp.tables[[1]]$chr[union.snps==1], ':', snp.tables[[1]]$pos[union.snps==1], sep='')
}
for(j in 1:7){
  non.refs[[1]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=0, thresh.rate=0.00)
  non.refs[[2]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=2, thresh.rate=0.05)
  non.refs[[3]][,j] <- nref.nuc.new(j, mask=union.snps==1, thresh.count=4, thresh.rate=0.10)
}

```

For comparison, I want to look at unfiltered SAMTools SNP calls. In complete opposition to the measures of consistency imposed above, I'm going to simply force this into the “non.refs” structure constructed above by imagining that any position called a SNP in any strain has its max nonref count on “A”, so any given position called a SNP in any strain will automatically be declared “consistent.” This will allow the tree-code, etc. given below to work in a uniform way (even though interpretation of the results is different.) Results will be jammed into a 4th component of the “non.refs” list; i.e., we have a 4th criterion:

- [[4]]: all called SNPs at a given position are considered “consistent.”

As this case was a late addition to the analysis, the commentary throughout this document has not necessarily been updated to reflect that this case is distinct from the first three.

```

for(j in 1:7){
  non.refs[[4]][,j] <- snp.tables[[j]]$snp[union.snps==1]
}

```

```

str(non.refs[[4]])

# num [1:474613, 1:7] 0 0 0 0 0 0 0 0 1 0 ...
# - attr(*, "dimnames")=List of 2
# ..$ : chr [1:474613] "Chr1:333" "Chr1:417" "Chr1:435" "Chr1:438" ...
# ..$ : chr [1:7] "1007" "1012" "1013" "1014" ...

```

“non.refs” indicates, among those positions in the union of all called SNPS having any non-reference read count above the thresholds listed above, the non-ref nucleotide having the highest read count in each strain. If, for a given position, the max of this code is the same as the min (among non-zero values), then every strain having over-threshold nonref reads in that position, in fact has most non-reference reads on the *same* nucleotide. These are defined as the “consistent” SNPs.

```

find.consistent <- function(nr){
  nr.max <- apply(nr, 1, max)
  nr.min <- apply(nr, 1, function(x){ifelse(max(x)==0, 0, min(x[x>0]))})
  return(nr.min == nr.max)
}
consistent <- lapply(non.refs, find.consistent)

```

4.2 Save them

```

# wrap this in a data structure to be cached:
Description <- [2753 chars quoted with '']

refined.snps <-
  list(Description=Description,

       Data=list(
         based.on.which.snp.tables=which.snp.tables(),
         number.union.snps=nusnps,
         number.intersection.snps=nisnps,
         non.ref.nucleotide=non.refs,
         consistent.snps=consistent),

```



```

Code=list(
  get.snps = function(strain, stringency=2){
    # return nusnps x 1 Bool vector of consistent SNPs @ specified stringency & strain
    return(refined.snps$Data$consistent.snps[[stringency]] &
      refined.snps$Data$non.ref.nucleotide[[stringency]][,strain] > 0)
  },
  get.snp.locs.char = function(strain, stringency=2){
    # return char vector of locations of consistent SNPs @ specified stringency & strain
    snps <- refined.snps$Code$get.snps(strain, stringency)
    return(names(snps)[snps])
  },
  get.snp.locs.df = function(strain, stringency=2){
    # return data frame (Chr/Pos) of locations of consistent SNPs @ specified stringency & strain
    snplist <- strsplit(refined.snps$Code$get.snp.locs.char(strain, stringency), ':', fixed=TRUE)
    # strsplit returns long list of 2-vectors, 1st=chr, 2nd=char position
    df <- data.frame(Chr=
      unlist(lapply(snplist,function(x){return(x[1])})),
      Pos=as.integer(unlist(lapply(snplist,function(x){return(x[2])}))),
      stringsAsFactors = FALSE)

    return(df)
  }
)

# dont't clobber existing .rda, but save if absent. (delete to re-save)
# result for trunc, unfiltered tables saved to "data" else "mycache"
if(which.snp.tables() == 'trunc-unfiltered'){
  rda.refined <- '../.../data/refined.snps-trunc-unfiltered.rda'
} else {
  rda.refined <- paste('../00common/mycache/refined.snps', which.snp.tables(), 'rda', sep='.')
}
if(file.exists(rda.refined)){
  cat('Pre-existing file', rda.refined, 'unchanged.\n')
} else {
  cat('Saving', rda.refined, '...\n')
  save(refined.snps, file=rda.refined, compress=TRUE)
  cat('Saved.\n')
}

# Saving ../.../data/refined.snps-trunc-unfiltered.rda ...Saved.

```

Knitr seems to be failing to format the long char string above, which says:

```

cat(refined.snps$Description)

# Contents of this .rda file:
#
# * Description: this text
#
# * Data -- 5 items defining refined SNPs, at 4 different stringency levels, as defined
#   in shared-snps.rnw:
#
# * based.on.which.snp.tables: {"Chr1","full","trunc"}-{"unfiltered","qfiltered"},
#   depending on which snp tables were used to build this data. ("trunc" = all Chrs.)
#
# * number.union.snps: the total number of SNPs (SAMtools calls) in the union of SNPs
#   across all 7 strains.
#
# * number.intersection.snps: similar, for the 7-way intersection.
#
#   nusnps/nisnps are easily recalculated from the data below, but their inclusion
#   may be convenient, e.g., to quickly see if the .rda represents the full genome
#   (nusnps=488848), or the chr 1 subset (nusnps=47499); (redundant with "based.on...";
#   numbers above are for unfiltered, perhaps slightly different if qfiltered)
#
# * non.ref.nucleotide: 4 arrays, each nusnps x 7, of values 0..4 (0..1 in the 4th

```

```
# array). In the 1st 3 arrays, 0 means the given position in the given strain did
# not have nonreference read counts above the corresponding filtering threshold,
# i.e., is NOT a refined SNP in that strain, whereas 1..4 mean that it did pass
# threshold, for A,C,G,T resp. In the 4th array, this value is just 1/0,
# indicating is/is not a called SNP in that strain.
#
# * consistent.snps: 4 Bool vectors of length nusnps flagging positions whose nonref
# nucs (wrt to the 4 filtering criteria) are deemed *consistent* across
# all 7 strains. For the 1st 3, this means all nonzero entries of non.ref.nuc
# are equal, i.e., nonref read counts passing threshold are on the SAME nonref
# nucleotide in all strains having over-threshold counts. Just for comparison
# and uniformity of data structures, the 4th is all TRUE, i.e., union of SNPs
# across all strains, without any regard for thresholds or consistency.
#
# In short, the refined SNPs according to our medium filtering criteria are
# strains/positions where consistent.snps[[2]]==TRUE and non.ref.nucleotide[[2]]>0.
#
# Rownames in both non.ref.nucs and consistent define location, e.g. "Chr1:333".
#
# * Code -- simple routines to extract refined SNPs in (potentially) convenient formats:
#
# * get.snps(strain, stringency=2)
# returns nusnps x 1 Bool vector of consistent SNPs @ specified stringency in
# given strain
#
# * get.snp.locs.char(strain, stringency=2)
# returns n x 1 char vector of locations of consistent SNPs @ specified stringency
# in given strain, e.g. "Chr1:1234", where n == sum(get.snps(...))
#
# * get.snp.locs.df(strain, stringency=2){
# As above, but returns data frame (char vector Chr, int vector Pos) with the same info.
```

```
str(consistent[[1]])

# Named logi [1:474613] TRUE FALSE TRUE FALSE TRUE TRUE ...
# - attr(*, "names")= chr [1:474613] "Chr1:333" "Chr1:417" "Chr1:435" "Chr1:438" ...
```

```
consistent.count <- unlist(lapply(consistent, sum)) ; consistent.count

# [1] 358872 468117 470970 474613

inconsistent.count <- consistent.count[4] - consistent.count; inconsistent.count

# [1] 115741 6496 3643 0

inconsistent.percent <- inconsistent.count/consistent.count[4]*100; inconsistent.percent

# [1] 24.3863948 1.3686941 0.7675727 0.0000000
```

I.e., of the 474613 positions in which a SNP is called, 358872 are consistent by my loose definition, and 470970 are consistent by my tightest definition. The increase in concordance supports the view that the loose definition is too loose. Perhaps misleadingly, these counts include positions that are “consistent SNPs” in only one strain; more below. (*TODO* I suspect, but have not yet systematically checked, that most of the rest are positions with low coverage and/or very low read counts on the mixture of non-reference nucleotides.)

4.3 Examples: Consistent

Here are a few (nonrandomly selected) prototypical consistent SNPs:

```
esnps <- names(consistent[[2]][consistent[[2]]])
esnps2 <- as.integer(unlist(lapply(strsplit(esnps[c(7,11:13,92)], ':', fixed=TRUE), function(x){x[2]})))
seecounts(esnps2, snp.tables=snp.tables)
```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	567	T										
# 2				1007	0	0	2	29	0	TRUE	FALSE		
# 3				1012	0	0	15	44	1	TRUE	FALSE		
# 4				1013	0	0	15	97	0	TRUE	FALSE		
# 5				1014	0	1	0	33	0	TRUE	FALSE		
# 6				1015	0	0	9	43	1	TRUE	FALSE		
# 7				3367	0	0	16	46	1	TRUE	FALSE		
# 8				1335	0	0	2	116	0	TRUE	FALSE		
# 9	Chr1	1053	A										
# 10				1007	39	0	0	5	0	TRUE	FALSE		
# 11				1012	55	0	0	12	0	TRUE	FALSE		
# 12				1013	17	1	0	40	0	TRUE	FALSE		
# 13				1014	25	0	0	5	0	TRUE	FALSE		
# 14				1015	38	0	1	20	1	TRUE	FALSE		
# 15				3367	13	0	0	9	0	TRUE	FALSE		
# 16				1335	71	1	0	46	1	TRUE	FALSE		
# 17	Chr1	1055	G										
# 18				1007	0	41	0	2	0	TRUE	FALSE		
# 19				1012	1	63	0	8	0	TRUE	FALSE		
# 20				1013	1	62	0	8	0	TRUE	FALSE		
# 21				1014	1	26	0	8	1	TRUE	FALSE		
# 22				1015	0	44	0	14	0	TRUE	FALSE		
# 23				3367	0	27	0	0	0	TRUE	FALSE		
# 24				1335	0	78	0	40	1	TRUE	FALSE		
# 25	Chr1	1176	G										
# 26				1007	2	67	0	0	0	FALSE	FALSE		
# 27				1012	1	68	0	0	0	FALSE	FALSE		
# 28				1013	29	73	0	0	0	FALSE	FALSE		
# 29				1014	1	52	0	0	0	FALSE	FALSE		
# 30				1015	4	103	0	0	0	FALSE	FALSE		
# 31				3367	11	8	0	0	1	FALSE	FALSE		
# 32				1335	1	206	0	0	0	FALSE	FALSE		
# 33	Chr1	8670	A										
# 34				1007	19	0	0	7	0	TRUE	FALSE		
# 35				1012	36	0	0	12	0	TRUE	FALSE		
# 36				1013	44	0	0	12	0	TRUE	FALSE		
# 37				1014	10	0	0	7	0	TRUE	FALSE		
# 38				1015	24	0	0	11	1	TRUE	FALSE		
# 39				3367	18	0	0	0	0	TRUE	FALSE		
# 40				1335	27	0	0	6	0	TRUE	FALSE		

4.4 Examples: Inconsistent

Here is a brief look at some *in*-consistent positions. E.g., Chr1:2013 shows nontrivial counts on 3 alleles in Wales, as do 2319, 3286, 5002, 5433, whereas 7878 shows a different alternate allele in Italy than in Wales.

```
unc <- names(consistent[[2]][!consistent[[2]]])
unc2 <- as.integer(unlist(lapply(strsplit(unc[1:10], ':', fixed=TRUE), function(x){x[2]})))
seecounts(unc2, snp.tables=snp.tables)
```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	2013	T										
# 2				1007	4	0	0	20	0	TRUE	FALSE		
# 3				1012	8	0	0	34	0	TRUE	FALSE		
# 4				1013	9	12	0	16	1	TRUE	FALSE		
# 5				1014	1	0	0	19	0	TRUE	FALSE		
# 6				1015	13	0	0	24	1	TRUE	FALSE		
# 7				3367	10	0	0	36	0	TRUE	FALSE		
# 8				1335	20	0	0	68	1	TRUE	FALSE		
# 9	Chr1	2319	C										
# 10				1007	0	29	22	0	1	TRUE	FALSE		
# 11				1012	0	54	26	0	1	TRUE	FALSE		
# 12				1013	19	19	18	0	1	TRUE	FALSE		
# 13				1014	0	25	19	0	1	TRUE	FALSE		

# 14				1015	0	54	29	0	1	TRUE	FALSE
# 15				3367	5	0	43	0	0	TRUE	FALSE
# 16				1335	0	132	48	0	1	TRUE	FALSE
# 17	Chr1	3286	T								
# 18				1007	4	0	1	17	0	TRUE	FALSE
# 19				1012	9	0	3	45	0	TRUE	FALSE
# 20				1013	39	1	38	12	1	TRUE	FALSE
# 21				1014	4	0	6	27	0	TRUE	FALSE
# 22				1015	11	0	7	37	0	TRUE	FALSE
# 23				3367	8	0	39	10	0	TRUE	FALSE
# 24				1335	15	0	4	75	0	TRUE	FALSE
# 25	Chr1	5002	T								
# 26				1007	0	15	0	12	0	TRUE	FALSE
# 27				1012	1	23	0	26	1	TRUE	FALSE
# 28				1013	21	11	0	39	0	TRUE	FALSE
# 29				1014	0	8	0	12	0	TRUE	FALSE
# 30				1015	0	19	0	16	1	TRUE	FALSE
# 31				3367	0	0	0	35	0	TRUE	FALSE
# 32				1335	0	57	0	60	0	TRUE	FALSE
# 33	Chr1	5433	G								
# 34				1007	0	50	0	3	0	TRUE	FALSE
# 35				1012	0	78	0	5	0	TRUE	FALSE
# 36				1013	18	47	0	14	1	TRUE	FALSE
# 37				1014	9	19	0	0	1	TRUE	FALSE
# 38				1015	7	63	0	2	0	TRUE	FALSE
# 39				3367	8	54	0	0	0	TRUE	FALSE
# 40				1335	6	109	0	4	0	TRUE	FALSE
# 41	Chr1	7858	C								
# 42				1007	0	0	48	0	0	TRUE	FALSE
# 43				1012	0	1	61	0	0	TRUE	FALSE
# 44				1013	0	0	131	10	0	TRUE	FALSE
# 45				1014	0	0	34	0	0	TRUE	FALSE
# 46				1015	0	0	74	0	0	TRUE	FALSE
# 47				3367	20	0	8	0	1	TRUE	FALSE
# 48				1335	0	0	120	0	0	TRUE	FALSE
# 49	Chr1	8914	A								
# 50				1007	23	0	0	2	0	TRUE	FALSE
# 51				1012	29	0	15	0	1	TRUE	FALSE
# 52				1013	25	0	6	0	0	TRUE	FALSE
# 53				1014	22	0	0	0	0	TRUE	FALSE
# 54				1015	31	0	5	2	0	TRUE	FALSE
# 55				3367	8	0	0	1	0	TRUE	FALSE
# 56				1335	68	0	7	0	0	TRUE	FALSE
# 57	Chr1	8974	C								
# 58				1007	0	2	6	0	0	TRUE	FALSE
# 59				1012	0	2	17	0	0	TRUE	FALSE
# 60				1013	10	22	4	0	1	TRUE	FALSE
# 61				1014	0	1	10	0	0	TRUE	FALSE
# 62				1015	0	2	15	0	0	TRUE	FALSE
# 63				3367	2	0	3	0	0	TRUE	FALSE
# 64				1335	0	11	49	0	0	TRUE	FALSE
# 65	Chr1	10099	T								
# 66				1007	17	0	0	29	0	TRUE	FALSE
# 67				1012	48	0	0	36	0	TRUE	FALSE
# 68				1013	0	2	6	68	0	TRUE	FALSE
# 69				1014	34	0	0	26	0	TRUE	FALSE
# 70				1015	41	0	0	38	0	TRUE	FALSE
# 71				3367	0	1	0	14	0	TRUE	FALSE
# 72				1335	55	0	0	68	1	TRUE	FALSE
# 73	Chr1	15154	A								
# 74				1007	25	0	0	0	0	FALSE	FALSE
# 75				1012	56	0	0	1	0	FALSE	FALSE
# 76				1013	10	0	38	10	1	FALSE	FALSE
# 77				1014	26	0	0	0	0	FALSE	FALSE
# 78				1015	37	0	0	0	0	FALSE	FALSE
# 79				3367	19	0	0	13	1	FALSE	FALSE
# 80				1335	70	0	0	3	0	FALSE	FALSE

4.5 Examples: Homozygous nonref

And at some *homozygous nonreference* positions (defined to be those with nonref fraction > 0.75):

```
hnr <- lapply(snp.tables, function(x){x$.match/x$.Cov < 0.25}) # find them
hnr <- lapply(hnr, function(x){ifelse(is.na(x), FALSE, x)}) # remove NA
unlist(lapply(hnr, sum)) # count per strain

# 1007 1012 1013 1014 1015 3367 1335
# 6619 7645 62072 440 3593 72356 558
```

Hmm, in L-clade, excluding the ref isolate (1335) this tracks time-in culture to some degree; Maybe many of these are in hemizygous regions. Next two chunks lifted from nc-snp to get tables for hemi-deletion.

```
cnv.chrononly <- load.cnv.tables('.../data/cnv.txt', chrs.only=TRUE)

str(cnv.chrononly)

# 'data.frame': 1956 obs. of 11 variables:
# $ strain : Factor w/ 7 levels "IT","tp1007",...: 3 3 3 3 3 3 3 3 3 ...
# $ chr : Factor w/ 65 levels "BD1_7","BD10_65",...: 38 38 38 38 38 38 38 38 38 ...
# $ start : int 10601 112001 215001 358901 536501 554801 673401 781801 806901 853201 ...
# $ end : int 13500 116500 221100 370300 538600 559300 685000 787400 811100 855600 ...
# $ length : int 2900 4500 6100 11400 2100 4500 11600 5600 4200 2400 ...
# $ filtered : logi FALSE FALSE FALSE TRUE FALSE FALSE ...
# $ type : Factor w/ 1 level "CNVnator": 1 1 1 1 1 1 1 1 1 1 ...
# $ cov_ratio: num 0.63738 1.54893 1.65381 0.00204 0.68486 ...
# $ dup_frac : num 0.41188 0.00908 0.01178 0.97997 0.0211 ...
# $ iStart : num 10601 112001 215001 358901 536501 ...
# $ iEnd : num 13500 116500 221100 370300 538600 ...

cnv.chrononly[c(1:4,nrow(cnv.chrononly)+c(-1,0)),] ## first/last few rows

# strain chr start end length filtered type cov_ratio dup_frac iStart iEnd
# 1 tp1012 Chr1 10601 13500 2900 FALSE CNVnator 0.63738000 0.41187900 10601 13500
# 2 tp1012 Chr1 112001 116500 4500 FALSE CNVnator 1.54893000 0.00907677 112001 116500
# 3 tp1012 Chr1 215001 221100 6100 FALSE CNVnator 1.65381000 0.01178470 215001 221100
# 4 tp1012 Chr1 358901 370300 11400 TRUE CNVnator 0.00204431 0.97997300 358901 370300
# 1955 tp1335 Chr24 259901 278000 18100 FALSE CNVnator 1.41458000 0.38091100 31264334 31282433
# 1956 tp1335 Chr24 286901 289800 2900 FALSE CNVnator 1.74941000 0.74228100 31291334 31294233
```

```
get.cnv.dels <- function(cov.thresh.lo = 0.0,
                        cov.thresh.hi = 0.8,
                        cnv,
                        snp.tables = NULL,
                        DEBUG = FALSE)
){
  # build list of 7 Bool vectors of genome length, with i-th == T iff
  # * i-th pos is 'NA' in genome seq (if snp.tables are provided), or
  # * in CNVnator call for coverage in half-open [cov.thresh.lo, hi), and
  # * not marked 'filtered' by CNVnator
  cnv.deletions <- vector(mode='list', 7) # make list of bool vectors
  if(is.null(snp.tables)){
    # if no tables, assume full
    t.len <- genome.length.constants()$genome.length.trunc
  } else {
    t.len <- nrow(snp.tables[[1]])
  }
  for(st in 1:7){
    if(is.null(snp.tables)){
      cnv.deletions[[st]] <- logical(t.len) # all F
    } else {
      cnv.deletions[[st]] <- is.na(snp.tables[[st]]$Pos[1:t.len]) # NA positions in genome
    }
  }
  strain.names <- c(paste('tp10', c('07', 12:15), sep=''), 'IT', 'tp1335')
```

```

names(cnv.deletions) <- strain.names
for(i in 1:nrow(cnv)){
  if(!cnv$filtered[i] &&
     cnv$cov_ratio[i] >= cov.thresh.lo &&
     cnv$cov_ratio[i] < cov.thresh.hi)
  {
    if(DEBUG){
      print(cnv[i,])
      print(as.character(cnv$strain[i]))
    }
    # following ASSUMES no CNVnator call crosses a chromosome bdry, & that
    # t.len ends at chr end (typically chr1 or chr24)
    if(cnv$iEnd[i] <= t.len){
      cnv.deletions[[as.character(cnv$strain[i])][cnv$iStart[i]:cnv$iEnd[i]]] <- TRUE
    }
  }
}
return(cnv.deletions)
}

# sanity check:
cnv.dels.38 <- get.cnv.dels(0.3, 0.8, cnv.chrononly, snp.tables = NULL)
unlist(lapply(cnv.dels.38, sum)) # does it match low.length.38 in tic ?

# tp1007 tp1012 tp1013 tp1014 tp1015 IT tp1335
# 1672500 1781500 1383600 1313700 988400 320900 1453000

# 1672500 1781500 1399400 1313700 988400 336500 1453000 <== low.length.38 from tic (circa page 8)
# 1672500 1781500 1399400 1313700 988400 336500 1453000 <== low.length.38 from tic (pg9, 6/28/17)
rm(cnv.dels.38)

```

Slight discrepancy in H-clade that I should hunt down, but basically OK. (hmm; maybe untrunc tbls.)

```

# the ones we want for the current analysis:
hemi.masks <- get.cnv.dels(0.3, 0.8, cnv.chrononly, snp.tables=snp.tables)

rbind(
  homnr      = unlist(lapply(hnr, sum)),
  hemi       = unlist(lapply(hemi.masks, sum)),
  homnr.unhemi = unlist(lapply(list(1,2,3,4,5,6,7), function(i){sum(hnr[[i]] & !hemi.masks[[i]]))}))
)

#           1007      1012      1013      1014      1015      3367      1335
# homnr      6619      7645      62072      440       3593      72356      558
# hemi      1834990 1940024 1527725 1472095 1134652 480817 1596965
# homnr.unhemi 4441      4220      59390      434       2683      71732      537

# based on the thought that hnr in 1335 may reflect errors in the ref seq,
# are they shared with others?
unlist(lapply(hnr, function(x){sum(x & hnr[[7]])})) # hnr shared with 1335

# 1007 1012 1013 1014 1015 3367 1335
# 271 306 311 243 332 330 558

# answer: around 300 in each strain, of 558 in NY, genomewide,
# so that seems like a plausibly important factor.

hnr.lclade <- hnr[[1]] | hnr[[2]] | hnr[[4]] | hnr[[5]] | hnr[[7]] # union over L-clade
sum(hnr.lclade) # count all in L-clade

# [1] 11010

sum(hnr[[3]] | hnr[[6]]) # present in H-clade

# [1] 104450

```

```

sum(hnr[[3]] & hnr[[6]]) # shared in H-clade

# [1] 29978

# look at a few in L-clade
w.hnr.l <- which(hnr.lclade)
seecounts(w.hnr.l[1:10], snp.tables=snp.tables)

#      chr    pos Ref Strain    A    G    C    T SNP  exon indel nrf rat
# 1  Chr1    5397    C
# 2      1007    0    0   24   27    1 TRUE FALSE
# 3      1012    0    0   34   40    1 TRUE FALSE
# 4      1013    0    0   12   42    0 TRUE FALSE
# 5      1014    1    0   30   28    1 TRUE FALSE
# 6      1015    0    0   33   35    1 TRUE FALSE
# 7      3367    0    0   20   38    1 TRUE FALSE
# 8      1335    0    0   29   98    1 TRUE FALSE
# 9  Chr1   20071    T
# 10     1007   22    0    0   15    1 FALSE FALSE
# 11     1012  109    0    0   41    1 FALSE FALSE
# 12     1013   28    0    0   33    1 FALSE FALSE
# 13     1014   76    0    0   29    1 FALSE FALSE
# 14     1015  130    0    0   28    1 FALSE FALSE
# 15     3367   27    0    0   28    0 FALSE FALSE
# 16     1335   95    0    0   57    0 FALSE FALSE
# 17 Chr1   25350    G
# 18     1007  104   31    0    0    1 FALSE FALSE
# 19     1012  171   53    0    0    1 FALSE FALSE
# 20     1013  209   87    1    0    0 FALSE FALSE
# 21     1014   19   32    0    0    0 FALSE FALSE
# 22     1015   91   44    0    0    1 FALSE FALSE
# 23     3367  397   94    0    0    0 FALSE FALSE
# 24     1335   80   64    0    0    0 FALSE FALSE
# 25 Chr1   26205    T
# 26     1007   50    0    0   20    1 FALSE FALSE
# 27     1012  104    0    0   33    1 FALSE FALSE
# 28     1013  224    0    0   69    1 FALSE FALSE
# 29     1014   23    0    1   16    1 FALSE FALSE
# 30     1015   88    0    0   33    1 FALSE FALSE
# 31     3367  143    0    0   41    1 FALSE FALSE
# 32     1335  196    0    0   67    1 FALSE FALSE
# 33 Chr1   90942    C
# 34     1007    0    0   15    0    0 FALSE FALSE
# 35     1012    0    0   33    0    0 FALSE FALSE
# 36     1013    0    0   46    0    0 FALSE FALSE
# 37     1014    0    0   16    0    0 FALSE FALSE
# 38     1015    0    0    7   25    1 FALSE FALSE
# 39     3367    0    0   56    0    0 FALSE FALSE
# 40     1335    0    0   70    0    0 FALSE FALSE
# 41 Chr1  149447    T
# 42     1007    0    0    0    1    0 FALSE FALSE
# 43     1012    0    1    0    0    0 FALSE FALSE
# 44     1013    0    0    0    1    0 FALSE FALSE
# 45     1014    0    0    0    8    0 FALSE FALSE
# 46     1015    0    0    0    2    0 FALSE FALSE
# 47     3367    0    0    0    1    0 FALSE FALSE
# 48     1335    0    1    0    1    0 FALSE FALSE
# 49 Chr1  149457 <NA>
# 50     1007 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 51     1012 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 52     1013 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 53     1014 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 54     1015 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 55     3367 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 56     1335 <NA> <NA> <NA> <NA>    0 FALSE FALSE
# 57 Chr1  156248    A
# 58     1007    2    0   39    0    0 FALSE FALSE
# 59     1012    7    0   67    0    0 FALSE FALSE

```

```
# 60      1013      5      0      53      0      0 FALSE FALSE
# 61      1014     11      0      13      0      1 FALSE FALSE
# 62      1015      6      0      44      0      0 FALSE FALSE
# 63     3367      9      0      66      0      0 FALSE FALSE
# 64     1335     62      0      31      0      1 FALSE FALSE
# 65 Chr1 176517      C      1007      0      0      0      1      0 TRUE FALSE
# 66      1012      0      0      2      0      0 TRUE FALSE
# 67      1013      0      0      4      0      0 TRUE FALSE
# 68      1014      0      0      6      0      0 TRUE FALSE
# 69      1015      0      0      0      0      0 TRUE FALSE
# 70     3367      0      0      4      0      0 TRUE FALSE
# 71     1335      0      0     11      0      0 TRUE FALSE
# 72 Chr1 193761      C      1007      0      0     20     14      1 FALSE FALSE
# 73      1012      0      0     19     31      1 FALSE FALSE
# 74      1013      0      1      4      6      1 FALSE FALSE
# 75      1014      0      0      9      4      1 FALSE FALSE
# 76      1015      0      0     28     39      1 FALSE FALSE
# 77     3367      0      0      7     11      0 FALSE FALSE
# 78     1335      0      0     10     43      1 FALSE FALSE
# 79
# 80
```

one of those is a little weird:

```
xx<-snp.tables[[1]][149457,]
for (i in 2:7){xx <- rbind(xx,snp.tables[[i]][149457,])}
row.names(xx)<-names(snp.tables)
# My guess is that Chr/Pos/Ref are left as NA if coverage is zero.
xx
```

#	chr	pos	snp	Chr	Pos	Ref	Cov	a	g	c	t	n	.match	exon	indel	
# 1007	Chr1	149457	0	<NA>	NA	<NA>	0	0	0	0	0	0	0	0	FALSE	FALSE
# 1012	Chr1	149457	0	<NA>	NA	<NA>	0	0	0	0	0	0	0	0	FALSE	FALSE
# 1013	Chr1	149457	0	<NA>	NA	<NA>	0	0	0	0	0	0	0	0	FALSE	FALSE
# 1014	Chr1	149457	0	Chr1	149457	G	1	0	0	0	0	0	0	1	FALSE	FALSE
# 1015	Chr1	149457	0	<NA>	NA	<NA>	0	0	0	0	0	0	0	0	FALSE	FALSE
# 3367	Chr1	149457	0	<NA>	NA	<NA>	0	0	0	0	0	0	0	0	FALSE	FALSE
# 1335	Chr1	149457	0	Chr1	149457	G	1	0	0	1	0	0	0	0	FALSE	FALSE

5 Table 1 stats

Here is a brief summary of per-strain SNP counts, pairwise overlaps, and other conveniently available stats, such as those shown in Table 1 of the paper.

```
snp.counts      <- matrix(NA,7,4)
snp.pctofny     <- matrix(NA,7,4)
snp.pctofself   <- matrix(NA,7,4)
snp.inter       <- matrix(NA,7,7)
snp.union       <- matrix(NA,7,7)
rownames(snp.counts)      <- names(snp.tables)
rownames(snp.pctofny)     <- names(snp.tables)
rownames(snp.pctofself)   <- names(snp.tables)
rownames(snp.inter)       <- names(snp.tables)
colnames(snp.inter)       <- names(snp.tables)
rownames(snp.union)       <- names(snp.tables)
colnames(snp.union)       <- names(snp.tables)
for(stringency in 1:4){
  cat('\nStringency', stringency, ifelse(stringency==4,'(i.e. raw SAMTools SNP calls)',''),
      '\n-----\n')
  for(i in 1:7){
    f.snps.i <- refined.snps$Code$get.snps(i, stringency)
    snp.counts[i,stringency] <- sum(f.snps.i)
    for(j in i:7){
      f.snps.j <- refined.snps$Code$get.snps(j, stringency)
      snp.inter[i,j] <- sum(f.snps.i & f.snps.j)
```



```

    snp.union[i,j] <- sum(f.snps.i | f.snps.j)
  }
}
snp.pctofny [,stringency] <- snp.inter[,7]/snp.counts[7,stringency]
snp.pctofself[,stringency] <- snp.inter[,7]/snp.counts[,stringency]
cat('Union Counts:\n'); print(snp.union)
cat('Intersect Counts:\n'); print(snp.inter)
cat('Intersect as percent of union:\n'); print(snp.inter/snp.union*100,digits=3)
}

#
# Stringency 1 :
# -----
# Union Counts:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 175110 185299 304446 191894 189265 297494 193810
# 1012      NA 180026 306192 195182 192148 299196 196561
# 1013      NA      NA 249044 302460 307109 316432 306430
# 1014      NA      NA      NA 168167 192630 295200 195041
# 1015      NA      NA      NA      NA 181549 300380 194151
# 3367      NA      NA      NA      NA      NA 237364 299559
# 1335      NA      NA      NA      NA      NA      NA 181546
# Intersect Counts:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 175110 169837 119708 151383 167394 114980 162846
# 1012      NA 180026 122878 153011 169427 118194 165011
# 1013      NA      NA 249044 114751 123484 169976 124160
# 1014      NA      NA      NA 168167 157086 110331 154672
# 1015      NA      NA      NA      NA 181549 118533 168944
# 3367      NA      NA      NA      NA      NA 237364 119351
# 1335      NA      NA      NA      NA      NA      NA 181546
# Intersect as percent of union:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 100 91.7 39.3 78.9 88.4 38.6 84.0
# 1012  NA 100.0 40.1 78.4 88.2 39.5 83.9
# 1013  NA      NA 100.0 37.9 40.2 53.7 40.5
# 1014  NA      NA      NA 100.0 81.5 37.4 79.3
# 1015  NA      NA      NA      NA 100.0 39.5 87.0
# 3367  NA      NA      NA      NA      NA 100.0 39.8
# 1335  NA      NA      NA      NA      NA      NA 100.0
#
# Stringency 2 :
# -----
# Union Counts:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 182700 189244 374013 193188 196130 364691 195659
# 1012      NA 186491 375922 195230 196927 366436 196736
# 1013      NA      NA 304293 361304 377929 407688 373319
# 1014      NA      NA      NA 149531 194388 351232 188751
# 1015      NA      NA      NA      NA 190302 368934 195492
# 3367      NA      NA      NA      NA      NA 290904 363768
# 1335      NA      NA      NA      NA      NA      NA 180187
# Intersect Counts:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 182700 179947 112980 139043 176872 108913 167228
# 1012      NA 186491 114862 140792 179866 110959 169942
# 1013      NA      NA 304293 92520 116666 187509 111161
# 1014      NA      NA      NA 149531 145445 89203 140967
# 1015      NA      NA      NA      NA 190302 112272 174997
# 3367      NA      NA      NA      NA      NA 290904 107323
# 1335      NA      NA      NA      NA      NA      NA 180187
# Intersect as percent of union:
#      1007    1012    1013    1014    1015    3367    1335
# 1007 100 95.1 30.2 72.0 90.2 29.9 85.5
# 1012  NA 100.0 30.6 72.1 91.3 30.3 86.4
# 1013  NA      NA 100.0 25.6 30.9 46.0 29.8
# 1014  NA      NA      NA 100.0 74.8 25.4 74.7
# 1015  NA      NA      NA      NA 100.0 30.4 89.5

```

```

# 3367 NA NA NA NA NA 100.0 29.5
# 1335 NA NA NA NA NA NA 100.0
#
# Stringency 3 :
# -----
# Union Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 172445 184829 367417 181202 191833 358522 189174
# 1012 NA 180855 371490 187550 193436 362568 192116
# 1013 NA NA 298841 340611 373801 406440 366421
# 1014 NA NA NA 106701 188560 330490 177650
# 1015 NA NA NA NA 185294 365440 191520
# 3367 NA NA NA NA NA 286507 357408
# 1335 NA NA NA NA NA NA 170198
# Intersect Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 172445 168471 103869 97944 165906 100430 153469
# 1012 NA 180855 108206 100006 172713 104794 158937
# 1013 NA NA 298841 64931 110334 178908 102618
# 1014 NA NA NA 106701 103435 62718 99249
# 1015 NA NA NA NA 185294 106361 163972
# 3367 NA NA NA NA NA 286507 99297
# 1335 NA NA NA NA NA NA 170198
# Intersect as percent of union:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 100 91.1 28.3 54.1 86.5 28.0 81.1
# 1012 NA 100.0 29.1 53.3 89.3 28.9 82.7
# 1013 NA NA 100.0 19.1 29.5 44.0 28.0
# 1014 NA NA NA 100.0 54.9 19.0 55.9
# 1015 NA NA NA NA 100.0 29.1 85.6
# 3367 NA NA NA NA NA 100.0 27.8
# 1335 NA NA NA NA NA NA 100.0
#
# Stringency 4 (i.e. raw SAMTools SNP calls) :
# -----
# Union Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 161103 176738 343873 171675 185741 336599 180313
# 1012 NA 166089 346766 176177 186459 339458 182312
# 1013 NA NA 247737 302322 352586 386037 339669
# 1014 NA NA NA 89184 179976 295574 162912
# 1015 NA NA NA NA 174701 345396 184068
# 3367 NA NA NA NA NA 240413 331982
# 1335 NA NA NA NA NA NA 153901
# Intersect Counts:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 161103 150454 64967 78612 150063 64917 134691
# 1012 NA 166089 67060 79096 154331 67044 137678
# 1013 NA NA 247737 34599 69852 102113 61969
# 1014 NA NA NA 89184 83909 34023 80173
# 1015 NA NA NA NA 174701 69718 144534
# 3367 NA NA NA NA NA 240413 62332
# 1335 NA NA NA NA NA NA 153901
# Intersect as percent of union:
#      1007 1012 1013 1014 1015 3367 1335
# 1007 100 85.1 18.9 45.8 80.8 19.3 74.7
# 1012 NA 100.0 19.3 44.9 82.8 19.8 75.5
# 1013 NA NA 100.0 11.4 19.8 26.5 18.2
# 1014 NA NA NA 100.0 46.6 11.5 49.2
# 1015 NA NA NA NA 100.0 20.2 78.5
# 3367 NA NA NA NA NA 100.0 18.8
# 1335 NA NA NA NA NA NA 100.0

vs.stringency <- cbind(snp.counts, matrix(NA,7,1), round(snp.counts[,1:3]/snp.counts[,4]*100,1))
colnames(vs.stringency) <- c('[[1]]', '[[2]]', '[[3]]', '[[4]]', '----', '[[1]]%', '[[2]]%', '[[3]]%')

# SNPs vs filtering stringency (raw counts and as % of [[4]]). Medium filter
# adds 10-20% in most cases. Big exception is Gyre, where low coverage,

```

```
# high err rate and SAMTools conservatism seemed to seriously undercall:
print(vs.stringency)

#      [[1]] [[2]] [[3]] [[4]] ---- [[1]]% [[2]]% [[3]]%
# 1007 175110 182700 172445 161103  NA  108.7  113.4  107.0
# 1012 180026 186491 180855 166089  NA  108.4  112.3  108.9
# 1013 249044 304293 298841 247737  NA  100.5  122.8  120.6
# 1014 168167 149531 106701  89184  NA  188.6  167.7  119.6
# 1015 181549 190302 185294 174701  NA  103.9  108.9  106.1
# 3367 237364 290904 286507 240413  NA   98.7  121.0  119.2
# 1335 181546 180187 170198 153901  NA  118.0  117.1  110.6

# Intersect NY as % of self (vs stringency):
print(snp.pctofself*100, digits=3)

#      [,1] [,2] [,3] [,4]
# 1007 93.0 91.5 89.0 83.6
# 1012 91.7 91.1 87.9 82.9
# 1013 49.9 36.5 34.3 25.0
# 1014 92.0 94.3 93.0 89.9
# 1015 93.1 92.0 88.5 82.7
# 3367 50.3 36.9 34.7 25.9
# 1335 100.0 100.0 100.0 100.0

# Intersect NY as % of NY (vs stringency):
print(snp.pctofny*100, digits=3)

#      [,1] [,2] [,3] [,4]
# 1007 89.7 92.8 90.2 87.5
# 1012 90.9 94.3 93.4 89.5
# 1013 68.4 61.7 60.3 40.3
# 1014 85.2 78.2 58.3 52.1
# 1015 93.1 97.1 96.3 93.9
# 3367 65.7 59.6 58.3 40.5
# 1335 100.0 100.0 100.0 100.0
```

Quick look at coverage. Are there any NA?:

```
nacount <- NULL
for(i in 1:4){
  if(!is.null(tset[[i]])){
    nacount <- rbind(nacount,
                     unlist(lapply(tset[[i]], function(x){sum(is.na(x$Cov))}))
  )
  rownames(nacount)[nrow(nacount)] <- names(tset)[i]
}
}
nacount

#      1007 1012 1013 1014 1015 3367 1335
# snp.tables.full.unfiltered    0    0    0    0    0    0    0
# snp.tables.full.qfiltered     0    0    0    0    0    0    0
```

Seemingly no. What's average in unq- vs q-filtered:

```
snp.tables.unqfil <- tset.picker(c(1,2), table.set = tset)
snp.tables.qfil   <- tset.picker(c(3,4), table.set = tset)
cov.unqfil <- unlist(lapply(snp.tables.unqfil, function(x){mean(x$Cov)}))
cov.qfil   <- unlist(lapply(snp.tables.qfil,   function(x){mean(x$Cov, na.rm=T)}))
cov.both <- rbind(cov.unqfil, cov.qfil, cov.qfil/cov.unqfil)
i <- 1
if(!is.null(snp.tables.unqfil)){
  rownames(cov.both)[i] <- which.snp.tables(snp.tables.unqfil)
  i <- i+1
}
if(!is.null(snp.tables.qfil)){
  rownames(cov.both)[i] <- which.snp.tables(snp.tables.qfil)
```

```

i <- i+1
}
if(i==3){
  rownames(cov.both)[i] <- 'Ratio'
}
cat('Mean Coverage:\n'); cov.both

# Mean Coverage:
#
#           1007           1012           1013           1014           1015           3367           1335
# trunc-unfiltered 37.0555484 70.8060724 69.6610432 33.1009373 61.5365159 64.0284488 107.7425968
# trunc-qfiltered  28.2750286 51.3249686 45.4036337 13.7261052 48.7880005 44.8042054 81.8823765
# Ratio           0.7630444  0.7248668  0.6517794  0.4146742  0.7928301  0.6997547  0.7599815

```

5.1 Table 1 Data

Throw together the conveniently-available Table 1 data, in Table 1 row order:

```

# if coverage unavailable, build NA vector
if(!is.null(cov.unqfil)){cov.unqfilv <- cov.unqfil} else {cov.unqfilv <- rep(NA,times=7)}
if(!is.null(cov.qfil )){cov.qfilv  <- cov.qfil } else {cov.qfilv  <- rep(NA,times=7)}
tlldata.df <- data.frame(
  id      = st.locs(1:7, id=T, loc=F, date=F),
  loc     = st.locs(1:7, id=F, loc=T, date=F),
  date    = st.locs(1:7, id=F, loc=F, date=T),
  cov.unq = cov.unqfilv,
  cov.q    = cov.qfilv,
  SNPs.4   = snp.counts[,4],
  SNPs.2   = snp.counts[,2],
  olap.ny.4 = snp.pctofny[,4]*100,
  olap.ny.2 = snp.pctofny[,2]*100
)
tlrow.order <- c(7,1,2,5,3,6,4)
print(tlldata.df[tlrow.order,],digits=3)

#           id           loc date cov.unq cov.q SNPs.4 SNPs.2 olap.ny.4 olap.ny.2
# 1335 CCMP1335      New York 1958  107.7  81.9 153901 180187    100.0    100.0
# 1007 CCMP1007      Virginia 1964   37.1  28.3 161103 182700     87.5     92.8
# 1012 CCMP1012    W. Australia 1965   70.8  51.3 166089 186491     89.5     94.3
# 1015 CCMP1015    Puget Sound 1985   61.5  48.8 174701 190302     93.9     97.1
# 1013 CCMP1013      Wales 1973   69.7  45.4 247737 304293     40.3     61.7
# 3367 CCMP3367      Italy 2007   64.0  44.8 240413 290904     40.5     59.6
# 1014 CCMP1014 N. Pacific Gyre 1971  33.1  13.7  89184 149531     52.1     78.2

```

6 Shared-SNPs P-Value

Text of the main paper quotes a “p-value” for the observed degree of SNP sharing in L-clade (and/or L-clade excluding Gyre) under a null model that these isolates were sampled from a population globally in Hardy-Weinberg equilibrium. Details of this analysis are as follows.

6.1 SNP Concordance

Arbitrarily pick one isolate, say, A , as the “template”. Arbitrarily pick a heterozygous (aka “SNP”) position in A . Let p_1 , and $q_1 = 1 - p_1$ be the frequencies in the overall population of the two nucleotides observed at that position in A . (Positions having 3 or 4 nucleotide variants segregating in the population are assumed to be negligibly rare.) Under the HWE null model, a second isolate B will also be heterozygous at the same position with probability $2p_1q_1 \leq 1/2$. Similarly, this position will be heterozygous in a third isolate C with the same probability, *independently*, and so on for isolates D and E . Overall, (assuming HWE) the probability that a heterozygous position in A is simultaneously heterozygous in the other 4 isolates is at most $1/2^4 = 1/16$. Continuing, suppose we pick a second heterozygous position in A , on a different chromosome with allele frequencies $p_2, q_2 = 1 - p_2$, say. Again assuming HWE, this

position will be a SNP in all of B, C, D and E with probability $(2p_2q_2)^4 \leq 1/16$, and this is independent of the first position, since segregation on different chromosomes is unlinked. Repeat this at 24 heterozygous positions in A , one per chromosome. Then, the number of five-way concordant positions observed should be dominated by the number observed when sampling from a binomial distribution with parameters $n = 24$ and $p = 1/16$, i.e., we expect at most $1/16 = 6.25\%$ of positions to agree, or at most $24/16 = 1.5$ five-way concordant positions in total. In sharp contrast, choosing CCMP 1014 (North Pacific Gyre) as the template, we see many more five-way concordant positions than predicted under these assumptions:

```
gyre.count <- sum(snp.tables[[4]]$snp)
# NOTE: what we now call "refined" SNPs were once called "filtered" SNPs and I have NOT tried
# to update variable names and annotation in the code below to reflect the terminology change...
# 'unfil.' => unfiltered for consistency; see below.
unfil.fiveway.count <- sum(snp.tables[[4]]$snp * i4.snps)
unfil.fiveway.percent <- unfil.fiveway.count / gyre.count * 100
unfil.p.value <- pbinom(floor(unfil.fiveway.count/gyre.count*24)-1, 24, 1/16, lower.tail = FALSE)
consistency.comparison <-
  data.frame(
    fiveway.count = unfil.fiveway.count,
    fiveway.percent = unfil.fiveway.percent,
    p.value = unfil.p.value
  )
consistency.comparison

#   fiveway.count fiveway.percent      p.value
# 1           70687          79.25973 4.142632e-19
```

Namely, 89184 positions are called as SNPs in CCMP1014, of which 70687 or 79.2597327% are also called as SNPs in *all four* other L-clade isolates. 79.2597327% of 24 is 19.0223358, and the probability of seeing 19 or more “Heads” in 24 flips of a biased coin with $P(\text{Heads}) \leq 1/16$, i.e., our p-value under the HWE null hypothesis, is at most: $4.1426317 \times 10^{-19}$ based on this simple binomial model. This is obviously strong evidence against the null hypothesis.

This analysis is potentially overly-simplistic in four respects, addressed below.

1. “ $2pq \leq 1/2$ ” is conservative. Neutral theory predicts that most variant nucleotides are rare in the population, so $2pq \ll 1/2$ is to be expected. This should make our quoted p-value very conservative.
2. Effect of Erroneous SNP calls. We base our analysis on *predicted* (by SAMTOOLS) heterozygous positions, not absolute-truth, which may affect our conclusions. However,
 - False negatives in A are irrelevant, since we never examine those positions. (This is the motivation for using CCMP1014 as the template; it has the lowest predicted SNP rate, likely due to a high false negative rate in that sequencing run. As noted elsewhere, it had the lowest coverage and lowest sequence quality of the 7 isolates, both of which impare SNP calling.)
 - False negatives in $BCDE$ make such positions appear *non*-concordant. For our purpose, this makes our statistic more conservative since it can only deflate a statistic that we argue is nevertheless unexpectedly large.
 - False positive calls in A are conservatively treated, as well: barring simultaneous false-positive calls in all of $BCDE$, such a position will appear non-concordant, again deflating the statistic. The *false* positive rates in B, C, D and E are unknown, but cannot exceed SAMTOOLS *total* positive rate, which is below 1% in all 7 isolates, suggesting a simultaneous $BCDE$ false positive rate $< 10^{-8}$, which will have a negligible effect.
 - A potentially more serious issue is a true positive in A aligned to false positives in BCD and/or E . (I.e., a position that is polymorphic in the population and heterozygous in A , under the HWE null model is likely to be homozygous for one of the two alleles in one or more of $BCDE$; false positive SNP calls in all of those isolates would make the site appear concordant, i.e., provide evidence against the null model.) However, (a) my impression is that SAMTOOLS is more prone to false negative calls than to false positive calls (see Section 4), and (b) we would need a high rate of false positives to turn a truly heterozygous but non-concordant A call into a false “concordant” call—I’d expect at most half (especially given point 1

above) of *BCDE* to be heterozygous, but all would need to be falsely declared heterozygous. Such a high false positive rate on *BCDE* seems unlikely (see previous bullet), and would likely be counterbalanced by a similarly increased rate of false positives on *A*, which, as noted, tend to deflate our statistic (previous bullet again).

- Systematic errors. If there were, say, a sequence-context-dependent bias in the DNA sequencing, mapping and/or SNP-calling that tended to suggest (or hide) a SNP at some position, we’re going to systematically over- (or under-) estimate concordant SNPs across isolates. The discordance of called SNPs between the L- and H-clades and within the H-clade suggests that this is not a major problem, but it is worth noting as a possibility.
3. Discordant nucleotides at “concordant” SNP positions. A “shared” SNP at a given position might be, say, G/C in one isolate vs T/C in another, reflecting an unexpected tri-allelic position in the population or a technical sequencing error. It is inappropriate to count such a “shared” SNP position as evidence against the null hypothesis, since it isn’t clear that it is truly shared. Instead, I will identify such inconsistent positions, based on the “stringency [[2]]” criteria established above, and treat each as non-concordant. I.e., a position will be considered to be a “5-way concordant SNP” if and only if it was called as a SNP by SAMTOOLS (independently) in all 5 L-clade isolates, *and* shows the same dominant non-reference nucleotide in all 5, according to criteria [[2]] above. As it turns out, this correction has a very minor effect on the resulting p-value:

```
# 'unfil.' => Ignoring "consistency"; 'fil.' => Filtering for "consistency":
fil.fiveway.count <- sum((snp.tables[[4]]$snp * i4.snps)[union.snps == 1] & consistent[[2]])
fil.fiveway.percent <- fil.fiveway.count / gyre.count * 100
fil.p.value <- pbinom(floor(fil.fiveway.count/gyre.count*24)-1, 24, 1/16, lower.tail = FALSE)
# append new stats to previous table for easy comparison
consistency.comparison <-
  rbind(consistency.comparison,
        data.frame(
          fiveway.count = fil.fiveway.count,
          fiveway.percent = fil.fiveway.percent,
          p.value = fil.p.value
        )
  )
rownames(consistency.comparison) <- c('unfiltered', 'consistency.filtered')
consistency.comparison

#               fiveway.count fiveway.percent      p.value
# unfiltered           70687         79.25973 4.142632e-19
# consistency.filtered    69915         78.39411 1.976512e-17
```

In particular, it removes 0.9% of five-way consistent positions (only 772 of 70687 positions), and still shows a highly significant p-value.

4. “ $P(E[X]) \neq E[P(X)]$ ”. I’m expressing this poorly, but finding the p-value based on the *expected* number of concordant positions is somewhat non-standard. A more typical set-up would use the *actual* value of some statistic, then calculate the probability of observing a value that extreme (or more extreme) under the null model. The fundamental problem is that we have thousands of SNPs, but I don’t see an easy way to use more than 24 of them at a time, because potential genetic linkage seemingly destroys statistical independence, which is key to most simple analyses. A somewhat more formal, but still non-standard, approach is the following. Suppose we randomly sample one SNP per chromosome and count the number X of them that are 5-way concordant. What I outlined above calculated the p-value based on $E[X]$, the expected value of X , i.e., $P(E[X])$. Alternatively, we can calculate $E[P(X)]$, the expected p-value. (They are not the same.) In effect, this averages the p-values that would be seen over many different randomly-sampled sets of 24 SNPs. This is not difficult to calculate. First, the probability that we would observe $0 \leq i \leq 24$ concordant positions in a sample of 24, given that 78.39% of positions are concordant follows this binomial distribution:

```
x.equals.i.distribution <- dbinom(0:24, 24, fil.fiveway.percent/100)
print(x.equals.i.distribution, digits=3)

# [1] 1.07e-16 9.33e-15 3.89e-13 1.04e-11 1.97e-10 2.86e-09 3.29e-08 3.07e-07 2.37e-06 1.53e-05
```

```
# [11] 8.31e-05 3.84e-04 1.51e-03 5.05e-03 1.44e-02 3.48e-02 7.11e-02 1.21e-01 1.71e-01 1.96e-01
# [21] 1.78e-01 1.23e-01 6.09e-02 1.92e-02 2.90e-03
```

Second, the p-value corresponding to $0 \leq i \leq 24$ observed concordant positions also follows a different binomial distribution:

```
p.val.of.x.equals.i <- c(1, pbinom(0:23, 24, 1/16, lower.tail = F))
print(p.val.of.x.equals.i, digits=3)

# [1] 1.00e+00 7.88e-01 4.48e-01 1.87e-01 5.95e-02 1.49e-02 3.01e-03 4.99e-04 6.90e-05 8.02e-06
# [11] 7.89e-07 6.60e-08 4.72e-09 2.87e-10 1.49e-11 6.59e-13 2.46e-14 7.66e-16 1.98e-17 4.14e-19
# [21] 6.88e-21 8.70e-23 7.88e-25 4.56e-27 1.26e-29
```

Finally, the expected (or “average”) p-value is just the weighted average of the latter values, weighted by the former:

```
e.of.p.of.x <- sum(x.equals.i.distribution * p.val.of.x.equals.i)
e.of.p.of.x

# [1] 6.939136e-10
```

This is still highly significant, but weaker than the $P(E[X])$ analysis, basically because $X < E[X]$ has a fair probability of occurring, and the corresponding p-value $P(X)$ rises rapidly as X declines.

Another way to look at the numbers:

```
pvdof <- data.frame(x.density=x.equals.i.distribution,
                    x.cdf=cumsum(x.equals.i.distribution),
                    pval.of.x=p.val.of.x.equals.i)
print(pvdof, digits=4)

#   x.density    x.cdf pval.of.x
# 1 1.071e-16 1.071e-16 1.000e+00
# 2 9.325e-15 9.432e-15 7.875e-01
# 3 3.891e-13 3.985e-13 4.476e-01
# 4 1.035e-11 1.075e-11 1.869e-01
# 5 1.972e-10 2.080e-10 5.950e-02
# 6 2.862e-09 3.070e-09 1.490e-02
# 7 3.289e-08 3.596e-08 3.010e-03
# 8 3.068e-07 3.428e-07 4.994e-04
# 9 2.366e-06 2.709e-06 6.899e-05
# 10 1.526e-05 1.797e-05 8.015e-06
# 11 8.306e-05 1.010e-04 7.887e-07
# 12 3.836e-04 4.846e-04 6.603e-08
# 13 1.508e-03 1.992e-03 4.716e-09
# 14 5.050e-03 7.042e-03 2.875e-10
# 15 1.440e-02 2.144e-02 1.493e-11
# 16 3.482e-02 5.626e-02 6.590e-13
# 17 7.107e-02 1.273e-01 2.456e-14
# 18 1.213e-01 2.487e-01 7.662e-16
# 19 1.712e-01 4.199e-01 1.977e-17
# 20 1.962e-01 6.161e-01 4.143e-19
# 21 1.780e-01 7.941e-01 6.877e-21
# 22 1.230e-01 9.170e-01 8.701e-23
# 23 6.085e-02 9.779e-01 7.884e-25
# 24 1.920e-02 9.971e-01 4.556e-27
# 25 2.903e-03 1.000e+00 1.262e-29
```

E.g., row 9 in that table says that the concordance rate (78%) is so high that a sample of 24 SNPs will almost always have 9 or more five-way concordant positions (probability of fewer is only 2.709e-06), while under the null model, seeing 9 or more is very unlikely (probability at most 6.899e-05). ***AM I OFF-BY-ONE INTERPRETING ROW 9 HERE??***

6.2 Notes

In earlier drafts, an analog of the above analysis was based on the concordance of *refined* SNPs. This now seems to me to be questionable, since the “refined” SNP calling makes SNPs called across L-clade non-independent. OTOH, the above analysis seems valid: SAMTOOLS was run on each isolate independently, and likewise “criterion [[2]]” is evaluated independently in each strain, and is being used here solely to remove SNP predictions, not to add them. “Systematic errors” as outlined above remain a potential problem, but again discordance with/within H-clade suggests that this is of limited concern.

For completeness, I did a similar analysis including a sample of H-clade comparisons: Gyre vs Italy, NY vs Italy, NY vs Italy+Wales, and of Italy vs Wales. As expected, none of these show a statistically significant p-value, although the $\approx 40\%$ concordance in the 2-way comparisons, while $< 1/2$ as predicted, is a bit higher than I expected based on “neutral theory implies many rare variants.” (I did not bother to include “criterion[[2]] filtering” in these calculations.)

```
# 'gi.twoway' => gyre vs italy 2-way concordance;
# 'ni.twoway' => new york vs italy 2-way concordance;
# not bothering with criterion[[2]] filtering
gi.twoway.count <- sum(snp.tables[[4]]$snp * snp.tables[[6]]$snp)
gi.twoway.percent <- gi.twoway.count / gyre.count * 100
gi.p.value <- pbinom(floor(gi.twoway.count/gyre.count*24)-1, 24, 1/2, lower.tail = FALSE)
ny.count <- sum(snp.tables[[7]]$snp)
ni.twoway.count <- sum(snp.tables[[7]]$snp * snp.tables[[6]]$snp)
ni.twoway.percent <- ni.twoway.count / ny.count * 100
ni.p.value <- pbinom(floor(ni.twoway.count/ny.count*24)-1, 24, 1/2, lower.tail = FALSE)
niw.threeway.count <- sum(snp.tables[[7]]$snp * snp.tables[[6]]$snp * snp.tables[[3]]$snp)
niw.threeway.percent <- niw.threeway.count / ny.count * 100
niw.p.value <- pbinom(floor(niw.threeway.count/ny.count*24)-1, 24, 1/4, lower.tail = FALSE)
it.count <- sum(snp.tables[[6]]$snp)
iw.twoway.count <- sum(snp.tables[[6]]$snp * snp.tables[[3]]$snp)
iw.twoway.percent <- iw.twoway.count / it.count * 100
iw.p.value <- pbinom(floor(iw.twoway.count/it.count*24)-1, 24, 1/2, lower.tail = FALSE)
consistency.comparison <-
  rbind(data.frame(
    fiveway.count = c(gi.twoway.count, ni.twoway.count, niw.threeway.count, iw.twoway.count),
    fiveway.percent = c(gi.twoway.percent, ni.twoway.percent, niw.threeway.percent, iw.twoway.percent),
    p.value = c(gi.p.value, ni.p.value, niw.p.value, iw.p.value)
  ))
colnames(consistency.comparison)[1:2] <- c('552232way.count', '552232way.percent') # old col names misleading
rownames(consistency.comparison)[3:6] <- c('gyre.vs.italy', 'new.york.vs.italy', # new rows
      'ny.vs.it.plus.wales', 'it.vs.wales')
consistency.comparison

#           552232way.count 552232way.percent      p.value
# unfiltered              70687          79.25973 4.142632e-19
# consistency.filtered    69915          78.39411 1.976512e-17
# gyre.vs.italy           34023          38.14922 9.242052e-01
# new.york.vs.italy       62332          40.50136 9.242052e-01
# ny.vs.it.plus.wales     35796          23.25911 7.533516e-01
# it.vs.wales             102113         42.47399 8.462719e-01
```

6.3 P-Value: The Bottom Line

So, what to say in the body of the paper? $E[P(X)]$ is highly significant, and conservative, but complex to explain. $P(E[X])$ is simpler to explain, but may be criticized as misleading if we aren’t very careful in that explanation. I’m slightly leaning towards the last option, but want to sleep on it and draft the key sentence or two before settling.

7 Sharing

The following analysis looks at the sharing patterns among the consistent SNPs. I assume that shared SNPs reflect shared ancestry, and that SNPs accumulate slowly over time. Then, in outline, the story is consistent with what we

have seen in other analyses—there seem to be 3 groups: 1013 (Wales) in one, 3367 (Italy) in another, and the other 5 in a third, with some hints as to the order of divergence. A caveat is that in a sexual population, non-shared SNPs do not immediately imply non-shared ancestry; they may merely reflect Hardy-Weinberg capturing a homozygous state in one isolate vs the other. (Or read errors, etc.) Thus, if we are right that the H-isolates retain sex, then the large number of “private” SNPs in H may be at least partially due to HWE.

Analysis is broken into cases based on how many strains share a particular SNP.

7.1 Code

To categorize SNPs by sharing patterns, first convert the 7-way consistent sharing pattern into a 7-bit binary number, and tabulate based on that:

```
# convert (n x 7) 0-1 matrix to n vector of 0-127
tobin <- function(x){
  bin <- integer(nrow(x)) # initialized to 0
  for(i in 1:7){
    bin <- bin*2 + as.integer(x[,i]>0)
  }
  return(bin)
}

# get full set of patterns
snp.pattern.all <- lapply(non.refs,tobin)
# prune to just the consistent ones
snp.pattern <- snp.pattern.all
for(i in 1:3){
  snp.pattern[[i]][!consistent[[i]]] <- NA
}

# analogous to built-in ``table'' but simpler. Count entries in an integer
# vector sharing values in a (smallish) range. Result is a 2-column matrix with
# the shared values in col 1 and count of occurrences of that value in col 2.
# Out-of-range values cause subscript error.
mytable <- function(vec, therange=range(vec,na.rm=T)){
  counts <- matrix(0,nrow=therange[2]-therange[1]+1,ncol=2,dimnames=list(NULL,c('val','count')))
  counts[1:nrow(counts),1] <- therange[1]:therange[2]
  for(i in 1:length(vec)){
    if(!is.na(vec[i])){
      counts[vec[i]-therange[1]+1,2] <- counts[vec[i]-therange[1]+1,2] + 1
    }
  }
  return(counts)
}

pattern.counts <- lapply(snp.pattern, function(x){mytable(x,c(0,127))})
```

To display the results, build a data frame whose i -th row, $0 \leq i \leq 127$ shows one of the 128 possible sharing patterns, with counts of the numbers of consistent, shared SNPs with that pattern according to criteria c1-c3.

```
tobitvec <- function(x){
  bitvec <- integer(7)
  for(i in 1:7){
    bitvec[i] <- x %% 2
    x <- x %% 2
  }
  return(bitvec)
}

flg <- function(x){
  return(ifelse(x==1,'X',''))
}

pat.summary <- function(listOfTbls){
  mydf <- data.frame(pat=0:127,sharedBy=NA,
```

```

        tp1007='',tp1012='',tp1013='',tp1014='',tp1015='',tp3367='',tp1335='',
        count1=NA,count2=NA,count3=NA,count4=NA,stringsAsFactors=F)

for(i in 1:128){
  bvec <- tobitvec(i-1)
  mydf[i,'sharedBy']=sum(bvec)
  mydf[i,'tp1007']=flg(bvec[1])
  mydf[i,'tp1012']=flg(bvec[2])
  mydf[i,'tp1013']=flg(bvec[3])
  mydf[i,'tp1014']=flg(bvec[4])
  mydf[i,'tp1015']=flg(bvec[5])
  mydf[i,'tp3367']=flg(bvec[6])
  mydf[i,'tp1335']=flg(bvec[7])
}

for(i in 1:length(listOfTbls)){
  tbl <- listOfTbls[[i]]
  if(!is.null(tbl)){
    mydf[,9+i] <- tbl[,2] ## count1/2/3/4 are columns 10/11/12/13 in mydf
    #for(j in 1:length(tbl)){
    #  k <- as.integer(rownames(tbl)[j]);
    #  mydf[k+1,9+i] <- tbl[j] ## count1/2/3 are columns 10/11/12
    #}
  }
}

mydf$pat <-as.octmode(mydf$pat) # display bit pattern in octal
return(mydf)
}

pat.summaries <- pat.summary(pattern.counts)

```

7.2 Sanity Checks

Some sanity checking: table sums equal to number of consistent positions?

```

all(consistent.count == apply(pat.summaries[,10:13],2,sum))

# [1] TRUE

```

More sanity checking: visually inspect a pattern with small counts, specifically pattern 12, i.e., consistent SNPs shared by only strains 1014 and 1015 (2nd and 3 rows from bottom, binary code $12 = 2^3 + 2^2$). There are only 10 such positions on Chr1. Chr1 2524239 has pattern 12 under criteria c1 and c2 but not c3; Chr1 1088766 has in c2 only. Both look good. Neither position is a *called* SNP except in 1015. However, all but 1 nonreference read agree with the called SNP (the exception being one read in Wales). Both 1014 and 1015 have at least 2 non-reference reads, comprising at least 5% of coverage, and in both strains, those reads are on the same non-reference base, satisfying criterion c2. The other strains have higher coverage and/or lower non-reference counts, so they do not satisfy c2. Position 2524239 also satisfies c1, but not c3, since 2 reads out of 35 is below the 10% threshold. (It is pattern 4 under c3, i.e., a SNP private to 1015.) Position 1088766 is also pattern 4 under c3 (2 reads out of 56 in 1335 is below both thresholds), and it is not consistent under c1, since the single A read in 1013 is discordant with the other non-reference reads.

```

unlist(lapply(snp.pattern,function(x){sum(x==12,na.rm=T)}))

# [1] 143 136 176 417

sp1 <- snp.pattern[[1]]==12
sp2 <- snp.pattern[[2]]==12
sp3 <- snp.pattern[[3]]==12
sp4 <- snp.pattern[[4]]==12
c(sum(sp1,na.rm=T), sum(sp2,na.rm=T), sum(sp3,na.rm=T), sum(sp4,na.rm=T))

# [1] 143 136 176 417

```

```
r1 <- rownames(non.refs[[1]])[which(sp1)]
r2 <- rownames(non.refs[[2]])[which(sp2)]
r3 <- rownames(non.refs[[3]])[which(sp3)]
r4 <- rownames(non.refs[[4]])[which(sp4)]
```

```
r2
```

```
# [1] "Chr1:1088766"      "Chr1:2524239"      "Chr2:713075"       "Chr2:1464209"
# [5] "Chr2:2406031"      "Chr2:2480466"      "Chr2:2480532"       "Chr2:2480838"
# [9] "Chr2:2481998"      "Chr2:2483322"      "Chr2:2488863"       "Chr2:2489189"
# [13] "Chr2:2490933"      "Chr2:2492886"      "Chr2:2492887"       "Chr2:2497794"
# [17] "Chr2:2500122"      "Chr2:2503000"      "Chr2:2507585"       "Chr2:2507680"
# [21] "Chr2:2510117"      "Chr2:2513923"      "Chr2:2515103"       "Chr2:2516669"
# [25] "Chr2:2516751"      "Chr2:2518558"      "Chr2:2518653"       "Chr2:2518980"
# [29] "Chr2:2519285"      "Chr2:2519288"      "Chr2:2519718"       "Chr2:2520984"
# [33] "Chr2:2521271"      "Chr2:2522648"      "Chr2:2524223"       "Chr2:2524439"
# [37] "Chr2:2525160"      "Chr2:2525463"      "Chr2:2527916"       "Chr2:2528472"
# [41] "Chr2:2528769"      "Chr2:2529076"      "Chr2:2529140"       "Chr2:2529186"
# [45] "Chr2:2529432"      "Chr2:2529684"      "Chr2:2530064"       "Chr2:2530216"
# [49] "Chr2:2530239"      "Chr2:2530294"      "Chr2:2530768"       "Chr2:2530896"
# [53] "Chr2:2531114"      "Chr2:2531285"      "Chr2:2531498"       "Chr2:2531567"
# [57] "Chr2:2532173"      "Chr2:2532365"      "Chr2:2533028"       "Chr2:2533171"
# [61] "Chr2:2533440"      "Chr2:2534441"      "Chr2:2535121"       "Chr2:2535122"
# [65] "Chr2:2535314"      "Chr2:2535493"      "Chr2:2535503"       "Chr2:2535509"
# [69] "Chr2:2535862"      "Chr2:2536242"      "Chr2:2537201"       "Chr2:2537864"
# [73] "Chr2:2537917"      "Chr2:2538072"      "Chr2:2538498"       "Chr2:2539318"
# [77] "Chr2:2543595"      "Chr2:2545615"      "Chr2:2545798"       "Chr2:2546865"
# [81] "Chr2:2546991"      "Chr2:2547055"      "Chr2:2547086"       "Chr2:2547120"
# [85] "Chr2:2547155"      "Chr2:2547212"      "Chr2:2547248"       "Chr2:2547318"
# [89] "Chr2:2547554"      "Chr2:2547938"      "Chr2:2547944"       "Chr2:2548131"
# [93] "Chr2:2549281"      "Chr2:2551574"      "Chr2:2551930"       "Chr2:2554708"
# [97] "Chr2:2554860"      "Chr2:2555005"      "Chr2:2555203"       "Chr2:2555820"
# [101] "Chr3:496665"       "Chr4:901220"       "Chr4:983962"        "Chr4:1086210"
# [105] "Chr4:1086589"      "Chr5:7509"         "Chr5:141375"        "Chr5:1397904"
# [109] "Chr6:1034519"      "Chr7:399475"       "Chr8:556556"        "Chr10:95217"
# [113] "Chr12:422344"      "Chr12:458461"      "Chr13:963939"       "Chr14:56058"
# [117] "Chr15:417704"      "Chr16a:39914"      "Chr16a:39917"       "Chr16a:394030"
# [121] "Chr17:461465"      "Chr18:673261"      "Chr19a_19:303090"   "Chr19a_19:308244"
# [125] "Chr19b_31:4468"    "Chr19b_31:138559"  "Chr19c_29:64170"    "Chr19c_29:64811"
# [129] "Chr19c_29:65720"   "Chr20:15766"       "Chr20:230994"       "Chr20:486431"
# [133] "Chr22:249009"      "Chr22:380816"      "Chr23:274291"       "Chr24:114599"
```

```
c1 <- as.integer(unlist(lapply(strsplit(r1[1:min(20,length(r1))],':',fixed=TRUE),function(x){x[2]})))
c2 <- as.integer(unlist(lapply(strsplit(r2[1:min(20,length(r2))],':',fixed=TRUE),function(x){x[2]})))
c3 <- as.integer(unlist(lapply(strsplit(r3[1:min(20,length(r3))],':',fixed=TRUE),function(x){x[2]})))
c4 <- as.integer(unlist(lapply(strsplit(r4[1:min(20,length(r4))],':',fixed=TRUE),function(x){x[2]})))
```

```
c1
```

```
# [1] 198498 914018 1317406 1481838 1501481 1878058 2145849 2388286 2524239 2718093 62676
# [12] 393166 458314 713075 1416054 2148271 2149651 2310069 2406031 2480466
```

```
c2
```

```
# [1] 1088766 2524239 713075 1464209 2406031 2480466 2480532 2480838 2481998 2483322 2488863
# [12] 2489189 2490933 2492886 2492887 2497794 2500122 2503000 2507585 2507680
```

```
c3
```

```
# [1] 371484 1210354 1886633 2264683 2898352 207186 903516 1264023 1276745 1464904 1464905
# [12] 2229060 2347253 2406031 2439655 2480532 2480838 2483322 2488863 2489189
```

```
c4
```

```
# [1] 518347 691730 767408 1049906 1390437 2072951 2254059 2254789 2264683 2823796 2898352
# [12] 2998868 77394 77407 155680 761325 968120 1182096 1222176 1264023
```

```
seecounts(c2,snp.tables=snp.tables)
```

#	chr	pos	Ref	Strain	A	G	C	T	SNP	exon	indel	nrf	rat
# 1	Chr1	1088766	G										
# 2				1007	0	32	1	0	0	FALSE	FALSE		
# 3				1012	0	39	1	0	0	FALSE	FALSE		
# 4				1013	1	74	0	0	0	FALSE	FALSE		
# 5				1014	0	26	2	0	0	FALSE	FALSE		
# 6				1015	0	38	9	0	1	FALSE	FALSE		
# 7				3367	0	36	1	0	0	FALSE	FALSE		
# 8				1335	0	54	2	0	0	FALSE	FALSE		
# 9	Chr1	2524239	C										
# 10				1007	0	0	37	0	0	TRUE	FALSE		
# 11				1012	0	0	47	0	0	TRUE	FALSE		
# 12				1013	0	0	62	0	0	TRUE	FALSE		
# 13				1014	0	0	33	2	0	TRUE	FALSE		
# 14				1015	0	0	11	15	1	TRUE	FALSE		
# 15				3367	0	0	41	0	0	TRUE	FALSE		
# 16				1335	0	0	95	0	0	TRUE	FALSE		
# 17	Chr1	713075	T										
# 18				1007	0	0	0	43	0	TRUE	FALSE		
# 19				1012	0	0	0	115	0	TRUE	FALSE		
# 20				1013	1	0	0	89	0	TRUE	FALSE		
# 21				1014	0	0	0	63	0	TRUE	FALSE		
# 22				1015	0	0	0	97	0	TRUE	FALSE		
# 23				3367	0	0	0	75	0	TRUE	FALSE		
# 24				1335	0	0	0	149	0	TRUE	FALSE		
# 25	Chr1	1464209	T										
# 26				1007	0	0	0	26	0	FALSE	FALSE		
# 27				1012	0	0	0	57	0	FALSE	FALSE		
# 28				1013	0	0	0	36	0	FALSE	FALSE		
# 29				1014	0	0	0	25	0	FALSE	FALSE		
# 30				1015	0	0	0	40	0	FALSE	FALSE		
# 31				3367	0	0	0	52	0	FALSE	FALSE		
# 32				1335	0	0	0	104	0	FALSE	FALSE		
# 33	Chr1	2406031	C										
# 34				1007	0	0	29	0	0	TRUE	FALSE		
# 35				1012	0	0	52	0	0	TRUE	FALSE		
# 36				1013	0	0	71	0	0	TRUE	FALSE		
# 37				1014	0	0	20	0	0	TRUE	FALSE		
# 38				1015	0	0	51	0	0	TRUE	FALSE		
# 39				3367	0	0	60	0	0	TRUE	FALSE		
# 40				1335	0	0	97	0	0	TRUE	FALSE		
# 41	Chr1	2480466	A										
# 42				1007	33	0	0	0	0	TRUE	FALSE		
# 43				1012	57	0	0	1	0	TRUE	FALSE		
# 44				1013	56	0	0	0	0	TRUE	FALSE		
# 45				1014	24	0	0	0	0	TRUE	FALSE		
# 46				1015	62	0	0	0	0	TRUE	FALSE		
# 47				3367	43	0	0	1	0	TRUE	FALSE		
# 48				1335	99	0	0	0	0	TRUE	FALSE		
# 49	Chr1	2480532	G										
# 50				1007	0	35	0	0	0	TRUE	FALSE		
# 51				1012	0	50	0	0	0	TRUE	FALSE		
# 52				1013	0	70	0	0	0	TRUE	FALSE		
# 53				1014	0	19	0	0	0	TRUE	FALSE		
# 54				1015	0	43	0	0	0	TRUE	FALSE		
# 55				3367	0	46	0	0	0	TRUE	FALSE		
# 56				1335	0	113	0	0	0	TRUE	FALSE		
# 57	Chr1	2480838	T										
# 58				1007	0	0	0	16	0	TRUE	FALSE		
# 59				1012	0	0	0	30	0	TRUE	FALSE		
# 60				1013	0	0	0	46	0	TRUE	FALSE		
# 61				1014	0	0	0	26	0	TRUE	FALSE		
# 62				1015	0	0	0	23	0	TRUE	FALSE		
# 63				3367	0	0	0	28	0	TRUE	FALSE		
# 64				1335	0	0	0	100	0	TRUE	FALSE		
# 65	Chr1	2481998	T										
# 66				1007	0	0	0	34	0	TRUE	FALSE		

# 67		1012	0	0	0	54	0	TRUE	FALSE
# 68		1013	0	0	0	90	0	TRUE	FALSE
# 69		1014	0	0	0	36	0	TRUE	FALSE
# 70		1015	0	0	0	59	0	TRUE	FALSE
# 71		3367	0	0	0	64	0	TRUE	FALSE
# 72		1335	0	0	0	100	0	TRUE	FALSE
# 73	Chr1 2483322	A							
# 74		1007	35	0	0	0	0	TRUE	FALSE
# 75		1012	59	1	0	0	0	TRUE	FALSE
# 76		1013	117	0	0	0	0	TRUE	FALSE
# 77		1014	53	0	0	0	0	TRUE	FALSE
# 78		1015	78	0	0	0	0	TRUE	FALSE
# 79		3367	71	0	1	0	0	TRUE	FALSE
# 80		1335	131	0	0	0	0	TRUE	FALSE
# 81	Chr1 2488863	C							
# 82		1007	0	0	29	0	0	FALSE	FALSE
# 83		1012	0	0	55	0	0	FALSE	FALSE
# 84		1013	0	0	49	0	0	FALSE	FALSE
# 85		1014	0	0	27	0	0	FALSE	FALSE
# 86		1015	0	0	48	0	0	FALSE	FALSE
# 87		3367	0	0	63	0	0	FALSE	FALSE
# 88		1335	0	0	90	0	0	FALSE	FALSE
# 89	Chr1 2489189	C							
# 90		1007	0	0	37	0	0	FALSE	FALSE
# 91		1012	0	0	88	0	0	FALSE	FALSE
# 92		1013	0	0	60	0	0	FALSE	FALSE
# 93		1014	0	0	45	0	0	FALSE	FALSE
# 94		1015	0	0	68	0	0	FALSE	FALSE
# 95		3367	0	0	39	0	0	FALSE	FALSE
# 96		1335	0	0	132	0	0	FALSE	FALSE
# 97	Chr1 2490933	G							
# 98		1007	0	35	0	0	0	FALSE	FALSE
# 99		1012	0	72	0	0	0	FALSE	FALSE
# 100		1013	0	60	0	0	0	FALSE	FALSE
# 101		1014	0	25	1	0	0	FALSE	FALSE
# 102		1015	0	44	0	0	0	FALSE	FALSE
# 103		3367	0	47	1	1	0	FALSE	FALSE
# 104		1335	0	71	0	0	0	FALSE	FALSE
# 105	Chr1 2492886	T							
# 106		1007	0	0	0	34	0	FALSE	FALSE
# 107		1012	0	0	0	98	0	FALSE	FALSE
# 108		1013	0	1	0	60	0	FALSE	FALSE
# 109		1014	0	0	0	37	0	FALSE	FALSE
# 110		1015	0	0	0	75	0	FALSE	FALSE
# 111		3367	0	0	0	73	0	FALSE	FALSE
# 112		1335	0	0	0	125	0	FALSE	FALSE
# 113	Chr1 2492887	G							
# 114		1007	0	33	0	0	0	FALSE	FALSE
# 115		1012	0	95	0	0	0	FALSE	FALSE
# 116		1013	0	59	0	0	0	FALSE	FALSE
# 117		1014	0	36	0	0	0	FALSE	FALSE
# 118		1015	0	72	0	0	0	FALSE	FALSE
# 119		3367	0	71	0	0	0	FALSE	FALSE
# 120		1335	0	125	0	0	0	FALSE	FALSE
# 121	Chr1 2497794	T							
# 122		1007	0	0	0	43	0	TRUE	FALSE
# 123		1012	0	0	0	77	0	TRUE	FALSE
# 124		1013	0	0	0	76	0	TRUE	FALSE
# 125		1014	0	0	0	25	0	TRUE	FALSE
# 126		1015	0	0	0	75	0	TRUE	FALSE
# 127		3367	0	0	0	63	0	TRUE	FALSE
# 128		1335	0	0	0	130	0	TRUE	FALSE
# 129	Chr1 2500122	A							
# 130		1007	33	0	0	0	0	FALSE	FALSE
# 131		1012	71	0	0	0	0	FALSE	FALSE
# 132		1013	58	0	0	0	0	FALSE	FALSE
# 133		1014	31	0	0	0	0	FALSE	FALSE

```

# 134      1015 48 0 0 0 0 FALSE FALSE
# 135      3367 46 0 0 0 0 FALSE FALSE
# 136      1335 74 0 0 0 0 FALSE FALSE
# 137 Chr1 2503000 T
# 138      1007 0 0 0 35 0 FALSE FALSE
# 139      1012 0 0 0 65 0 FALSE FALSE
# 140      1013 0 0 0 83 0 FALSE FALSE
# 141      1014 0 0 0 33 0 FALSE FALSE
# 142      1015 0 0 0 54 0 FALSE FALSE
# 143      3367 0 0 0 56 0 FALSE FALSE
# 144      1335 0 0 0 65 0 FALSE FALSE
# 145 Chr1 2507585 A
# 146      1007 40 0 0 0 0 TRUE FALSE
# 147      1012 65 0 0 0 0 TRUE FALSE
# 148      1013 56 0 0 0 0 TRUE FALSE
# 149      1014 31 0 0 0 0 TRUE FALSE
# 150      1015 47 0 0 0 0 TRUE FALSE
# 151      3367 78 0 1 0 0 TRUE FALSE
# 152      1335 118 0 0 0 0 TRUE FALSE
# 153 Chr1 2507680 A
# 154      1007 30 0 0 0 0 FALSE FALSE
# 155      1012 82 0 0 0 0 FALSE FALSE
# 156      1013 62 0 0 0 0 FALSE FALSE
# 157      1014 32 1 0 0 0 FALSE FALSE
# 158      1015 70 0 0 0 0 FALSE FALSE
# 159      3367 87 0 0 0 0 FALSE FALSE
# 160      1335 125 1 0 0 0 FALSE FALSE

```

Position 1088766, however, in a good example of the situation that motivated this analysis—one strain has a G/C SNP and 5 of the other 6 strains have nonreference reads consistent with that SNP. Although, excluding 1015, the nonreference read counts are not high enough to justify a SNP call in any strain considered in isolation, the fact that they *consistently* agree with the 1015 SNP suggests that they are real. One alternative hypothesis is that there is some sequence-dependent bias at this locus that favors misreading a G as a C. On the other hand, one could equally well posit a shared SNP, and a locus-dependant bias that *supresses* C reads, explaining the unbalanced readout that we observe. However, it is hard to reconcile either view with the significant strain-specific patterns that we see in the shared SNPs (as seen below). I think a more likely explanation is that (a) there are some number of relatively rare SNPs present in each of the sampled populations, (b) some of these SNPs happened to be present in one or two cells of the roughly 5-10 cells that we believe constituted the founding population of the culture grown for sequencing, and (c) stochastic effects during culture growth and during sequencing may have further perturbed the apparent frequency of each variant, but the bottom line is that the above-threshold presence of consistent non-reference reads is evidence for shared SNPs at the population level (and the proportions of such reads represent estimates of the population-level frequencies of the variants, albeit a noisy estimate at any specific position).

An aside: I was curious to see whether there is any consistent pattern to positions that are called consistent SNPs in all but Italy, so I repeated the above, basically. My summary is that coverage in Italy tends to be below average in these positions, but otherwise they don't stand out. For the record:

```

abit <- snp.pattern[[2]]==125
abit[is.na(abit)]<-F
sum(abit)

# [1] 14648

rabit <- rownames(non.refs[[2]])[which(abit)]
rabits <- rabit[1:20]
cabit <- as.integer(unlist(lapply(strsplit(rabits,':',fixed=TRUE),function(x){x[2]})))
cabit

# [1] 1244 1575 6485 7181 7220 7661 8144 8208 8518 8552 8567 8670 8685 14361 15254
# [16] 15280 16103 17587 18904 25546

seecounts(cabit,snp.tables=snp.tables)

#      chr   pos Ref Strain   A   G   C   T SNP  exon indel nrf rat
# 1  Chr1 1244   G
# 2      1007 3 30 0 0 0 TRUE FALSE
# 3      1012 5 54 0 0 0 TRUE FALSE

```

# 4			1013	15	47	0	0	1	TRUE	FALSE
# 5			1014	3	30	0	0	0	TRUE	FALSE
# 6			1015	21	68	0	0	1	TRUE	FALSE
# 7			3367	0	10	0	0	0	TRUE	FALSE
# 8			1335	108	108	0	0	1	TRUE	FALSE
# 9	Chr1	1575	G							
# 10			1007	26	11	0	0	0	TRUE	FALSE
# 11			1012	49	27	0	0	0	TRUE	FALSE
# 12			1013	19	28	0	0	0	TRUE	FALSE
# 13			1014	15	17	0	0	0	TRUE	FALSE
# 14			1015	46	42	0	0	1	TRUE	FALSE
# 15			3367	0	11	0	0	0	TRUE	FALSE
# 16			1335	37	99	0	0	0	TRUE	FALSE
# 17	Chr1	6485	G							
# 18			1007	26	20	0	0	0	TRUE	FALSE
# 19			1012	33	39	0	0	0	TRUE	FALSE
# 20			1013	54	48	0	0	0	TRUE	FALSE
# 21			1014	13	10	0	0	0	TRUE	FALSE
# 22			1015	34	41	0	0	1	TRUE	FALSE
# 23			3367	0	42	0	0	0	TRUE	FALSE
# 24			1335	71	69	0	0	0	TRUE	FALSE
# 25	Chr1	7181	G							
# 26			1007	0	37	31	0	0	TRUE	FALSE
# 27			1012	0	66	36	0	0	TRUE	FALSE
# 28			1013	0	30	86	0	0	TRUE	FALSE
# 29			1014	0	19	8	0	0	TRUE	FALSE
# 30			1015	0	44	33	0	1	TRUE	FALSE
# 31			3367	0	33	0	0	0	TRUE	FALSE
# 32			1335	0	94	78	0	0	TRUE	FALSE
# 33	Chr1	7220	C							
# 34			1007	17	0	26	6	0	TRUE	FALSE
# 35			1012	45	0	31	14	0	TRUE	FALSE
# 36			1013	112	1	41	14	0	TRUE	FALSE
# 37			1014	16	1	16	2	0	TRUE	FALSE
# 38			1015	66	0	26	7	1	TRUE	FALSE
# 39			3367	0	0	24	0	0	TRUE	FALSE
# 40			1335	68	0	51	25	0	TRUE	FALSE
# 41	Chr1	7661	T							
# 42			1007	0	0	10	14	0	TRUE	FALSE
# 43			1012	0	0	7	24	0	TRUE	FALSE
# 44			1013	0	0	32	23	1	TRUE	FALSE
# 45			1014	0	0	8	11	0	TRUE	FALSE
# 46			1015	0	0	6	41	0	TRUE	FALSE
# 47			3367	0	0	0	8	0	TRUE	FALSE
# 48			1335	0	0	8	42	0	TRUE	FALSE
# 49	Chr1	8144	G							
# 50			1007	10	16	0	1	0	TRUE	FALSE
# 51			1012	19	28	0	0	1	TRUE	FALSE
# 52			1013	63	67	0	0	0	TRUE	FALSE
# 53			1014	7	12	0	0	0	TRUE	FALSE
# 54			1015	18	28	0	0	0	TRUE	FALSE
# 55			3367	0	7	0	0	0	TRUE	FALSE
# 56			1335	17	58	0	0	1	TRUE	FALSE
# 57	Chr1	8208	G							
# 58			1007	0	15	0	8	1	TRUE	FALSE
# 59			1012	0	28	0	16	0	TRUE	FALSE
# 60			1013	0	24	0	63	0	TRUE	FALSE
# 61			1014	0	15	0	4	0	TRUE	FALSE
# 62			1015	0	25	0	13	1	TRUE	FALSE
# 63			3367	0	9	0	1	0	TRUE	FALSE
# 64			1335	0	49	0	21	1	TRUE	FALSE
# 65	Chr1	8518	T							
# 66			1007	0	0	20	18	1	FALSE	FALSE
# 67			1012	0	0	45	30	1	FALSE	FALSE
# 68			1013	0	0	57	75	1	FALSE	FALSE
# 69			1014	0	0	10	32	0	FALSE	FALSE
# 70			1015	0	0	41	18	1	FALSE	FALSE
# 71			3367	0	0	0	11	0	FALSE	FALSE
# 72			1335	0	0	120	71	1	FALSE	FALSE
# 73	Chr1	8552	G							
# 74			1007	3	13	0	0	0	TRUE	FALSE
# 75			1012	21	31	0	1	0	TRUE	FALSE
# 76			1013	33	35	0	0	1	TRUE	FALSE
# 77			1014	7	15	0	0	0	TRUE	FALSE
# 78			1015	14	22	0	0	0	TRUE	FALSE
# 79			3367	0	28	0	0	0	TRUE	FALSE
# 80			1335	27	59	0	0	0	TRUE	FALSE
# 81	Chr1	8567	A							
# 82			1007	16	18	0	0	1	TRUE	FALSE
# 83			1012	34	35	0	0	1	TRUE	FALSE

# 84			1013	66	75	0	0	1	TRUE	FALSE
# 85			1014	9	4	0	0	0	TRUE	FALSE
# 86			1015	17	31	0	0	1	TRUE	FALSE
# 87			3367	29	0	0	0	0	TRUE	FALSE
# 88			1335	59	44	0	0	1	TRUE	FALSE
# 89	Chr1	8670	A							
# 90			1007	19	0	0	7	0	TRUE	FALSE
# 91			1012	36	0	0	12	0	TRUE	FALSE
# 92			1013	44	0	0	12	0	TRUE	FALSE
# 93			1014	10	0	0	7	0	TRUE	FALSE
# 94			1015	24	0	0	11	1	TRUE	FALSE
# 95			3367	18	0	0	0	0	TRUE	FALSE
# 96			1335	27	0	0	6	0	TRUE	FALSE
# 97	Chr1	8685	G							
# 98			1007	7	16	0	0	0	TRUE	FALSE
# 99			1012	12	37	0	0	0	TRUE	FALSE
# 100			1013	18	30	0	0	1	TRUE	FALSE
# 101			1014	5	32	0	0	0	TRUE	FALSE
# 102			1015	11	35	0	0	1	TRUE	FALSE
# 103			3367	0	12	0	0	0	TRUE	FALSE
# 104			1335	5	45	0	0	0	TRUE	FALSE
# 105	Chr1	14361	A							
# 106			1007	29	7	0	0	0	FALSE	FALSE
# 107			1012	54	6	0	0	0	FALSE	FALSE
# 108			1013	28	12	0	0	1	FALSE	FALSE
# 109			1014	22	2	1	0	0	FALSE	FALSE
# 110			1015	51	9	0	0	0	FALSE	FALSE
# 111			3367	12	1	0	0	0	FALSE	FALSE
# 112			1335	64	8	0	0	0	FALSE	FALSE
# 113	Chr1	15254	T							
# 114			1007	11	0	0	22	1	FALSE	FALSE
# 115			1012	28	0	0	53	1	FALSE	FALSE
# 116			1013	39	0	0	66	1	FALSE	FALSE
# 117			1014	3	0	0	14	1	FALSE	FALSE
# 118			1015	18	0	0	39	1	FALSE	FALSE
# 119			3367	0	0	0	89	0	FALSE	FALSE
# 120			1335	15	0	0	63	1	FALSE	FALSE
# 121	Chr1	15280	T							
# 122			1007	0	14	0	32	1	FALSE	FALSE
# 123			1012	0	31	0	53	1	FALSE	FALSE
# 124			1013	0	6	0	102	0	FALSE	FALSE
# 125			1014	0	3	1	40	0	FALSE	FALSE
# 126			1015	0	22	1	51	1	FALSE	FALSE
# 127			3367	0	0	0	74	0	FALSE	FALSE
# 128			1335	0	26	0	109	1	FALSE	FALSE
# 129	Chr1	16103	A							
# 130			1007	12	0	14	0	1	FALSE	FALSE
# 131			1012	50	0	19	0	1	FALSE	FALSE
# 132			1013	29	0	15	0	1	FALSE	FALSE
# 133			1014	28	0	2	0	0	FALSE	FALSE
# 134			1015	37	0	10	0	1	FALSE	FALSE
# 135			3367	41	0	0	0	0	FALSE	FALSE
# 136			1335	56	0	12	0	0	FALSE	FALSE
# 137	Chr1	17587	A							
# 138			1007	22	2	0	0	0	FALSE	FALSE
# 139			1012	62	6	0	1	0	FALSE	FALSE
# 140			1013	22	12	0	0	1	FALSE	FALSE
# 141			1014	22	2	0	0	0	FALSE	FALSE
# 142			1015	29	3	0	0	0	FALSE	FALSE
# 143			3367	20	1	0	0	0	FALSE	FALSE
# 144			1335	82	11	0	0	0	FALSE	FALSE
# 145	Chr1	18904	T							
# 146			1007	0	5	0	34	0	FALSE	FALSE
# 147			1012	0	4	0	39	0	FALSE	FALSE
# 148			1013	0	9	0	21	0	FALSE	FALSE
# 149			1014	0	3	0	21	0	FALSE	FALSE
# 150			1015	0	9	0	48	0	FALSE	FALSE
# 151			3367	0	5	0	96	0	FALSE	FALSE
# 152			1335	0	27	0	73	1	FALSE	FALSE
# 153	Chr1	25546	A							
# 154			1007	31	0	0	14	1	FALSE	FALSE
# 155			1012	64	0	0	22	1	FALSE	FALSE
# 156			1013	20	0	0	50	1	FALSE	FALSE
# 157			1014	22	0	1	18	1	FALSE	FALSE
# 158			1015	64	0	0	18	1	FALSE	FALSE
# 159			3367	73	0	0	0	0	FALSE	FALSE
# 160			1335	80	0	0	5	0	FALSE	FALSE

More sanity: there are 83 sites on Chr1 shared by zero strains in the tightest condition. (I.e., SAMTOOLS called

it a SNP, but the read counts/proportions fall below our 3rd threshold). Are they due to low coverage? Seemingly yes:

```

zp3 <- snp.pattern[[3]] == 0
zr3 <- rownames(non.refs[[3]])[which(zp3)]
zc3 <- as.integer(unlist(lapply(strsplit(zr3[1:min(100, length(zr3))], ':', fixed=TRUE), function(x){x[2]})))
zc3

#      [1]      91284      127986      161271      196862      196864      199166      282391      289344      289363      314132      314661
# [12]      438976      447253      475823      501830      501975      504462      652889      657955      692139      709443      762174
# [23]      826899      856950      875379      913014      938651      967184      1036942      1100300      1113225      1181146      1203203
# [34]     1210360     1212223     1224082     1270250     1270251     1348311     1431628     1473437     1516083     1526912     1628300
# [45]     1637082     1686331     1736789     1763837     1782580     1967158     2024930     2075603     2098145     2110716     2194162
# [56]     2242316     2258647     2261176     2325671     2376777     2432898     2441781     2498706     2550796     2554565     2581374
# [67]     2614631     2619528     2659281     2675254     2691279     2703771     2737914     2744068     2802553     2842231     2846930
# [78]     2906880     2931365     2948653     2957936     3014028     3016252           31184      101081      109502      195069      198570
# [89]      208189      278516      292413      297200      320853      349833      357243      357245      403824      418951      422130
# [100]      459508

seecounts(zc3[1:5], snp.tables=snp.tables)

#      chr      pos Ref Strain      A G      C T SNP      exon indel nrf rat
# 1  Chr1      91284      T
# 2              1007      0 0      0 17      0 FALSE FALSE
# 3              1012      0 0      0 38      0 FALSE FALSE
# 4              1013      2 0      0 13      0 FALSE FALSE
# 5              1014      0 0      0 20      0 FALSE FALSE
# 6              1015      0 0      0 35      0 FALSE FALSE
# 7              3367      3 0      0 12      1 FALSE FALSE
# 8              1335      0 0      0 47      0 FALSE FALSE
# 9  Chr1     127986      A
# 10             1007     47 0      0 0      0 TRUE FALSE
# 11             1012     92 0      0 0      0 TRUE FALSE
# 12             1013     19 1      0 0      0 TRUE FALSE
# 13             1014     73 0      0 0      0 TRUE FALSE
# 14             1015     83 0      0 0      0 TRUE FALSE
# 15             3367     13 3      0 0      1 TRUE FALSE
# 16             1335    160 0      0 0      0 TRUE FALSE
# 17  Chr1     161271      A
# 18             1007     31 0      0 0      0 TRUE FALSE
# 19             1012     47 0      0 0      0 TRUE FALSE
# 20             1013     18 3      0 0      0 TRUE FALSE
# 21             1014     30 0      0 0      0 TRUE FALSE
# 22             1015     59 0      0 0      0 TRUE FALSE
# 23             3367      8 3      0 0      1 TRUE FALSE
# 24             1335    102 0      0 0      0 TRUE FALSE
# 25  Chr1     196862      C
# 26             1007      0 0     10 0      0 FALSE FALSE
# 27             1012      0 0     22 0      0 FALSE FALSE
# 28             1013      0 0      8 2      0 FALSE FALSE
# 29             1014      0 0     14 0      0 FALSE FALSE
# 30             1015      0 0     18 0      0 FALSE FALSE
# 31             3367      1 0      4 3      1 FALSE FALSE
# 32             1335      0 0     18 0      0 FALSE FALSE
# 33  Chr1     196864      T
# 34             1007      0 0      0 11      0 FALSE FALSE
# 35             1012      0 0      1 23      0 FALSE FALSE
# 36             1013      3 0      0 8      1 FALSE FALSE
# 37             1014      0 0      0 12      0 FALSE FALSE
# 38             1015      1 0      1 19      0 FALSE FALSE
# 39             3367      3 0      0 4      1 FALSE FALSE
# 40             1335      0 0      1 19      0 FALSE FALSE

```

7.3 Main Analysis

Turning to the main analysis, there is a large increase in the number of consistent positions between the loose and medium stringency levels; medium and tight are similar in most respects. The likely interpretation is that the loose

criterion is including many “SNPs” induced by read errors, and that either of the tighter criteria are successfully filtering them out. In the interest of simplicity, the narrative below will focus on the shared SNPs at the medium stringency level (the “count2” column in the data frame), although the numbers for all three (sometimes all 4) are displayed. Also note that the prose and some comments in the code were based on the Chr1 analysis, and so may occasionally be off-target for the whole-genome data.

```
# Show a subset of pat.summaries, optionally with totals of count_i in last row, and optionally
# aggregating low-count rows as ``Other''
#
# sharedBy=c(2,4) selects SNPs shared by 2 or 4 strains,
# subset=as.octmode('35') select those with sharing pattern a subset (optionally proper) of this
# split=as.octmode('14') additionally restricts to patterns stradling split/subset minus split
# c2.thresh=42 suppresses printout of rows with count2 < 42
# restrict.to=c(0,42,127) restrict to these 3 rows
showgroup <- function(p.summ=pat.summaries, sharedBy=0:7, subset=127, split=NULL, proper.subset=F,
                      total=T, c2.thresh=0, fourteenth=F, restrict.to=NULL){
  # pick just those bit patterns that are subsets of 'subset'
  pick <- bitwAnd(0:127,bitwNot(subset))==0
  if(proper.subset){
    pick[subset+1] <- F
  }
  if(!is.null(split)){ # AND that stradle left/right subtrees
    cosplit <- bitwAnd(subset,bitwNot(split))
    pick <- pick & bitwAnd(0:127,split)!=0 & bitwAnd(0:127,cosplit)!=0
  }
  # and have desired shareBy counts
  pick <- pick & (p.summ$sharedBy %in% sharedBy)
  # and are among the set of interest
  if(!is.null(restrict.to)){
    pick <- pick & (0:127 %in% restrict.to)
  }
  # find rows with low counts
  pick.low <- pick & (p.summ$count2 < c2.thresh)
  # now show them
  show <- p.summ[pick & ! pick.low,]
  # rename columns just to narrow the printouts
  colnames(show) <- c('Pat','ShrBy','1007', '1012', '1013', '1014', '1015', '3367', '1335',
                     'count1', 'count2', 'count3','count4')
  show[,1] <- format(show[,1]) # convert octal col to char so can override in last row(2)
  nlow <- sum(pick.low)
  if(nlow > 0){
    n <- nrow(show)+1
    lows <- apply(p.summ[pick.low,10:13],2,sum)
    show[n,10:13] <- lows
    show[n,1:9] <- ''
    row.names(show)[n] <- 'Other'
    if(fourteenth){
      # do this: add 14th col just to hold this comment:
      show <- cbind(show, ' '=, stringsAsFactors=F)
      show[n,14] <- paste('(', nlow, 'rows w/ c2 < ', c2.thresh, ')')
    } else {
      ## or this (looks a bit funky, but fits across page without line-wrap):
      show[n,1:8] <- c('(', nlow, 'rows', 'w/', 'c2', '<', c2.thresh, ')')
    }
  }
  if(total){
    n <- nrow(show)+1
    tots <- apply(show[,10:13],2,sum)
    show[n,10:13] <- tots
    show[n,1:9] <- ''
    row.names(show)[n] <- 'Total'
    if(ncol(show)==14){show[n,14]<-''}
  }
  return(show)
}
```

First, are there any SNPs that are not “consistent SNPs?” Yes, a few in c3. As noted above, they seem to be mainly

low-coverage positions.

```
showgroup(pat.summaries,0,total=F) # chr1 totals: 0 0 83
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 1	0	0								0	2	1164	0

Next, look at completely shared SNPs, those found in all 7 strains.

```
showgroup(pat.summaries,7,total=F) # Chr1 count1 = 8593, count2 = 7054, count3 = 4790 c4=1641
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 128	177	7	X	X	X	X	X	X	X	82193	67223	46524	15186

I.e., of the 468117 consistent positions, 67223 or 14.4% are shared by all 7 strains.

Next look at singletons, aka private SNPs—SNPs that are called in one strain and no other strain has a significant number of non-ref reads at that position. Presumably these are variants that arose in a given population after it separated from the others.

```
showgroup(pat.summaries,1) # chr1 totals: 9669 18865 19670 23574
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 2	001	1							X	199	620	1157	2260
# 3	002	1						X		41774	84335	88149	105614
# 5	004	1					X			921	2070	2578	4608
# 9	010	1								208	559	714	1231
# 17	020	1				X				47772	93481	96798	113191
# 33	040	1		X						285	611	1031	2450
# 65	100	1	X							121	321	542	2005
# Total										91280	181997	190969	231359

The import of shared/private SNPs changes between sexual and asexual populations. Presumably asexuals slowly gain and rarely lose private SNPs; shared ones predate separation of the lineages. In sexual lineages, however, SNPs may be rather freely “gained” or “lost,” merely by recombination (converting between homo- and heterozygous in the sample we sequenced). Thus, the low private counts for the 5 L-isolates compared to the large count of het positions overall suggest that (a) they are asexual, and (b) none of them has been isolated from the others for very long (if at all). Conversely, the high counts for Italy and Wales suggest that (a) if asexual, they have been separated from each other and from the rest for a long time, but (b) if sexual, there is little surprise: we have ≈ 160 K SNPs shared between the two (90K just in those two (below), plus 70K shared by all 7), and ≈ 90 K additional positions that are het in one but not the other. These are close to, but not exactly equal to, the 1:2:1 ratios we would naively expect from two samples of a single HWE population. The most parsimonious explanation seems to be that the H-clade is sexual, but perhaps some het positions private to each population separates them.

Aside: counts of “consistent” SNPs minus these singletons yields count of shared SNPs:

```
singlets <- apply(pat.summaries[pat.summaries$sharedBy==1,10:13],2,sum)
rbind(consistent=consistent.count,singlets=singlets,shared=consistent.count-singlets)
```

#	count1	count2	count3	count4
# consistent	358872	468117	470970	474613
# singlets	91280	181997	190969	231359
# shared	267592	286120	280001	243254

The slightly higher count of shared positions in the medium case further supports this choice for subsequent analysis.

Next look at consistent SNPs shared between just a pair of isolates.

```
showgroup(pat.summaries,2) # chr 1 counts: 7641 9549 9472 6924
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 4	003	2						X	X	2266	281	432	587
# 6	005	2					X		X	210	410	854	1407
# 7	006	2					X	X		1282	119	261	590

# 10	011	2				X			X	358	535	384	827
# 11	012	2				X		X		2060	158	59	93
# 13	014	2				X	X			143	136	176	417
# 18	021	2			X				X	2445	154	260	402
# 19	022	2			X			X		55300	87584	84944	58009
# 21	024	2			X		X			1406	178	371	625
# 25	030	2			X	X				2257	180	59	93
# 34	041	2		X					X	31	85	230	368
# 35	042	2		X				X		1429	117	224	394
# 37	044	2		X			X			72	215	895	1809
# 41	050	2		X		X				31	25	54	105
# 49	060	2		X	X					1651	98	247	388
# 66	101	2	X						X	19	33	75	314
# 67	102	2	X					X		887	98	126	351
# 69	104	2	X				X			39	105	356	1196
# 73	110	2	X			X				19	20	50	150
# 81	120	2	X		X					1009	117	116	309
# 97	140	2	X	X						591	1150	1281	2144
# Total										73505	91798	91454	70578

I.e., of the 91798 paired SNPs, 87584 or 95.4% are found between Italy and Wales, with comparatively few shared between any other pairs (only).

SNPs shared among exactly 3 isolates are relatively rare. (The 5 trios containing both Italy and Wales predominate in the loose set, probably because they share many pairs that become triples with the addition of a few read errors.)

showgroup(pat.summaries,3) # chr 1 counts: 1438 294 671 1034													
#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 8	007	3					X	X	X	152	146	278	557
# 12	013	3					X	X	X	350	253	183	338
# 14	015	3					X	X	X	776	1050	757	1389
# 15	016	3					X	X	X	109	62	65	152
# 20	023	3			X				X	3322	481	838	533
# 22	025	3			X		X		X	197	168	333	522
# 23	026	3			X		X	X		1794	395	771	789
# 26	031	3			X	X			X	361	255	178	361
# 27	032	3			X	X		X		2845	237	112	86
# 29	034	3			X	X	X			144	116	125	219
# 36	043	3		X					X	175	70	104	133
# 38	045	3		X			X		X	116	409	1369	1656
# 39	046	3		X			X	X		103	124	367	604
# 42	051	3		X		X			X	46	55	79	126
# 43	052	3		X		X		X		101	11	27	22
# 45	054	3		X		X	X			22	79	171	292
# 50	061	3		X	X				X	169	67	98	115
# 51	062	3		X	X			X		2161	289	447	469
# 53	064	3		X	X		X			114	148	390	601
# 57	070	3		X	X	X				123	21	18	24
# 68	103	3	X					X	X	83	25	51	143
# 70	105	3	X				X		X	35	113	283	805
# 71	106	3	X				X	X		42	63	127	377
# 74	111	3	X			X			X	16	9	19	139
# 75	112	3	X			X		X		55	12	16	26
# 77	114	3	X			X	X			8	38	69	365
# 82	121	3	X		X				X	87	19	29	73
# 83	122	3	X		X			X		1325	191	220	354
# 85	124	3	X		X		X			66	81	167	400
# 89	130	3	X		X	X				73	12	9	27
# 98	141	3	X	X					X	65	109	208	519
# 99	142	3	X	X				X		342	419	477	755
# 101	144	3	X	X			X			193	1079	2430	4432
# 105	150	3	X	X		X				47	44	78	238
# 113	160	3	X	X	X					306	359	421	712
# Total										15923	7009	11314	18353

Four-way sharing is more common, but dominated by the coastal (i.e., non-Gyre) L-clade isolates. This is likely a

reflection of the strong 5-way sharing among the L-clade, from which the Gyre commonly drops out due to the lower coverage/higher error rate in that sequencing run.

```
showgroup(pat.summaries,4) # chr 1 counts: 564 1346 2552 3479
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 16	017	4				X	X	X	X	344	361	250	564
# 24	027	4			X		X	X	X	441	602	1120	771
# 28	033	4			X	X		X	X	1174	1062	725	306
# 30	035	4			X	X	X		X	495	529	373	503
# 31	036	4			X	X	X	X		381	367	287	211
# 40	047	4		X			X	X	X	93	178	485	708
# 44	053	4		X		X		X	X	46	38	36	56
# 46	055	4		X		X	X		X	480	709	750	971
# 47	056	4		X		X	X	X		17	50	65	88
# 52	063	4		X	X			X	X	325	167	265	194
# 54	065	4		X	X		X		X	88	208	528	582
# 55	066	4		X	X		X	X		263	432	944	851
# 58	071	4		X	X	X			X	30	17	14	28
# 59	072	4		X	X	X		X		158	50	35	31
# 61	074	4		X	X	X	X			28	51	60	116
# 72	107	4	X				X	X	X	36	64	105	330
# 76	113	4	X			X		X	X	18	8	8	66
# 78	115	4	X			X	X		X	103	141	138	604
# 79	116	4	X			X	X	X		5	8	19	101
# 84	123	4	X		X			X	X	162	85	77	124
# 86	125	4	X		X		X		X	41	79	142	283
# 87	126	4	X		X		X	X		124	214	297	425
# 90	131	4	X		X	X			X	17	5	5	52
# 91	132	4	X		X	X		X		86	27	19	38
# 93	134	4	X		X	X	X			15	24	25	143
# 100	143	4	X	X				X	X	46	55	86	190
# 102	145	4	X	X			X		X	3312	9777	18140	23189
# 103	146	4	X	X			X	X		121	455	958	1795
# 106	151	4	X	X		X			X	33	77	71	220
# 107	152	4	X	X		X		X		27	21	24	67
# 109	154	4	X	X		X	X			896	1611	1430	1738
# 114	161	4	X	X	X				X	72	88	102	207
# 115	162	4	X	X	X			X		1552	1955	1909	1014
# 117	164	4	X	X	X		X			165	603	1060	1752
# 121	170	4	X	X	X	X				21	26	21	69
# Total										11215	20144	30573	38387

Five-way sharing is much more common, and is strongly dominated by the 5 L-clade isolates.

```
showgroup(pat.summaries,5) # chr 1 counts: 3969 5047 4624 6125
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 32	037	5			X	X	X	X	X	2087	1987	1386	620
# 48	057	5		X		X	X	X	X	227	231	205	324
# 56	067	5		X	X		X	X	X	422	685	1836	1151
# 60	073	5		X	X	X		X	X	155	89	72	38
# 62	075	5		X	X	X	X		X	200	185	233	328
# 63	076	5		X	X	X	X	X		108	158	188	128
# 80	117	5	X			X	X	X	X	42	40	35	241
# 88	127	5	X		X		X	X	X	133	253	399	482
# 92	133	5	X		X	X		X	X	56	31	16	52
# 94	135	5	X		X	X	X		X	116	121	96	235
# 95	136	5	X		X	X	X	X		41	71	63	106
# 104	147	5	X	X			X	X	X	1372	3155	5536	10001
# 108	153	5	X	X		X		X	X	35	32	30	96
# 110	155	5	X	X		X	X		X	33045	38232	26997	30602
# 111	156	5	X	X		X	X	X		492	681	566	735
# 116	163	5	X	X	X			X	X	271	263	302	316
# 118	165	5	X	X	X		X		X	1875	3928	6825	9715
# 119	166	5	X	X	X		X	X		621	1958	3252	2688
# 122	171	5	X	X	X	X			X	30	29	18	70

# 123	172	5	X	X	X	X		X	105	95	59	86
# 125	174	5	X	X	X	X	X		567	789	656	782
# Total									42000	53013	48770	58796

Six-way sharing is also common, with the sets *excluding* Gyre, Italy, or Wales having the most mutually-shared SNPs.

```
showgroup(pat.summaries,6) # chr 1 counts: 4166 4741 5312 4722
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 64	077	6		X	X	X	X	X	X	920	872	917	485
# 96	137	6	X		X	X	X	X	X	394	325	275	333
# 112	157	6	X	X			X	X	X	13257	11725	8245	12202
# 120	167	6	X	X	X		X	X	X	8614	16443	28402	15091
# 124	173	6	X	X	X	X		X	X	140	93	74	114
# 126	175	6	X	X	X	X	X		X	17128	14648	10156	12697
# 127	176	6	X	X	X	X	X	X		2303	2825	2133	1032
# Total										42756	46931	50202	41954

8 Trees

So, overall, the picture looks like a long shared history (67223 7-way shared positions), followed by a split of the 5 L-isolates from the 2 H-isolates, then a long shared history in the 5 (38232 quintuples), in parallel with a long shared history in H- (87584 pairs), then separate histories in Italy and Wales (>84335 “private” SNPs in each, although again if they are sexual, many of these just reflect HWE), and very limited differentiation among the 5 L-isolates.

Branch lengths of course depend on filtering criteria used (and, of course, full vs Chr1 differ by about a factor of 10), but the tree *topology* appears to be fairly stable. Various versions are drawn below, exactly to explore how robust this story is. I think we should go with “medium stringency” SNP filtering (based on un-qfiltered reads).

NOTE: Much of this analysis make less sense for q-filtered read data, since (a) the point of the SNP filtering was to try to correct for noise in the raw reads, which may (or may not; haven’t looked closely, yet) be largely fixed by qfiltering (e.g., “loose” or no SNP filtering may be more appropriate, post-q-filtering, esp. if we had re-run SAMTools to call SNPs based on the q-filtered reads), and (b) tree topology *does* appear to change, in that Gyre’s coverage has been so sharply reduced by qfiltering that it clearly stands aside from the others (and that’s confirmed by bootstrap), but this also seems to be clearly a technical rather than a biological artifact. SO, code below will run on q-filtered data, but *is not tuned to it*. Likewise, most comments in the prose below were made to describe the un-q-filtered data, and *are misleading and in some cases flatly wrong* for qfiltered data, but it doesn’t seem worthwhile to bother with a rewrite...

Trees are coded in newick format, which doesn’t seem to tolerate line-breaks; print with line-wrap:.

```
# wrap a long char string across multiple lines in printout
cat.hardwrap <- function(str,width=80){
  while(nchar(str)>width){
    cat(substr(str,1,width),'\n')
    str <- substr(str,width+1,nchar(str))
  }
  cat(str,'\n')
}
```

Trees are built as follows. Code for drawing, especially, is specific to the topology of the medium tree, and placement of some of the figure elements have been hand-optimized for this case; drawings for the other variants will not be as pretty.

```
# set up for tree figs

# the newick parser in ape seems to be confused by commas and parens in
# tip names, and blanks are not allowed, so replace by *, <, >, _, resp.
newick.name <- function(name){
  name <- gsub(' ', '_', name, fixed=TRUE)
  name <- gsub(',', '*', name, fixed=TRUE)
  name <- gsub('(', '<', name, fixed=TRUE)
```

```

name <- gsub(')', '>', name, fixed=TRUE)
return(name)
}
# undo above changes
newick.name.undo <- function(name){
#name <- gsub('_', ' ', name, fixed=TRUE) # unnecessary; ape plot routine handles this one
name <- gsub('*', ' ', name, fixed=TRUE)
name <- gsub('<', '(', name, fixed=TRUE)
name <- gsub('>', ')', name, fixed=TRUE)
return(name)
}

# make a newick string from tree; see it below
# 'pre' is prefixed to ccmpid; 'nb' optionally included;
# 'alt' can be used instead of pre/ccmp/nb/where for less formal labeling
# 'newstyle'=T => new node label: [nb_]where[(pre-less-id)]
# 'newstyle'=F => old node label: [nb_][pre id]where
newickize <- function(tree,pre='CCMP',nb=TRUE,alt=F,newstyle=TRUE){
  if(is.null(tree$where)){
    # not a leaf; paste together newick from subtrees
    sub1 <- newickize(tree$sub1,pre=pre,nb=nb,alt=alt,newstyle=newstyle)
    sub2 <- newickize(tree$sub2,pre=pre,nb=nb,alt=alt,newstyle=newstyle)
    new <- paste('(', sub1, ',', sub2, ')', sep='')
    if(!is.null(tree$length)){
      # internal node, add length
      return(paste(new, ':', tree$length, sep=''))
    } else {
      # top level; escape blanks and add trailing ';'
      return(paste(gsub(' ', '_', new), ';', sep=''))
    }
  } else {
    # a leaf; build label and branch length
    if(alt){
      # label is just alt; if alt omitted, default to where
      new <- newick.name(ifelse(is.null(tree$alt), tree$where, tree$alt))
    } else {
      if(newstyle){
        # new node label = [nb_]where[(pre-less-id)]
        new <- ifelse(nb && !is.null(tree$nb), paste(tree$nb, '_', sep=''), '')
        new <- newick.name(paste(new, tree$where, sep=''))
        new <- ifelse(is.null(tree$id), new, paste(new, '_', tree$id, ')', sep=''))
        new <- newick.name(new)
      } else {
        # old style node label = [nb_][pre id]where
        new <- ifelse(nb && !is.null(tree$nb), paste(tree$nb, '_', sep=''), '')
        new <- ifelse(is.null(tree$id), new, paste(new, pre, tree$id, '_', sep=''))
        new <- newick.name(paste(new, tree$where, sep=''))
      }
    }
    #add length to either
    new <- paste(new, ':', tree$length, sep='')
  }
  return(new)
}

# Make a tree as nested lists, **based on the chr1, count2 topology**, but using any of the counts.
# Internal nodes have subtrees sub1/2 and length
# Root has sub1/2, but no length
# Leaves have where, length, optionally, id, alt, nb. (Omit id for 'outgroup'. Use 'alt' for less formal
# labeling in cartoon version; it defaults to 'where'. Use 'nb' to add abcde annotations for legend.)
# The single parameter v is any of the 4 count vectors contained in pat.summaries (most conveniently
# indexed in octal). E.g., make.tree(pat.summaries[, 'count2']) reproduces the count2 tree.
# (This was previously built by hand-pasting the edge lengths; tree.by.hand is retained in appendix
# for comparison, & its counts are in comments below).
#
make.tree <- function(v){
  pat.count <- function(pat, pat.counts=v){return(pat.counts[1+strtoi(pat,8)])}
  thetree <-
    list(
      sub1 = list(
        sub1 = list(
          sub1 = list(id=3367, length=pat.count('002'), where='Venice, Italy', alt='Venice'), #8813
          sub2 = list(id=1013, length=pat.count('020'), where='Wales, UK'), #9652
          length=pat.count('022')), #9365
        sub2 = list(
          sub1 = list(
            sub1 = list(
              sub1 = list(id=1007, length=pat.count('100'), nb='e', where='Virginia, USA'), #30
              sub2 = list(id=1012, length=pat.count('040'), nb='d', where='Perth, W. Australia', alt='Perth'), #61
              length=pat.count('140')), #19
            )
          )
        )
    )

```

```

    sub2 = list(
      sub1 = list(id=1015, length=pat.count('004'), nb='c', where='Washington, USA', alt='Puget Sound'), #207
      sub2 = list(id=1335, length=pat.count('001'), nb='b', where='New York, USA', alt='NY'), #41
      length=pat.count('005'), #18
      length=pat.count('145'), #1005
      sub2 = list(id=1014, length=pat.count('010'), nb='a', where='N. Pacific Gyre'), #61
      length=pat.count('155'), #3912
      length=pat.count('177'), #7054
      sub2 = list(length=0, where='outgroup')
    )
  return(thetree)
}

```

Code to plot a tree given newick description. Again, code is somewhat general, but has some specializations tied to the medium-stringency, full-genome, un-filtered data.

```

# run following 2 lines after an R upgrade
# update.packages()
# install.packages("ape")
library(ape)
show.tree <- function(newick.str=newick.medium,
  col.edge = 'darkblue', lwd.edge = 2,
  col.elabel='darkblue', cex.elabel=0.8, font.elabel=3,
  col.arrow = 'red', lwd.arrow=1.5, cex.arrow = 0.9, font.arrow = 4,
  col.clade = 'black', lwd.clade=1, cex.clade = 1.0, font.clade = 3,
  col.legend = 'beige', cex.legend=0.8,
  col.tip = 'darkblue', font.tip = 4,
  plusx=FALSE, pltdebug=FALSE, total.snps=consistent.count[2],
  straight.arrow=FALSE){

  ####
  #
  # ADJUST NEWICK & GET LENGTHS, COORDINATES
  #
  newick.str.noout <- sub('outgroup','_',newick.str) # Hide outgroup ('_' prints as blank)
  the.tree <- read.tree(text=newick.str.noout)

  ## nasty hack: ape's newick parser seems to be confused by commas, () in tip labels, so
  ## newickize replaced them by '*<>'; before plotting, I want to convert them back, and hope
  ## this doesn't break anything else... And if a revised version of ape changes the internal
  ## representation of a tree, this may need to be redone.
  the.tree$tip.label <- newick.name.undo(the.tree$tip.label)

  # extract branch lengths as char string of comma-separated numbers via pattern matching hack:
  # lengths always preceded by colon
  lengths.ch <- strsplit(paste(newick.str, ':'), '^[^0-9][^:]*:')[[1]]

  # then convert to ints, dropping empty string at front
  lengths.int <- scan(what=integer(), quiet=T, sep=',', text=lengths.ch[-1])

  # then to data frame with named rows; a.g are terminal branches; others are internal.
  # a.e match legend in plot; f/g = wales/italy. lengths appear in postfix order of
  # newick tree, and ape draws the 1st of them at the bottom of the plot.
  lmed <- data.frame(lengths=lengths.int,
    row.names=c('g', 'f', 'fg', 'e', 'd', 'de', 'c', 'b', 'bc', 'bcde', 'a', 'abcde', 'all', 'out'))

  # extract counts needed for legend:
  #leg.counts <- c( 61, 41, 207, 61, 30, 1005, 18, 19) #by hand, medium chr1
  leg.counts <- lmed[c('a', 'b', 'c', 'd', 'e', 'bcde', 'bc', 'de'), 1]
  discord <- total.snps - sum(lmed$lengths)

  #tree.labels <- list( ## x,y,text; coords are all picked by eye
  # 3000, 3.62, paste(lmed['all', 1], 'shared by 7', sep='\n'), # 7054
  # 8900, 5.75, paste(lmed['abcde', 1], 'by 5', sep='\n'), # 3912
  # 12000, 1.50, paste(lmed['fg', 1], 'shared by 2', sep='\n'), # 9365
  # 21000, 2.00, paste(lmed['f', 1], 'only\nin Wales'), # 9652
  # 21000, 1.00, paste(lmed['g', 1], 'only\nin Italy'), # 8813
  # 11500, 4.50, '*')
  # automating x-placement, below; retain above for comparison...
  tip <- integer(7) # x coords of tree tips
  tip[1] <- sum(lmed[c('all', 'fg', 'g'), 1])
  tip[2] <- sum(lmed[c('all', 'fg', 'f'), 1])
  tip[3] <- sum(lmed[c('all', 'abcde', 'bcde', 'de', 'e'), 1])
  tip[4] <- sum(lmed[c('all', 'abcde', 'bcde', 'de', 'd'), 1])
  tip[5] <- sum(lmed[c('all', 'abcde', 'bcde', 'bc', 'c'), 1])
  tip[6] <- sum(lmed[c('all', 'abcde', 'bcde', 'bc', 'b'), 1])
  tip[7] <- sum(lmed[c('all', 'abcde', 'a'), 1])

  inode <- integer(5) # x coords of (some) internal nodes

```



```

inode[1] <- 0 # root
inode[2] <- lmed['all',1] # lca of all
inode[3] <- sum(lmed[c('all','fg'),1]) # lca H-clade
inode[4] <- sum(lmed[c('all','abcde'),1]) # lca L-clade
inode[5] <- sum(lmed[c('all','abcde','bcde'),1]) # lca L-clade, nonGyre
tree.labels <- list( ## x,y,text; y coords partially picked by eye
  sum(inode[c(1,2)])/2, 3.62, paste(lmed['all',1], 'shared by 7', sep='\n'), # 7054
  sum(inode[c(2,4)])/2, 5.75, paste(lmed['abcde',1], 'by 5', sep='\n'), # 3912
  sum(inode[c(2,3)])/2, 1.50, paste(lmed['fg',1], 'shared by 2', sep='\n'), # 9365
  (inode[3]+tip[2])/2, 2.00, paste(lmed['f',1], 'only\nin 1013'), # 9652
  (inode[3]+tip[1])/2, 1.00, paste(lmed['g',1], 'only\nin 3367'), # 8813
  sum(inode[c(4,5)])/2, 4.35, '*' )

tree.labels <- list( ## x,y,text; y coords partially picked by eye
  sum(inode[c(1,2)])/2, 3.62, paste(lmed['all',1], 'in 7', sep='\n'), # 7054
  sum(inode[c(2,4)])/2, 5.75, paste(lmed['abcde',1], 'in 5', sep='\n'), # 3912
  sum(inode[c(2,3)])/2, 1.50, paste(lmed['fg',1], 'in 2', sep='\n'), # 9365
  (inode[3]+tip[2])/2, 2.00, paste(lmed['f',1], 'only\nin 1013'), # 9652
  (inode[3]+tip[1])/2, 1.00, paste(lmed['g',1], 'only\nin 3367'), # 8813
  sum(inode[c(4,5)])/2, 4.35, '*' )

####
#
# BOGUS PLOT
#
# a messy bit: need string widths to set xlim; but strwidth needs x-scale so must plot first.
# M plot completely invisible, overlay 2nd plot via par(new=F...) .
#
# PROVISIONALLY set x.lim here at about 30% wider than tree; fine tune it for the real plot
# based on strwidth(tip labels) below.
#
provisional.tree.x.lim <- 1.3 * max(tip) # <== PROVISIONAL plot width
plot(0,0, type='n', bty='n', xaxt='n', yaxt='n', xlab='', ylab='', xlim=c(0,provisional.tree.x.lim), ylim=c(0,7))

tiplabel.x <- integer(7)
for(i in 1:7){
  # see warning above about internals of the tree; labels have '_', printed as ' '.
  tiplabel.x[i] <- tip[i]+strwidth(gsub('_', ' ', the.tree$tip.label[i], fixed=T), font=font.tip)
}

# visually show tip coords & max x to debug placement issues
plt.debug <- function(tree.x.lim, tip, tiplabel.x, spx=NULL, spy=NULL){
  if(plt.debug){ # F to hide/T to show debug
    cat('Tip labels:', paste(the.tree$tip.label, sep=', collapse='/'), '\n')
    axis(2) # useful only for placing labels
    for(i in 1:7){
      points(c(tip[i], tiplabel.x[i]), c(i,i)) # debug: do I have right tip coordinates?
    }
    lines(rep(tree.x.lim, 2), c(0,7)) # where is right edge?
    if(!is.null(spx)){
      points(spx, spy) # show spline control points, for tweaking
    }
  }
}

plt.debug(provisional.tree.x.lim, tip, tiplabel.x)

label.end.H <- max(tiplabel.x[1:2])
label.end.L <- max(tiplabel.x[3:7])
clade.dx <- strwidth('x') # space between clade marker line and its label
xdel <- 3*clade.dx # space between labeled clade tips and marker line

tree.x.lim <- 1.03*(max(tiplabel.x)+xdel) # <== FINAL plot width
tree.y.lim <- 7
if(plt.debug){cat('Plot width hacking:', provisional.tree.x.lim, tree.x.lim, tree.x.lim/1.03/max(tip), clade.dx)}

par(new=T) # I.e., NOT starting a new plot

####
#
# REAL PLOT
#
plot(the.tree,
  x.lim = c(0, tree.x.lim),
  y.lim = c(0, tree.y.lim),
  font=font.tip, label.offset=100, # bold-italic, nudged slightly right
  tip.color=col.tip, edge.color=col.edge,
  edge.width=lwd.edge,

```

```

    edge.lty=c(1,1,1,1, 1 ,1,1,1,1,1,1,1,0)      # 5th is bottleneck edge; 14th is outgroup
  )
  lines(00+c(0,0),c(3.5,6),col='white',lwd=6)      # Hide vertical line to outgroup
  axis(1, pos=0.25, at=seq(0,25,by=5)*10^round(log10(max(tip)/25)))

  if(pltdebug){text(tip[1]+100, 1.0, 'Venice, Italy (3367)', adj=0, font=font.tip)}

  ####
  #
  # BOTTLENECK ANNOTATION
  #
  # spline/ellipse control points (spy/y) & tweaks thereto (dx/y)
  dx <- 0.01 * tree.x.lim
  dy <- .04
  spx <- c(7400, 7400, 9900, 10500) # by eye, chr1, for comparison
  spx <- c(inode[2]+dx,inode[2]+dx,inode[4]-3*dx,inode[4]-dx)
  spy <- c( 3.8,  3.9,  5.6-dy,  5.6-dy)

  plt.debug(tree.x.lim, tip, tiplabel.x, spx, spy)

  if(T){
    #ellipse version, defined by rect thru 2 middle pts of spx/y
    spf<-function(x){
      ifelse(x <= spx[2], spy[1],
             ifelse(x >= spx[3], spy[4],
                    spy[2]+(spy[3]-spy[2])*sqrt(pmax(0,1-((x-spx[3])/(spx[3]-spx[2]))^2))))
    }
  } else {
    # spline version
    spf <- splinefun(spx,spy,method='hyman')
  }
  serx <- seq(spx[1],spx[length(spx)],length.out=50)
  sery <- spf(serx)
  tailx <- spx[1]
  taily <- spy[1]
  headx <- spx[4]
  heady <- spy[4]
  textx <- (headx+tailx)/2+(headx-tailx)*(-.01)
  texty <- (heady+taily)/2+(heady-taily)*(-.10)
  bottle.txt <- "inbreeding\nLoH / LoS"
  if(!straight.arrow){
    arrows(headx,heady,headx+tree.x.lim*1e-3,heady, length=.1,col=col.arrow,lwd=lwd.arrow)
    lines(rev(serx), rev(sery), lty=c(5,1),col=col.arrow, lwd=lwd.arrow)
    textangle <- 66
    textadj <- c(0,0)
  } else {
    # Tweak positioning slightly; visualize a rectangle from 7-node to base of L-clade;
    # center text, rotated, on diagonal towards L-clade; ditto the straight arrow.
    llx <- inode[2] # the aforementioned rectangle
    urx <- inode[4]
    lly <- 3.62
    ury <- 5.75
    # rect(llx,lly,urx,ury) # show rect for debug
    textx <- (llx+urx)/2 # center text
    texty <- (lly+ury)/2
    textangle <- atan(grconvertY(ury-lly,to='dev')/grconvertX(urx-llx,to='dev'))*360/(2*pi)
    textadj <- c(0.50, 0.43) #tweak position; ".5" = center in x , ".43" raises, THEN rotate.
    alpha <- .78 # fraction along diag at which arrow begins
    beta <- .95 # ... and ends
    arrows((1-alpha)*llx + alpha*urx,
           (1-alpha)*lly + alpha*ury,
           (1-beta)*llx + beta*urx,
           (1-beta)*lly + beta*ury, length=.1,col=col.arrow,lwd=lwd.arrow,angle=25)
  }
  if(T){
    text(textx, texty, bottle.txt, srt=textangle, font=font.arrow, cex=cex.arrow,
         col=col.arrow, adj=textadj)
  } else {
    # experiment at wrapping text along curved path; unpretty, but retain for now, maybe revisit
    bottlec <- strsplit(bottle,split=NULL)[[1]]
    for(i in 1:length(bottlec)){
      text(xser[i],yser[i],bottlec[i], srt=65, font=4, cex=.7, col=col.arrow)
    }
  }

  ####
  #
  # CLADE ANNOTATION
  #
  clade.L.x <- label.end.L + xdel

```

```

clade.H.x <- label.end.H + xdel
dy <- .33
lines(rep(clade.L.x,2),c(3-dy,7+dy),lwd=lwd.clade,col=col.clade)
lines(rep(clade.H.x,2),c(1-dy,2+dy),lwd=lwd.clade,col=col.clade)
text(clade.L.x+clade.dx,5.0,'L-clade',srt=90,font=font.clade,cex=cex.clade,col=col.clade)
text(clade.H.x+clade.dx,1.5,'H-clade',srt=90,font=font.clade,cex=cex.clade,col=col.clade)

####
#
# LEGEND
#
# parameter plusx controls whether we try to annotate b/c (+) and d/e (x) sharing in tree; I think
# it looks cluttered, rather than adding clarity, so I vote no, but code is here, in case. "Logic,"
# if any, for my symbol choice is that + overlaid on x looks like the * at the next level; this
# analogy is more visible if we use pch 3/4/8 rather than Courier or Helvetica chars, but probably
# should use same in both tree & legend, which will take a modicum of additional work.
legend.text <- c('a: only in 1014 ',
                'b: only in 1335 ',
                'c: only in 1015 ',
                'd: only in 1012 ',
                'e: only in 1007 ',
                'x: shared by bcde',
                paste(ifelse(plusx,'+:',' '), 'shared by b/c '),
                paste(ifelse(plusx,'x:',' '), 'shared by d/e ')
)

legend.text <- c('a: only in 1014 ',
                'b: only in 1335 ',
                'c: only in 1015 ',
                'd: only in 1012 ',
                'e: only in 1007 ',
                'x: in bcde ',
                paste(ifelse(plusx,'+:',' '), 'in bc '),
                paste(ifelse(plusx,'x:',' '), 'in de '),
                'Discordant SNPs '
)

legend.text <- paste(legend.text,format(c(leg.counts,discord),width=4),sep=' - ')
legend.text <- paste(legend.text,' ') # add a little more right margin in box
opar <- par(family='mono',cex=cex.legend)
legend('topright', legend=legend.text, cex=cex.legend, inset=c(0.05,0), bg=col.legbox, box.col=col.legbox)
par(opar)
if(plusx){
  points(tree.labels[[16]],tree.labels[[17]]+.14,pch=8,col=col.elabel)
  points(tree.labels[[16]]+200,tree.labels[[17]]+1,pch=3,col=col.elabel)
  points(tree.labels[[16]]+200,tree.labels[[17]]-1,pch=4,col=col.elabel)
}

####
#
# EDGE LENGTHS
#
for(i in seq(1,length(tree.labels)-ifelse(plusx,5,2),by=3)){
  if(F){ # T for \n in edge labels; F to remove (except "by 5")
    text(tree.labels[[i]], tree.labels[[i+1]], tree.labels[[i+2]])
  } else {
    # points(tree.labels[[i]], tree.labels[[i+1]], pch=3,col='green') # for debugging
    text(tree.labels[[i]], tree.labels[[i+1]], sub('\n([~z])',' \\1', tree.labels[[i+2]]),
         pos=3, offset=.4, font=font.elabel, col=col.elabel,cex=cex.elabel)
  }
}
}
if(FALSE){#for debug convenience
pdf(paperfig.path, width=8,height=5,onefile=TRUE,family='Helvetica',fonts='Courier',pointsize=10)
show.tree(newick.medium, total.snps=consistent.count[2], pltdebug=F, straight.arrow=T)
dev.off()
}

```

```

caption <- function(stringency,which.tables=which.snp.tables(string.val=F)){
  caption.where <- '(UNKNOWN genome subset).'
  if(which.tables[1]=='Chr1') {caption.where <- 'on Chr1.'}
  if(which.tables[1]=='full') {caption.where <- 'genome-wide.'}
  if(which.tables[1]=='trunc'){caption.where <- 'all Chrs.'}
  cap.stringency <- c(
    'loose SNP filters.',
    'medium SNP filters.',
    'strict SNP filters.',
    'unfiltered SNPs.')
}

```

```
cap <- paste('Tree based on', which.tables[2], 'reads and', cap.stringency[stringency],
            '``Lengths`` are numbers of shared/private SNPs', caption.where)
return(cap)
}
```

Trees based on all four SNP filtering criteria are shown below. Their topologies are exactly the same, although the branch lengths are different. In all four, the length of the branch labeled “*” is probably inflated by lower coverage and higher error rate in 1014, which may mask further legitimate sharing between it and the other L-isolates. The branch lengths among the other 4 are too short for their topology to be convincing without a more rigorous analysis (e.g., a bootstrap test), but detail there is irrelevant to the story.

My sense is that the “medium” version is the best for the paper, made here and shown in Fig 1. In theory, this should look exactly like Fig 3, but something is apparently different between Knitr and direct-to-pdf. (Increasing fig.width in Knitr’s chunk headers from 8 (as in the pdf call below) to 9 helps somewhat, but probably still best to make the paper fig directly rather than via Knitr.)

```
###
#
# MAKE PROTOTYPE PDF FOR PAPER, *AND* SAVE DATA NEEDED TO BUILD IT
#
w.s.t. <- which.snp.tables()
if(w.s.t. == 'trunc-unfiltered'){
  rda.Description <- 'This .rda contains data to generate Fig 3; see shared.snps.rnw for details.'
  save(rda.Description, w.s.t., pat.summaries, consistent.count, file='Fig3-data.rda')
  paperfig.path <- paste('figs-mine/paperfig-medium-tree-', w.s.t., '--Fig3proto.pdf', sep='')
} else {
  paperfig.path <- paste('figs-mine/paperfig-medium-tree-', w.s.t., '.pdf', sep='')
}
pdf(paperfig.path, width=8,height=5,onefile=TRUE,family='Helvetica',fonts='Courier',pointsize=10)
newick.medium <- newickize(make.tree(pat.summaries[, 'count2']))
show.tree(newick.medium, total.snps=consistent.count[2], pltdebug=F, straight.arrow=T)
dev.off()

# pdf
# 2
```

```
# fig.paths for knitr chunks below; .h for "hand-made" trees; plain for automatic chrl/full versions
myfigpath <- paste(getwd(), '/figs-knitr/newick-', which.snp.tables(), '-', sep='')
myfigpath.h <- paste(getwd(), '/figs-knitr/newick-', sep='')
```

Figure 2, i.e., criteria [[1]]:

```
newick.loose <- newickize(make.tree(pat.summaries[, 'count1']))
show.tree(newick.loose, total.snps=consistent.count[1])
```

Figure 3, i.e. [[2]]:

```
# newick.medium <- newickize(tree.by.hand)
# simple.newick.medium <- newickize(tree.by.hand, alt=TRUE)
newick.medium <- newickize(make.tree(pat.summaries[, 'count2']))
simple.newick.medium <- newickize(make.tree(pat.summaries[, 'count2']), alt=TRUE)
show.tree(newick.medium, total.snps=consistent.count[2])
```

Figure 4, i.e. [[3]]:

```
newick.strict <- newickize(make.tree(pat.summaries[, 'count3']))
show.tree(newick.strict, total.snps=consistent.count[3])
```

Figure 5, i.e. [[4]]:

```
newick.unfiltered <- newickize(make.tree(pat.summaries[, 'count4']))
show.tree(newick.unfiltered, total.snps=consistent.count[4])
```

Some other versions of the trees are included in the appendix.

Counts for all tree edges in the medium tree:

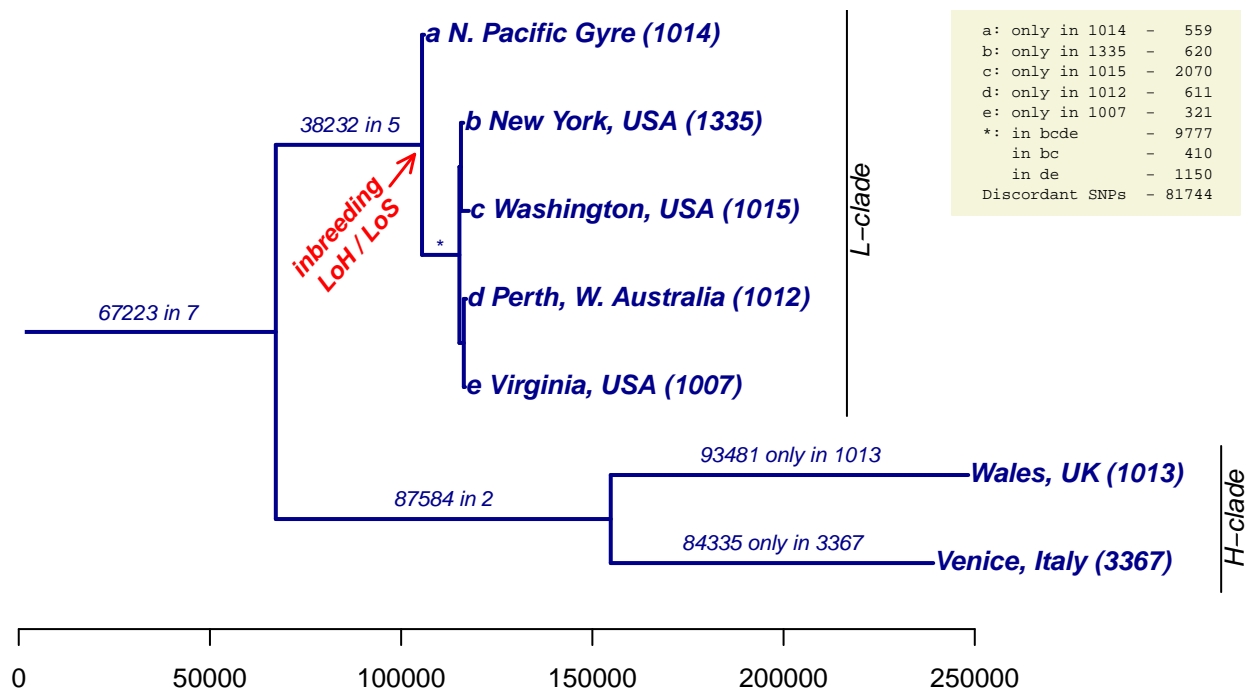


Figure 1: Proposed fig. for paper: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs.




Figure 2: Tree based on unfiltered reads and loose SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs.




Figure 3: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs.




Figure 4: Tree based on unfiltered reads and strict SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs.




Figure 5: Tree based on unfiltered reads and unfiltered SNPs. “Lengths” are numbers of shared/private SNPs all Chrs.

```
#pat.summaries[c(128,110,102,6,97,19,9,2,5,33,65,17,3),]
tree.edges <- c(128,110,102,6,97,19,9,2,5,33,65,17,3)-1
non.edges <- setdiff(0:127, tree.edges)
sg.edges <- showgroup(restrict.to=tree.edges) ; sg.edges
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 2	001	1							X	199	620	1157	2260
# 3	002	1						X		41774	84335	88149	105614
# 5	004	1					X			921	2070	2578	4608
# 6	005	2					X		X	210	410	854	1407
# 9	010	1				X				208	559	714	1231
# 17	020	1			X					47772	93481	96798	113191
# 19	022	2			X			X		55300	87584	84944	58009
# 33	040	1		X						285	611	1031	2450
# 65	100	1	X							121	321	542	2005
# 97	140	2	X	X						591	1150	1281	2144
# 102	145	4	X	X			X		X	3312	9777	18140	23189
# 110	155	5	X	X		X	X		X	33045	38232	26997	30602
# 128	177	7	X	X	X	X	X	X	X	82193	67223	46524	15186
# Total										265931	386373	369709	361896

Counts for the top 10 discordant patterns, i.e., SNPs whose sharing pattern does not match any of the bifurcations in the tree:

```
tenth <- sort(showgroup(restrict.to=non.edges)[-length(non.edges)+1, 'count2'], decreasing=T)[10]
sg.non.edges <- showgroup(restrict.to=non.edges, c2.thresh = tenth) ; sg.non.edges
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 32	037	5			X	X	X	X	X	2087	1987	1386	620
# 104	147	5	X	X			X	X	X	1372	3155	5536	10001
# 109	154	4	X	X		X	X			896	1611	1430	1738
# 112	157	6	X	X		X	X	X	X	13257	11725	8245	12202
# 115	162	4	X	X	X			X		1552	1955	1909	1014
# 118	165	5	X	X	X		X		X	1875	3928	6825	9715
# 119	166	5	X	X	X		X	X		621	1958	3252	2688
# 120	167	6	X	X	X		X	X	X	8614	16443	28402	15091
# 126	175	6	X	X	X	X	X		X	17128	14648	10156	12697
# 127	176	6	X	X	X	X	X	X		2303	2825	2133	1032
# Other	(105 rows	w/	c2	< 1611)				43236	21509	31987	45919
# Total										92941	81744	101261	112717

And percent of discordant SNPs:

```
nsge <- nrow(sg.edges)
discordv <- consistent.count - sg.edges[nsge,c('count1','count2','count3','count4')] ; discordv
```

#	count1	count2	count3	count4
# Total	92941	81744	101261	112717

```
discordv.pct <- round(discordv/consistent.count*100,1) ; discordv.pct
```

#	count1	count2	count3	count4
# Total	25.9	17.5	21.5	23.7

In short, the sharing pattern observed at 81744 or 17.5% of the 468117 medium-stringency consistent SNPs positions observed across all 7 isolates are discordant with the medium tree. (The strict tree has slightly more.)

A majority of the discordant SNPs fall into one of three patterns: 6-way sharing excluding Gyre (likely a technical artifact since the low coverage in Gyre reduces our power to detect SNPs there), or 6-way sharing excluding one of the two H-isolates (likely a reflection of sexuality in the H-clade—SNP positions in a population in Hardy-Weinberg equilibrium are fairly likely to be homozygous for the reference allele in a given individual).

```
third.biggest <- sort(showgroup(pat.summaries,6)[-8, 'count2'], decreasing=T)[3]
big.three <- showgroup(pat.summaries,6,c2.thresh = third.biggest); big.three
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 112	157	6	X	X		X	X	X	X	13257	11725	8245	12202
# 120	167	6	X	X	X		X	X	X	8614	16443	28402	15091
# 126	175	6	X	X	X	X	X		X	17128	14648	10156	12697
# Other	(4 rows	w/	c2	< 11725)				3757	4115	3399	1964
# Total										42756	46931	50202	41954

```
big.three.frac <- sum(big.three[1:3, 'count2'])/discordv$count2; big.three.frac
```

```
# [1] 0.5237816
```


I.e., 52.4% of discordant SNPs fall into one of these three categories.

Out of curiosity: what is the ratio of full genome to Chr 1 branch lengths. Except for the shortest few, generally $\approx 10\times$, as expected given the length of Chr 1:

```
# (vectors derived by editing Newick strings, and in that order)
print(
  c(Italy=86155, Wales=95697, IW=89598, Virg=330, Aust=632, VA=1296,
    Puget=2113, NY=658, PNY=480, four=10059, Gyre=568, five=39517, all=69526) /
  c(Italy=8813, Wales=9652, IW=9365, Virg=30, Aust=61, VA=19,
    Puget=207, NY=41, PNY=18, four=1005, Gyre=61, five=3912, all= 7054),
  digits=3)

# Italy Wales   IW Virg Aust   VA Puget   NY   PNY   four   Gyre   five   all
# 9.78 9.91 9.57 11.00 10.36 68.21 10.21 16.05 26.67 10.01 9.31 10.10 9.86

round(genome.length.constants()$genome.length.trunc / genome.length.constants()$chr1.length, digits=4)

# [1] 10.2879
```

9 Semi-Automated Tree-Building

Slightly formalizing the process above: Look for the bifurcation of the 7 strains that maximizes the number of shared SNPs *within* each side of the partition while minimizing the number and fraction of SNPs that are shared by subsets that include at least one strain on each side of the partition. The 2/5 split is the winner, with 6418 SNPs in conflict with that partition (16% of the 39842 SNPs not shared by all 7; Chr1 data). The runner-up places the Gyre in a group by itself (7079 = 18% in conflict).

```
treepart <- function(p.summ=pat.summaries, root=127, verbose=T, stringency='count2'){
  root.shared <- p.summ[root+1,stringency]
  df<-NULL
  for(i in 1:floor(root/2)){
    if(bitwAnd(i,root)==i && i < root-i){
      l1 <- showgroup(p.summ,subset=i,split=NULL,proper.subset=F,total=T)
      l <- l1[nrow(l1),stringency]
      r1 <- showgroup(p.summ,subset=root-i,split=NULL,proper.subset=F,total=T)
      r <- r1[nrow(r1),stringency]
      c1 <- showgroup(p.summ,subset=root,split=i,proper.subset=T,total=T)
      c <- c1[nrow(c1),stringency]
      df <- rbind(df, data.frame(pat=i,left=l,right=r,both=l+r,cross=c,all=l+r+c,ratio=c/(l+r+c),
                                best=' ',stringsAsFactors=F))
    }
  }
  df$pat<-as.octmode(df$pat)
  maxl <- which.max(df$left)
  maxr <- which.max(df$right)
  maxb <- which.max(df$both)
  minc <- which.min(df$cross)
  minr <- which.min(df$ratio)
  df$best[c(maxl,maxr,maxb,minc,minr)] <- '<'
  df$best[maxl] <- paste(df$best[maxl], 'L') # max Left
  df$best[maxr] <- paste(df$best[maxr], 'R') # max Right
  df$best[maxb] <- paste(df$best[maxb], 'B') # max Both (L+R)
  df$best[minc] <- paste(df$best[minc], 'C') # min Cross
  df$best[minr] <- paste(df$best[minr], 'O') # min ratio (Cross/(Left+Right+Cross))
  if(verbose){
    same <- all(maxl==c(maxr,maxb,minc,minr))
    cat('root:', format(as.octmode(root),width=3),
        '; shared:', root.shared,
        '. max l', format(as.octmode(df$pat[maxl]),width=3),
        ', max r', format(as.octmode(df$pat[maxr]),width=3),
        ', max both', format(as.octmode(df$pat[maxb]),width=3),
        ', min cross', format(as.octmode(df$pat[minc]),width=3),
        ', min ratio', format(as.octmode(df$pat[minr]),width=3),
        '. \nAll the same?:', same,
        '\n')
```

```

    cat('\n')
  }
  return(df)
}

```

```

treepart()

# root: 177 ; shared: 67223 . max l 077 , max r 010 , max both 022 , min cross 022 , min ratio 022 .
# All the same?: FALSE
#
# pat left right both cross all ratio best
# 1 01 622 287930 288552 112344 400896 0.2802323
# 2 02 84337 177213 261550 139346 400896 0.3475864
# 3 03 85238 104349 189587 211309 400896 0.5270918
# 4 04 2072 277815 279887 121009 400896 0.3018464
# 5 05 3102 272625 275727 125169 400896 0.3122231
# 6 06 86526 99183 185709 215187 400896 0.5367651
# 7 07 87983 97026 185009 215887 400896 0.5385112
# 8 10 561 318586 319147 81749 400896 0.2039157 < R
# 9 11 1716 279366 281082 119814 400896 0.2988655
# 10 12 85054 116885 201939 198957 400896 0.4962808
# 11 13 86743 100618 187361 213535 400896 0.5326444
# 12 14 2767 273729 276496 124400 400896 0.3103049
# 13 15 5382 271127 276509 124387 400896 0.3102725
# 14 16 87441 97314 184755 216141 400896 0.5391448
# 15 17 91097 96139 187236 213660 400896 0.5329562
# 16 20 93483 163824 257307 143589 400896 0.3581702
# 17 21 94257 94798 189055 211841 400896 0.5284188
# 18 22 265402 60429 325831 75065 400896 0.1872431 < B C O
# 19 23 266938 8065 275003 125893 400896 0.3140291
# 20 24 95731 90188 185919 214977 400896 0.5362413
# 21 25 97083 87903 184986 215910 400896 0.5385686
# 22 26 268164 4255 272419 128477 400896 0.3204746
# 23 27 271026 2732 273758 127138 400896 0.3171346
# 24 30 94222 106813 201035 199861 400896 0.4985358
# 25 31 95786 91283 187069 213827 400896 0.5333727
# 26 32 266536 17109 283645 117251 400896 0.2924724
# 27 33 270177 5553 275730 125166 400896 0.3122156
# 28 34 96722 88331 185053 215843 400896 0.5384015
# 29 35 100443 87053 187496 213400 400896 0.5323076
# 30 36 269979 2931 272910 127986 400896 0.3192499
# 31 37 278873 2084 280957 119939 400896 0.2991773
# 32 40 613 281626 282239 118657 400896 0.2959795
# 33 41 1318 271381 272699 128197 400896 0.3197762
# 34 42 85065 101681 186746 214150 400896 0.5341784
# 35 43 86121 97440 183561 217335 400896 0.5421231
# 36 44 2898 271190 274088 126808 400896 0.3163115
# 37 45 4422 267334 271756 129140 400896 0.3221284
# 38 46 87593 96322 183915 216981 400896 0.5412401
# 39 47 89792 94692 184484 216412 400896 0.5398208
# 40 50 1197 272887 274084 126812 400896 0.3163214
# 41 51 2492 269354 271846 129050 400896 0.3219039
# 42 52 85818 97951 183769 217127 400896 0.5416043
# 43 53 87755 96355 184110 216786 400896 0.5407537
# 44 54 3697 267827 271524 129372 400896 0.3227071
# 45 55 7570 266129 273699 127197 400896 0.3172818
# 46 56 88673 94747 183420 217476 400896 0.5424749
# 47 57 94104 93921 188025 212871 400896 0.5309881
# 48 60 94192 92195 186387 214509 400896 0.5350739
# 49 61 95118 88106 183224 217672 400896 0.5429638
# 50 62 266517 6162 272679 128217 400896 0.3198261
# 51 63 268442 3251 271693 129203 400896 0.3222856
# 52 64 96803 87269 184072 216824 400896 0.5408485
# 53 65 98924 85505 184429 216467 400896 0.5399580
# 54 66 270198 2099 272297 128599 400896 0.3207790
# 55 67 274929 902 275831 125065 400896 0.3119637
# 56 70 94977 88805 183782 217114 400896 0.5415719

```

```
# 57 71 96765 87113 183878 217018 400896 0.5413324
# 58 72 267758 3674 271432 129464 400896 0.3229366
# 59 73 271987 2498 274485 126411 400896 0.3153212
# 60 74 97970 85715 183685 217211 400896 0.5418138
# 61 75 103426 84756 188182 212714 400896 0.5305965
# 62 76 272458 976 273434 127462 400896 0.3179428
# 63 77 285417 323 285740 115156 400896 0.2872466 < L
```

Comparing the 5/2 split to the second-place NPG/rest split (below), the former has fewer pattern instances in conflict with the split (6418 vs 7079), as well as somewhat more random distribution of the conflicting patterns (92 vs 62 rows), whereas the 1/6 split has the majority of its conflicts (3912 of 7079, or 55%) concentrated in one pattern—the 5 NE strains. Collectively, these seem to favor the 5/2 split as the correct “history.”

```
showgroup(pat.summaries,split=strtoi('022'), subset=127, proper.subset=T, c2.thresh=100)
```

#	Pat	ShrBy	1007	1012	1013	1014	1015	3367	1335	count1	count2	count3	count4
# 4	003	2						X	X	2266	281	432	587
# 7	006	2					X	X		1282	119	261	590
# 8	007	3					X	X	X	152	146	278	557
# 11	012	2				X		X		2060	158	59	93
# 12	013	3				X		X	X	350	253	183	338
# 16	017	4				X	X	X	X	344	361	250	564
# 18	021	2		X					X	2445	154	260	402
# 20	023	3		X				X	X	3322	481	838	533
# 21	024	2		X			X			1406	178	371	625
# 22	025	3		X			X		X	197	168	333	522
# 23	026	3		X			X	X		1794	395	771	789
# 24	027	4		X			X	X	X	441	602	1120	771
# 25	030	2		X		X				2257	180	59	93
# 26	031	3		X	X				X	361	255	178	361
# 27	032	3		X	X			X		2845	237	112	86
# 28	033	4		X	X			X	X	1174	1062	725	306
# 29	034	3		X	X	X				144	116	125	219
# 30	035	4		X	X	X			X	495	529	373	503
# 31	036	4		X	X	X	X			381	367	287	211
# 32	037	5		X	X	X	X	X	X	2087	1987	1386	620
# 35	042	2		X				X		1429	117	224	394
# 39	046	3		X			X	X		103	124	367	604
# 40	047	4		X			X	X	X	93	178	485	708
# 48	057	5		X		X	X	X	X	227	231	205	324
# 51	062	3		X	X			X		2161	289	447	469
# 52	063	4		X	X			X	X	325	167	265	194
# 53	064	3		X	X		X			114	148	390	601
# 54	065	4		X	X		X		X	88	208	528	582
# 55	066	4		X	X		X	X		263	432	944	851
# 56	067	5		X	X		X	X	X	422	685	1836	1151
# 62	075	5		X	X	X	X		X	200	185	233	328
# 63	076	5		X	X	X	X	X		108	158	188	128
# 64	077	6		X	X	X	X	X	X	920	872	917	485
# 81	120	2	X		X					1009	117	116	309
# 83	122	3	X		X			X		1325	191	220	354
# 87	126	4	X		X		X	X		124	214	297	425
# 88	127	5	X		X		X	X	X	133	253	399	482
# 94	135	5	X		X	X	X		X	116	121	96	235
# 96	137	6	X		X	X	X	X	X	394	325	275	333
# 99	142	3	X	X				X		342	419	477	755
# 103	146	4	X	X			X	X		121	455	958	1795
# 104	147	5	X	X			X	X	X	1372	3155	5536	10001
# 111	156	5	X	X		X	X	X		492	681	566	735
# 112	157	6	X	X		X	X	X	X	13257	11725	8245	12202
# 113	160	3	X	X	X					306	359	421	712
# 115	162	4	X	X	X			X		1552	1955	1909	1014
# 116	163	5	X	X	X			X	X	271	263	302	316
# 117	164	4	X	X	X		X			165	603	1060	1752
# 118	165	5	X	X	X		X		X	1875	3928	6825	9715
# 119	166	5	X	X	X		X	X		621	1958	3252	2688
# 120	167	6	X	X	X		X	X	X	8614	16443	28402	15091

```
# 125 174 5 X X X X X 567 789 656 782
# 126 175 6 X X X X X X 17128 14648 10156 12697
# 127 176 6 X X X X X X 2303 2825 2133 1032
# Other ( 38 rows w/ c2 < 100 ) 5050 1815 2294 5023
# Total 89393 75065 90025 94037
```

```
showgroup(pat.summaries, split=strtoi('010'), subset=127, proper.subset=T, c2.thresh=100)
```

```
# Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 10 011 2 X 358 535 384 827
# 11 012 2 X X 2060 158 59 93
# 12 013 3 X X 350 253 183 338
# 13 014 2 X X 143 136 176 417
# 14 015 3 X X X 776 1050 757 1389
# 16 017 4 X X X X 344 361 250 564
# 25 030 2 X X 2257 180 59 93
# 26 031 3 X X X 361 255 178 361
# 27 032 3 X X X 2845 237 112 86
# 28 033 4 X X X X 1174 1062 725 306
# 29 034 3 X X X 144 116 125 219
# 30 035 4 X X X X 495 529 373 503
# 31 036 4 X X X X 381 367 287 211
# 32 037 5 X X X X X 2087 1987 1386 620
# 46 055 4 X X X X X 480 709 750 971
# 48 057 5 X X X X X 227 231 205 324
# 62 075 5 X X X X X 200 185 233 328
# 63 076 5 X X X X X 108 158 188 128
# 64 077 6 X X X X X 920 872 917 485
# 78 115 4 X X X X X 103 141 138 604
# 94 135 5 X X X X X 116 121 96 235
# 96 137 6 X X X X X 394 325 275 333
# 109 154 4 X X X X 896 1611 1430 1738
# 110 155 5 X X X X X 33045 38232 26997 30602
# 111 156 5 X X X X X 492 681 566 735
# 112 157 6 X X X X X 13257 11725 8245 12202
# 125 174 5 X X X X X 567 789 656 782
# 126 175 6 X X X X X X 17128 14648 10156 12697
# 127 176 6 X X X X X X 2303 2825 2133 1032
# Other ( 33 rows w/ c2 < 100 ) 1755 1270 1424 3544
# Total 85766 81749 59463 72767
```

Below is the full summary of shared SNPs that do *not* directly correspond to tree splits, e.g. deep coalescence, independent coincident mutations, false positives/false negatives in the shared SNP calls, loss of SNPs in hemizygous regions, etc. (Additionally, SAMTools' SNP calls exclude positions it judges to be homozygous, and I think it operates without regard to the reference sequence, so homozygous nonreference positions, while rare except in IT/Wales, often are not called SNPs by SAMTools, but are relevant for this analysis. Provided the position is called a SNP in some other isolate, the consistency filtering we've done above should recover it, but this is still worth keeping in mind when examining the data.)

First, here are SNPs that “coalesce” on the branch from the LCA of bcde, i.e., shared among some nonempty, proper subset of bcde other than bc or de. There are 8 such patterns: any of the 4 choose 3 trios plus any of the 4 pairs having exactly one of bc.

```
sg4 <- showgroup(pat.summaries, subset=strtoi('0145'), split=5, proper.subset = F)
sg4
```

```
# Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 34 041 2 X 31 85 230 368
# 37 044 2 X X 72 215 895 1809
# 38 045 3 X X 116 409 1369 1656
# 66 101 2 X X 19 33 75 314
# 69 104 2 X X 39 105 356 1196
# 70 105 3 X X 35 113 283 805
# 98 141 3 X X 65 109 208 519
# 101 144 3 X X X 193 1079 2430 4432
# 102 145 4 X X X X 3312 9777 18140 23189
```

```
# Total 3882 11925 23986 34288

sg4n <- nrow(sg4)
sg4pct <- round(sg4$count2[sg4n-1]/sg4$count2[sg4n]*100,1)
sg4pct

# [1] 82
```

So, of the 11925 SNPs found only in bcde, 82% have a sharing pattern consistent with the given tree structure.

Similarly, we analyze patterns relative to the root of the L-clade (14 patterns—any nonempty proper subset of bcde together with a):

```
sg5 <- showgroup(pat.summaries, subset=strtoi('0155'), split=8, proper.subset = F)
sg5

# Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 10 011 2 X 358 535 384 827
# 13 014 2 X X 143 136 176 417
# 14 015 3 X X 776 1050 757 1389
# 41 050 2 X 31 25 54 105
# 42 051 3 X 46 55 79 126
# 45 054 3 X X 22 79 171 292
# 46 055 4 X X 480 709 750 971
# 73 110 2 X X 19 20 50 150
# 74 111 3 X X 16 9 19 139
# 77 114 3 X X 8 38 69 365
# 78 115 4 X X 103 141 138 604
# 105 150 3 X X 47 44 78 238
# 106 151 4 X X X 33 77 71 220
# 109 154 4 X X X 896 1611 1430 1738
# 110 155 5 X X X 33045 38232 26997 30602
# Total 36023 42761 31223 38183

sg5n <- nrow(sg5)
sg5pct <- round(sg5$count2[sg5n-1]/sg5$count2[sg5n]*100,1)
```

I.e., of the 42761 SNPs found only in abcde, 89.4% have a sharing pattern consistent with the given tree structure.

Finally, how many SNPs have patterns inconsistent with the 5-2 split, i.e., include at least one strain on each side of the 5-2 split, but not shared by all 7?

```
sg7 <- showgroup(pat.summaries, subset=127, split=strtoi('022'), proper.subset=F)
sg7

# Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 4 003 2 X X 2266 281 432 587
# 7 006 2 X X 1282 119 261 590
# 8 007 3 X X 152 146 278 557
# 11 012 2 X X 2060 158 59 93
# 12 013 3 X X 350 253 183 338
# 15 016 3 X X 109 62 65 152
# 16 017 4 X X 344 361 250 564
# 18 021 2 X 2445 154 260 402
# 20 023 3 X X 3322 481 838 533
# 21 024 2 X X 1406 178 371 625
# 22 025 3 X X 197 168 333 522
# 23 026 3 X X 1794 395 771 789
# 24 027 4 X X 441 602 1120 771
# 25 030 2 X X 2257 180 59 93
# 26 031 3 X X 361 255 178 361
# 27 032 3 X X 2845 237 112 86
# 28 033 4 X X 1174 1062 725 306
# 29 034 3 X X X 144 116 125 219
# 30 035 4 X X X 495 529 373 503
# 31 036 4 X X X 381 367 287 211
# 32 037 5 X X X X 2087 1987 1386 620
```

# 35	042	2		X				X		1429	117	224	394
# 36	043	3		X				X	X	175	70	104	133
# 39	046	3		X			X	X		103	124	367	604
# 40	047	4		X			X	X	X	93	178	485	708
# 43	052	3		X		X		X		101	11	27	22
# 44	053	4		X		X		X	X	46	38	36	56
# 47	056	4		X		X	X	X		17	50	65	88
# 48	057	5		X		X	X	X	X	227	231	205	324
# 49	060	2		X	X					1651	98	247	388
# 50	061	3		X	X				X	169	67	98	115
# 51	062	3		X	X			X		2161	289	447	469
# 52	063	4		X	X			X	X	325	167	265	194
# 53	064	3		X	X		X			114	148	390	601
# 54	065	4		X	X		X		X	88	208	528	582
# 55	066	4		X	X		X	X		263	432	944	851
# 56	067	5		X	X		X	X	X	422	685	1836	1151
# 57	070	3		X	X	X				123	21	18	24
# 58	071	4		X	X	X			X	30	17	14	28
# 59	072	4		X	X	X		X		158	50	35	31
# 60	073	5		X	X	X		X	X	155	89	72	38
# 61	074	4		X	X	X	X			28	51	60	116
# 62	075	5		X	X	X	X		X	200	185	233	328
# 63	076	5		X	X	X	X	X		108	158	188	128
# 64	077	6		X	X	X	X	X	X	920	872	917	485
# 67	102	2	X					X		887	98	126	351
# 68	103	3	X					X	X	83	25	51	143
# 71	106	3	X				X	X		42	63	127	377
# 72	107	4	X				X	X	X	36	64	105	330
# 75	112	3	X			X		X		55	12	16	26
# 76	113	4	X			X		X	X	18	8	8	66
# 79	116	4	X			X	X	X		5	8	19	101
# 80	117	5	X			X	X	X	X	42	40	35	241
# 81	120	2	X		X					1009	117	116	309
# 82	121	3	X		X				X	87	19	29	73
# 83	122	3	X		X			X		1325	191	220	354
# 84	123	4	X		X			X	X	162	85	77	124
# 85	124	3	X		X		X			66	81	167	400
# 86	125	4	X		X		X		X	41	79	142	283
# 87	126	4	X		X		X	X		124	214	297	425
# 88	127	5	X		X		X	X	X	133	253	399	482
# 89	130	3	X		X	X				73	12	9	27
# 90	131	4	X		X	X			X	17	5	5	52
# 91	132	4	X		X	X		X		86	27	19	38
# 92	133	5	X		X	X		X	X	56	31	16	52
# 93	134	4	X		X	X	X			15	24	25	143
# 94	135	5	X		X	X	X		X	116	121	96	235
# 95	136	5	X		X	X	X	X		41	71	63	106
# 96	137	6	X		X	X	X	X	X	394	325	275	333
# 99	142	3	X	X				X		342	419	477	755
# 100	143	4	X	X				X	X	46	55	86	190
# 103	146	4	X	X			X	X		121	455	958	1795
# 104	147	5	X	X			X	X	X	1372	3155	5536	10001
# 107	152	4	X	X		X		X		27	21	24	67
# 108	153	5	X	X		X		X	X	35	32	30	96
# 111	156	5	X	X		X	X	X		492	681	566	735
# 112	157	6	X	X		X	X	X	X	13257	11725	8245	12202
# 113	160	3	X	X	X					306	359	421	712
# 114	161	4	X	X	X				X	72	88	102	207
# 115	162	4	X	X	X			X		1552	1955	1909	1014
# 116	163	5	X	X	X			X	X	271	263	302	316
# 117	164	4	X	X	X		X			165	603	1060	1752
# 118	165	5	X	X	X		X		X	1875	3928	6825	9715
# 119	166	5	X	X	X		X	X		621	1958	3252	2688
# 120	167	6	X	X	X		X	X	X	8614	16443	28402	15091
# 121	170	4	X	X	X	X				21	26	21	69
# 122	171	5	X	X	X	X			X	30	29	18	70
# 123	172	5	X	X	X	X		X		105	95	59	86

```
# 124 173 6 X X X X X X 140 93 74 114
# 125 174 5 X X X X X 567 789 656 782
# 126 175 6 X X X X X X 17128 14648 10156 12697
# 127 176 6 X X X X X X 2303 2825 2133 1032
# 128 177 7 X X X X X X 82193 67223 46524 15186
# Total 171586 142288 136549 109223
```

```
sg7n <- nrow(sg7)
sg7pct <- round(sg7$count2[sg7n-1]/sg7$count2[sg7n]*100,1)
sg7pct

# [1] 47.2
```

A more compact version of that table, showing only the larger counts:

```
thresh <- signif(.02 * sg7$count2[sg7n],1)
thresh

# [1] 3000

showgroup(pat.summaries, subset=127, split=stoi('022'), proper.subset=F, c2.thresh = thresh)

# Pat ShrBy 1007 1012 1013 1014 1015 3367 1335 count1 count2 count3 count4
# 104 147 5 X X X X X 1372 3155 5536 10001
# 112 157 6 X X X X X 13257 11725 8245 12202
# 118 165 5 X X X X X 1875 3928 6825 9715
# 120 167 6 X X X X X 8614 16443 28402 15091
# 126 175 6 X X X X X 17128 14648 10156 12697
# 128 177 7 X X X X X 82193 67223 46524 15186
# Other ( 87 rows w/ c2 < 3000 ) 47147 25166 30861 34331
# Total 171586 142288 136549 109223
```

So, of the 142288 SNPs found both in the L- and H-clades, 47.2% have a sharing pattern consistent with the given tree structure, i.e., are found in all 7 isolates. Among the others, three patterns dominate—(i) the 6-way pattern excluding the Gyre is the largest, plausibly explained by 7-way sharing from which the Gyre drops out due to low coverage/high error rate, (ii) the 6-way excluding Italy, and (iii) ditto for Wales. Origin of the later two cases is unclear, but may partly reflect Hardy-Weinberg—some positions that are *population-level* SNPs in those isolates will be homozygous-reference in the CCMP founder cell for IT or Wales. If I take the 7-way shared SNP count (69526) as a surrogate approximating the number of population-level SNPs in either IT or Wales that are shared with the L-clade, then I might expect, based on HWE, roughly half that number to be lost (become homozygous) in IT, and a similar number in Wales. However, the observed counts of these positions are lower by $\approx 20K$ than I might have guessed from HWE, perhaps suggesting that IT and Wales are distinct populations, each with a pool of many thousand private polymorphisms.

In aggregate:

```
untreelike <-
  sg7$count2[sg7n]-sg7$count2[sg7n-1] +
  sg5$count2[sg5n]-sg5$count2[sg5n-1] +
  sg4$count2[sg4n]-sg4$count2[sg4n-1]
untreelike

# [1] 81742

consistent.count[2]

# [1] 468117

unpct <- round(untreelike/consistent.count[2]*100,1)
unpct

# [1] 17.5
```

I.e., 81742 or 17.5% of the 468117 consistent SNPs identified (by criterion 2) across all 7 isolates are discordant with the assumed tree.

# 16	017	4				X	X	X	X	NA	403	NA	NA
# 17	020	1			X					NA	93546	NA	NA
# 19	022	2			X			X		NA	87716	NA	NA
# 20	023	3			X			X	X	NA	481	NA	NA
# 24	027	4			X		X	X	X	NA	627	NA	NA
# 28	033	4			X	X		X	X	NA	1037	NA	NA
# 30	035	4			X	X	X		X	NA	503	NA	NA
# 32	037	5			X	X	X	X	X	NA	1974	NA	NA
# 33	040	1		X						NA	594	NA	NA
# 38	045	3		X				X		NA	406	NA	NA
# 46	055	4		X		X	X		X	NA	703	NA	NA
# 55	066	4		X	X		X	X		NA	425	NA	NA
# 56	067	5		X	X		X	X	X	NA	674	NA	NA
# 64	077	6		X	X	X	X	X	X	NA	886	NA	NA
# 97	140	2	X	X						NA	1130	NA	NA
# 99	142	3	X	X				X		NA	424	NA	NA
# 101	144	3	X	X			X			NA	1037	NA	NA
# 102	145	4	X	X			X		X	NA	9812	NA	NA
# 103	146	4	X	X			X	X		NA	456	NA	NA
# 104	147	5	X	X			X	X	X	NA	3197	NA	NA
# 109	154	4	X	X		X	X			NA	1658	NA	NA
# 110	155	5	X	X		X	X		X	NA	38094	NA	NA
# 111	156	5	X	X		X	X	X		NA	684	NA	NA
# 112	157	6	X	X		X	X	X	X	NA	11714	NA	NA
# 115	162	4	X	X	X			X		NA	1973	NA	NA
# 117	164	4	X	X	X		X			NA	586	NA	NA
# 118	165	5	X	X	X		X		X	NA	3981	NA	NA
# 119	166	5	X	X	X		X	X		NA	1939	NA	NA
# 120	167	6	X	X	X		X	X	X	NA	16361	NA	NA
# 125	174	5	X	X	X	X	X			NA	762	NA	NA
# 126	175	6	X	X	X	X	X		X	NA	14747	NA	NA
# 127	176	6	X	X	X	X	X	X		NA	2758	NA	NA
# 128	177	7	X	X	X	X	X	X	X	NA	67423	NA	NA
# Other	(88 rows	w/	c2	<	400)			NA	9895	NA	NA
# Total										NA	468117	NA	NA

Tree partition analysis (and how to pluck out only the best rows based on 3 smallest cross counts and “best” criteria):

```
tp <- treepart(boot.summaries, root=127) ; tp

# root: 177 ; shared: 67423 . max l 077 , max r 010 , max both 022 , min cross 022 , min ratio 022 .
# All the same?: FALSE
#   pat  left  right  both  cross  all  ratio  best
# 1  01    656 287705 288361 112333 400694 0.2803461
# 2  02   84134 177289 261423 139271 400694 0.3475745
# 3  03   85067 104362 189429 211265 400694 0.5272477
# 4  04    2148 277617 279765 120929 400694 0.3017989
# 5  05    3226 272454 275680 125014 400694 0.3119937
# 6  06   86410 99164 185574 215120 400694 0.5368685
# 7  07   87904 96999 184903 215791 400694 0.5385431
# 8  10     539 318600 319139 81555 400694 0.2035344
# 9  11    1710 279265 280975 119719 400694 0.2987791
# 10 12   84808 117015 201823 198871 400694 0.4963164
# 11 13   86509 100637 187146 213548 400694 0.5329453
# 12 14    2827 273636 276463 124231 400694 0.3100396
# 13 15    5517 271004 276521 124173 400694 0.3098948
# 14 16   87294 97317 184611 216083 400694 0.5392719
# 15 17   91056 96122 187178 213516 400694 0.5328655
# 16 20   93546 163624 257170 143524 400694 0.3581885
# 17 21   94342 94582 188924 211770 400694 0.5285080
# 18 22  265396 60374 325770 74924 400694 0.1869856 < B C O
# 19 23  266950 8063 275013 125681 400694 0.3136583
# 20 24   95866 89864 185730 214964 400694 0.5364792
# 21 25   97225 87575 184800 215894 400694 0.5388002
# 22 26  268241 4162 272403 128291 400694 0.3201720
# 23 27  271124 2635 273759 126935 400694 0.3167879
# 24 30   94278 106731 201009 199685 400694 0.4983479
# 25 31   95861 91068 186929 213765 400694 0.5334869
# 26 32  266498 17167 283665 117029 400694 0.2920658
# 27 33  270129 5536 275665 125029 400694 0.3120311
# 28 34   96847 88094 184941 215753 400694 0.5384483
```

```
# 29 35 100593 86777 187370 213324 400694 0.5323863
# 30 36 270023 2888 272911 127783 400694 0.3189042
# 31 37 278960 2012 280972 119722 400694 0.2987866
# 32 40 594 281586 282180 118514 400694 0.2957718
# 33 41 1330 271336 272666 128028 400694 0.3195156
# 34 42 84841 101737 186578 214116 400694 0.5343629
# 35 43 85927 97505 183432 217262 400694 0.5422143
# 36 44 2979 271061 274040 126654 400694 0.3160866
# 37 45 4543 267221 271764 128930 400694 0.3217667
# 38 46 87474 96343 183817 216877 400694 0.5412534
# 39 47 89710 94700 184410 216284 400694 0.5397735
# 40 50 1156 272926 274082 126612 400694 0.3159818
# 41 51 2463 269362 271825 128869 400694 0.3216145
# 42 52 85545 98025 183570 217124 400694 0.5418699
# 43 53 87487 96444 183931 216763 400694 0.5409689
# 44 54 3750 267776 271526 129168 400694 0.3223607
# 45 55 7685 266057 273742 126952 400694 0.3168303
# 46 56 88509 94786 183295 217399 400694 0.5425562
# 47 57 94032 93945 187977 212717 400694 0.5308714
# 48 60 94228 92098 186326 214368 400694 0.5349918
# 49 61 95167 87908 183075 217619 400694 0.5431052
# 50 62 266451 6240 272691 128003 400694 0.3194532
# 51 63 268395 3266 271661 129033 400694 0.3220238
# 52 64 96941 86988 183929 216765 400694 0.5409739
# 53 65 99093 85214 184307 216387 400694 0.5400305
# 54 66 270234 2044 272278 128416 400694 0.3204840
# 55 67 275014 839 275853 124841 400694 0.3115619
# 56 70 95013 88695 183708 216986 400694 0.5415255
# 57 71 96812 86957 183769 216925 400694 0.5413732
# 58 72 267656 3757 271413 129281 400694 0.3226427
# 59 73 271878 2538 274416 126278 400694 0.3151482
# 60 74 98097 85506 183603 217091 400694 0.5417875
# 61 75 103598 84516 188114 212580 400694 0.5305295
# 62 76 272437 970 273407 127287 400694 0.3176663
# 63 77 285475 288 285763 114931 400694 0.2868299 < L
```

```
otp <- order(tp[, 'cross'])[1:3] # 3 smallest 'cross' counts
btp <- which(tp[, 'best'] != '') # 'best' by Left/Right/Both/Cross/ratio
toptp <- unique(c(otp, btp, 18, 8)) # above, plus 5/2, 6/1 splits
print(tp[toptp,]) # show the winners
```

```
# pat left right both cross all ratio best
# 18 22 265396 60374 325770 74924 400694 0.1869856 < B C O
# 8 10 539 318600 319139 81555 400694 0.2035344 < R
# 1 01 656 287705 288361 112333 400694 0.2803461
# 63 77 285475 288 285763 114931 400694 0.2868299 < L
```

Now repeat the above nboot times, and summarize results:

```
nboot <- params$nboot # default from params set in section 2
nboot <- ((nboot+2) %/% 4) * 4 + 1 # summary is cleaner if n mod 4 == 1, so int median/quartiles
cat('***\n*** Doing', nboot, 'bootstrap replicates.\n***\n')

# ***
# *** Doing 101 bootstrap replicates.
# ***

bcor <- numeric(nboot)
b52cross <- integer(nboot)
b61cross <- integer(nboot)
brev <- logical(nboot)
for(i in 1:nboot){
  boot.sample <- sample(0:127, n2, replace=T, prob=pattern.counts[[2]][,2])
  boot.count <- mytable(boot.sample, c(0, 127))
  boot.counts <- list(NULL, boot.count, NULL) # dummy list with just c2 summaries
  boot.summaries <- pat.summary(boot.counts)
  tp <- treepart(boot.summaries, root=127, verbose=F)
  bcor[i] <- cor(pattern.counts[[2]][,2], boot.counts[[2]][,2]) # just curious - how correlated are they?
  b52cross[i] <- tp[18, 'cross']
  b61cross[i] <- tp[8, 'cross']
  brev[i] <- (b52cross[i] > b61cross[i])
```

```

if(brev[i]){
  # show the unexpected ones; probably breaks w/ cache
  otp <- order(tp[, 'cross'])[1:3]
  btp <- which(tp[, 'best'] != '')
  totp <- unique(c(otp, btp, 18, 8))
  print(tp[totp,])
}
}
# summarize:
corsummary <- t(as.matrix(c(summary(bcor), sd=sd(bcor))))
row.names(corsummary) <- 'bcor'
bdelta <- b61cross-b52cross
brevp <- 100*brev # make it percent reversed instead of logical
thesummary <- rbind(summary(b52cross), summary(b61cross), summary(c(bdelta)), summary(brevp))
row.names(thesummary) <- c('b52cross', 'b61cross', 'b61-b52', '% rev')
thesummary <- cbind(thesummary, sd=c(sd(b52cross), sd(b61cross), sd(bdelta), sd(brevp)))

```

SUMMARY: In 101 bootstrap replicates, we saw 0 samples with the 6/1 split having fewer conflicts than the 5/2 split, and the minimum separation between them was ≈ 20 sigma, hence highly statistically significant.

```

# 'opt' hacking is trying to force knitr to show more digits of bcor in summary, as Rstudio does, but
# it still fails... Bottom line is the correlation seems to be > .999 in all samples, rounds to 1.0,
# as seen in this run of 1001 samples cut/paste from Rstudio:
#           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.      sd
# bcor      " 0.9998" " 0.9999" " 0.9999" " 0.9999" " 1" " 1" " 0.00003462"
# > max(bcor)
# [1] 0.9999915
o.opts <- options(digits=7,width=127)
format(rbind(corsummary, thesummary), scientific=F, digits=4, drop0trailing=T)

#           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
# bcor      " 0.999979714" " 0.999991809" " 0.999994085" " 0.999993422" " 0.999996278" " 0.9999
# b52cross  "74410"      "74883"      "75078"      "75043.148514851" "75192"      "75542"
# b61cross  "81022"      "81564"      "81761"      "81748.564356436" "81937"      "82333"
# b61-b52   " 5857"      " 6524"      " 6694"      " 6705.415841584" " 6859"      " 7499"
# % rev     " 0"        " 0"        " 0"        " 0"        " 0"        " 0"
#           sd
# bcor      " 0.000003645"
# b52cross  " 231.908446855"
# b61cross  " 255.187594363"
# b61-b52   " 297.27442767"
# % rev     " 0"

options(o.opts)

```

Based on this, it is reasonable to claim that we are confident that the tree topology is as shown in the earlier figures, with the exception of the exact order of the splits with the 4 NE coastal isolates.

10 Notes

This section is a random brain dump of limitations of the current analysis, ideas for improvements, etc. In the main, these may not be worth doing, unless we see significant holes or get pushed by reviewers, etc, but I wanted to catalog before we forget them.

Noise: Various sources of “noise” in the data:

1. Read errors, low read depth — perhaps fixed by medium/strict thresholding
2. Deep coalescence
3. Skew because 1335 is the reference. (Julie notes we could partially fix this by remapping based on discovered SNPs, tho that wouldn't fix gross misassembly in 1335, e.g. collapsed or misordered tandem duplicates, or segments missing in 1335 that are present in one or more other strains, etc.; much harder to fix those, let's just hope they are rare...)

4. Varying error rates and sequencing depth among the 7. E.g., plausibly the 1000 SNPs shared by 4 but not by Gyre are a result of lower read depth (we missed a SNP that is actually present) and/or higher error rates (causing a position to appear inconsistent in gyre) in the gyre data. I can't think of a way to correct for this effect. It might be possible, perhaps by simulation, to estimate the size of the effect and see whether it could explain ≈ 1000 SNPs.
5. Varying numbers of founder cells in the sequencing cultures. (Again, I made some attempts at modeling this, but nothing very satisfactory yet.)
6. Tri-allelic positions where stochastic fluctuation in sequence sampling promotes the rare allele to prominence. (Julie replies: "isn't this the same as more than one founder cell? If they are diploid there should only ever be two alleles, unless there were random and very rare, thus unlikely, trisomy events?" I agree, but it is a concrete example of an effect of multiple founders that might be important. Not sure this is the most important such effect...)
7. Gaps/indels - alignments are likely to be of lower quality in the vicinity of an indel, so, maybe lower coverage/more SNPs. We ignored them. Does this add any systematic bias? e.g. if one strain had more indels than another, would this confound other analyses? unclear. Julie suggested a paper titled "Barking up the wrong tree-length: yada yada yada gap penalties"; maybe relevant?

Other Items/Potential To Dos:

1. any spacial structure to various sub-classes?
2. after top level split, should I reanalyze halves of partition in isolation? said another way, I think the tree-building is sensible, but not sure it's optimal.
3. if we believe no sex, then I think gain of SNP should be more common than loss of SNP, since the later can only happen by (a) mutation reverting to reference, (b) second mutation matching nonreference, (c) homologous repair (look for blocks of LOH), or (d) false negative e.g. from low read depth. Does tree-building appropriately weight the gain vs loss cases? (Does it even care?)
4. should we weight coding and/or nonsynonymous SNPs more heavily? Julie says "you do not want to weight the coding or nonsynonymous/coding SNPs because for time you want the more clock-like neutral mutations." I.e., I got this backwards. Maybe should redo tree based on noncoding SNPs only.
5. We could also do an actual parsimony analysis based on 2-state model (homozygous-ref vs not), but I'm not quite sure how to handle this in a mixed sex/nosex case.
6. Might be interesting to look at sharing just within (shared?) deserts. Given tree model above and expectation that bottleneck followed split of H- from L-clades, I would expect little or no sharing of L-clade desert SNPs with H-clade; sharing between It/Wales might suggest "desert" is actually a region under strong purifying selection (e.g. a gene); sharing/non-sharing within L-clade deserts might suggest more about evo history of the 5.

11 Appendix: Old Trees, etc.

Tangents, old stuff of historical interest at best, etc..

11.1 HWE Sharing

Tangent: As a function of nonref allele freq, assuming HWE, what is probability that nonref allele will be seen in k strains, $0 \leq k \leq 4$ (Fig 6).

```
myfigpath.h <- paste(getwd(), '/figs-knitr/', sep='')
```




Figure 6: Sharing Probability

```

p <- (0:20)/20
q <- 1-p
r <- 2*p*q+p^2
plot(p, 1*q^0*r^4, type='b', pch='4', ylab="share prob")
points(p, 4*q^2*r^3, type='b', pch='3')
points(p, 6*q^4*r^2, type='b', pch='2')
points(p, 4*q^6*r^1, type='b', pch='1')
points(p, 1*q^8*r^0, type='b', pch='0')

```

11.2 Old Tree Stuff

All based on un-q-filtered reads.

The first pass at the tree analysis was the Chr1 tree, *loose criteria* (c1); it is rendered via <http://iubio.bio.indiana.edu/treeapp/treeprint-form.html> as Fig 7, and in newick format is:

```

newick.chr1.loose <- '(((tp3367_Italy:4551,tp1013_Wales:4954):5920,(((tp1007_Virginia:10,tp1012_Australia:29):9,
cat.hardwrap(newick.chr1.loose)

# (((tp3367_Italy:4551,tp1013_Wales:4954):5920,(((tp1007_Virginia:10,tp1012_Austra
# lia:29):9,(tp1015_Puget_Sound:90,tp1335_NY:13):11):320,tp1014_Gyre:22):3484):859
# 3,outgroup:0);

```

Chr 1 tree based on *medium criteria* (c2) has exactly the same topology is, although the branch lengths are different. As noted earlier, the length of the branch labeled “*” is probably inflated by lower coverage and higher error rate in 1014, which may mask further legitimate sharing between it and the other L-isolates. The branch lengths among the other 4 are too short for its topology to be convincing without a more rigorous analysis (e.g., a bootstrap test).

Chr1 tree, medium criteria, in newick format:

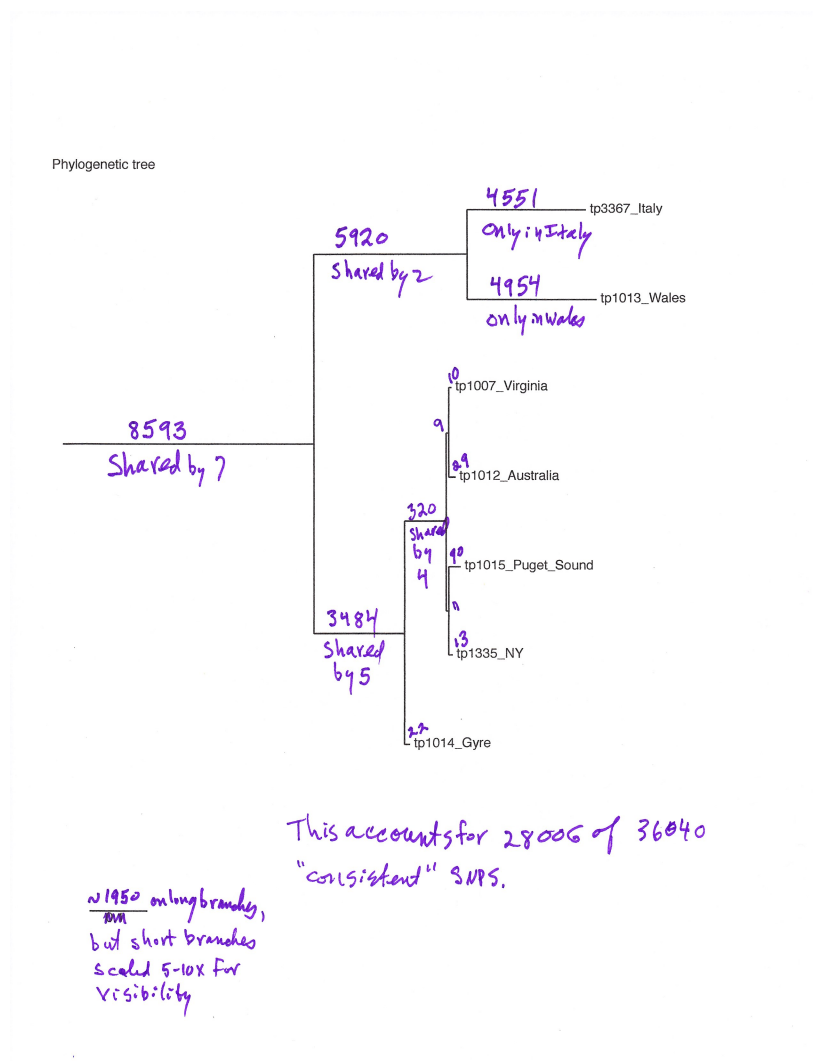


Figure 7: Inferred Tree, based on Chr1, un-q-filtered reads, loose criteria. (Note: to visually resolve the edges among the 5, their lengths were scaled by 5x – 10x in this figure, but not in the newick description shown in the text.)

```
newick.chr1.med <- '(((tp3367_Italy:8813,tp1013_Wales:9652):9365,(((e_tp1007_Virginia:30,d_tp1012_Australia:61):19,(c_tp1015_Puget_Sound:207,b_tp1335_NY:41):18):1005,a_tp1014_Gyre:61):3912):7054,outgroup:0);
cat.hardwrap(newick.chr1.med)

# (((tp3367_Italy:8813,tp1013_Wales:9652):9365,(((e_tp1007_Virginia:30,d_tp1012_Australia:61):19,(c_tp1015_Puget_Sound:207,b_tp1335_NY:41):18):1005,a_tp1014_Gyre:61):3912):7054,outgroup:0);
```

NOTE: In early code, tree was not being recalculated; it was defined by constants in the following code chunk, hand-copied from the analysis above.

```
# tree parameters as nested lists
# Internal nodes have subtrees sub1/2 and length
# Root has sub1/2, but no length
# Leaves have where, length, optionally, id, alt, nb. (Omit id for 'outgroup'. Use 'alt' for less formal
# labeling in cartoon version; it defaults to 'where'. Use 'nb' to add abcde annotations for legend.)
# This hand-made version is now subsumed by make.tree; retained for comparison
tree.by.hand <-
  list(
    sub1 = list(
      sub1 = list(
        sub1 = list(id=3367, length=8813, where='Venice, Italy', alt='Venice'),
        sub2 = list(id=1013, length=9652, where='Wales, UK'),
        length=9365),
      sub2 = list(
        sub1 = list(
          sub1 = list(
            sub1 = list(id=1007, length=30, nb='e', where='Virginia, USA'),
            sub2 = list(id=1012, length=61, nb='d', where='Perth, W. Australia', alt='Perth'),
            length=19),
          sub2 = list(
            sub1 = list(id=1015, length=207, nb='c', where='Washington, USA', alt='Puget Sound'),
            sub2 = list(id=1335, length=41, nb='b', where='New York, USA', alt='NY'),
            length=18),
          length=1005),
        sub2 = list(id=1014, length=61, nb='a', where='N. Pacific Gyre'),
        length=3912),
        length=7054),
      sub2 = list(length=0, where='outgroup')
    )
  )

# historical, format example, and debug help:
oldwick.medium <- '(((CCMP3367_Italy:8813,CCMP1013_Wales:9652):9365,(((e_CCMP1007_Virginia:30,d_CCMP1012_Australia:61):19,(c_CCMP1015_Puget_Sound:207,b_CCMP1335_NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
# with simpler labeling for cartoon
simple.oldwick.medium <- '(((Italy:8813,Wales:9652):9365,(((Virginia:30,Australia:61):19,(Puget:207,NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
cat.hardwrap(oldwick.medium)

# (((CCMP3367_Italy:8813,CCMP1013_Wales:9652):9365,(((e_CCMP1007_Virginia:30,d_CCMP1012_Australia:61):19,(c_CCMP1015_Puget_Sound:207,b_CCMP1335_NY:41):18):1005,a_CCMP1014_NPG:61):3912):7054,outgroup:0);

cat.hardwrap(simple.oldwick.medium)

# (((Italy:8813,Wales:9652):9365,(((Virginia:30,Australia:61):19,(Puget:207,NY:41):18):1005,Gyre:61):3912):7054,outgroup:0);
```

Two other versions of the tree, for possible use in figs in the main paper.

Figure 8: [** as of 10/4/2015, this fig and next have stray stars on virginia, wales labels; probably due to hacking with commas in newick; not worth fixing unless we resurrect these trees for some purpose, but if so, see use of newick.name.undo in show.tree as probable fix. **]

```
tree.scale <- ifelse(which.snp.tables(string.val=F)[1]=='Chr1', 1, 10)
tree.x.lim <- 3e4 * tree.scale
the.simple.tree <- read.tree(text=simple.newick.medium)
plot(the.simple.tree, x.lim = tree.x.lim)
axis(1)
```

Figure 9:

```
plot(the.simple.tree, x.lim = tree.x.lim)
axis(1, (0:4)*7000*tree.scale, (0:4)*7000*tree.scale)
```

At some much earlier point, Tony ran the whole-genome version of the then-current code above, and manually entered tree branch lengths/legend for the resulting tree, shown in Fig 10. Code above can now automatically generate




Figure 8: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs. (no edge labels, nolegend)




Figure 9: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs all Chrs. (no edge labels, no legend, short scale bar)

such a tree, but retain the following for comparison. The basic story seems clear—same topology and branch lengths scaled by about 10x, which is completely reasonable given that Chr1 is about 10% of the genome. Note that this tree is not being recalculated; it is defined by constants in the following code chunk.

```
fullgenome.newick.medium <- '(((3367_Italy:86155,1013_Wales:95697):89598,((e_1007_VA:330,d_1012_Australia:632):1296,(c_1015_WA:2
cat.hardwrap(fullgenome.newick.medium)

# (((3367_Italy:86155,1013_Wales:95697):89598,((e_1007_VA:330,d_1012_Australia:63
# 2):1296,(c_1015_WA:2113,b_1335_NY:658):480):10059,a_1014_NPG:568):39517):69526,o
# utgroup:0);

legend.text <- c('a: only in 1014 ',
                'b: only in 1335 ',
                'c: only in 1015 ',
                'd: only in 1012 ',
                'e: only in 1007 ',
                '*: shared by bcde',
                '  shared by b/c ',
                '  shared by d/e ')
)
fullgenome.tree.x.lim <- 300000
fullgenome.counts <- c( 568, 658, 2113, 632, 330, 10059, 480, 1296 )
fullgenome.legend.text <- paste(legend.text, format(fullgenome.counts,width=5), sep=' - ')
fullgenome.tree.labels <- list( ## x,y,text
  41000,3.63,'69526\nshared by 7',
  90000,5.75,'39517\nby 5 (*)',
  115000,1.5, '89598\nshared by 2',
  210000,2.0, '95697 only\nin Wales',
  210000,1.0, '86155 only\nin Italy',
  113500,4.6, '*' )
```

Figure 10:

```
library(ape)
the.fullgenome.tree <- read.tree(text=fullgenome.newick.medium)
plot(the.fullgenome.tree, x.lim = fullgenome.tree.x.lim)
axis(1) # ; axis(2) useful only for placing labels
opar <- par(family='mono',cex=.8)
legend('topright', legend=fullgenome.legend.text)
par(opar)
for(i in seq(1,length(fullgenome.tree.labels)-2,by=3)){
  text(fullgenome.tree.labels[[i]], fullgenome.tree.labels[[i+1]], fullgenome.tree.labels[[i+2]])
}
```




Figure 10: Tree based on unfiltered reads and medium SNP filters. “Lengths” are numbers of shared/private SNPs genome-wide. (By-hand legacy version)

